Orchestron User Documentation

Compiled from markdown documentation

Generated on 12 November 2026 at 23:47

Table of Contents

Orchestron User Documentation 4
Getting Started 5
Instrument Design 7
Patch Toolbar and Instrument Tabs 9
Opcode Catalog and Integrated Documentation 13
Graph Editor 20
Input Formula Assistant 27
GEN Table Editor (GEN Meta-Opcode) 31
If, Switch and Drumset 37
Monophonic Legato 39
Defining Performance Controls 41
Runtime Panel and Compilation Workflow 43
Instrument Import / Export and CSD Export 48
Supported Opcodes 50
Performance 62
Instrument Rack and Engine Transport 65
Audio mixer and routing 69
Using Performance Controls 78
MIDI Setup and Inputs 82
Piano Rolls 84
MIDI Controllers 88
Multitrack Arranger 91
Patterns and Reusable Phrases 111
Sequencer Timing 113
Melodic Sequencers and Step Editing 114
Drummer Sequencers 118
Score Sequencers 122
Arpeggiators 126
Controller Sequencers 129
Instrument Automation Clips 133
Live Status and Troubleshooting 136
Performance Import / Export 139
Supplied Instruments and Effects 142
Supplied Instruments 143
Physical Steel-String Guitar 145
Ambient Techno Instruments 148
HSB Mono 152
Enough and Logan 153
Goa Trance FM Lead 155

Standard Effects . . . 156
Tube Overdrive . . . 157
Stereo Delay . . . 159
Stereo Flanger . . . 161
Autopanner . . . 163
Parametric Equalizer . . . 165
Stereo Widener, Bit Crusher, and Harmonizer . . . 167
Granular Capture: Stutter, Freeze and Grains . . . 169
Stereo Moog Filter . . . 172
Stereo Auto Wah Wah . . . 174
Trance Gate . . . 176
Configuration . . . 179
GUI Language and Integrated Help . . . 180
Audio Engine Settings (Config Page) . . . 184
Browser-Clock Latency . . . 190
Saving, Workspace Recovery and Defaults . . . 193
Appendix: Optional Authoring Tools . . . 194
Agent CLI . . . 195
Orchestron Patch Creator Skill . . . 196
Orchestron Performance Creator Skill . . . 197
Orchestron Composition Creator Skill . . . 199
Sung Melody Transcription Skill . . . 204

Orchestron User Documentation

Navigation: Next

Orchestron lets you build instruments, arrange music and perform it in your browser. You can start with a built-in instrument, import a supplied sound, or design your own graph of Csound opcodes.

Start here

Follow Getting Started to make a sound and save a short performance. No external MIDI keyboard or existing instrument library is required. If Orchestron is not running yet, use the installation and startup instructions.

The app has three working areas:

- **Instrument Design** creates, edits and saves individual instruments and effects.
- **Perform** assembles saved instruments into a rack, mixes them and plays notes, patterns and songs.
- **Config** adjusts audio-engine and browser playback settings. Language and help are available throughout the app.

Read Instrument Design before Performance when learning the whole app. For ready-made sounds, go directly to Supplied Instruments and Effects. Configuration and the optional authoring tools can be consulted when needed.

Manual contents

- Getting Started
- Instrument Design
- Performance
- Supplied Instruments and Effects
- Configuration
- Appendix: Optional Authoring Tools

Musical terms

Term	Meaning in Orchestron
Instrument or patch	A saved graph that generates or processes sound. “Patch” is used in some controls and file formats.
Rack instance	One use of a saved instrument in a performance, with its own channel, controls and mixer settings.
Performance	A saved setup containing instruments, routes, mixer settings and musical material.
Track	A lane in the arranger assigned to a source of notes or controller values.
Pattern	Musical content such as a melody, drum rhythm or controller curve.
Clip	A window that places patterns or phrases on the song timeline; trimming the window retains the underlying content.
Arrangement	The song order and timing of clips across tracks.
Audition	Temporary playback used to hear an instrument or musical draft before saving it.
Workspace	The current editing context. Some workspace state is recovered after reload, but it is not a substitute for saving library entries.

For backups and sharing, see saving and recovery, instrument files, and performance files.

Navigation: Next

Getting Started

Navigation: Up | Previous | Next

This walkthrough creates a playable instrument and a short song using only the browser. Start with Orchestron open, your speakers or headphones at a modest level, and the GUI language set to English if you want labels to match this manual. An empty library is fine.

Create your first instrument

1. Open **Instrument Design** and choose **New**.
2. Choose **Playable instrument**, then **Create**. This built-in starter contains an oscillator, an envelope and a stereo output; it does not require a saved template.
3. Enter **First Sound** as the patch name. Keep the default Melody instrument type.
4. Choose **Compile**. A successful compile shows that the graph is valid; it does not save or start it.

See the patch toolbar if you want to import an existing sound instead.

Hear and save it

1. Open **Audition** in the graph header.
2. Choose **Isolated**, leave the test MIDI note at **60**, and choose **Prepare audition**. When preparation finishes, choose **Stop and audition**. The preview sends the test note automatically.
3. Choose **Close audition · transport stays stopped** when finished. Audition restores the previous setup with transport stopped.
4. Choose **Save** to add First Sound to your instrument library.

The draft audition guide explains each preview control and how to audition a changed instrument in a performance.

Build a short performance

1. Open **Perform**, choose **New**, and enter **First Performance** as the performance name.
2. Choose **Add Instrument**. With only one saved instrument, First Sound is selected automatically; otherwise select it in the patch picker. Set its MIDI channel to **1**.
3. Open **Mixer**. Check that First Sound reaches **Master**, with both strips unmuted. The built-in starter receives this route automatically.
4. Under **Add track** in the arranger, select **Melodic**. Orchestron creates a four-beat clip and opens its first pattern editor.
5. Check that the track sends to MIDI channel **1**. Leave the default 4/4 meter, sixteenth-note subdivision and octave **4**. In the first three step positions, replace **Rest** with **C**, **D**, and **G**; leave the remaining steps as rests.
6. Choose **Play edit** to hear the draft, then **Stop edit**, **Save**, and **Close editor**.
7. Choose **Play** in the arranger. It starts stopped instruments automatically and plays the saved pattern inside its clip.
8. Choose **Stop** to finish. Instruments that Play started are stopped again; an engine that was already running stays running.

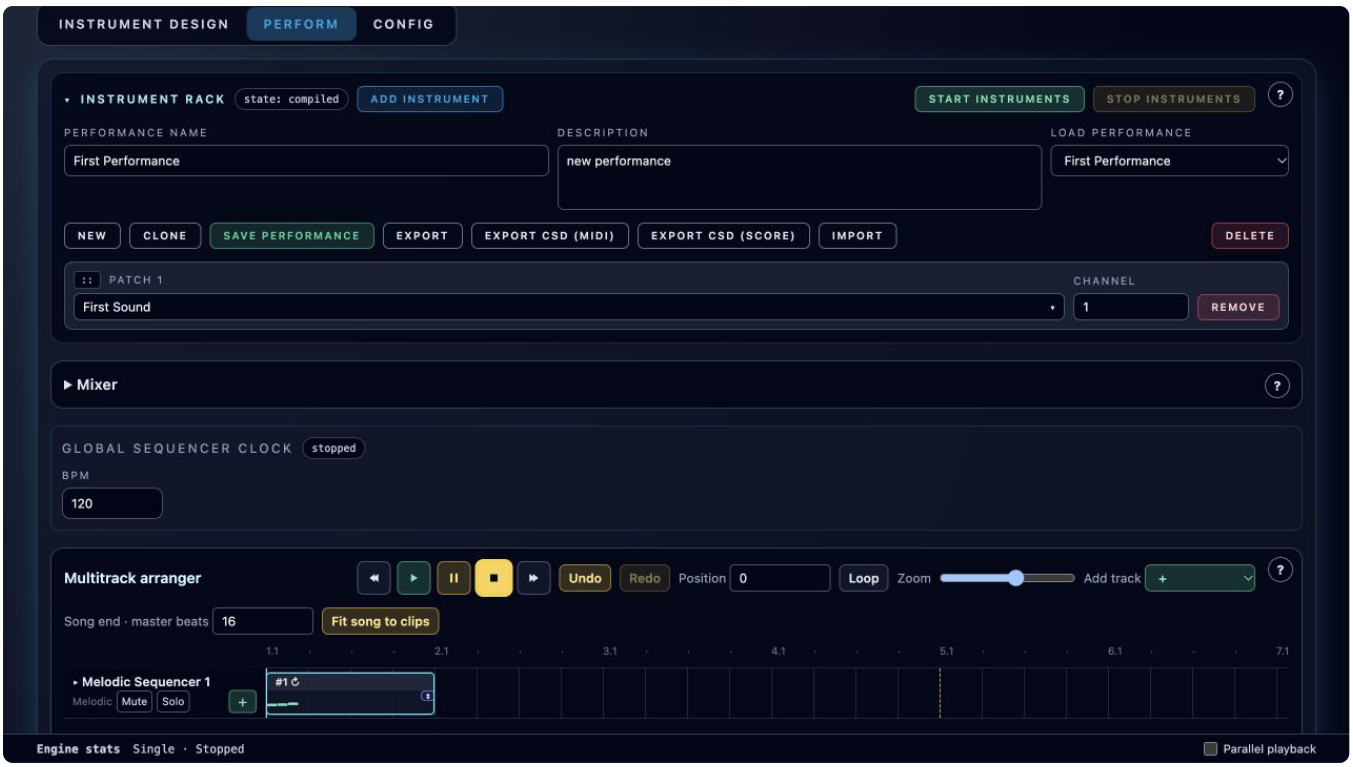

The saved starter performance. Mixer is collapsed to leave room for the arranger.

Read melodic step editing for note entry and the arranger for clip editing. Pattern length and the clip's window are separate: resizing a clip does not erase notes.

Save and reopen

1. Save any open pattern or structure draft. **Play edit** alone does not commit changes.
2. Choose **Save Performance**.
3. Choose **New**, then select **First Performance** under **Load Performance**. Confirm that First Sound and the melodic clip return, and choose arranger **Play** again.
4. Use **Export** for a portable, editable backup of the performance and its instruments.

Reload recovery is useful, but use explicit saves for work you intend to keep. See Saving, Workspace Recovery and Defaults.

If you hear no sound

- Check your device volume and browser audio permissions; interact with the page if the browser has not enabled audio yet.
- Confirm that the instrument compiled and was saved as a normal instrument, not a template.
- Check that the track or piano roll sends to the rack instrument's MIDI channel.
- Check the instrument strip and Master routing, levels and mute controls.
- Check that your clip contains a saved pattern and that the playhead passes through it.
- Open the Runtime or Perform status display for an error or audio-queue warning. See Live Status and Troubleshooting.

Navigation: Up | Previous | Next

Instrument Design

Navigation: Up | Previous | Next

This chapter covers the complete **Instrument Design** workflow on the `Instrument Design` page.

Start with a built-in playable instrument or import a supplied sound. You can work with an empty library. If you are new to the app, complete Getting Started first.

Create or load an instrument, connect its graph, compile and audition it, then save it. The following guides explain each step and the more advanced authoring tools.

Chapter Contents

- Patch Toolbar and Instrument Tabs
- Opcode Catalog and Integrated Documentation
- Graph Editor
- Input Formula Assistant
- GEN Table Editor (GEN Meta-Opcode)
- If, Switch and Drumset
- Monophonic Legato
- Defining Performance Controls
- Runtime Panel and Compilation Workflow
- Instrument Import / Export and CSD Export
- Supported Opcodes

What You Can Do Here

- Build instruments visually from Csound opcodes using the graph editor.
- Maintain multiple instrument tabs (parallel drafts or different patches).
- Use localized integrated help and opcode-level documentation without leaving the app.
- Choose built-in instrument, effect and output starters and define grouped stereo audio interfaces.
- Compile the current graph, inspect generated ORC, and audition drafts in isolation with compact performance controls and a mini keyboard.
- Define advanced input-combine formulas when multiple signals feed the same input.
- Configure function tables with the `GEN` meta-opcode (including `GEN01` audio-file tables and `GENpadsynth`).
- Export instruments as Orchestron bundle files and as `.csd`.

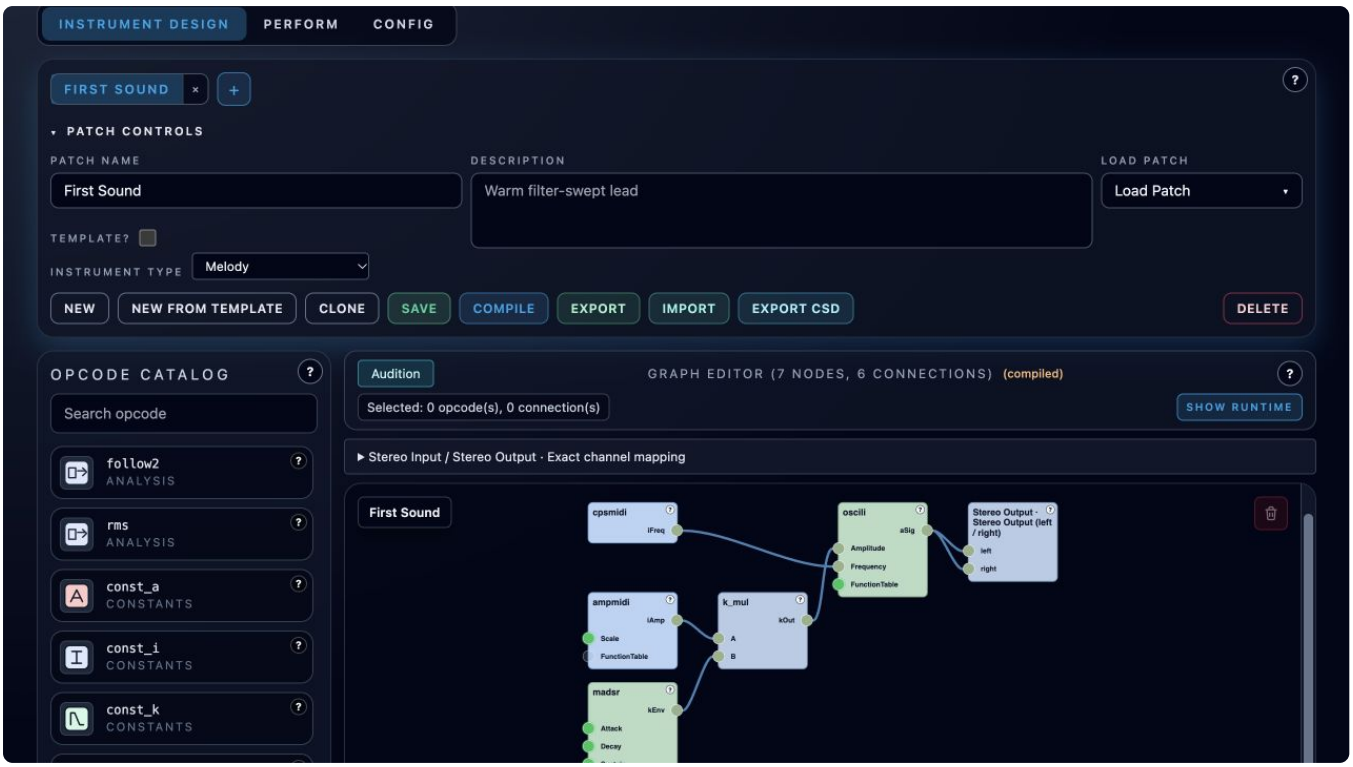

Instrument Design with the saved starter graph. Open Show runtime to inspect compilation output.

Recommended Workflow

1. Create or load a patch in the patch toolbar.
2. Add opcodes from the catalog (click or drag-and-drop).
3. Connect signals and set constant values.
4. If needed, define an Input Formula on a socket with multiple inputs.
5. Compile and inspect the runtime panel output.
6. Save the patch.
7. Export a `.csd` or an Orchestron instrument bundle.

Important Behavior

- The compile status badge is per active instrument tab/patch snapshot (`compiled`, `pending changes`, `errors`).
- Runtime starts collapsed to give more graph space; open it with `Show runtime`. Collapse `Patch controls` to hide metadata and file actions while keeping instrument tabs visible. Both choices survive view switches until browser reload.
- Saving a normal patch performs compile validation first; a failing compile prevents a bad save.
- Saving a patch marked `Template?` skips compile validation so incomplete starter graphs can be stored and reused. Template patches remain loadable in Instrument Design but are not selectable as Perform rack instruments.

Navigation: Up | Previous | Next

Patch Toolbar and Instrument Tabs

Navigation: Up | Previous | Next

The patch toolbar is the main control surface for patch metadata, patch library operations, and instrument-level file actions.

Collapse Patch Controls

Use Patch controls to collapse metadata and load/save/compile/import/export actions. Instrument tabs, including add and close, stay visible. Controls start expanded; your choice survives view switches until browser reload.

Patch Toolbar Layout

The toolbar includes:

- Instrument tabs (multiple editable patch workspaces)
- Patch metadata fields (`Patch Name` , `Description`)
- Instrument Type selector
- Grouped, searchable patch picker (`Load Patch`)
- Action buttons (`New` , `Clone` , `Delete` , `Save` , `Compile` , `Export` , `Import` , `Export CSD`)

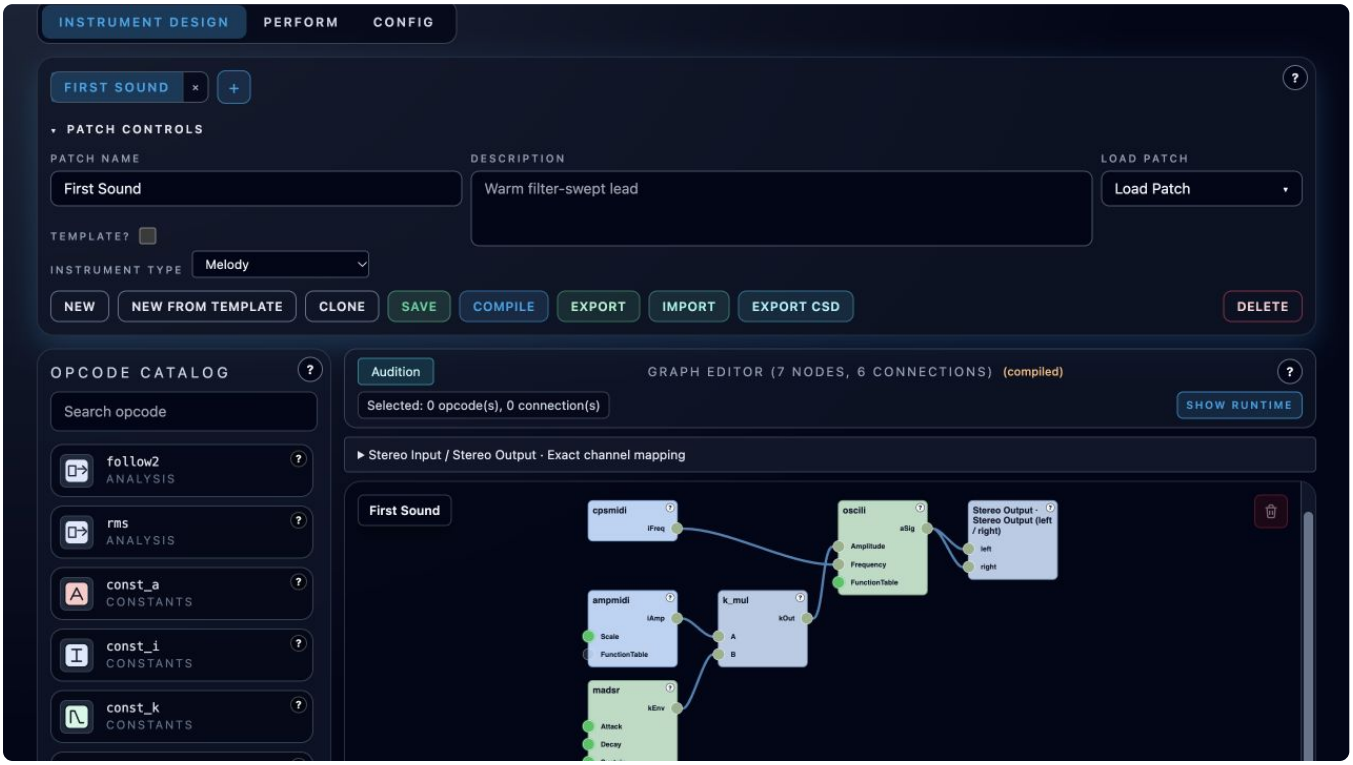

The saved First Sound starter: Instrument Type, library actions and Audition are available above the graph.

Instrument Tabs

- Use the `+` button to add a new instrument tab.
- Each tab stores its own editable patch snapshot (graph + metadata).
- Tabs can be switched without losing the current edit state.
- Tabs can be closed with the `x` button.
- Tabs are persisted in app state (see Configuration chapter), so drafts can be restored after reload.

Patch Metadata

- `Patch Name` is the display name used in the patch library and performance rack dropdowns for normal patches.

- `Description` is a three-line patch note field and accepts up to 2048 characters. Hover over the field to see the full description in a tooltip.
- `Template?` marks the patch as a reusable starter graph. Template patches show a `TEMPLATE` token beside their name and are excluded from Perform rack instrument choices.
- `Instrument Type` offers Percussion, Melody, Bass, Effects / Noise, and Continuous (always-on). Choosing Continuous enables always-on scheduling; every other type uses MIDI activation. Continuous patches run when added to the Perform rack and started. They may generate audio without an inlet. Receiving audio requires an actual inlet, independently of activation. Musical role is separate interface metadata.
- Metadata updates affect the current tab immediately, but they are not stored in the patch library until you save.

Patch Library Loading

- `Load Patch` opens collapsed type groups with counts. Expand a group to see patch names in alphabetical order; template patches retain their marker. Patches already open in instrument tabs are excluded.
- Search matches names and descriptions across all groups, ignoring case and surrounding spaces. Enter at least **four characters**; results update after **500 ms without typing**. Matching groups expand automatically. Shortening the search below four characters restores browsing immediately.
- Use Tab or Up/Down to move through controls and Enter to expand a group or load a patch. Enter in an active search loads its first result. Escape or clicking outside closes the picker; closing resets search and group expansion.
- Loading replaces the active tab with the selected saved patch.

Actions

New

- Opens the built-in creation chooser: Playable instrument, Drumset, Audio effect, Audio Output, or Empty patch.
- Creates an unsaved draft with the selected starter graph and default engine settings.
- This draft is not yet stored in the patch library until `Save` is used.

New from Template

- Opens the same chooser with the five built-in starters plus any saved templates.
- Creates a new unsaved normal patch with the chosen template's graph and description.
- Built-in choices remain available when the library has no saved templates.

Clone

- Creates a new saved patch by duplicating the current patch graph and metadata.
- Preserves the `Template?` flag and Instrument Type from the source patch.
- If the name already exists, Orchestron generates a `(... copy)` style name.
- The cloned patch is loaded automatically after creation.

Delete

- Deletes the currently saved patch from the patch library.
- Requires confirmation.
- After delete, Orchestron loads another available patch or creates a new draft if none remain.

Compile

- Compiles the current instrument graph into generated Csound ORC/CSD artifacts.
- Updates compile status badge state.
- Does not save the patch or start audio automatically. Use Audition in the graph header to hear an unsaved draft.

Save

- Performs compile validation first for normal patches.
- If compile succeeds, saves the patch to the instrument library.

- If compile fails, save is blocked and the patch remains unsaved.
- Template patches skip compile validation so incomplete starter graphs can be saved.
- Always-on patches still save as normal patch definitions; they run continuously only after they are added to a started Perform rack. A continuous source need not be an effect or have audio inlets.

Export / Import

- `Export` writes an Orchestron instrument definition bundle (`.orch.instrument.json` or `.orch.instrument.zip`).
- `Import` loads an instrument definition bundle and supports conflict handling (overwrite / rename / skip).
- Details are in Instrument Import / Export and CSD Export.

Export CSD

- Triggers compile and downloads the compiled `.csd` file.
- Useful for running the generated instrument outside Orchestron.

Practical Notes

- `Compile` is the fast iteration action while designing.
- `Save` is the persistence action for the patch library.
- `Export` is the sharing/transfer action for Orchestron-specific instrument bundles.
- `Export CSD` is the interoperability action for raw Csound usage.
- The integrated `?` help for this toolbar now focuses on tab-local draft state versus saved library state, plus the difference between `Compile`, `Save`, `Clone`, and export actions.

Built-in creation choices

New and New from template offer Playable instrument (MIDI pitch/velocity, sine oscillator and madsr), Audio effect (stereo pass-through), Audio Output (stereo input to direct output), and Empty patch. These choices work without saved templates. Guided instruments connect to the fixed Master when added to a performance.

Choose a starter

Choice	Type, activation and audio interface	Next step
Playable instrument	Melody; MIDI notes; named stereo output; pitch, velocity, sine oscillator and envelope	Save, add to Perform, and play its MIDI channel.
Drumset	Percussion; MIDI notes; named stereo output; analog drum voices	Save, add to Perform, and sequence its drum notes.
Audio effect	Continuous; named stereo input and output, initially pass-through	Add processing between the input and output before using it as an effect.
Audio Output	Continuous; stereo input to Direct Audio Output (`outs`)	Save for use as a custom output processor.
Empty patch	Melody; empty graph for custom design	Build and compile the required signal path.

Role describes the patch's audio interface; Instrument Type determines whether it runs continuously or on MIDI notes. Audio effect and Audio Output starters use Continuous; use Effects / Noise for MIDI-triggered sound effects. Every performance has a fixed internal Master. It has no editable patch or rack assignment; add processing through its mixer inserts.

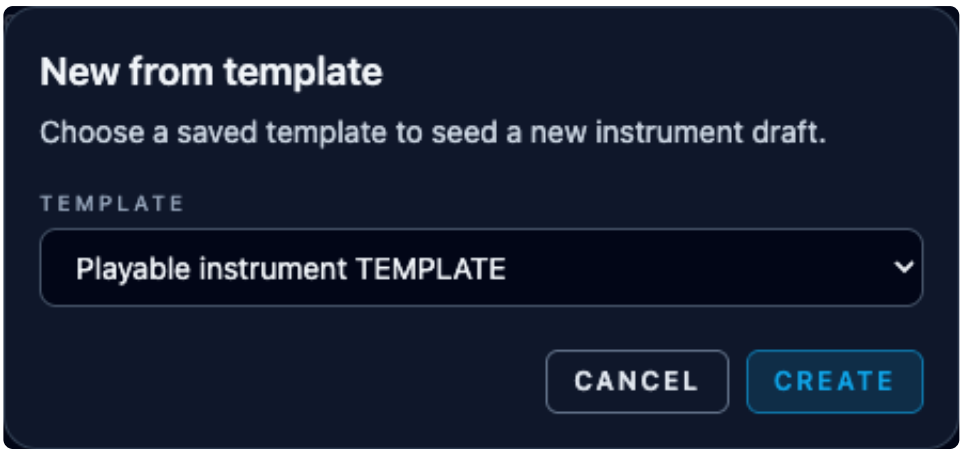

New opens the starter chooser even when no saved templates exist.

Older patches and saved metadata

Existing always-on patches become Continuous without changing playback. Other older patches are classified from their names first, then descriptions: drum/percussion terms precede bass terms, followed by melodic and sound-effect terms. Unmatched patches default to Melody. Review the inferred Instrument Type and change it when needed; later name or description edits do not reclassify a patch.

Types are saved with the patch and retained in clones, templates, restored tabs, and instrument/performance JSON or ZIP bundles. New playable and empty patches default to Melody, drumsets to Percussion, and effect/output starters to Continuous.

Navigation: Up | Previous | Next

Opcode Catalog and Integrated Documentation

Navigation: Up | Previous | Next

The opcode catalog is the source list of all supported building blocks you can place in the graph editor.

What The Catalog Supports

- Search by opcode name
- Search by category
- Search by description text
- Search by tags
- Click to add an opcode to the graph
- Drag-and-drop an opcode into the graph canvas at a chosen position

Search Behavior

The search field filters on these opcode fields:

- `name`
- `category`
- `description`
- `tags`

This makes it easy to find opcodes by function (for example `filter`, `midi`, `reverb`, `soundfont`) even if you do not remember the exact Csound opcode name.

Opcode catalog filtered by search text.

Adding Opcodes

Click To Add

- Clicking an opcode row inserts the node into the graph using the app's add-node behavior.
- This is fastest when exact placement is not important yet.

Drag-and-Drop To Place

- Drag an opcode from the catalog and drop it into the graph canvas.
- The graph editor computes the drop location and places the node there.
- While dragging over the graph, the canvas highlights to confirm drop targeting.

Visual Information In The Catalog

Each catalog entry shows:

- Opcode icon (backend-served icon asset)
- Opcode name
- Category label

These same categories influence node coloring in the graph editor for quick visual grouping.

Integrated Opcode Documentation (? on node)

After placing an opcode in the graph, use the node's `?` button to open the opcode documentation modal.

The opcode documentation modal provides:

- Localized description (English/German/French/Spanish)
- Category and syntax/template summary
- Input/output port descriptions (including optional/default/accepted types)
- Tags
- Direct link to the Csound reference page (`Open Csound Reference`)

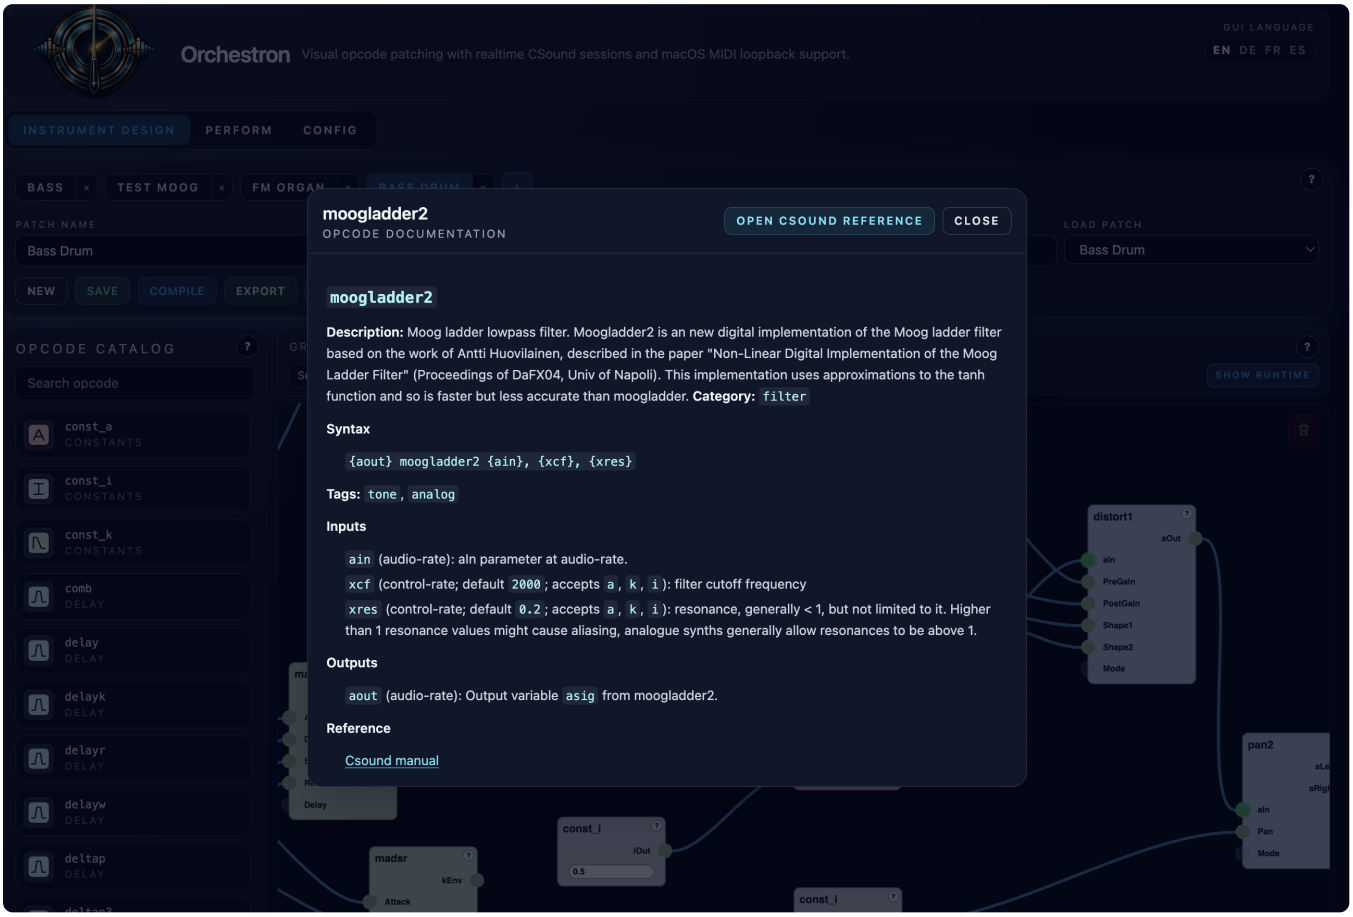

Integrated opcode documentation modal opened from a graph node.

Monophonic Legato

Use `midi_legato` for touching or overlapping notes that share a voice. Its live pitch and velocity outputs drive the sound while `madsr` follows the phrase. Gaps restore articulation. The node is opt-in and belongs in the main graph.

Atone High-Pass Filters

Search for `atone` or `highpass` in the `filter` category:

- atone filters an audio signal (`asig`) and produces audio (`aout`). Its `khp` cutoff accepts control-rate or init-rate values and starts at 200 Hz.
- atonek filters a control signal (`ksig`) and produces control output (`kout`) for modulation. Its `khp` cutoff starts at 10 Hz. Both inputs accept control-rate or init-rate values.
- atonex cascades multiple `atone` stages for a sharper cutoff. Its audio input/output are `asig` / `aout`; `xhp` accepts audio-rate, control-rate, or init-rate cutoff values and starts at 200 Hz. The optional init-rate `inumlayer` sets the stage count and defaults to 4.

Each filter exposes optional init-rate `iskip`: 0 clears its internal state at initialization, while a nonzero value retains previous state. The default is 0. Setting only `iskip` on `atonex` retains the default four layers. Cutoff defaults are editor starting values; Csound requires a cutoff argument.

Connect the audio outputs of `atone` and `atonex` to downstream audio nodes or both `outs` inputs. Connect `atonek` to a control input such as oscillator amplitude or frequency. The node's `?` help documents every port in English, German, French, and Spanish.

ZDF and Custom IIR Filters

Search for `zdf` in the `filter` category. All five nodes accept audio input `ain`; `xcf` accepts init, control, or audio rate and starts at 1200 Hz. The two-pole and ladder filters also expose `xq` (0.5–25, default 1) at the same rates.

- `zdf_1pole`: 6 dB/oct, with `kmode` 0 = low-pass, 1 = high-pass, 2 = allpass.
- `zdf_1pole_mode`: simultaneous `alp`, `ahp` outputs, in that order.
- `zdf_2pole`: 12 dB/oct, with `kmode` 0 = low-pass, 1 = high-pass, 2 = band-pass, 3 = unity-gain band-pass, 4 = notch, 5 = allpass, 6 = peak.
- `zdf_2pole_mode`: simultaneous `alp`, `abp`, `ahp` outputs, in that order.
- `zdf_ladder`: four-pole (24 dB/oct) Moog ladder low-pass; Q = 25 reaches self-oscillation.

`kmode` accepts control or init rate and defaults to 0. Every ZDF node has optional init-rate `istor`: 0 clears filter memory, nonzero retains it. Setting `istor` alone preserves other defaults. Cutoff and Q defaults are editor starting values.

`zfilter2` uses custom IIR coefficients with control/init-rate `kdamp` and `kfreq`. Positive damping adjustments lengthen ringing and negative ones shorten it; positive warp shifts pole frequencies upward and negative warp shifts them downward. Both start at 0 (unchanged poles); use values strictly between -1 and 1.

Edit its inline **Coefficients** field as a comma-separated list: all numerator coefficients starting at b0, followed by denominator coefficients starting at a1. The leading denominator coefficient a0 is implicitly 1. **NumeratorCount** (`im`) counts numerator coefficients including b0; **DenominatorCount** (`in`) counts denominator coefficients. Connect `const_i` nodes to change these counts. Defaults are 3 and 2, with `0.06745527, 0.13491055, 0.06745527, -1.1429805, 0.4128016`, a stable second-order low-pass filter.

You can instead connect init-rate sources to **Coefficients** in coefficient order. Each connection supplies one coefficient, replacing the entire inline list; these inputs are not summed, and the list does not accept a combine formula. Connection order survives save/export. Use separate upstream init-rate calculations for individual coefficients.

Supply exactly `im + in` coefficients. The editor supports 1–51 numerator coefficients and 1–49 denominator coefficients for Csound 6 compatibility: zero poles and 50 poles encounter native bugs in Csound 6.18. Invalid literal counts fail compilation; invalid connected counts stop the affected voice before the native filter initializes. Coefficients remain the responsibility of the filter design: count checks do not establish stability.

Compression, Equalization and Cyclic Tonality

- `compress2`, in **dynamics**, processes `aasig` using peak detection on the separate `acsig` sidechain. Connect the same source to both for ordinary compression, or a different source for ducking. Gate and knee levels use dBFS. Editor defaults are -90 dB gate, -52/-30 dB knees, ratio 3, attack 0.01 s, release 0.1 s and lookahead 0.05 s. Lookahead adds audio latency. The detector and ratio controls accept control or init rate; lookahead is init rate.
- `pareq`, in **filter**, offers **Mode** 0 = peaking, 1 = low shelf, 2 = high shelf. **Gain** is linear: 1 is neutral, less than 1 cuts, greater than 1 boosts. Frequency, gain and Q accept control or init rate; their editor defaults are 1000 Hz, 1 and approximately 0.7071. Optional init-rate **SkipInit** defaults to 0; setting it alone retains peaking mode.
- `hsboscil`, in **oscillator**, separates cyclic **Tonality** from **Brightness**, both measured in octaves relative to **BaseFrequency** (initially 220 Hz). Ramp tonality with fixed brightness for a Risset glissando. Amplitude, tonality and brightness accept control or init rate. **Octaves** is an optional init-rate count from 2 to 10, default 3; **Phase** defaults to 0, with -1 skipping initialization.

`hsboscil` needs two function tables. **WaveformTable** defaults to the built-in sine table 1. **OctaveTable** must be supplied: add a **GEN** node with routine **19**, normalization disabled, table size **1024**, and a GEN19 partial row with **Partial number 1**, **Strength 0.5**, **Phase 270°**, and **DC offset 0.5**; connect its `ift` output to **OctaveTable**. Connect a second GEN node to **WaveformTable** for a custom waveform. Table numbers use init-rate ports, matching GEN's `ift` output. They are distinct from the `f` signals used by spectral nodes.

Each node produces mono audio; connect its output to both `outs` inputs for centered stereo. The `?` help describes all ports in English, German, French and Spanish. Prefilled required parameters are editor starting values; optional parameters retain Csound defaults.

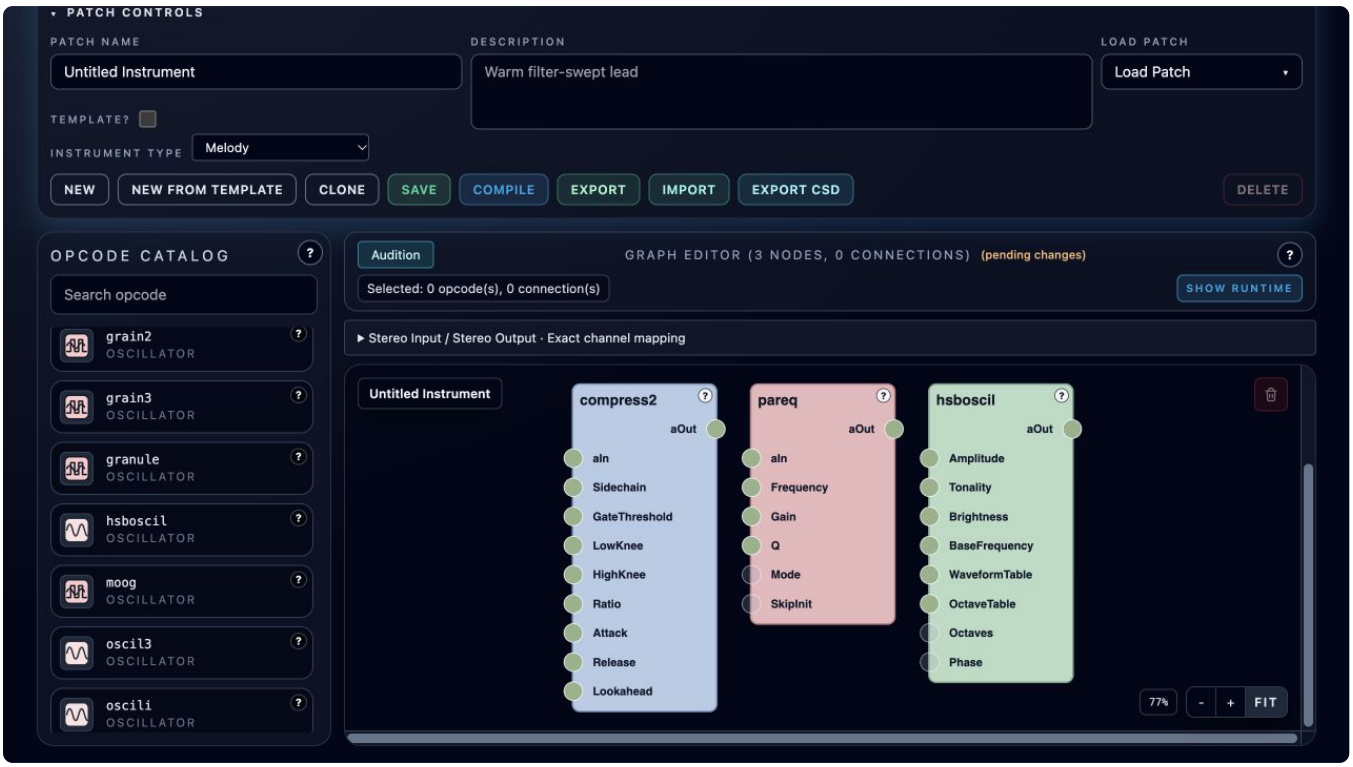

Prepared Piano

Search for `prepiano` in **physical_modeling**. The node exposes the stereo form: connect `al` to the left output and `ar` to the right output. Connect `cpsmidi` to init-rate `ifreq` for note pitch; use downstream gain or envelopes for volume shaping.

The starting values follow the Csound example: 60 Hz, three strings, 10 cents detuning, stiffness 1, a 3-second decay to –30 dB, high-frequency loss 0.002, pivoting ends, hammer mass 1, hammer frequency 5000 Hz, initial hammer position –0.01, strike position 0.09, strike velocity 20, scan frequency 0 Hz and stereo scan-frequency spread 0.1 Hz. These are editor starting values. Use two or three strings with Csound 6; its implementation divides the detuning spread by the string count minus one.

Only `kbcl` and `kbcr` accept live control-rate changes (or init-rate constants): **1 = clamped**, **2 = pivoting**, **3 = free**. Changing boundaries during a note can glitch. Other inputs are init rate. Keep stiffness positive for new voices; negative stiffness skips native initialization and is intended for continuing an already initialized model.

Optional `irattles` and `irubbers` default to **0**, disabling that preparation. To add one, create a GEN02 value-list table, disable normalization, and connect its init-rate `ift` output to the corresponding table port. Use a table size large enough for a count followed by four values per object (`1 + 4 × count` ; size 8 holds one object):

Preparation	Values after the count	Example value list
Rattle	Position along the string, mass-density ratio, frequency, vertical length	`1, 0.6, 10, 100, 0.001`
Rubber	Position along the string, mass-density ratio, frequency, loss	`1, 0.7, 50, 500, 1000`

Either table can be used alone. When only a rubber table is connected, the rattle argument stays 0. These examples render with the 60 Hz starting settings; retune stiffness and preparations when changing register. The native physical model can become unstable with other combinations, including these rubber settings at 220 Hz with stiffness 1. Start with low output gain when experimenting. The node's `?` help covers every port in English, German, French and Spanish.

STK Instruments

Search for `stk` to find the 27 Synthesis Toolkit instruments in the `physical_modeling` , `fm` , and `oscillator` categories. They include strings, winds, percussion, voices, organs, and electric pianos. Each node produces one mono audio signal;

connect it to both `outs` inputs for centered stereo output.

- `ifrequency` and `iamplitude` are init-rate inputs: frequency in Hz and note amplitude from 0 to 1. Use `cpsmidi` for note pitch; apply envelopes or gain after the audio output for continuous level changes.
- `STKDrummer` uses the frequency to select a drum sample. Connect `cpsmidi` and play MIDI drum notes. `STKDrummer`, `STKPlucked`, and `STKSitar` have no documented controller inputs.
- The other nodes expose the manual's controller number/value pairs in order. The controller number is prefilled with the documented number. Set or connect its paired `kv` value to enable that control. Both ports accept init-rate or control-rate sources, so controller values can change during a note.
- Unset controller values omit the entire pair and retain STK's instrument defaults. You can enable a later pair without filling earlier pairs. For example, set only `kv7` to `3` on `STKBandedWG` to send instrument preset controller `16` with value `3`.
- Most values range from 0 to 127, but preset selectors use model-specific ranges. `STKShakers` uses controller **1071**, not a MIDI CC in the 0–127 range, for its instrument selection. Check the node's `?` help for each controller's meaning.

Runtime requires the Csound STK plugin (`stkopd`) and its rawwave data on the machine running the Csound backend. Set `RAWWAVE_PATH` to the installed rawwaves directory before starting the backend; some Csound builds do not register any STK opcodes without it. For example, an Apple Silicon Homebrew installation of STK uses `RAWWAVE_PATH=/opt/homebrew/opt/stk/share/stk/rawwaves/`. Catalog availability and successful graph compilation do not install this runtime dependency.

Context Help (`?` in page sections)

In addition to opcode-level docs, the app provides context help buttons for page sections (toolbar, catalog, graph, runtime, sequencer panels, config panels). These open integrated markdown help content in the currently selected GUI language.

For the opcode catalog specifically, the integrated help now focuses on search/add workflow and on the distinction between catalog help and the node-level opcode documentation modal.

See also:

- GUI Language and Integrated Help
- Supported Opcodes for the full opcode index table

Additive Resynthesis with adsyn

`adsyn` resynthesizes sinusoidal partials from a Csound `hetro` analysis file. Find it in **oscillator**, add it to the graph, and open its configuration button to upload the analysis. The selected filename is saved with the instrument. Use **Clear Asset** to remove a selection or upload another file to replace it.

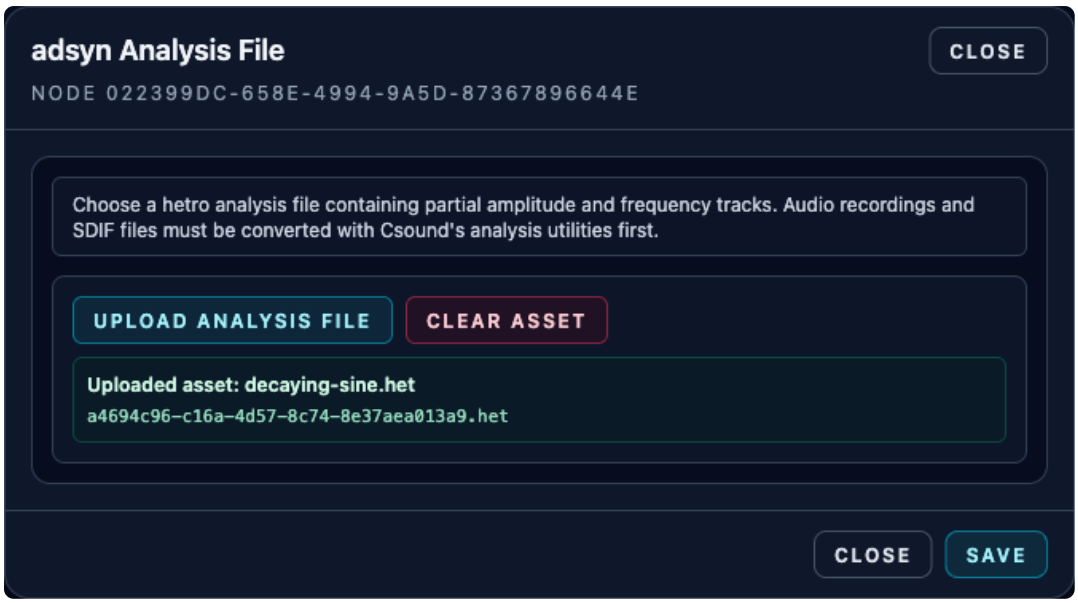

The three control-rate inputs also accept init-rate connections and formulas:

Input	Default	Meaning
`kamod`	1	Amplitude multiplier
`kfmod`	1	Frequency multiplier; 2 raises every partial by an octave
`ksmod`	1	Analysis playback speed; 2 traverses the envelopes twice as fast

Connect the mono audio output `ares` to your envelope/effects chain and then an audio output. Add your own release envelope for MIDI articulation. `adsyn` can be placed in the main graph or inside an If/Switch synthesis case. The analysis file is chosen when each note initializes; the three controls can change at control rate during playback.

Upload analysis data, not a WAV recording. Generate it outside Orchestron with Csound hetro, for example:

```
csound -U hetro -f440 -h16 source.wav source.het
```

Set `-f` to the approximate fundamental of the source. SDIF data must first be converted with Csound's `sdif2ad` utility. The file picker allows arbitrary extensions because Csound analysis files need not be named `.het`.

Orchestron accepts portable HETRO text and legacy native-endian binary analysis. Malformed tracks are rejected before compilation or CSD export. Orchestron resolves uploaded/imported files from its asset storage; raw paths and numeric Csound filename aliases are not supported. Instrument/performance ZIPs carry the analysis file, and both performance CSD modes include it under `assets/`. No analysis-generation tool or piano conversion is included.

Navigation: Up | Previous | Next

Graph Editor

Navigation: Up | Previous | Next

The graph editor is the central visual patching surface for instrument design.

Core Graph Editing Features

- Typed input/output ports (audio, control, init, string, function-table)
- Node placement and repositioning (drag nodes on the canvas)
- Connection creation between compatible ports
- Connection and node selection highlighting
- Deletion of selected nodes/connections via explicit delete action
- Zoom controls and fit-to-graph navigation
- Node category color coding
- Inline constant value editing for `const_a`, `const_i`, `const_k`, `const_s`

Canvas Navigation

The graph editor supports pan/zoom navigation and includes a zoom HUD in the lower-right corner:

- Zoom percentage display
- `−` zoom out
- `+` zoom in
- `Fit` to fit the current graph into view

If the graph is empty, `Fit` resets toward a centered default view.

Horizontal and vertical scrollbars along the canvas edges let you reach every part of a large patch, including expanded If/Switch case frames. They remain available when **Patch controls** is collapsed and adjust when Runtime or the browser window changes the canvas size. Drag a scrollbar, or focus it with **Tab** and use the arrow keys for that direction. Scrolling moves the view without moving nodes or clearing selection. The scrollbars stay synchronized with background panning, zoom, and `Fit`.

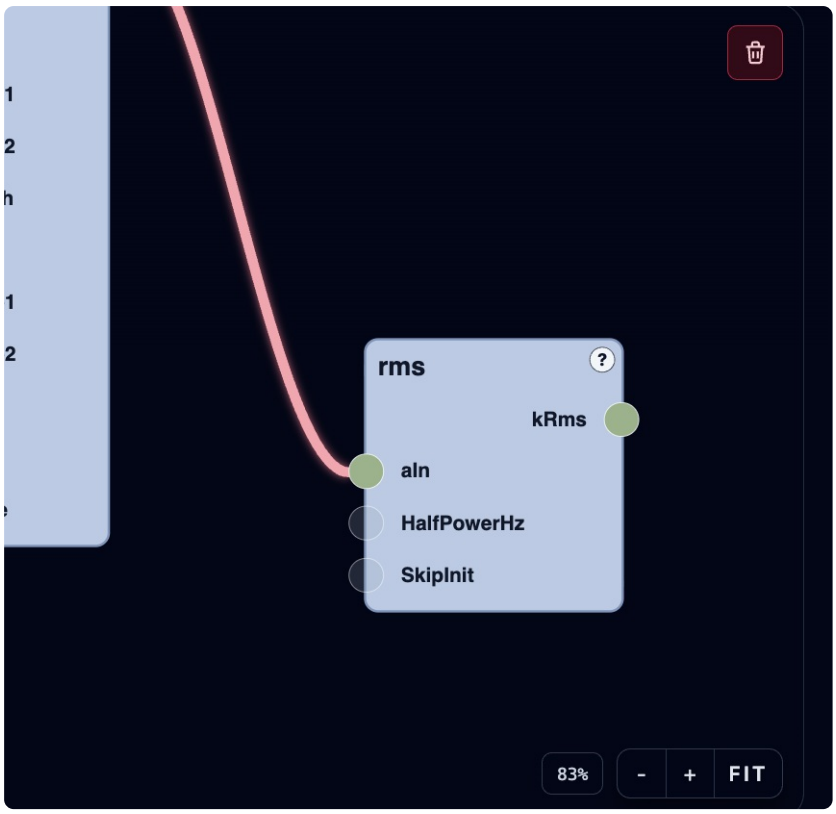

Graph editor detail showing selection state, delete button, and zoom controls.

Node Appearance and Visual Cues

Category Coloring

Nodes are color-coded by category (for example generator, filter, MIDI, constants, etc.) to help visually scan signal flow.

Optional Inputs

Optional inputs are visually distinct from required inputs.

Stored Values and Formulas

Bright-green input sockets have an active stored value or saved formula. Zero, false, and explicitly stored values equal to the catalog default count as stored values. A connection overriding a stored value is not highlighted unless it also has a formula. Inherited catalog defaults remain unhighlighted.

Hover over an input to see its source and value. Double-click an input, with or without cables, to inspect or edit its effective value in the Input Formula Assistant. The dialog distinguishes **Stored value**, **Saved formula**, **Connected signals**, **Catalog default**, and **Unset**. Opening or cancelling it leaves the patch unchanged.

Reset input removes both the saved formula and the stored value, preserving connections. The dialog previews the resulting connection, default, or unset state before you reset. Required inputs with no remaining source produce the existing compile diagnostic.

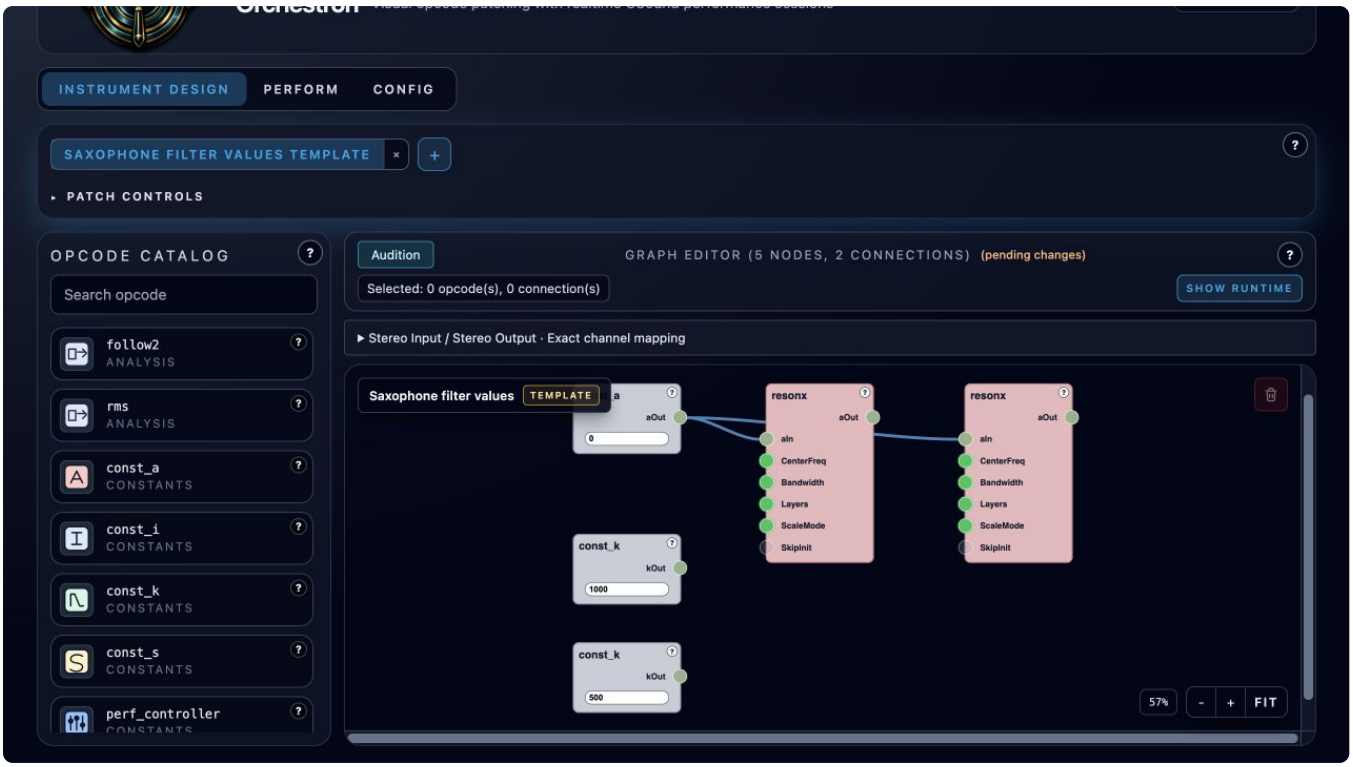

The two filters retain their stored CenterFreq, Bandwidth, Layers and ScaleMode values. SkipInit is unhighlighted because its zero comes from the catalog.

Adding Nodes

You can add nodes in two ways:

- Click from the opcode catalog
- Drag an opcode from the catalog and drop it into the graph canvas

Drag-and-drop placement is preferable for larger patches because it preserves your layout intent.

Connections

Creating Connections

- Connect outputs to compatible inputs.
- Compile the patch to check all connections together and see any remaining errors.

Selecting Connections

- Connections can be selected directly (not just nodes).
- Selected connections receive a visual highlight.
- Ctrl-based accumulation allows multi-select of connections and nodes.

Connection Deletion Is Intentional

Orchestron prevents accidental edge removal during socket interaction.

- Connections are not removed by casual socket clicks/drags.
- To remove a connection, select it explicitly and use the delete action.

Selection Behavior

- Click nodes to select.

- Ctrl-click accumulates node selections.
- Click connections to select them.
- Ctrl-click accumulates/deselects connection selections.
- Background click clears selection (when Ctrl is not held).

Delete Selected Elements

The graph editor overlay includes a trash/delete button (top-right of the graph canvas area).

- Disabled when nothing is selected
- Enabled when at least one node or connection is selected

Deletion flow:

- If one item is selected, deletion can happen directly.
- If multiple items are selected, Orchestron opens a confirmation dialog listing all affected items (nodes and connections) before deletion.

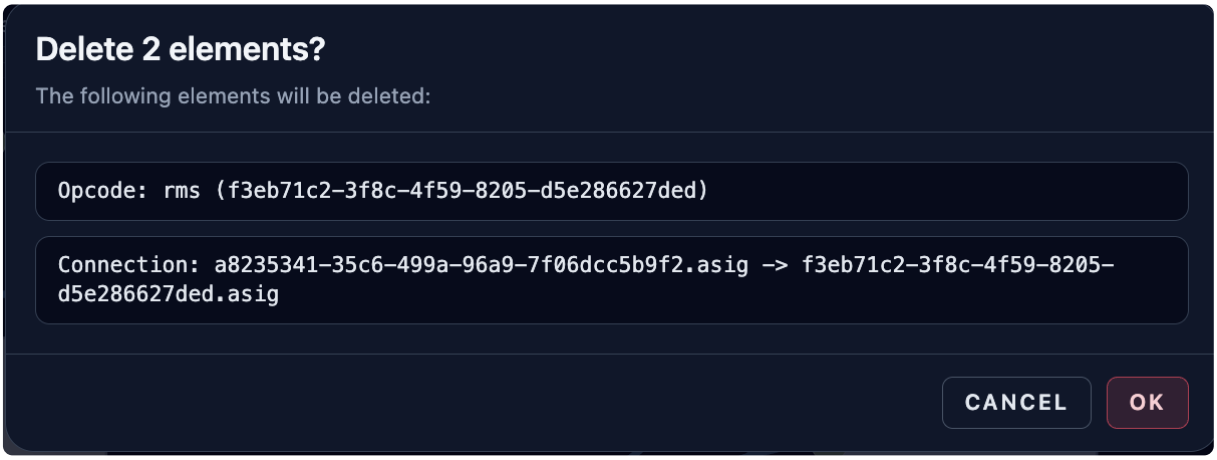

Delete-selection confirmation dialog listing selected graph items.

Node-Specific Controls

Constant Nodes

The constant opcodes (`const_a` , `const_i` , `const_k` , `const_s`) provide inline editable values directly on the node. Numeric fields keep focus while you type complete values, including a leading decimal point (for example `.25`), negative numbers and scientific notation. `const_s` accepts lowercase letters, digits, and underscores, must begin with a lowercase letter, and is limited to 50 characters.

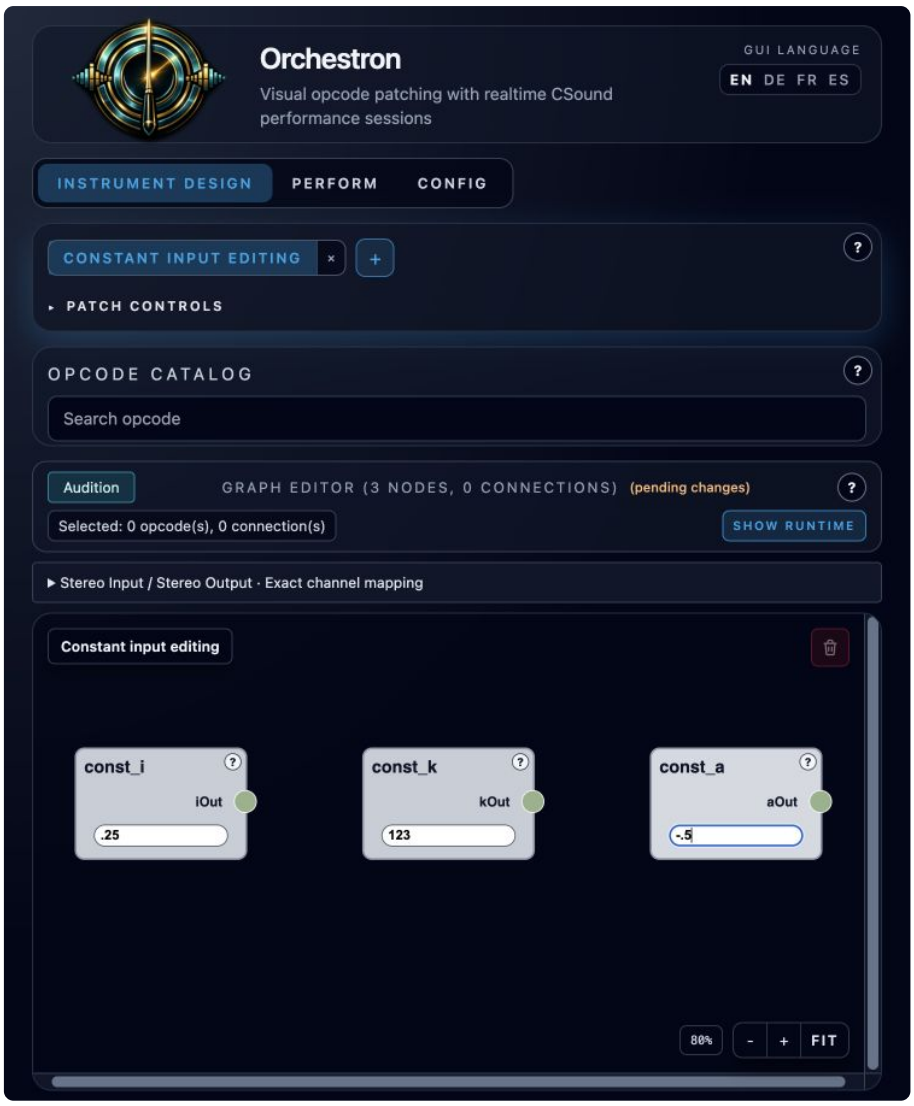

Custom IIR Coefficients

The `zfilter2` node has an inline comma-separated coefficient field. Its Coefficients socket also accepts ordered init-rate connections that replace the field. See ZDF and custom IIR filters for coefficient order, counts, and Csound compatibility limits.

Documentation Button (`?`)

Nodes with documentation expose a `?` button to open the opcode documentation modal.

GEN Button (for `GEN` node)

The `GEN` meta-opcode node shows a `GEN` button that opens the specialized GEN table editor.

Layout Persistence

- Node positions are stored in the patch graph.
- The graph structure (nodes + connections) is part of the patch data.
- The visual layout and related UI metadata (for example input formulas, GEN node configuration) are stored in `graph.ui_layout`.

Compile Implications

- Compilation order is derived from graph dependencies.

- Invalid wiring, missing required inputs, or invalid formulas surface compile diagnostics.
- Graph compile status is reflected in the instrument page header (`compiled` , `pending changes` , `errors`).

Integrated Help Focus

The graph editor `?` help now concentrates on graph-specific behavior: persisted layout, deliberate deletion, constant-node inline editing, and the Input Formula Assistant workflow for combining multiple signals on one input.

If and Switch at note start

Add **If** or **Switch** from **Control flow**. Drag catalog nodes directly into expanded cases, then patch each case into its managed **Case Result**. Frames measure the complete node bodies, automatically grow to the right and bottom and retain their dimensions, and keep the result in a separate bottom-right row. Drag the **Case Result title or unused body** to resize or shrink the frame; complete member bounds sets its minimum. Other selected nodes stay still. This also works for Silence. Escape, cancellation or focus loss restores the previous dimensions and positions. Sizes persist through reload and export/import; bright purple outlines remain about two screen pixels thick at every zoom. The block remains above the first case; top and left movement limits keep members clear of the header.

Ordinary main-graph drags into a case transfer the selection. Ordinary case-member drags rearrange the current case; hold **Alt/Option at drag start** to transfer to another case or the main graph. The cyan/red border shows whether a drop is allowed. Every connection and formula binding must lead to another moved node or an endpoint already in the destination scope. Invalid drops list the connections and restore the whole group; Escape cancels. Existing shared main-graph inputs remain valid and can be wired after placement. A valid drop into Silence activates Synthesis. Dropping onto a collapsed block expands it without choosing a case.

The common output carries only the selected case. Collapse preserves selection, positions, viewport, shared input wiring and formulas. **Configure branches** retains **Move selected nodes** with the same transfer validation, plus case editing, Silence and reviewed mono/stereo changes. The built-in **Drumset** template demonstrates overlapping voices.

See If, Switch and Drumset for the complete workflow and case membership rules.

Stereo blocks and audio interfaces

Stereo Input and Stereo Output are collapsible views of ordinary paired inleta/outleta nodes. Expand them to edit the underlying opcodes and connections. Exact port names, stored input values and input formulas survive collapsing, expansion, saving and export. Direct Audio Output (outs) stays distinct.

Route an instrument’s Stereo Output through the performance mixer to Master for playback.

The interface panel describes musical role, stable group IDs, display names, exact member ports, purpose and designated main input/output. Stereo blocks and their mappings share the same creation and deletion operations. Advanced mono/custom groups remain metadata views; changing a stereo group to another layout exposes the ordinary nodes without deleting them.

Create, rename and remove a stereo interface

1. Click **Stereo Input** or **Stereo Output** in the opcode catalog, drag either onto the canvas, or expand **Stereo Input / Stereo Output · Exact channel mapping** and use its matching add button. Each action creates two channels and their mapping. Connect the new audio ports as needed.
2. Set Role and review the group's display name, purpose and designated Main input/Main output. New pairs use `left` / `right` or an available `busN.left` / `busN.right` pair. Channel names are unique within each direction.
3. Edit **Left** and **Right**, then choose **Apply** to review and confirm the rename. **Cancel** discards the draft. Renaming changes the mapping and underlying opcode names together. A string constant supplying an old name remains available to its other connections.
4. **Show underlying nodes** exposes the pair; **Collapse** returns to its grouped view. Delete a collapsed block or choose **Remove** on its mapping to remove both channels, attached connections and affected input formulas. The confirmation lists all affected elements. Deleting only one expanded channel removes the mapping and leaves the other as an ordinary opcode. Deleting a connection alone retains the mapping.

5. Deleting the designated main group clears its selection without promoting another group. Select a remaining group explicitly when needed.
6. Compile and save before assigning the patch in Perform. Check Audio Mixer and Routing after any external port changes.

Upgrade or repair an existing instrument

- **Group existing channels:** choose Stereo Input or Stereo Output, select the existing Left/Right nodes and review the change. This preserves the channels' exact names, IDs, positions, audio wiring and formulas. Channel names become part of the stereo block's configuration: values from connected `const_s` nodes are stored directly on the underlying channels, and their naming connections are removed. Dedicated naming constants disappear from the canvas. Constants still used by other nodes or formulas remain available to those consumers; editing or deleting them no longer changes the grouped names. Edit the block's **Left**/**Right** fields to rename its channels. Channels already used by another mapping, duplicate names and unresolved expressions must be repaired first.
- **Convert Direct Audio Output:** select one `outs` node and review its replacement with a named stereo outlet pair. Each side retains its incoming wires, literal value and input formula. Missing inputs remain unconnected. Other direct outputs are left intact. This changes the destination model: connect the named ports in the performance mixer to hear them.
- Ungrouped old patches are not converted automatically. Previously saved valid stereo groups receive the same channel-name cleanup when opened, including groups that are currently expanded. Stale mappings show a reason and **Repair** / **Remove mapping**. Repair can bind a chosen existing pair or explicitly **Create missing channels**. Removing an unresolved mapping removes only metadata, preserving existing nodes.
- Conversion and renaming preserve role and activation. Previews list affected current-performance routes for explicit repair in the mixer after saving; other saved performances remain unchanged.

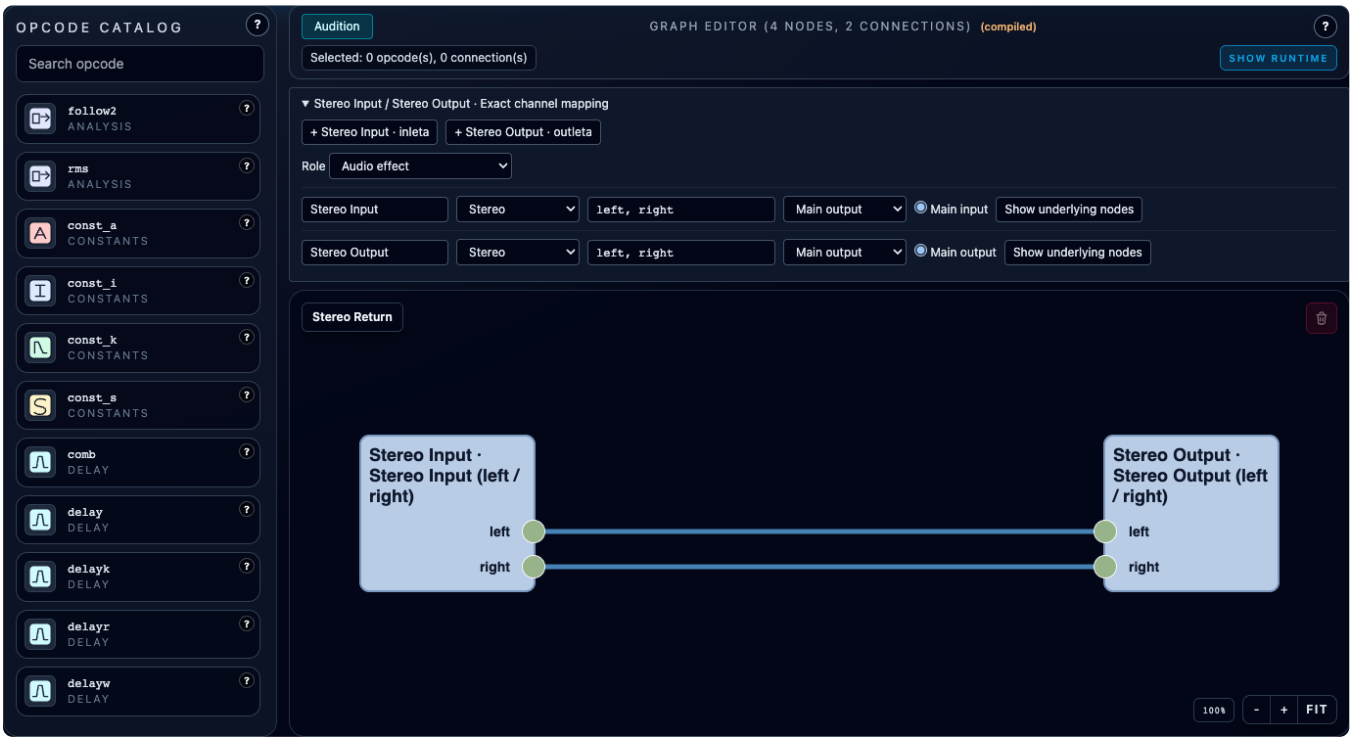

A continuous pass-through effect with explicit left/right groups and collapsed Stereo Input and Stereo Output blocks.

Navigation: Up | Previous | Next

Input Formula Assistant

Navigation: Up | Previous | Next

The Input Formula Assistant shows the effective value of an input and lets you edit a constant or define how incoming signals are combined.

When To Use It

Most commonly, you use it when **multiple outputs are connected to the same input**.

Examples:

- Mix two modulation sources before feeding a filter cutoff input
- Scale one source more than another (`in1 * 0.8 + in2 * 0.2`)
- Apply utility transforms (`abs(in1)`, `dbamp(in1)`) before an input

It is also useful for advanced cases with only one or even zero connected inputs:

- One input: reshape or scale the signal (`ceil(in1)`, `in1 * 0.5`)
- Zero inputs: edit a stored constant (`650`, `-0.25`, `1e-7`) or provide an expression (`ampdb(6)`)

How To Open It

- Double-click an **input socket** in the graph editor.
- The assistant opens for that exact target node + target input port.

What The Assistant Shows

Effective Value and Source

The initial field follows the compiler's precedence:

Source	Initial contents	Bright-green socket
Saved formula	Saved expression	Yes
Connected signals	Connection token or sum	No
Stored value	Stored constant	Yes
Catalog default	Inherited default	No
Unset	Empty field	No

A stored value overridden by a connection or formula appears separately as an **Inactive stored fallback**. The **After reset** line previews what will control the input once both saved layers are removed. Boolean constants are editable as their numeric equivalents, `0` and `1`.

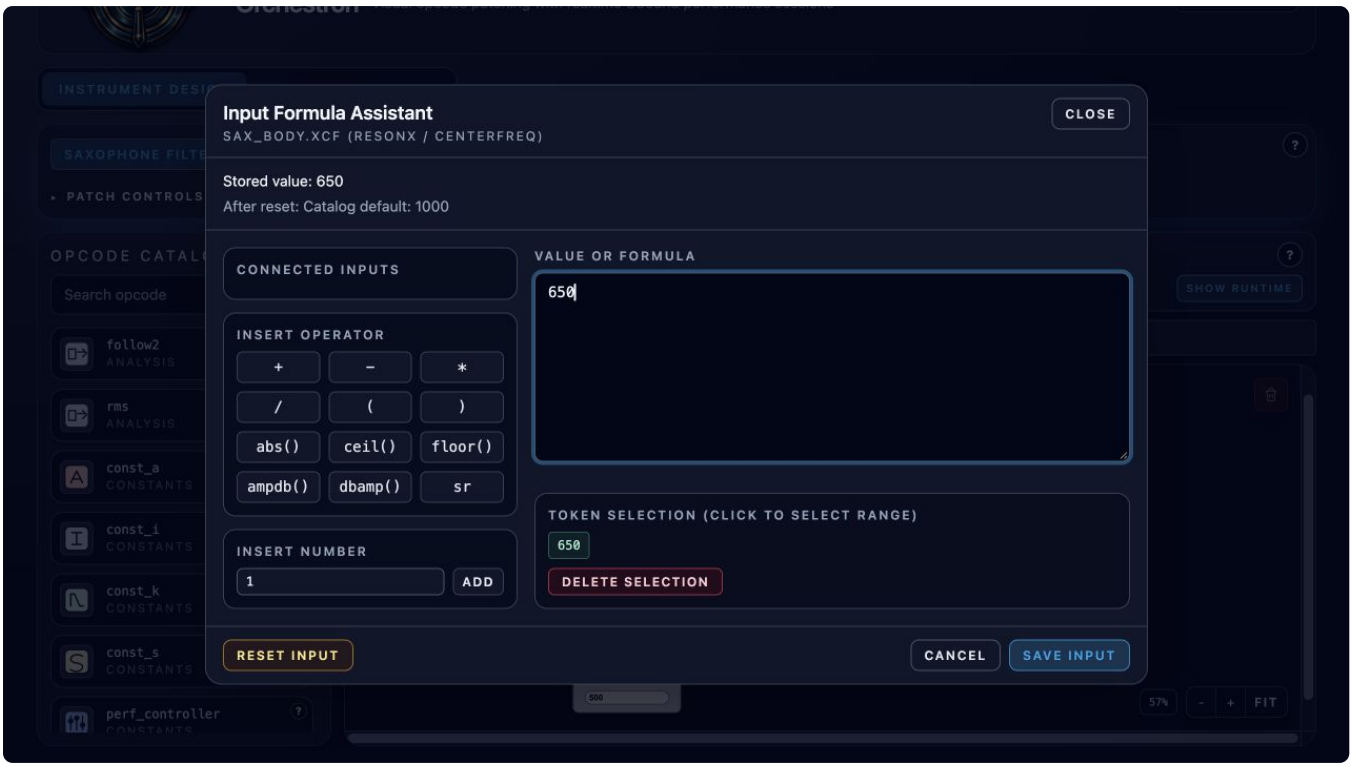

Connected Inputs (Token Mapping)

The left panel lists inbound connections to the target input and assigns tokens such as:

- `in1`
- `in2`
- `in3`

Each token shows source information (node/opcode/port) so you can see what signal you are referencing.

Operator and Function Buttons

The assistant provides insertion buttons for:

- Operators: `+`, `−`, `*`, `/`
- Grouping: `(`, `)`
- Unary functions: `abs()`, `ceil()`, `floor()`, `ampdb()`, `dbamp()`
- Built-in literals: `sr` (configured audio sample rate)
- Custom numeric literals via the number entry field

Formula Text Area

You can type directly into the formula field.

Token Selection / Delete Helper

The assistant tokenizes the current expression and displays clickable tokens.

- Click a token to select that exact range in the text area.
- Use `Delete Selection` to remove the selected token/range quickly.

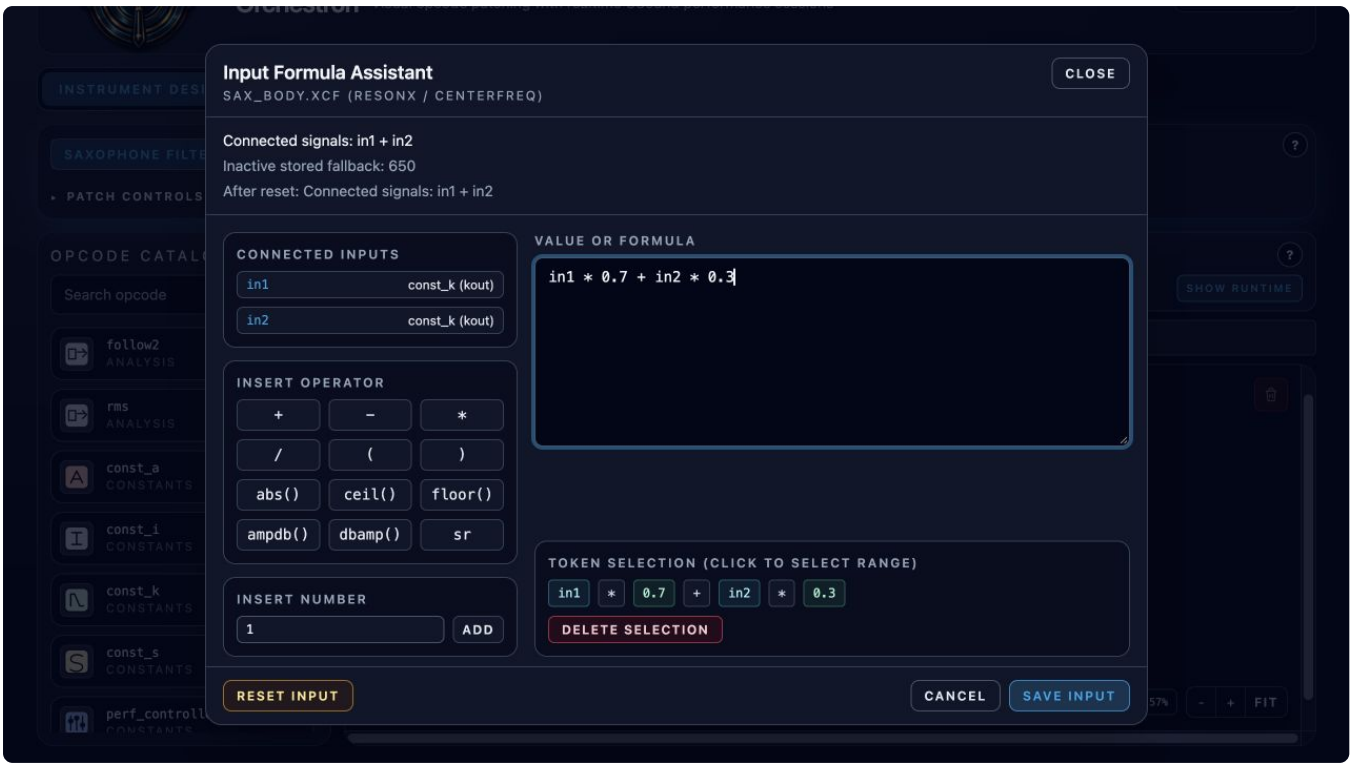

Input Formula Assistant with connected input tokens and editable formula.

Validation Rules

The formula must be syntactically valid before it can be saved.

Validation includes:

- Non-empty expression
- Valid numeric literals
- Supported characters/operators only
- Balanced parentheses
- Valid token ordering
- Known identifiers only (`inN` tokens, supported built-in literals like `sr` , or supported unary function names)

If validation fails:

- The formula text area border changes to an error style
- A list of validation errors is shown
- `Save input` is disabled

Save / Reset / Cancel

- **Save input** stores an unconnected finite numeric literal directly as the input's value, removing any previous formula. Scientific notation is accepted for these constants. Arithmetic expressions and values entered on connected inputs use formula metadata and replace the obsolete stored value. The formula language itself is unchanged.
- **Reset input** removes both the stored value and saved formula for this input. Connections stay in place; otherwise the catalog default applies, or the input becomes unset. A required input without a remaining source fails compilation with a missing-input diagnostic.
- **Cancel** / **Close** discards dialog edits. Opening the dialog alone does not modify the patch.

Specialized Inputs

`maxalloc.icount` accepts only a numeric literal. String inputs use literal text editing. Disconnect a literal-only or string input before editing its literal value. Stereo channel names remain managed by Stereo Input / Stereo Output configuration, and

`zfilter2` coefficients retain their dedicated inline list editor. Stored values remain visible in socket tooltips and highlighting for these inputs too.

Collapsed Stereo Output blocks preserve the values and formulas on their underlying left/right audio inputs. Collapsed control-flow captures retain their existing edit restrictions.

What Happens During Compile

- Inputs resolve in this order: saved formula → connected signals → stored value → catalog default → unset.
- If multiple connections target one input and **no** formula is defined, Orchestron falls back to summing the inbound signals.
- If a formula exists, the compiler resolves the token bindings and compiles the expression.
- Formula validation also happens backend-side, so invalid formulas fail compile with diagnostics.

Formula Examples

Weighted Mix

```
in1 * 0.7 + in2 * 0.3
```

Gain Conversion

```
dbamp(in1)
```

Rectified Modulation

```
abs(in1)
```

Constant Override

```
0.5
```

Sample-Rate-Aware Expression

```
sr * 0.5
```

Navigation: Up | Previous | Next

GEN Table Editor (GEN Meta-Opcode)

Navigation: Up | Previous | Next

The `GEN` meta-opcode is a specialized table generator node that opens a dedicated editor for building `ftgen / ftgenonce` lines visually.

Why The GEN Meta-Opcode Exists

Csound function-table generation is powerful but argument-heavy. The GEN editor provides a structured UI for common GEN routines and a preview of the generated `ftgen` / `ftgenonce` line.

Opening The GEN Editor

- Add the `GEN` opcode to the graph.
- Click the `GEN` button on the node card.

Editor Overview

The editor has two main areas:

- Left side: generation mode, routine selection, parameters
- Right side: preview (`Effective GEN` , flattened args, rendered line) and notes

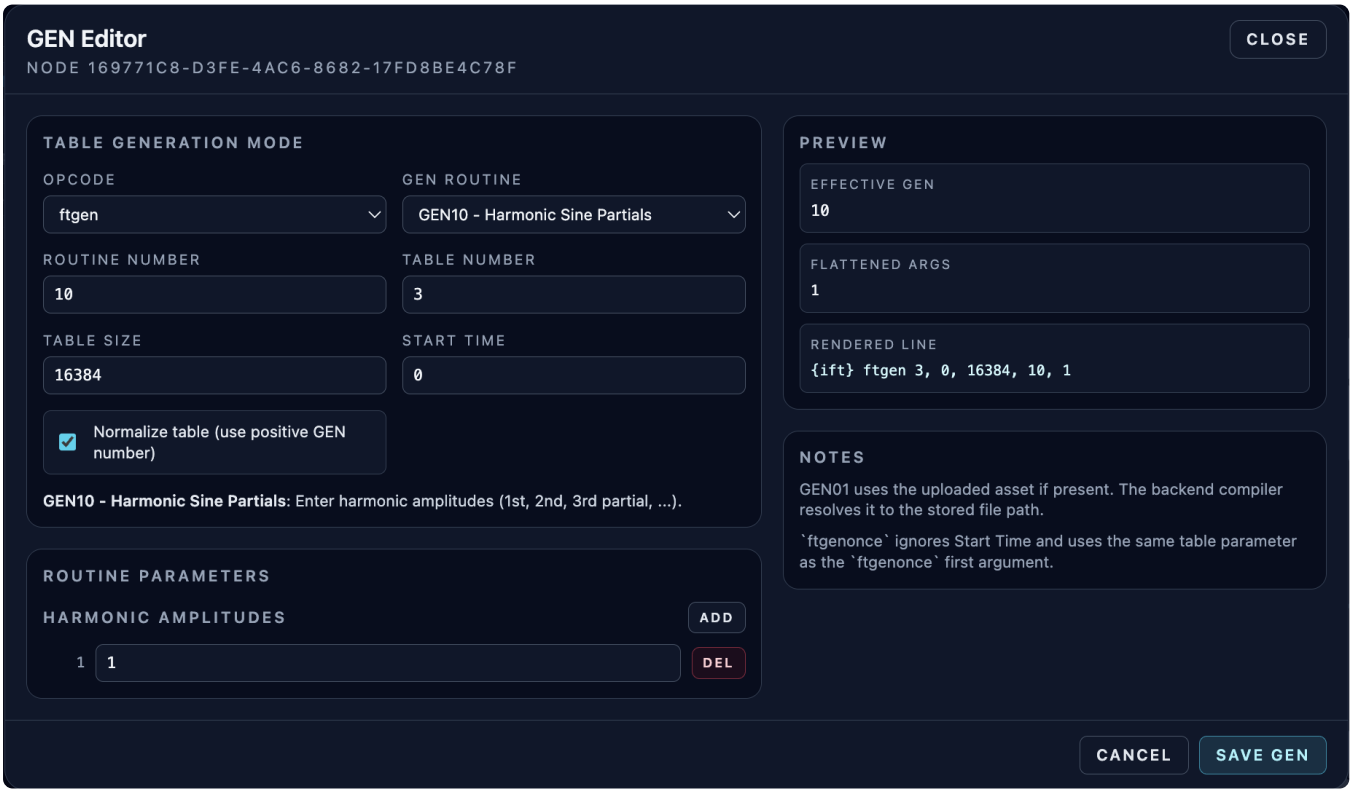

GEN editor in structured routine mode with live preview panel.

Table Generation Mode (Top Section)

Opcode Mode

Choose which Csound table opcode style to generate:

- `ftgen`
- `ftgenonce`

Routine Selection

You can choose from structured editors for common routines and a named routine option:

- `GEN01` - Audio File
- `GEN02` - Value List
- `GEN07` - Segments
- `GEN08` - Cubic Spline Segments
- `GEN10` - Harmonic Sine Partials
- `GEN11` - Harmonic Cosine Partials
- `GEN17` - Step Table From x/y Pairs
- `GEN19` - Sinusoidal Partials with DC Offset
- `GEN20` - Window / Distribution Function
- `GENpadsynth` (named routine)
- Custom routine number / raw arguments

Core Table Fields

- `Routine Number`
- `Table Number`
- `Table Size`
- `Start Time` (disabled for `ftgenonce` UI mode)
- `Normalize table` (positive GEN number vs negative GEN number behavior)

The backend validates GEN table settings before storage and again before Csound orchestra emission. `Table Size` can be `0` only for `GEN01`; other routines require a positive size. The maximum table size is `4194304` samples, and raw or structured GEN argument lists are capped at 512 entries. Raw argument text is limited to 8192 characters, with each token limited to 512 characters.

Supported Routine Editors (Structured UI)

Routine	Purpose	Main Parameters In Editor
`GEN10`	Harmonic sine partial amplitudes	Harmonic amplitude list
`GEN11`	Harmonic cosine partials	harmonic count, lowest harmonic, multiplier
`GEN02`	Literal value list	Numeric value list
`GEN07`	Segments	Start value + repeated length/value rows
`GEN08`	Cubic spline segments	Start value + repeated length/value rows
`GEN17`	Step table from x/y pairs	x/y pairs (editable list)
`GEN19`	Sinusoidal partials with independent phase and DC offset	repeated partial number / strength / phase / offset rows
`GEN20`	Window / distribution	window type, max, optional parameter (for selected windows)
`GEN01`	Load audio file into table	uploaded asset, skip time, format, channel

GEN19: Sinusoidal Partials with DC Offset

Choose **GEN19** and add a row for each sinusoid. Each row emits four arguments in order: **Partial number**, **Strength**, **Phase (°)**, **DC offset**. Partial numbers are positive frequency ratios, including fractions; strengths may be negative. The offset is added after strength scaling. Phase uses degrees.

The initial row is `1, 1, 0, 0`. For an `hsboscil` octave window, use `1, 0.5, 270, 0.5`, table size `1024`, and disable normalization. Use a power-of-two table size or that size plus one. Up to 128 partials fit the 512-argument limit. Both `ftgen`

and `ftgenonce` support these rows; **Normalize table** selects `19` or `-19`.

Existing GEN19 patches with raw arguments keep their raw editor and values on save/reload. Structured rows take precedence if a configuration also contains old raw arguments. See the Csound GEN19 reference.

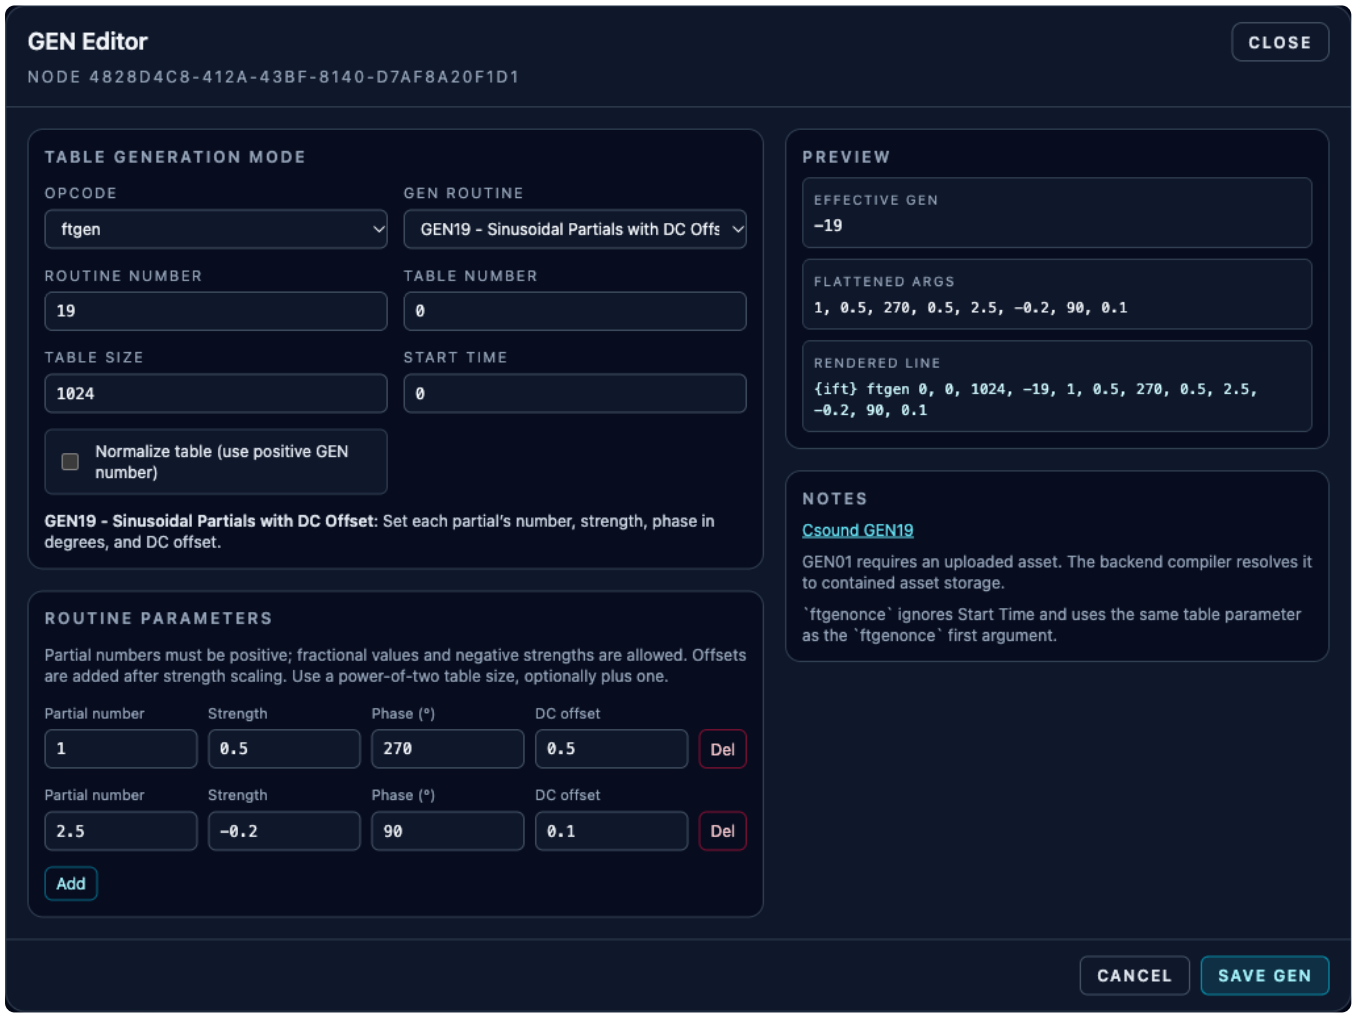

`GENpadsynth` (Named Routine)

Orchestron supports a named GEN routine for padsynth (`GENpadsynth`).

- Select the `GENpadsynth` routine option.
- Enter padsynth parameters and partial amplitude/frequency pairs in the raw arguments editor.
- Argument order follows the Csound `GENpadsynth` documentation.

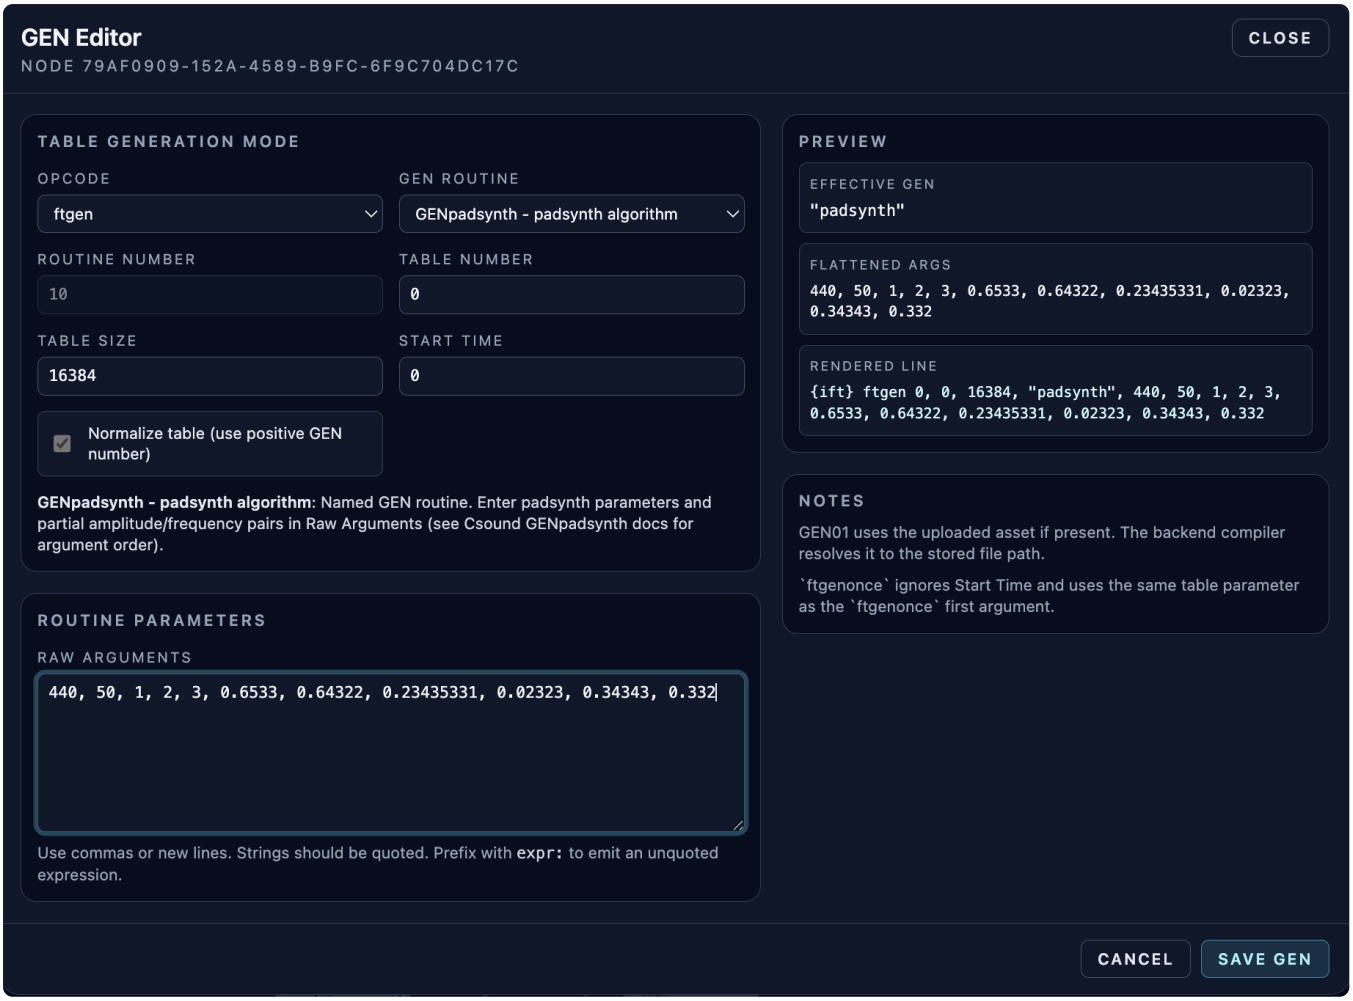

GENpadsynth named routine using raw arguments input.

Custom / Raw Routine Arguments

For routines without a dedicated structured editor:

- Use the raw arguments editor.
- Separate tokens by commas or new lines.
- Quote string literals manually (for example `"file.wav"`).
- Prefix an argument with `expr:` to force raw Csound expression rendering without quotes.
- Keep raw argument lists within the accepted limits: 512 tokens total, 512 characters per token, and 8192 characters for the raw text field.

Examples:

```
1, 0.5, expr:1024*2, "file.wav"
```

GEN01 Audio File Workflow (Important)

The `GEN01` editor loads sound files through uploaded backend assets:

Uploaded Asset

- Click `Upload Audio File`
- Select an audio file (UI accepts common formats such as WAV/AIFF/FLAC/MP3/OGG and `audio/*`)
- Orchestron stores the uploaded asset on the backend and references it in the GEN node config

Benefits:

- Instrument/performance exports can include the audio file automatically
- Imports can restore the asset automatically when using ZIP bundles
- Runtime compilation resolves the asset inside the configured backend asset directory
- Raw absolute paths and parent-traversal paths are rejected before Csound starts

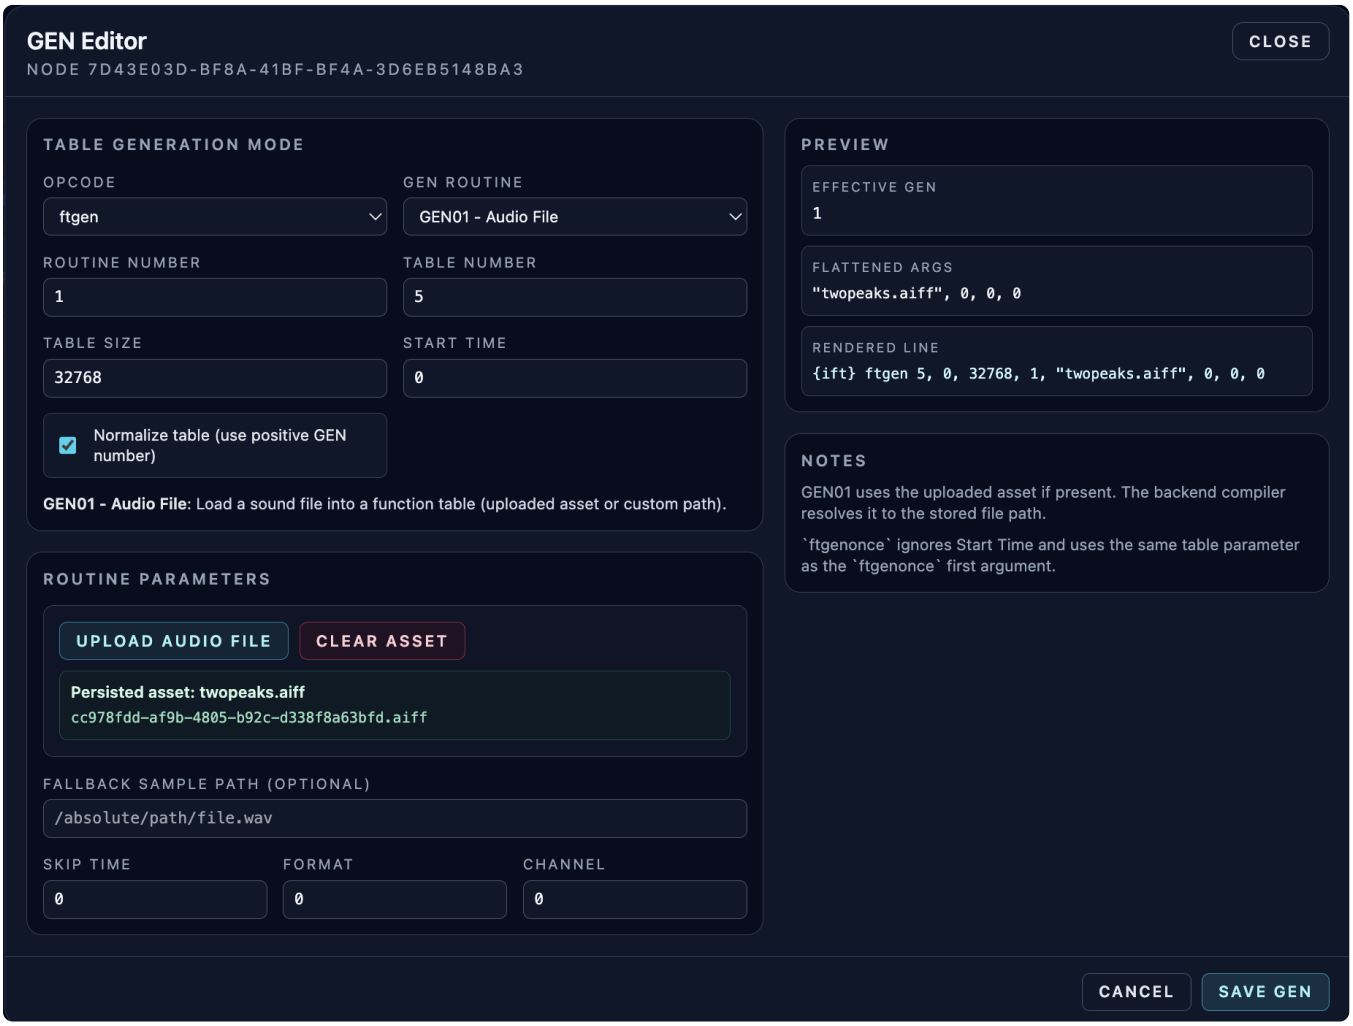

GEN01 table setup using an uploaded backend audio asset.

Additional GEN01 Parameters

- `Skip Time`
- `Format`
- `Channel`

Upload Limits and Notes

- Default backend upload limit for GEN audio assets is 64 MiB per file (`VISUALCSOUND_GEN_AUDIO_ASSET_MAX_BYTES`)
- Persistent generated-asset storage is also capped by total bytes (`VISUALCSOUND_GEN_AUDIO_ASSETS_MAX_TOTAL_BYTES`, 1 GiB by default) and asset count (`VISUALCSOUND_GEN_AUDIO_ASSETS_MAX_COUNT`, 1024 by default)
- Garbage collection removes unreferenced generated assets older than `VISUALCSOUND_GEN_AUDIO_ASSET_GC_MIN_AGE_SECONDS` (24 hours by default)
- The backend streams GEN audio uploads and rejects requests once they exceed the configured cap
- ZIP bundle imports also enforce request, member-count, JSON, and total-uncompressed-size caps before importing bundled assets
- Deployments should set a matching maximum request size at the ASGI server, reverse proxy, or load balancer

- The preview/note panel explains that uploaded assets are required for GEN01 sample loading

Preview Panel

The preview section shows:

- `Effective GEN` (including normalization sign / named routine behavior)
- Flattened argument list
- Rendered `ftgen` / `ftgenonce` line preview

Use this to verify the exact Csound line before compile.

Editor Notes Worth Knowing

- `ftgenonce` behavior differs from `ftgen` and the editor shows notes about the differences.
- `GEN01` asset-backed generation is resolved by the backend compiler to contained stored asset paths.
- `GEN01` asset dependencies are automatically bundled in ZIP exports when referenced.

Save Behavior

- `Save GEN` writes the GEN node configuration into patch UI metadata (`graph.ui_layout`).
- Closing without saving leaves the previous GEN node settings unchanged.

Navigation: Up | Previous | Next

If, Switch and Drumset

Navigation: Up | Previous | Next

Choose a Synthesis Case at Note Start

If and **Switch**, in the **Control flow** catalog category, are editor constructs that generate Csound `if/elseif/else/endif`. Each synthesis voice chooses one case when it starts. Only that case initializes and performs; selection remains fixed until the voice ends. Different notes, including repetitions of the same note, have independent voices.

1. Add If or Switch. New blocks produce stereo audio. If compares two init-rate inputs using `==`, `!=`, `<`, `<=`, `>` or `>=`. Switch compares its init-rate selector with unique numeric case values. For drums, connect `notnum.inote` to `selector`. Input connections and the existing arithmetic Input Formula Assistant take precedence over literal values in **Configure branches**.
2. Expand the cases on the canvas and drag opcodes from the catalog into a case. A valid drop into**Silence** changes it to **Synthesis** and exposes its result inputs. Ordinary canvas drops create main-graph nodes. Dropping onto a collapsed block expands it; drop again into a specific case.
3. Connect each synthesis case to its managed **Case Result** audio inputs. Several sources can sum at an input, or use an input formula. Connect the block's common output to shared processing or a named Stereo Output. Every non-silent case needs a value for every result channel, even if that case is rarely selected.
4. False and Default initially use **Silence**. Silence assigns zero to every output and contains no synthesis nodes. Changing an existing case to Silence previews the nodes, connections and formulas that will be removed. Switch cases can be renamed, reordered, added and deleted; Default remains last and cannot be deleted. If keeps its two ordered True/False cases.
5. Collapse to see the condition, case summary, shared inputs and common output. Node IDs, selection, positions and the viewport survive collapse, save and reload. Captured input formulas are retained; expand to edit them. Format changes show an impact preview: mono preserves `left` and removes `right` wiring only after confirmation; switching back to stereo leaves `right` unconnected. Block/case deletion removes owned content and associated formulas and configuration together.

Arrange and transfer nodes

The block stays above and left-aligned with the first case, with room for its complete body and buttons. Case frames measure every visible parameter row at any zoom. They automatically grow to the right and bottom and retain their size when nodes move inward, are removed, or become smaller; **Case Result** stays in a separate row at the bottom-right. Later cases move down or up without changing their internal arrangement.

Drag the **Case Result title or unused body** to resize a branch in both dimensions, including shrinking it. This also works for the Silence marker. The pointer displacement sets the requested size; complete node bounds, padding and the reserved result row limit how small it can become. Other selected nodes stay still and selection is preserved. Sockets, formula editing and help buttons keep their usual actions. Release commits the size and managed positions together; **Escape**, cancellation or focus loss restores the full prior layout. Ordinary dragging retains the largest frame size reached during the gesture. Bright purple borders stay about two screen pixels wide at every zoom. Sizes survive collapse, block movement, reordering, save/reload, copies and export/import.

- Drag main-graph nodes into an expanded case to transfer them. Dragging existing case members normally only rearranges them within that case. Their top and left movement limits keep them clear of the header; selected groups keep their relative spacing. Moving the block moves all contents together.
- Hold **Alt/Option before starting the drag** to transfer case members to another case or the main graph. Keep the desired group selected. **Move selected nodes** in Configure branches is the alternative, with identical validation.
- A transfer requires every connection and formula binding touching a moved node to end at another moved node or a node already in the destination scope. A cyan border marks an allowed target; a red border and diagnostic list refused connections and unselected endpoints. Nothing is automatically selected, disconnected or deleted. An isolated ordinary node can always transfer.
- This restriction applies to **membership transfers**. Existing shared main-graph inputs remain valid; you can deliberately connect them after placing the nodes. Case-local values still leave through Case Result only. Main-graph-only constructs, nested blocks and independent Case Result transfers are refused.

- Release to commit positions and membership together. A refused drop, **Escape**, pointer cancellation or loss of focus restores the entire group and clears highlighting. A refused drop leaves Silence unchanged.

Play the Drumset template

Choose **New from template → Drumset** for 17 analog percussion voices: **35/36** kicks, **37** side stick, **38** snare, **39** clap, **40** electric snare, **41/43/45/47/48/50** toms, **42/46** closed/open hi-hat, **49** crash, **51** ride and **54** tambourine. Other notes select Default Silence. Velocity scales finite envelopes; branch-local `xtratim` preserves tails after short note-offs. Shared BD snappiness, shared hi-hat tuning and snare tuning default to **1**, preserving the previous sounds. Each added voice has one characteristic control. Save, assign a MIDI channel in Perform and route Stereo Output through the mixer. Controls affect newly started notes.

Processing after the block runs **once per voice**. Put effects in the performance mixer when they should process the **sum of the whole kit**. For velocity-dependent timbre, feed `ampmidi` into an If selector and place the alternative sounds in its cases.

Continuous switching, nested blocks, numeric results, raw jumps, crossfades, MIDI learn, ranges and choke groups are outside this release. Keep `outs`, `inleta`, `outleta`, `sfload`, `maxalloc` and `midi_legato` in the main graph. GEN nodes and duration controls may belong to cases. Cycles and illegal scope crossings are compile errors, and all cases are validated.

Saving and sharing

Save the instrument to retain its branches, case sizes and formulas. Native export/import and performance exports preserve the selected synthesis behavior. For programmatic authoring, see the control-flow technical reference.

Navigation: Up | Previous | Next

Monophonic Legato

Navigation: Up | Previous | Next

Phrase Behavior

Add one **midi_legato** node from **MIDI** to the main graph of a MIDI instrument. Connect **Frequency (Hz)** to control-rate oscillator frequencies and **Velocity** to amplitude and blowing-pressure formulas. At least one **madsr** phrase envelope is required; compilation fails without it. Keep **cpsmidi** and **ampmidi** for initial values needed by init-rate inputs.

- Touching and overlapping notes share one synthesis voice. Pitch changes immediately without glide. Envelopes, oscillator phase, resonances and vibrato continue.
- The last note played has priority. Releasing it returns to the most recently held note, including its velocity.
- Velocity is MIDI velocity divided by 128, matching `ampmidi 1`. Its first value is immediate; changes have a **10 ms half-time**.
- A gap starts a new phrase, even during a release tail. Initialization contours and performance-controller values refresh then. Connected notes retain phrase-initial settings and do not repeat attack bursts or pressure overshoots.
- All events within the existing Csound control boundary settle together. There is **no extra legato grace period**. Live resolution follows `ksmps`; CSD exports use `ksmps=1`.
- `madsr` retains delay, attack, decay, sustain and release semantics, including explicit `ireltim`. A new phrase applies a 2 ms amplitude declick to interrupt an old release; this is not pitch glide.
- `maxalloc` limits synthesis voices; internal note collectors remain available. Each rack instance has separate state. Live MIDI, sequencers and MIDI/SCORE CSD exports use the same phrase logic. Stopping the instrument engine clears its phrase state; MIDI panic releases the phrase. Stopping a sequencer releases its own notes without stopping other devices or the rack engine.

This is an **Orchestron virtual opcode**, not a native Csound opcode. Use it once, outside If/Switch cases, on a MIDI-activated instrument. Native release/turnoff opcodes other than `madsr` are not supported in a legato graph. Patches without this node retain their normal voice behavior. Recompile/restart a running instrument after adding it.

Steel-string guitar example

The supplied pair includes a legato guitar and a polyphonic guitar. The legato patch rounds its pitch slide to equal-tempered semitones. Connected notes preserve the decaying string without another pluck; the phrase's initial velocity sets its pluck strength. A gap starts a new pluck and reads the current performance controls. The Slide setting is a smoothing half-time, so larger intervals take longer to reach their final fret. The matching polyphonic patch gives each note its own pluck for chords and manually timed strums. Both guitars use `wgpluck2` with velocity-sensitive pluck position and adjustable damping. Since that model's pitch is fixed at initialization, a period-synchronous delay readout transposes the ringing string. The guitar rounds the target pitch before this readout, with no pitch smoothing afterward and no new excitation at fret changes. Its default Slide half-time is 0.12 seconds so intermediate frets have more distinct dwell times. Reset an existing rack override to adopt this default.

Expressive tenor saxophone example

Tenor Saxophone — Expressive Legato uses one continuous STK reed model with velocity-sensitive breath, reed response, filtering and loudness. Play concert A♭2–E5 (MIDI 44–76); there is no automatic tenor transposition. Route **Stereo Output → Master** in Perform.

The balanced acoustic revision retunes the reed's blow position, aperture, stiffness and pressure together, then adds body at 850 Hz and presence at 2,400 Hz. Normal velocities have a richer upper-harmonic spectrum, with restrained breath noise and output gain compensated for the stronger excitation. The example guide provides level-matched before/after previews; the four controls keep their IDs, ranges and defaults so existing performance overrides remain valid.

Velocities 1–95 cover normal playing, 96–114 progressively add growl, and 115–127 blend in a sustained overblown octave squeal. Hard connected notes can provoke the squeal; softer notes and held-note fallback restore their own velocity. Gaps restart a 10–35 ms tongue contour followed by 80 ms settling; release is 100 ms. Tone, Growl, Squeal and Vibrato are 0–1 performance controls, with defaults 0.5, 0.35, 0.65 and 0.2. Zero Growl/Squeal disables its added effect. Knob changes apply

at the next phrase, while note velocity remains live throughout legato.

The native model fixes pitch at initialization, so period-synchronous delay readouts transpose the same reed signal for legato and overblowing. The readout can add up to two initial note periods of delay, and wide leaps can color the timbre. The graph stays editable using existing opcodes. Csound STK and its rawwave data are required and included in Orchestron's Docker image. The example guide includes native import instructions, control translations and WAV previews.

Further reading

For implementation and authoring references, see the legato technical notes.

Navigation: Up | Previous | Next

Defining Performance Controls

Navigation: Up | Previous | Next

Add controls to an instrument when performers should be able to adjust its sound without editing the graph. Each rack instance can retain different values.

Define a Controller

Add `perf_controller` from Constants in Instrument Design. Click its gear button and configure:

Type	Configuration	Perform control
Float	Minimum, Maximum, Default, linear/logarithmic Scale, Label, Controller Group	Existing knob with exact numeric entry
Integer	Whole-number Minimum, Maximum and Default, Label, Controller Group; linear only	Up triangle, editable number, down triangle
Bool	On/Off Default, Label, Controller Group	On/off toggle, output `0` or `1`

The Type selector initially chooses **Float**. Existing patches without a stored type also use Float. Float defaults remain minimum `0`, maximum `1`, default `0.5`, linear scale and label `Parameter`. Integer definitions without numeric fields use minimum `0`, maximum `1`, default `0`. Bool defaults to **Off** and uses a fixed internal `0–1` linear range.

The fields are fixed values, without input sockets or formulas. Float numbers must be finite, minimum smaller than maximum, and default within the range. Logarithmic minimum must be positive. Integer has the same ordering requirements and accepts only safe whole numbers from `-9007199254740991` to `9007199254740991`. Labels must contain non-whitespace text and be no more than 128 characters long.

Changing Type to Integer rounds the numeric fields to the nearest integer, with ties toward positive infinity (`-1.5` becomes `-1`). If rounding makes minimum and maximum equal, correct the range before saving. Changing to Bool maps a default below `0.5` to Off and any default at least `0.5` to On. Integer and Bool always use linear scale.

For a logarithmic attack-time control, choose Float, minimum `0.001`, maximum `5`, default `0.01`, and label `Attack (s)`. Connect `iout` to the envelope's I-rate attack input, then save the patch. All three types use the same numeric `iout` socket and existing compatible connections. Continuous instruments automatically use live updates when every connected input accepts K-rate.

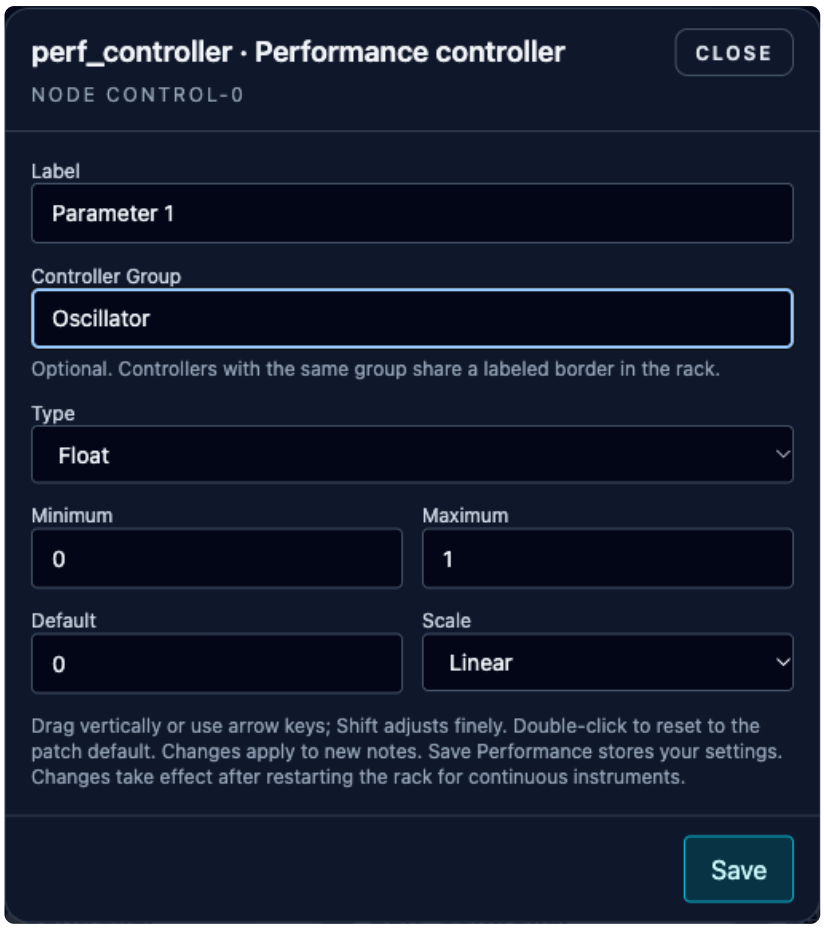

Organize controls into groups

Set **Controller Group** in the gear dialog to organize related controls, such as Oscillator, Filter or Envelope. The optional field accepts up to 128 characters. Names are case-sensitive; surrounding spaces are ignored. Leave it blank for ungrouped controls.

Matching group names appear inside one labeled border in Perform. Groups and controls follow their order in the patch. Group names survive native import/export and do not affect values, MIDI assignments or sound. Renaming a group does not require restarting the rack.

Use the controls in Perform

Save the instrument, add it to a rack, and adjust its controls. See Using Performance Controls for MIDI assignments, saving and the Live / Requires restart labels.

Navigation: Up | Previous | Next

Runtime Panel and Compilation Workflow

Navigation: Up | Previous | Next

The Runtime panel supports live testing of the current patch and exposes compile/runtime diagnostics for the unified browser-clock runtime.

Runtime Panel Sections

- `MIDI Input` selector
- `Compile Output` (generated ORC)
- `Browser Audio` status for browser-clock mode
- `Session Events` log (recent backend/runtime events)

Runtime panel showing MIDI input selection, generated ORC, active browser-owned PCM runtime, queue diagnostics and recent session events.

MIDI Input Binding

Use the `MIDI Input` dropdown to bind an external MIDI input to the active runtime session.

- This is how external hardware/DAW MIDI reaches the running instrument session.
- `internal:loopback` is always available as the built-in app-only loopback.
- Internal sequencers, piano rolls, and manual controller lanes always go straight to the session engine and do not depend on this binding.

- The selected input is session-specific and also reflected on the Performance page status footer.

See MIDI Setup and Inputs for OS-level MIDI setup guidance.

Compile Output (Generated ORC)

After compiling a patch, the panel shows the generated ORC text.

This is useful for:

- Verifying the compiled signal chain
- Understanding how formulas/GEN/meta-opcodes are rendered
- Debugging compile issues with concrete generated code

A compilable patch must contain at least one audio output sink: either `outs` for direct stereo output or `outleta` for named audio routing into another instrument. This allows source instruments to feed always-on effects without also writing directly to the final stereo output.

Session Events

The event list shows recent session events (most recent first, limited window in the UI).

Typical uses:

- Confirm that start/stop events happened
- Check runtime errors and session state
- Diagnose issues during live testing

Browser Audio

The runtime now always uses `browser_clock` audio mode, so the Runtime panel can show:

- Browser PCM queue/runtime status (`connecting`, `live`, `error`)
- Error message when the controller socket or AudioWorklet path fails
- No `<audio>` player, because playback is owned directly by the browser `AudioContext`

See Browser-Clock Latency for the browser-clock mode guide and tuning workflow.

Runtime Panel Collapse / Show

Runtime starts collapsed. Open it with Show runtime in the graph header and close it with Hide. Your choice survives view switches until browser reload; compilation and playback continue while it is hidden.

To maximize graph editor space:

- Click `Hide` in the Runtime panel to collapse it
- Use `Show runtime` in the graph header to bring it back

This is especially useful on smaller displays or complex patches.

Compile Status Badge (Graph Header)

The instrument page header above the graph displays a compile-state badge for the current patch/tab snapshot:

- `compiled` - the current graph snapshot matches the last successful compile
- `pending changes` - the graph changed since the last successful compile
- `errors` - the last compile for this exact snapshot failed

This helps you avoid assuming a graph is live-valid after edits.

Recommended Test Loop

1. Edit the graph.
2. Click **Audition** to compile and start an isolated preview.
3. Play the mini keyboard and adjust the compact performance controls.
4. Click **Stop audition** when finished.
5. Save the instrument when you are happy with it.

The integrated ? help for this panel now focuses on session-specific external MIDI binding, generated ORC inspection, runtime event feedback, and browser-clock PCM playback.

Audition a draft

Audition is a toggle in the graph header. Click it to compile the current draft and start an isolated preview. While it is running, the button reads **Stop audition**, and compact performance controls and a mini keyboard appear below it. Click again to stop and hide both. The graph stays accessible throughout. Compile remains available to validate the graph without starting audio.

Compilation finishes before any current performance is stopped. A compilation error leaves that performance running and appears beside the audition controls. During preparation, click **Preparing… · Cancel** to cancel the pending start.

Play and adjust the sound

- **Float** controls have a label and tiny knob. Drag vertically or use the arrow keys; hold Shift for fine adjustments. Hover to see the value, or double-click to reset to the patch default.
- **Integer** controls have a label and number field. Type a whole number or use Up/Down; Home/End selects the limits. Enter or leaving the field applies an entry; Escape cancels. Fractional entries are rejected and whole numbers are clamped to the configured range.
- **Bool** controls have a label and small on/off toggle.

- Controls appear in one flat row that wraps as needed. There are no controller groups, MIDI assignments or automation controls in audition.
- Hold a piano key to play, then release it to stop the note. Focused keys also respond to Space or Enter. The– and + buttons shift the two-octave keyboard. It uses fixed velocity, with no scale, mode or automatic chords. Starting audition does not hold a test note automatically.

Controls start at the patch defaults. Changes are temporary and do not change the library patch or a performance. They survive stopping and restarting audition for the same instrument, and reset when you leave Instrument Design or switch instrument tabs. Changing a controller range or type reconciles its temporary value on the next start.

Note instruments use changed values on the next note or phrase. Continuous instruments update live when all connected inputs support it; other settings require stopping and restarting audition. Hover over a control for its timing behavior. See Using Performance Controls.

Effects and continuous instruments

An isolated effect automatically receives the built-in playable test instrument. Play the mini keyboard to send that source through the effect. Audition needs a stereo main output; effects also need a stereo main input compatible with the built-in source. Direct Audio Output is supported.

Continuous generators start immediately. If there is no note-triggered instrument in the preview, the mini keyboard remains visible but disabled.

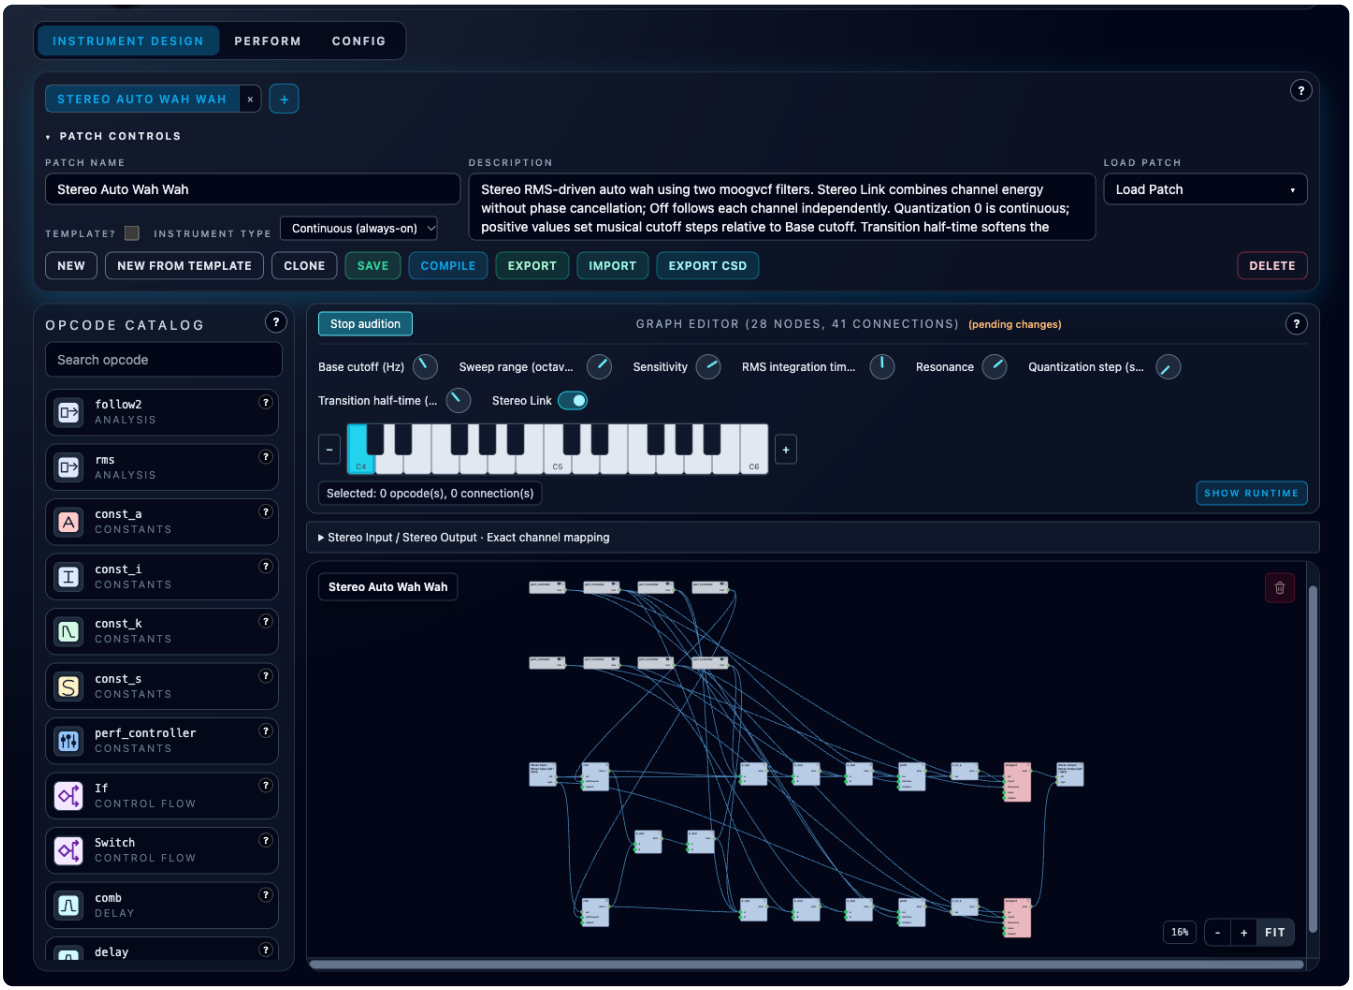

Stop or change the draft

Stopping audition releases held notes, removes the temporary session and restores the original performance playhead with transport stopped. Restart that performance explicitly when ready. Switching instruments or leaving Instrument Design also ends audition. Temporary patches, test sources and routes are never saved into the library.

If you edit the graph during audition, a **Draft changed** message reminds you to stop and restart audition to hear the updated draft. Performance controls continue to address the running snapshot until then.

Audition notes for branching instruments

The mini keyboard initially shows the octave containing the first MIDI-number Switch case (36 for Drumset); other patches start at 60. Play 38 or 42 to test other drums, or an unmatched note to check Default Silence. Use the octave buttons to reach other pitches. See the Branching guide.

Navigation: Up | Previous | Next

Instrument Import / Export and CSD Export

Navigation: Up | Previous | Next

This page documents all user-facing import/export options available from the Instrument Design page.

Instrument Export Types

1. Orchestron Instrument Bundle (`Export` button)

The `Export` button creates an Orchestron instrument definition bundle containing the patch graph and metadata.

File extensions:

- `.orch.instrument.json` (JSON bundle)
- `.orch.instrument.zip` (ZIP bundle)

JSON vs ZIP Export Selection

Orchestron chooses the format automatically:

- JSON export when no referenced uploaded GEN01, `sfload`, or `adsyn` assets are present
- ZIP export when the instrument references uploaded GEN01 audio assets or `sfload` SoundFont assets or `adsyn` analysis files (so audio dependencies can be included)

2. Csound CSD Export (`Export CSD` button)

The `Export CSD` button compiles the patch and downloads a `.csd` file.

Use this when you want to run the generated instrument outside Orchestron in a Csound workflow.

What Is In An Instrument Bundle

An instrument bundle contains:

- patch name
- patch description
- template flag
- instrument type (Percussion, Melody, Bass, Effects / Noise, or Continuous)
- complete editable graph, layout and engine settings

If a ZIP bundle is required, it also contains referenced uploaded audio/SoundFont/analysis assets under an `audio/` folder in the archive.

Instrument Import (`Import` button)

The import button accepts:

- standalone instrument bundle JSON
- standalone instrument bundle ZIP
- performance bundle JSON/ZIP (to extract included patch definitions)

This means you can import patches directly from a performance export without manually splitting files. Instrument Type is preserved for both new and overwritten patches, including imports renamed during conflict resolution. Older bundles without a type receive an inferred type; always-on patches become Continuous and unmatched MIDI patches default to Melody.

Conflict Resolution (Patch Name Already Exists)

When importing one or more patch definitions and a patch name already exists, Orchestron opens a conflict dialog.

Per patch, you can choose:

- `Overwrite` (update existing patch)
- `Rename` (disable overwrite and provide a new target name)
- `Skip` (patch imports only)

The dialog validates name collisions before proceeding. Cancelling leaves the library and audio storage unchanged. After confirmation, the selected patches and their new audio assets are saved together; a failed import rolls back the whole batch. Invalid definitions and duplicate source IDs are rejected before anything is stored.

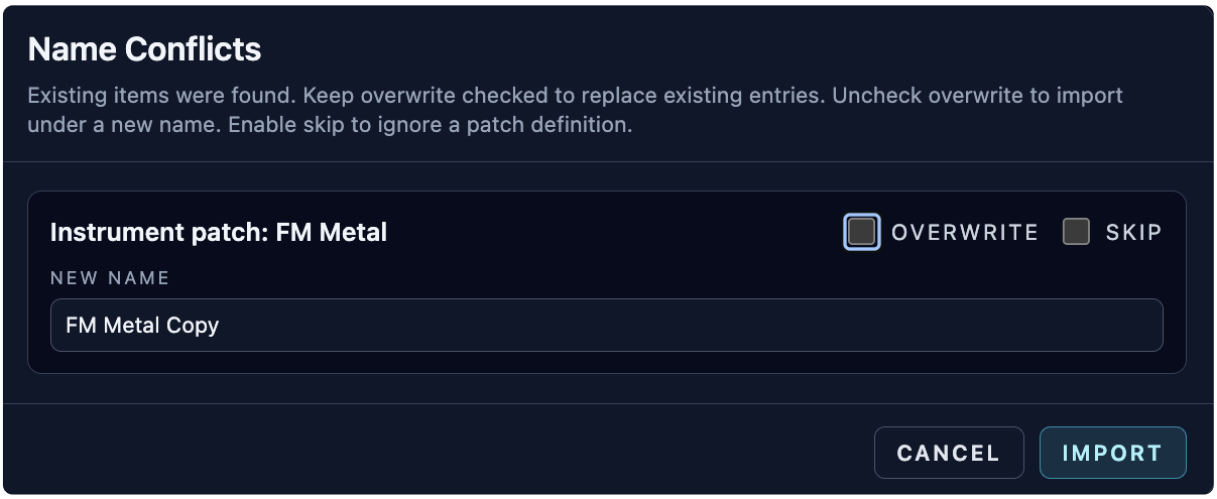

Patch import conflict dialog with overwrite/rename/skip decisions.

Audio Dependency Handling

For ZIP imports that include referenced GEN01 audio or `sfload` SoundFont assets or `adsyn` analysis files:

- Orchestron validates the archive structure
- Stores the selected instruments' audio files after confirmation
- Restores the patch graph with the asset references intact
- Rejects archives that exceed import limits before unpacking large members

Import limits and compatibility

Keep the downloaded archive intact: missing audio assets prevent import. Default limits are 256 MiB for the compressed file, 8 MiB for its instrument/performance definitions, 512 ZIP entries, and 256 MiB total unpacked content. An installation may use different limits; the import error identifies the rejected input.

Instrument Type survives export/import. Older bundles infer Continuous for always-on instruments and Melody for unmatched MIDI instruments. For format and deployment details, see the technical reference.

Instrument Bundle vs Performance Bundle

Use the right export for the right task:

- Instrument bundle: share one patch design
- Performance bundle: share a full performance setup (instrument rack + sequencers + controllers + piano rolls + optional patch definitions)

See Performance Import / Export for performance bundles.

Navigation: Up | Previous | Next

Supported Opcodes

Navigation: Up | Previous | Next

This reference lists **199** user-selectable opcodes and editor constructs in the Instrument Design catalog. Use the catalog search to find an opcode and its **?** button to open integrated help.

How to Use This Reference

- Use the table for quick discovery (category, signal I/O, one-line purpose).
- In the app, click the `?` on an opcode node to open the integrated, localized documentation with port-level details and a Csound reference link.
- The `Inputs / Outputs` columns show port count and signal-rate shorthand.
- Signal legend: `a` = audio-rate, `k` = control-rate, `i` = init-rate, `S` = string, `f` = phase-vocoder frame signal.

Category Counts

Category	Count
analysis	2
constants	5
control_flow	2
delay	12
distortion	5
dynamics	3
envelope	13
filter	27
fm	17
math	3
midi	10
mixer	3
modulation	11
noise	3
oscillator	22
output	1
physical_modeling	31
reverb	4
routing	4
soundfont	3
spectral	9
tables	4
utility	5

Opcode Index

analysis

Opcode	Inputs	Outputs	Short Description
follow2	3 (a, k, k)	1 (a)	Envelope follower with independent attack and release controls.
rms	3 (a, i, i)	1 (k)	RMS envelope follower for an audio signal.

constants

Opcode	Inputs	Outputs	Short Description
const_a	-	1 (a)	Audio-rate constant value.
const_i	-	1 (i)	Init-rate constant value.
const_k	-	1 (k)	Control-rate constant value.
const_s	-	1 (S)	String constant value.
perf_controller	-	1 (i)	Virtual per-instance performance setting with a labeled rack knob.

control_flow — editor constructs

These nodes generate structured Csound conditionals; they are not ordinary opcode calls. If, Switch and Drumset documents note-start selection and the managed Case Result boundary.

Construct	Inputs	Outputs	Short Description
If	2 (i, i)	1 or 2 (a)	Select True/False by an init-rate comparison for each voice.
Switch	1 (i)	1 or 2 (a)	Select a named numeric case or mandatory Default for each voice.

delay

Opcode	Inputs	Outputs	Short Description
comb	4 (a, k, i, i)	1 (a)	Comb filter / feedback delay.
delay	3 (a, i, i)	1 (a)	Simple non-interpolating audio delay line.
delayk	3 (k, i, i)	1 (k)	Control-rate delay line.
delayr	2 (i, i)	1 (a)	Read tap from a classic delay-line memory buffer.
delayw	1 (a)	-	Write into the delay-line memory buffer.
deltap	1 (k)	1 (a)	Read a delay tap using linear interpolation.
deltap3	1 (k)	1 (a)	Read a delay tap using cubic interpolation.
flanger	4 (a, k, k, i)	1 (a)	Flanger effect with delay modulation and feedback.
stereo_capture	13 (a, a, k, k, k, k, i, i, i, i, i, i, i)	2 (a, a)	Orchestron stereo stutter, freeze and granular processor.
vcomb	6 (a, k, k, i, i, i)	1 (a)	Variable-time comb reverb with controllable loop time.

Opcode	Inputs	Outputs	Short Description
vdelay3	4 (a, k, i, i)	1 (a)	Variable delay line with cubic interpolation.
vdelayxs	6 (a, a, a, i, i, i)	2 (a, a)	Stereo variable delay with high-quality sinc interpolation.

distortion

Opcode	Inputs	Outputs	Short Description
clip	4 (a, i, i, i)	1 (a)	Signal clipper with selectable transfer curves.
distort1	6 (a, k, k, k, k, i)	1 (a)	Waveshaping distortion with configurable transfer curve.
fold	2 (a, k)	1 (a)	Artificial foldover effect for audio signals.
powershape	3 (a, k, i)	1 (a)	Power-law waveshaper for controllable nonlinear distortion.
mirror	3 (a, k, k)	1 (a)	Reflect an audio signal between lower and upper thresholds.

dynamics

Opcode	Inputs	Outputs	Short Description
compress2	9 (a, a, k, k, k, k, k, k, i)	1 (a)	Compressor, limiter, expander and gate with an audio sidechain and dBFS thresholds.
dam	6 (a, k, i, i, i, i)	1 (a)	Dynamic amplitude processor (downward compressor/noise suppressor).
limit	3 (a, k, k)	1 (a)	Hard clamp limiter.

envelope

Opcode	Inputs	Outputs	Short Description
adsr	4 (i, i, i, i)	1 (k)	Control-rate ADSR envelope.
envlpxr	8 (k, i, i, i, i, i, i, i)	1 (k)	Control-rate envelope with optional release scaling controls.
expon	3 (i, i, i)	1 (k)	Control-rate exponential segment between two points.
expseg	9 (i, i, i, i, i, i, i, i, i)	1 (k)	Control-rate exponential breakpoint envelope generator.
expsega	9 (i, i, i, i, i, i, i, i, i)	1 (a)	Audio-rate exponential breakpoint envelope generator.
expsegr	5 (i, i, i, i, i)	1 (k)	Control-rate exponential envelope with release segment.
linenr	4 (k, i, i, i)	1 (k)	Control-rate envelope with note-release stage.
linseg	7 (i, i, i, i, i, i, i)	1 (k)	Control-rate linear breakpoint envelope generator.
linsegr	5 (i, i, i, i, i)	1 (k)	Control-rate linear breakpoint envelope with release segment.

Opcode	Inputs	Outputs	Short Description
madsr	6 (i, i, i, i, i, i)	1 (k)	MIDI release-sensitive ADSR envelope.
mxadsr	6 (i, i, i, i, i, i)	1 (k)	Extended MIDI release-sensitive ADSR envelope.
transegr	7 (i, i, i, i, i, i, i)	1 (k)	Control-rate transition envelope with release-triggered final segment.
linseg_a	7 (i, i, i, i, i, i, i)	1 (a)	Audio-rate linear breakpoint envelope; holds its final value.

filter

Opcode	Inputs	Outputs	Short Description
atone	3 (a, k, i)	1 (a)	Audio-rate high-pass filter complementary to tone.
atonek	3 (k, k, i)	1 (k)	Control-rate high-pass filter complementary to tonek.
atonex	4 (a, k, i, i)	1 (a)	Cascaded atone high-pass filters with a sharper cutoff.
butterbp	4 (a, k, k, i)	1 (a)	Second-order Butterworth band-pass filter.
butterbr	4 (a, k, k, i)	1 (a)	Second-order Butterworth band-reject filter.
butterhp	3 (a, k, i)	1 (a)	Second-order Butterworth high-pass filter.
butterlp	3 (a, k, i)	1 (a)	Second-order Butterworth low-pass filter.
diode_ladder	6 (a, k, k, i, i, i)	1 (a)	Diode ladder low-pass filter model.
exciter	5 (a, k, k, k, k)	1 (a)	Harmonic exciter that adds controlled upper partials.
fofilter	5 (a, k, k, k, i)	1 (a)	Formant filter.
mode	4 (a, k, k, i)	1 (a)	Mass-spring-damper resonator for acoustic body modes.
moogladder	3 (a, k, k)	1 (a)	Moog ladder low-pass filter.
moogladder2	3 (a, k, k)	1 (a)	Nonlinear Moog-style ladder filter with audio-rate modulation support.
moogvcf	5 (a, k, k, i, i)	1 (a)	Moog ladder voltage-controlled filter emulation.
pareq	6 (a, k, k, k, i, i)	1 (a)	Parametric equalizer with peaking, low-shelf and high-shelf modes.
resonx	6 (a, k, k, i, i, i)	1 (a)	Multi-layer resonant band-pass filter.
rezzy	5 (a, k, k, i, i)	1 (a)	Resonant low-pass or high-pass filter.
skf	5 (a, k, k, i, i)	1 (a)	Sallen-Key low-pass or high-pass filter.

Opcode	Inputs	Outputs	Short Description
statevar	5 (a, k, k, i, i)	4 (a, a, a, a)	State-variable filter with simultaneous HP/LP/BP/BR outputs.
tbvcf	6 (a, k, k, k, k, i)	1 (a)	TB-303 style voltage-controlled filter model.
vclpf	4 (a, k, k, i)	1 (a)	Virtual-analog low-pass filter.
zdf_1pole	4 (a, k, k, i)	1 (a)	One-pole zero-delay feedback filter with selectable low-pass, high-pass, or allpass output.
zdf_1pole_mode	3 (a, k, i)	2 (a, a)	One-pole zero-delay feedback filter with simultaneous low-pass and high-pass outputs.
zdf_2pole	5 (a, k, k, k, i)	1 (a)	Two-pole zero-delay feedback filter with seven selectable response modes.
zdf_2pole_mode	4 (a, k, k, i)	3 (a, a, a)	Two-pole zero-delay feedback filter with simultaneous low-pass, band-pass, and high-pass outputs.
zdf_ladder	4 (a, k, k, i)	1 (a)	Four-pole zero-delay feedback Moog ladder low-pass filter.
zfilter2	6 (a, k, k, i, i, i)	1 (a)	Custom IIR filter with control-rate pole damping and frequency warping.

fm

Opcode	Inputs	Outputs	Short Description
STKBeeThree	12 (i, i, k, k, k, k, k, k, k, k, k, k)	1 (a)	STK four-operator FM organ inspired by Hammond organs.
STKFMVoices	12 (i, i, k, k, k, k, k, k, k, k, k, k)	1 (a)	STK FM singing-voice instrument with vowel and spectral-tilt controls.
STKHevyMetl	12 (i, i, k, k, k, k, k, k, k, k, k, k)	1 (a)	STK four-operator FM instrument for metallic timbres.
STKPercFlut	12 (i, i, k, k, k, k, k, k, k, k, k, k)	1 (a)	STK four-operator FM percussive flute instrument.
STKRhodey	12 (i, i, k, k, k, k, k, k, k, k, k, k)	1 (a)	STK four-operator FM electric piano inspired by Rhodes pianos.
STKTubeBell	12 (i, i, k, k, k, k, k, k, k, k, k, k)	1 (a)	STK four-operator FM tubular-bell instrument.
STKWurley	12 (i, i, k, k, k, k, k, k, k, k, k, k)	1 (a)	STK four-operator FM electric piano inspired by Wurlitzer pianos.
crossfmi	9 (k, k, k, k, k, i, i, i)	2 (a, a)	Interpolating crossed frequency-modulation oscillator pair.
crossfmpmi	9 (k, k, k, k, k, i, i, i)	2 (a, a)	Interpolating crossed frequency/phase-modulation oscillator pair.
crosspmi	9 (k, k, k, k, k, i, i, i)	2 (a, a)	Interpolating crossed phase-modulation oscillator pair.

Opcode	Inputs	Outputs	Short Description
fmb3	10 (k, k, k, k, k, i, i, i, i)	1 (a)	B3 organ FM model.
fmbell	12 (k, k, k, k, k, k, i, i, i, i, i, i)	1 (a)	Bell FM model.
fmmetal	11 (k, k, k, k, k, k, i, i, i, i, i)	1 (a)	Metallic FM model.
fmpercfl	11 (k, k, k, k, k, k, i, i, i, i, i)	1 (a)	Percussive flute FM model.
fmrhode	11 (k, k, k, k, k, k, i, i, i, i, i)	1 (a)	Rhodes electric piano FM model.
fmvoice	11 (k, k, k, k, k, k, i, i, i, i, i)	1 (a)	Voice-like FM model.
fmwurlie	11 (k, k, k, k, k, k, i, i, i, i, i)	1 (a)	Wurlitzer electric piano FM model.

math

Opcode	Inputs	Outputs	Short Description
a_mul	2 (a, a)	1 (a)	Multiply two audio signals.
k_mul	2 (k, k)	1 (k)	Multiply two control signals.
tanh	1 (a)	1 (a)	Hyperbolic tangent function for audio waveshaping and limiting.

midi

Opcode	Inputs	Outputs	Short Description
ctrl7	5 (i, i, k, k, i)	1 (k)	Explicit-channel MIDI CC reader, including continuous instruments.
initc7	3 (i, i, i)	-	Initialize a shared MIDI CC once at instrument startup with a normalized value.
ampmidi	2 (i, i)	1 (i)	Generate init-rate amplitude from MIDI note velocity.
ampmidicurve	3 (k, k, k)	1 (k)	Map MIDI velocity to gain using dynamic range and curve exponent.
ampmidid	2 (k, i)	1 (k)	Map MIDI velocity to amplitude using a decibel range.
cpsmidi	-	1 (i)	Read active MIDI note pitch as cycles-per-second.
midi_legato	-	2 (k, k)	Opt-in monophonic phrases, immediate pitch and smoothed velocity.
midi_note	1 (i)	2 (k, k)	Extract MIDI note frequency and velocity amplitude.
midictrl	3 (i, i, i)	1 (k)	Read a MIDI controller value with optional scaling.
notnum	-	1 (i)	Read the MIDI note number for the current note event.

mixer

Opcode	Inputs	Outputs	Short Description
mix2	2 (a, a)	1 (a)	Mix two audio signals.

Opcode	Inputs	Outputs	Short Description
ntrpol	3 (a, a, k)	1 (a)	Linear crossfade between two audio signals.
pan2	3 (a, k, i)	2 (a, a)	Stereo panner.

modulation

Opcode	Inputs	Outputs	Short Description
jitter	3 (k, k, k)	1 (k)	Control-rate segmented random line generator.
lfo	3 (k, k, i)	1 (k)	Low-frequency oscillator for control-rate modulation.
portk	3 (k, k, i)	1 (k)	Control-rate portamento/slew limiter with optional init value.
random	2 (k, k)	1 (k)	Control-rate random value generator between a minimum and maximum.
randomh	5 (k, k, k, i, i)	1 (k)	Control-rate random sample-and-hold generator with update frequency control.
randomi	5 (k, k, k, i, i)	1 (k)	Control-rate interpolated random generator with update frequency control.
release	-	1 (k)	Return 1 during note release phase and 0 otherwise.
samphold	4 (k, k, i, i)	1 (k)	Control-rate sample-and-hold processor.
swept_gate	6 (k, i, i, i, i, k)	2 (a, a)	Orchestron swept triangle gate with LFO and sample-and-hold random panning.
vibr	4 (k, k, i, i)	1 (k)	Simple vibrato control oscillator with table lookup.
vibrato	9 (k, k, k, k, k, k, k, k, i)	1 (k)	Randomized vibrato generator.

noise

Opcode	Inputs	Outputs	Short Description
noise	4 (k, k, i, i)	1 (a)	Variable-color random audio noise.
pinker	-	1 (a)	Pink noise generator.
pinkish	2 (a, i)	1 (a)	Pinkening filter for shaping white noise toward a pink spectrum.

oscillator

Opcode	Inputs	Outputs	Short Description
STKMoog	12 (i, i, k, k, k, k, k, k, k, k, k, k)	1 (a)	STK sampled-wave synthesizer with swept formant filters.
STKSimple	10 (i, i, k, k, k, k, k, k, k, k)	1 (a)	STK wavetable and noise synthesizer with filters and an ADSR envelope.

Opcode	Inputs	Outputs	Short Description
adsyn	3 (k, k, k)	1 (a)	Resynthesize sinusoidal partials from an uploaded hetro analysis file.
fof	15 (k, k, k, k, k, k, k, k, i, i, i, i, i, i, i)	1 (a)	Formant/granular source using sinusoid bursts.
fof2	15 (k, k, k, k, k, k, k, k, i, i, i, i, k, k, i)	1 (a)	FOF source with per-grain phase indexing and glissando.
foscili	7 (k, k, k, k, k, i, i)	1 (a)	Audio-rate FM oscillator with harmonic ratios.
gbuzz	7 (k, k, k, k, k, i, i)	1 (a)	Generalized buzz oscillator with controllable harmonics.
grain	10 (k, k, k, k, k, k, i, i, i, i)	1 (a)	Classic granular synthesis oscillator with table-based grains.
grain2	9 (k, k, k, i, k, i, i, i, i)	1 (a)	Easy-to-use granular synthesis texture generator.
grain3	13 (k, k, k, k, k, k, i, k, i, k, k, i, i)	1 (a)	Granular oscillator with independent frequency and phase modulation controls.
granule	22 (k, i, i, i, i, i, i, i, i, i, k, i, k, i, i, i, i, i, i, i, i, i)	1 (a)	Multi-voice granular processor with independent gap and grain-size controls.
hsboscil	8 (k, k, k, i, i, i, i, i)	1 (a)	Oscillator with independent cyclic tonality and octave brightness controls.
moog	9 (k, k, k, k, k, k, i, i, i)	1 (a)	Mini-Moog style synthesizer model source.
oscil3	4 (k, k, i, i)	1 (a)	Cubic-interpolating oscillator with low distortion.
oscili	3 (k, k, i)	1 (a)	Classic interpolating oscillator.
poscil3	4 (k, k, i, i)	1 (a)	High-precision cubic interpolating oscillator.
syncphasor	3 (k, a, i)	2 (a, a)	Audio-rate normalized phase generator with sync trigger I/O.
vco	5 (k, k, i, k, i)	1 (a)	Band-limited voltage-controlled oscillator.
vco2	6 (k, k, i, k, k, i)	1 (a)	Improved anti-aliased analog-style oscillator.
voice	8 (k, k, k, k, k, k, i, i)	1 (a)	Formant-based vocal source synthesizer.
vosim	8 (k, k, k, k, k, k, i, i)	1 (a)	Vocal formant synthesis using overlapping sinusoidal bursts.
phasor	2 (k, i)	1 (a)	Audio-rate normalized phase ramp; accepts audio-rate frequency.

output

Opcode	Inputs	Outputs	Short Description
outs	2 (a, a)	-	Stereo output sink.

physical_modeling

Opcode	Inputs	Outputs	Short Description
STKBandedWG	16 (i, i, k, k, k, k, k, k, k, k, k, k, k, k, k, k)	1 (a)	STK banded waveguide model of bowed bars, glasses, and bowls.
STKBlowBotl	10 (i, i, k, k, k, k, k, k, k, k)	1 (a)	STK blown-bottle model with a Helmholtz resonator and jet excitation.
STKBlowHole	12 (i, i, k, k, k, k, k, k, k, k, k, k)	1 (a)	STK clarinet model with register and tone holes.
STKBowed	12 (i, i, k, k, k, k, k, k, k, k, k, k)	1 (a)	STK bowed-string waveguide instrument.
STKBrass	12 (i, i, k, k, k, k, k, k, k, k, k, k)	1 (a)	STK brass waveguide instrument with lip and slide controls.
STKClarinet	12 (i, i, k, k, k, k, k, k, k, k, k, k)	1 (a)	STK single-reed clarinet waveguide instrument.
STKDrummer	2 (i, i)	1 (a)	STK drum synthesizer using sampled raw waves and one-pole filters.
STKFlute	12 (i, i, k, k, k, k, k, k, k, k, k, k)	1 (a)	STK flute waveguide instrument with jet and breath controls.
STKMandolin	12 (i, i, k, k, k, k, k, k, k, k, k, k)	1 (a)	STK mandolin model with paired strings and sampled body responses.
STKModalBar	16 (i, i, k, k, k, k, k, k, k, k, k, k, k, k, k, k)	1 (a)	STK modal percussion instrument with bar and bell presets.
STKPlucked	2 (i, i)	1 (a)	STK Karplus-Strong plucked-string instrument.
STKResonate	12 (i, i, k, k, k, k, k, k, k, k, k, k)	1 (a)	STK noise-driven resonator with pole and zero controls.
STKSaxofony	16 (i, i, k, k, k, k, k, k, k, k, k, k, k, k, k, k)	1 (a)	STK saxophone-like reed waveguide instrument.
STKShakers	14 (i, i, k, k, k, k, k, k, k, k, k, k, k, k)	1 (a)	STK stochastic percussion and environmental-sound models.
STKSitar	2 (i, i)	1 (a)	STK plucked-string sitar model.
STKStifKarp	8 (i, i, k, k, k, k, k, k)	1 (a)	STK stiff-string model with pickup, sustain, and stretch controls.
STKVoicForm	12 (i, i, k, k, k, k, k, k, k, k, k, k)	1 (a)	STK four-formant voice synthesizer with voiced and unvoiced excitation.
STKWhistle	12 (i, i, k, k, k, k, k, k, k, k, k, k)	1 (a)	STK whistle model with a pea and fipple modulation.
dripwater	8 (k, i, i, i, i, i, i, i)	1 (a)	Stochastic dripping-water physical model source.
marimba	11 (k, k, i, i, i, k, k, i, i, i, i)	1 (a)	Physical model of a marimba bar and resonator.
pluck	6 (k, k, i, i, i, i)	1 (a)	Karplus-Strong plucked-string model.

Opcode	Inputs	Outputs	Short Description
prepiano	17 (i, i, i, i, i, i, k, k, i, i, i, i, i, i, i, i, i)	2 (a, a)	Stereo prepared-piano physical model with detuned strings, hammer controls, rattles and rubbers.
sekere	5 (i, i, i, i, i)	1 (a)	PhISEM sekere shaker model.
sleighbells	8 (k, i, i, i, i, i, i, i)	1 (a)	PhISEM sleigh bells model.
wgbow	8 (k, k, k, k, k, k, i, i)	1 (a)	Waveguide bowed-string model.
wgbowedbar	9 (k, k, k, k, k, i, i, i, i)	1 (a)	Waveguide bowed-bar model.
wgclar	10 (k, k, k, i, i, k, k, k, i, i)	1 (a)	Waveguide clarinet model.
wgflute	10 (k, k, k, i, i, k, k, k, i, i)	1 (a)	Waveguide flute model.
waveguide_string	10 (a, i, i, i, i, i, i, i, i, k)	2 (a, a)	Virtual dual-polarization string with local sample feedback and compensated dispersion.
wgpluck2	5 (i, k, i, k, k)	1 (a)	Waveguide plucked string model with pick and reflection controls.
wguide2	4 (a, k, k, k)	1 (a)	Two-point waveguide resonator.

reverb

Opcode	Inputs	Outputs	Short Description
freeverb	6 (a, a, k, k, i, i)	2 (a, a)	Stereo Freeverb processor with room size and high-frequency damping controls.
platerev	9 (i, i, k, i, i, i, i, a, a)	2 (a, a)	Physical-model plate reverb with stereo output taps.
reverb2	4 (a, k, k, i)	1 (a)	Schroeder reverb processor.
reverbsc	7 (a, a, k, k, i, i, i)	2 (a, a)	8-delay-line stereo FDN reverb based on Sean Costello's design.

routing

Opcode	Inputs	Outputs	Short Description
inleta	1 (S)	1 (a)	Receive an audio-rate signal from a named instrument inlet port.
inletk	1 (S)	1 (k)	Receive a control-rate signal from a named instrument inlet port.
outleta	2 (S, a)	-	Send an audio-rate signal to a named instrument outlet port.
outletk	2 (S, k)	-	Send a control-rate signal to a named instrument outlet port.

soundfont

Opcode	Inputs	Outputs	Short Description
sfinstr3	8 (i, i, k, k, i, i, i, i)	2 (a, a)	Play a SoundFont instrument as stereo audio with cubic interpolation.
sfload	-	1 (i)	Load a SoundFont2 file and return a file handle.

Opcode	Inputs	Outputs	Short Description
sfplay3	8 (i, i, k, k, i, i, i, i)	2 (a, a)	Play a SoundFont preset as stereo audio with cubic interpolation.

spectral

Opcode	Inputs	Outputs	Short Description
pvsanal	7 (a, i, i, i, i, i, i)	1 (f)	Phase-vocoder analysis from an audio input to an fsig stream.
pvscale	5 (f, k, k, k, k)	1 (f)	Transpose spectral audio by multiplying its frequencies by a ratio.
pvshift	6 (f, k, k, k, k, k)	1 (f)	Shift fsig partial frequencies by a fixed amount in Hz.
pvsmooth	3 (f, k, k)	1 (f)	Smooth fsig amplitude and frequency trajectories with lowpass filtering.
pvsmorph	4 (f, f, k, k)	1 (f)	Morph between two fsig streams by amplitude and frequency interpolation.
pvsosc	8 (k, k, k, i, i, i, i, i)	1 (f)	Generate oscillator spectra directly as an fsig stream.
pvsvoc	5 (f, f, k, k, k)	1 (f)	Cross-synthesize fsig amplitudes and excitation frequencies.
pvswarp	7 (f, k, k, k, k, k, k)	1 (f)	Warp and shift the spectral envelope of an fsig stream.
pvsynth	2 (f, i)	1 (a)	Resynthesize audio from an fsig stream via overlap-add.

tables

Opcode	Inputs	Outputs	Short Description
GEN	-	1 (i)	Routine-aware function table generator (meta-opcode) that renders ftgen/ftgenonce from a specialized editor, including GEN19 sinusoidal partials with phase and DC offset.
ftgen	12 (i, i, i, i, i, i, i, i, i, i, i, i)	1 (i)	Create a function table at init time using a GEN routine.
ftgenonce	12 (i, i, i, i, i, i, i, i, i, i, i, i)	1 (i)	Generate a function table once and reuse it across instances.
vco2init	6 (i, i, i, i, i, i)	1 (i)	Precompute band-limited wavetable sets for vco2 oscillators.

utility

Opcode	Inputs	Outputs	Short Description
downsamp	2 (a, i)	1 (k)	Convert an audio signal to control-rate by downsampling.

Opcode	Inputs	Outputs	Short Description
k_to_a	1 (k)	1 (a)	Interpolate control signal to audio-rate.
maxalloc	1 (i)	-	Limit simultaneous instrument allocations for the compiled instrument.
upsamp	1 (k)	1 (a)	Convert a control signal to audio-rate by repeating k-values.
xtratim	1 (i)	-	Extend current note duration by an additional init-time amount.

See Physical steel-string guitar for the editable excitation, string and body example.

Navigation: Up | Previous | Next

Performance

Navigation: Up | Previous | Next

Use **Perform** to turn saved instruments into a song or live setup. Start with at least one saved playable instrument, or follow Getting Started. Your first task is to add that instrument to the rack and play it from a track on the same MIDI channel.

The performance description field displays three lines with automatic text wrapping and preserves manual line breaks. Longer descriptions scroll vertically within the field. Hover over the description to see its full text in a tooltip.

New performances start in the multitrack arranger with one piano keyboard. Use **Add track** to create melodic, Score, drummer, controller or arpeggiator tracks, then add clips and patterns to build the song.

Chapter Contents

- Instrument Rack and Engine Transport
- Audio mixer and routing
- Using Performance Controls
- MIDI Setup and Inputs
- Piano Rolls
- MIDI Controllers
- Multitrack Arranger
- Patterns and Reusable Phrases
- Sequencer Timing
- Melodic Sequencers and Step Editing
- Drummer Sequencers
- Score Sequencers
- Arpeggiators
- Controller Sequencers
- Instrument Automation Clips
- Live Status and Troubleshooting
- Performance Import / Export

What You Can Do Here

- Assemble an instrument rack from saved patches and assign MIDI channels
- Mix audio with dB faders, pan/balance, mute/solo, meters, inserts and pre/post sends
- Route instruments and returns through Master or explicit direct paths
- Start/stop the instrument engine session
- Create multiple melodic sequencers, drummer sequencers, and controller sequencers
- Add arpeggiators that turn held notes into routed arpeggiated instrument output
- Edit reusable patterns and phrases, assemble song order in the arranger, and audition drafts in context
- Arrange multiple track timelines in the multitrack arranger with shared transport and loop-range playback
- Copy song sections across all or selected lanes, overwrite destinations, or insert time and shift later material together
- Perform live with piano rolls and manual MIDI controller knobs
- Save/load/clone/delete performances
- Import/export performance bundles (with optional patch definitions)
- Export offline Csound render ZIPs either with `.csd + .mid` playback or with inline Csound score playback

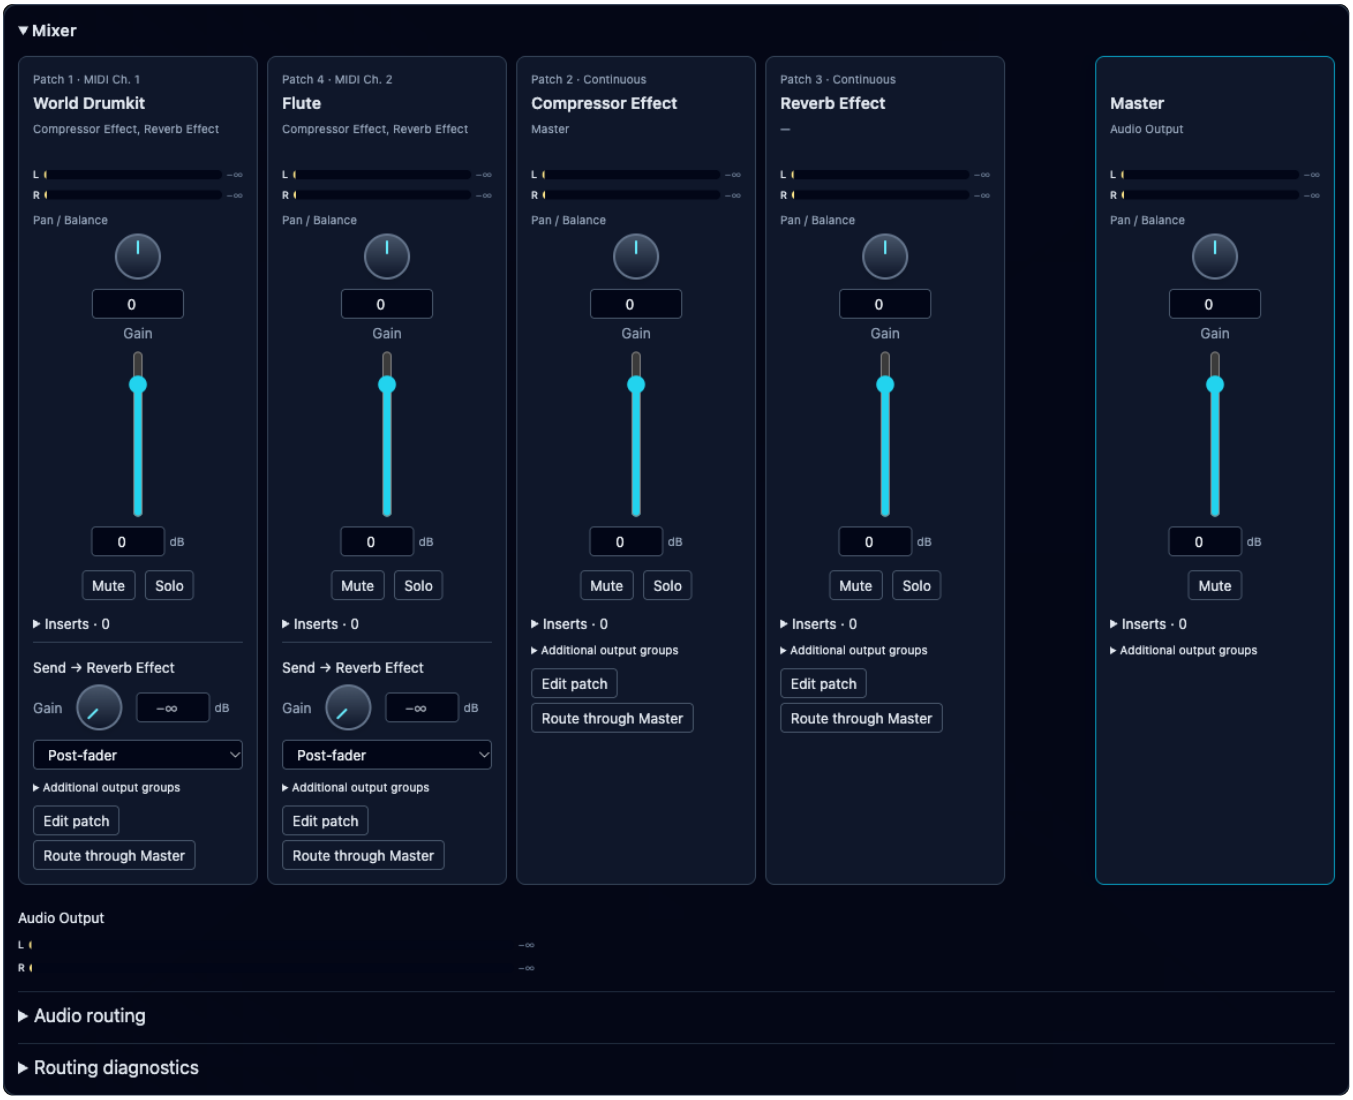

Live “Ratchet drums demo” mixer with World Drumkit and Flute before the continuous effect strips, followed by pinned Master. Playback is stopped. See the following chapters for rack and sequencer controls.

Practical Workflow

1. Save at least one patch in Instrument Design.
2. Add instruments to the performance rack and assign channels.
3. Open Mixer, check the destination of each strip, and connect effects/returns. Guided instruments get a neutral Master automatically; existing direct routes stay direct.
4. Use Routing diagnostics → Check routing, then Start Instruments. Adjust existing mixer controls during playback.
5. Choose **Add track**, add clips and patterns, then edit their musical content. Save each draft with the desired occurrence scope.
6. Save the performance to retain the rack, routes and current mix.
7. Export the performance bundle for backup/sharing, or use `Export CSD (MIDI) / Export CSD (SCORE)` when you need an offline Csound render package.

Tracks start compact. Expand a track to work with nested patterns, groups and supergroups. Double-click content to edit it and use **Play edit** to audition the draft. Lane Mute/Solo is temporary and independent of the instrument mixer; see lane controls.

Sequencer timing

Use Meter, Pattern length and Subdivision for melodic, drum and controller sequencers. A one-bar 4/4 pad with three steps per beat displays twelve steps in four labeled beat groups. Playback speed is under Advanced timing. See Sequencer timing

for the full guide.

Device Names

Melodic sequencers, drummer sequencers, controller sequencers, arpeggiators, piano rolls, and MIDI controllers all support custom names. Click the pen immediately to the right of a device name to edit it. The input receives focus with its current name selected.

- **Save** or Enter applies the name; **Cancel** or Escape discards the draft. Clicking elsewhere does not save.
- Use a nonempty name of at most 65 characters, without angle brackets. Names must be unique across device types in the performance, ignoring case and surrounding spaces.
- Renaming keeps the music, routing and playback state. Names are retained by Save Performance and native exports. Switching performances or removing the device discards an unfinished name edit.

Collapsible Panels

Click a panel title or focus it and press Enter or Space to collapse or expand its contents. Header actions, status and help remain available. Panels collapse independently, start expanded and remember your choices across view switches until browser reload. Adding a device opens its group. Collapsing preserves edits and does not stop playback. The shared tempo stays outside the sequencer groups.

The Instrument Rack, Mixer, keyboard and controller panels can be collapsed to make room for the song. The arranger is visible before adding a track. Click a track header to expand its content; double-click a pattern to open its editor. These layout choices are not saved into performances or exports.

Draft names and controller values, validation feedback, arpeggiator preset drafts, clip selections, open containers, arranger selections, range clipboard and undo history, zoom and scroll positions survive collapse/expand and switching away from Perform. Nested Mixer sections remember their open state when a parent collapses. New, Load and Import clear retained editor state after successfully replacing the workspace, even when reloading the same performance. Saving does not clear it. Removed devices and invalid selections are discarded; restored scroll positions are limited to the current content. This temporary editor state is not stored in performance files or browser storage and does not survive a page reload.

Collapsing closes temporary menus and pickers, cancels active drag previews and releases manually held piano notes. Already-applied edits remain, and existing input save/commit rules still apply. Sequenced notes continue playing. Collapsing reduces visual work and can help on a busy display.

The collapsed rack shows channel and patch name in one horizontally scrolling row. See Compact Rack for continuous instruments, long names, and keyboard scrolling.

Navigation: Up | Previous | Next

Instrument Rack and Engine Transport

Navigation: Up | Previous | Next

The Instrument Rack is the top section of the Perform page and controls the live session, instrument assignments, and performance metadata.

Adaptive Expanded Rack

Perform fills the available browser width with normal outer margins. Instrument cards automatically share rows or span the full rack according to their controller content. The layout reduces vertical space while keeping the saved instrument order and the divider between note-triggered and continuous instruments.

Optional **Controller Group** names in Instrument Design give related controls a thin labeled border. Groups remain expanded; cards and groups wrap without shrinking the controls. Narrow windows use a single column and a wrapping header. See Performance Controllers for grouping rules.

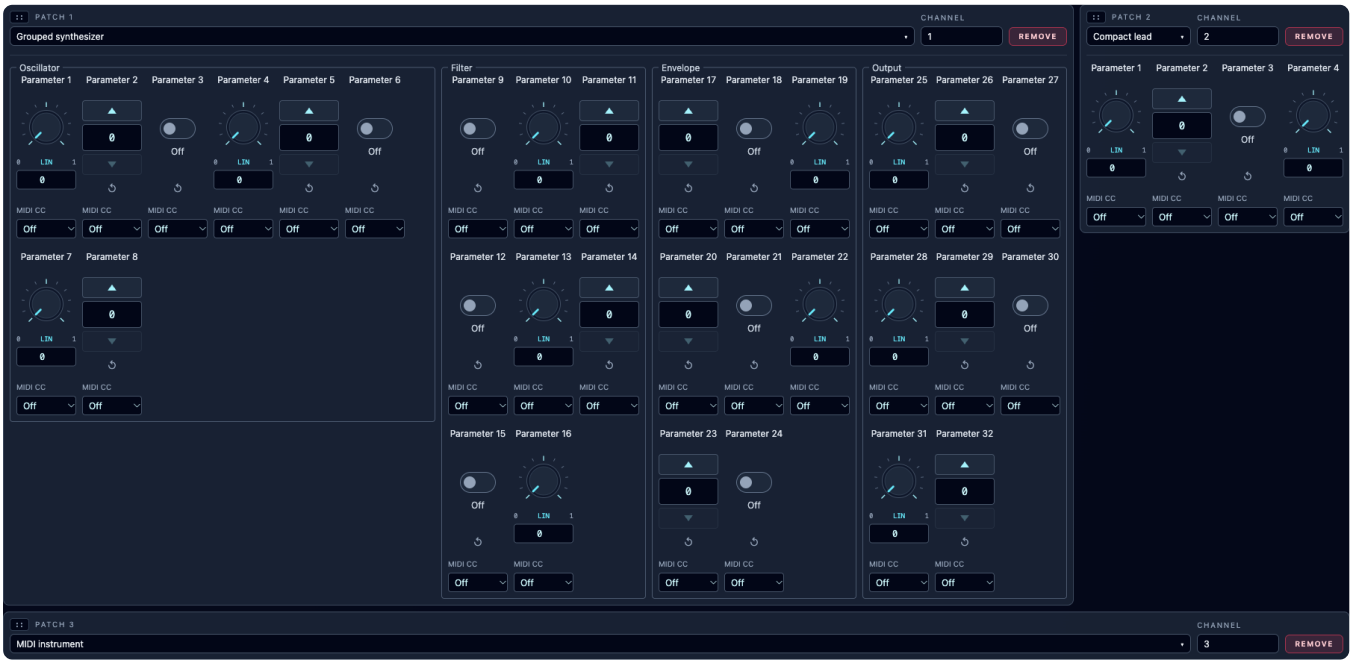

Compact Rack

When collapsed, the rack keeps Add Instrument and Start/Stop available. A single horizontally scrolling row shows MIDI channel and patch name for every instance, including duplicates; continuous instruments show Continuous instead of a channel. Long names have full-name tooltips. Note-triggered instruments appear first, followed by continuous instruments, with a thin vertical divider when both groups are present. The row follows the saved rack order; expand the rack to reorder instruments. Focus the row and use the arrow keys to scroll. Metadata, file actions and detailed controls return when expanded.

The header contains rack status, Add Instrument, Start/Stop, and help in both views. Collapse choices survive view switches until browser reload.

Performance Metadata and Library Actions

The rack includes fields and actions for the current performance:

- `Performance Name`
- `Description`
- `Load Performance` dropdown
- `Save Performance`

- `Clone`
- `Delete`
- `Export`
- `Export CSD (MIDI)`
- `Export CSD (SCORE)`
- `Import`

These actions operate on the performance configuration (instrument rack + audio routing/mixer + sequencers + controllers + piano rolls), not on individual patch definitions.

- `Export` writes an Orchestron `.orch.json` / `.orch.zip` performance bundle for backup, sharing, and re-import.
- `Export CSD (MIDI)` creates a Csound package with a MIDI file; `Export CSD (SCORE)` embeds the events in the Csound score. See Performance Import / Export for choosing a format, samples, rendering and release tails.

Instrument Assignments (Rack Slots)

Each rack entry selects a saved patch and a MIDI channel (1–16), or displays Continuous activation. Up to 64 user patch instances are supported, including processors. Generated mixer stages do not occupy rack slots or MIDI channels. Use distinct MIDI channels for note-triggered instances.

Instrument Order

Note-triggered instruments appear above continuous (always-on) instruments. A thin horizontal line separates the groups when both are present.

With the instrument engine stopped, drag a card's `::` handle onto another card in the same group. An insertion line shows the destination: the upper half inserts before the target, and the lower half inserts after it. You can also focus the handle and use `↑/↓` to move one position within its group. Instruments cannot move across the divider. Press **Escape** to cancel a drag; collapsing the rack or starting the engine also cancels it.

The compact rack and mixer strips follow this order. Master stays pinned at the right of the mixer, and insert processors remain inside their owner's insert chain. Moving a rack card changes display order without changing MIDI channels, routing, insert processing order, or controller and mixer values. Autosave, Save/Load, Clone, and native export/import preserve the instrument order.

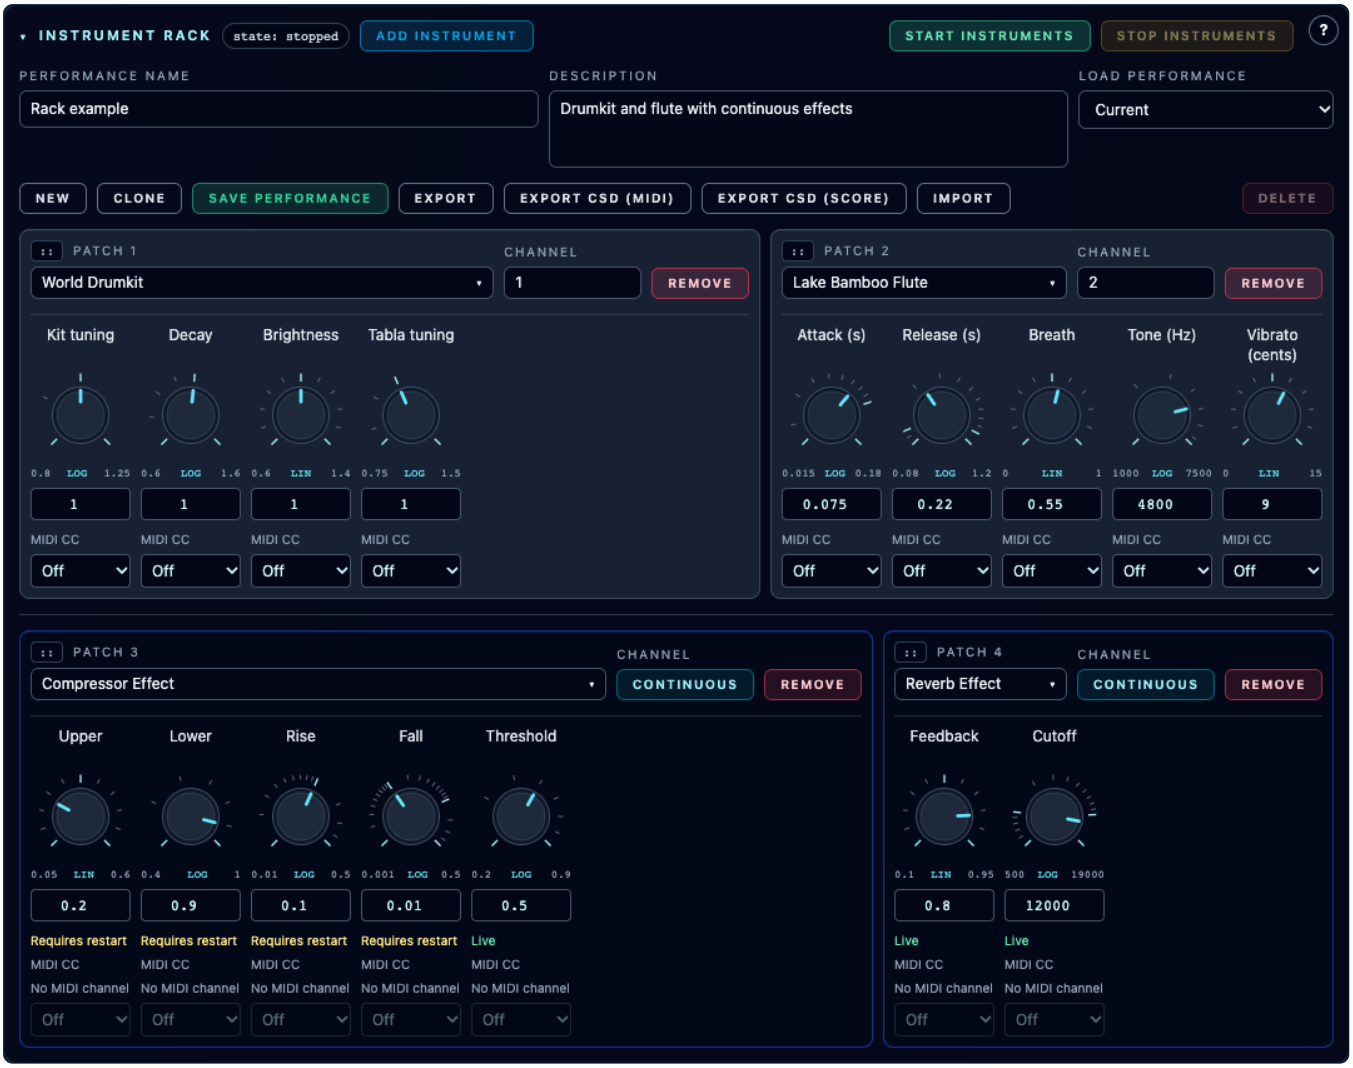

World Drumkit and Lake Bamboo Flute sit above the thin divider, with continuous compressor and reverb below. The :: handles reorder instruments within each group while the engine is stopped.

Instrument Controls

Patches containing `perf_controller` show five compact tuning knobs per standard rack-card row; extras wrap below. They remain editable while playing. New notes use updated values. Continuous controls marked **Live** update during playback. **Requires restart** identifies I-rate or mixed connections. Both CSD export modes use initialization values for continuous controls. Open the rack's **?** help and look for **Performance Controllers** for details. Save Performance stores explicit overrides.

The separate audio mixer below the rack provides audio faders, pan/balance knobs, mute/solo, pre/post sends, inserts and meters. The old numeric Level field has been removed. Mixer gain affects held notes and external MIDI audio without changing note velocity.

While instruments run, adding/removing assignments, changing patch/channel assignments, and changing routing topology are locked. Existing mixer controls stay live. Use **Stop to edit routing** before changing connections or insert order.

Add Instrument

- Use `Add Instrument` to create another rack slot.
- Open the slot's patch picker to browse instruments in collapsed type groups: Percussion, Melody, Bass, Effects / Noise, and Continuous (always-on). Expand a group to see its instruments alphabetically.
- Search names and descriptions across all types by entering at least four characters; results appear after a 500 ms typing pause. Clear the search to return to browsing. Template patches are excluded from the rack picker.
- This enables multi-instrument performances driven by different MIDI channels.

- Continuous effects run as explicit rack instances. Adding an insert creates a dedicated instance automatically; simply saving an effect in the library does not start it.
- The button is unavailable while the engine is running; stop instruments before changing rack assignments.

Rack Transport (Instrument Engine Control)

The rack transport controls the instrument runtime session itself.

Buttons:

- `Start Instruments`
- `Stop Instruments`

`Start Instruments / Stop Instruments`

These start/stop the underlying instrument engine session.

Global song transport lives in the multitrack arranger. **Play** starts or resumes, **Pause** holds the song cursor, and **Stop** returns to the loop start or song start. Rewind and forward move one master beat. Play starts instruments when needed; Stop stops instruments started by Play, while instruments that were already running remain running. **Stop Instruments** stops the engine. Use **Play edit** to audition a pattern or phrase.

Session State Badge

The rack shows current session state (localized label), for example:

- running
- stopped / idle

This is the state of the instrument engine session, not just the sequencer transport.

Error Banner

If engine start/stop or transport actions fail, the Perform page shows an error banner below the rack transport section.

Tips

- Save the performance after significant changes (rack assignments, sequencers, piano rolls, controller mappings).
- Use distinct MIDI channels for note-triggered rack instances.
- Use `sname` labels on source `outleta` nodes to name the available matrix channels. The label can be stored directly on the node or supplied by a direct `const_s` connection.

Navigation: Up | Previous | Next

Audio mixer and routing

Navigation: Up | Previous | Next

The collapsible mixer sits below the instrument rack. Select a strip to edit its outputs and sends. The pinned Master controls audio routed into it. A **Direct Audio Output** badge identifies paths that bypass Master; these still respond to their own strip's fader, balance, mute and solo.

Collapsing Mixer suspends its visual contents and meter subscriptions while audio controls and meter collection keep running. Nested sections suspend independently and remember their expansion state when their parent closes. Reopening restores routing selections, editor feedback and strip scroll position, and meters show the latest snapshot. See Collapsible Panels for state lifetime and gesture cleanup.

Use the top-right **?** for integrated help in English, German, French or Spanish, following the selected GUI language. The button remains available while Mixer is collapsed and opens help without changing the mix or expanding the panel.

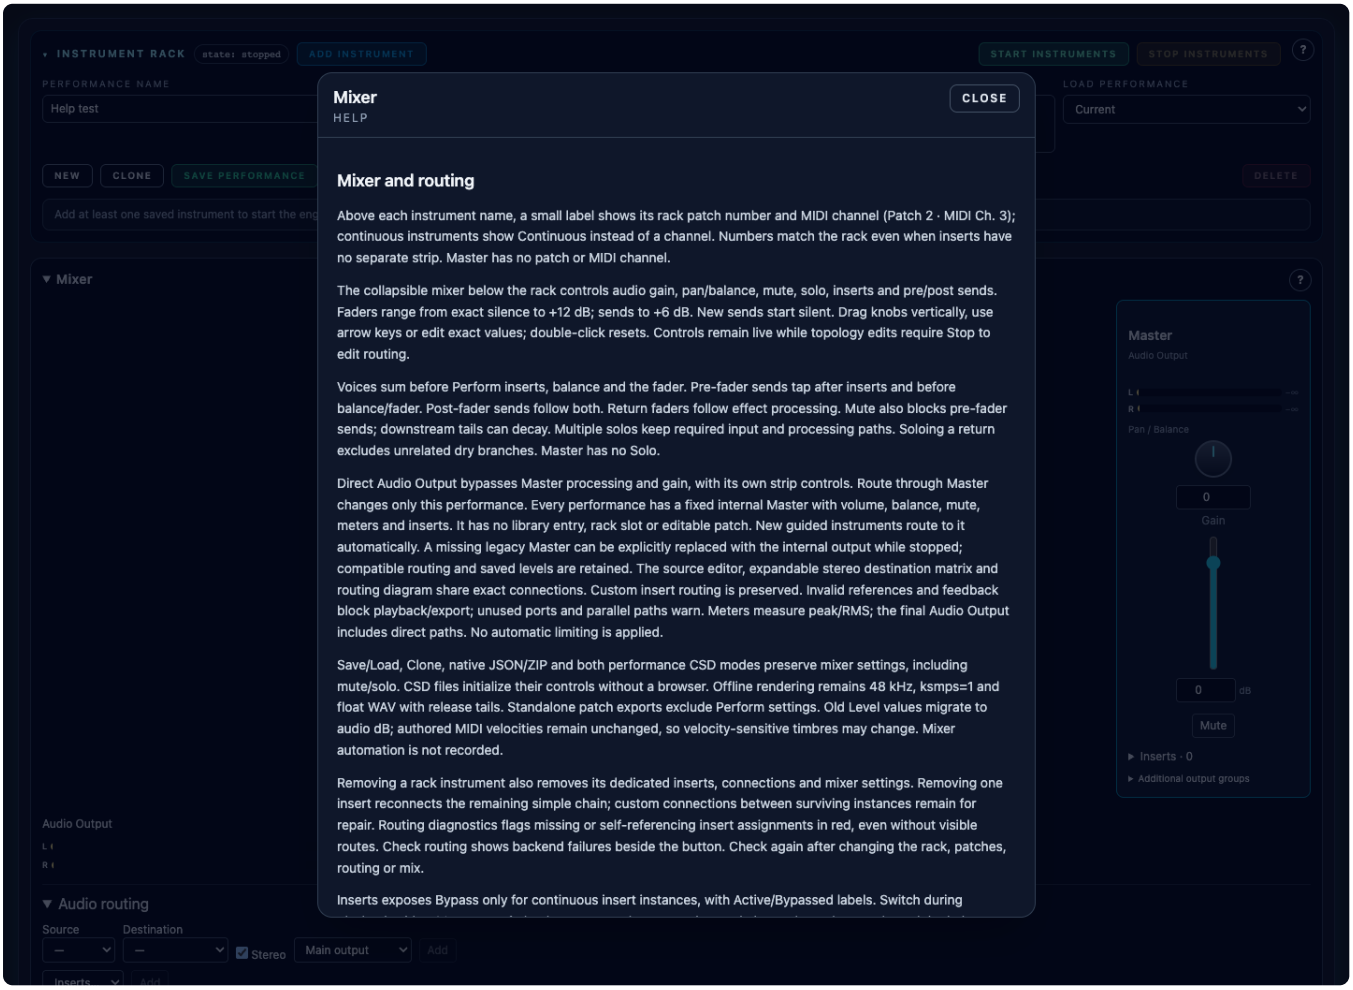

Build your first mix

1. In Instrument Design, choose **New → Playable instrument**, name the draft and **Save** it. Create and save an **Audio effect** when you need a processor. The built-in effect is a stereo pass-through starter; add your processing inside its graph.
2. In Perform, use **Add Instrument** and select the saved instrument. Set its MIDI channel. Guided instruments connect to the fixed internal Master automatically. For existing patches, inspect their destination labels and use **Route through Master** if needed.
3. Expand **Mixer** below the rack. Above each instrument name, a small label shows its rack patch number and MIDI channel, for example **Patch 2 · MIDI Ch. 3**. Continuous instruments and effects show **Patch 4 · Continuous**. Patch numbers match the rack even when insert processors have no separate mixer strip. Each strip header selects its

source for the routing editor. The Master strip stays pinned at the right; scroll horizontally when the rack has many strips. Strips follow the instrument rack's order, with note-triggered instruments before continuous returns. Stop the instruments and drag the rack's `::` handles to reorder within either group. This changes display order only; insert processors remain in their owner's insert chain.

4. Start with 0 dB gain and centered Pan / Balance. Build routes and insert chains while stopped, then open **Routing diagnostics → Check routing**.
5. Click **Start Instruments** and play via a piano roll, sequencer or MIDI input. Change gain, balance, mute/solo and existing send settings while listening. Watch strip meters and the final Audio Output meter.
6. Use **Save Performance** to retain this mix. Stop the engine before adding/removing routes, replacing assignments or reordering inserts.

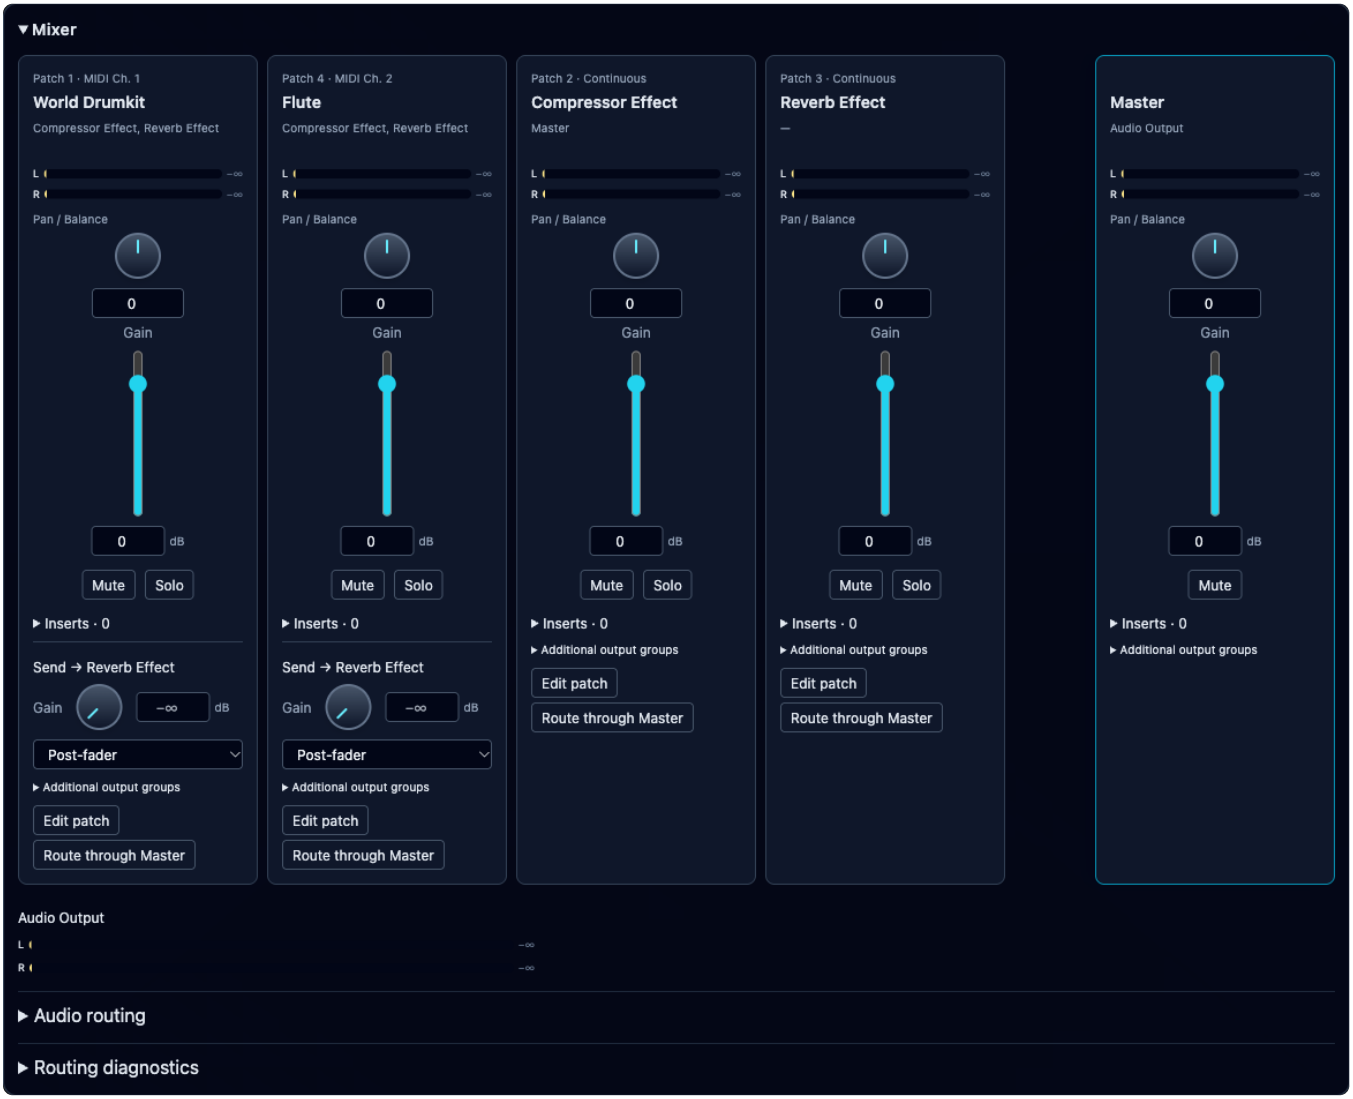

Live "Ratchet drums demo" mixer, stopped: the strips show their rack patch numbers and MIDI channels or Continuous above their names, with Master pinned at the right. World Drumkit and Flute have post-fader send controls for Reverb Effect; both sends are set to silence (–∞ dB). This performance has no insert processors.

Signal flow

Path	Processing order
Instrument	Summed voices → Perform inserts → pan/balance → fader → destination
Pre-fader send	After inserts → send amount → destination
Post-fader send	After pan/balance and fader → send amount → destination
Return	Summed inputs → effect patch → inserts → balance → return fader

Path	Processing order
Master	Incoming routes → internal stereo bus → inserts → balance → Master fader → Audio Output

Effects drawn inside an instrument graph run per voice. Perform inserts receive summed voices. Each inserted effect is a dedicated rack instance. Reordering an insert changes explicit connections; if custom connections no longer form a simple chain, **Custom routing** preserves those connections.

Additional output groups retain their own port names and mappings. They share the owning strip's gain and mute; inserts process the designated main output. Custom routes can express other processing arrangements.

Inserts versus shared returns

An **insert** processes one strip in series after its voices are summed. Select that strip, choose a saved continuous patch with audio inputs and outputs in the **Inserts…** selector below the routing controls, then click the adjacent **Add**. Open **Inserts** on the strip to inspect its chain. Click a processor name to edit its patch, ↑ to move it earlier, or **×** to remove that insert instance. Reusing one saved effect creates separate processing instances; an insert does not appear as a separate mixer strip.

Removing an instrument from the rack also removes its dedicated inserts, including any inserts on those processors, plus their connections and mixer settings. Removing a single insert through the rack or its **×** button reconnects the remaining simple chain. Custom connections between surviving instances are preserved for inspection and repair; unrelated returns, sends and Master controls remain unchanged.

A **return** is a continuous rack instance that can receive sends from several sources and has its own mixer strip. Add it with **Add Instrument**, connect its output to Master, then add sends from the source strips. For a parallel reverb/delay arrangement, design the return patch to output the desired wet signal; the starter pass-through would merely add another copy of the dry signal.

Compare inserts with Bypass

Open **Inserts** on the owning mixer strip. Each continuous insert has a **Bypass** button and an **Active** or **Bypassed** label. You can switch while instruments are running. Reordering and removing inserts still require stopping playback. The same controls appear for nested inserts and for inserts with custom routing.

Bypass passes that effect’s main input to its main output, using the declared channel order. A 20 ms crossfade suppresses abrupt transitions. The selected effect’s output, including its tails, fades out; downstream effects keep processing and can retain their own tails. The bypassed processor keeps receiving audio and running internally, so reactivation retains its current state. Live bypass does not reduce CPU use.

For example, bypassing Distortion in `Instrument → EQ → Distortion → Reverb` lets Reverb receive the EQ output directly. Other inserts, including inserts nested inside the bypassed processor, retain their own states. Strip gain/balance and mute/solo still apply. Additional effect outputs have no dry substitute and fade to silence. Bypass requires matching main input/output channel counts; an incompatible mapping disables the button. Bypass does not infer or compensate a patch’s internal processing latency.

The control is available only on continuous instances assigned as inserts, not on standalone generators, shared returns or note instruments. New inserts default to Active. Each instance’s state survives Save/Load, Clone, autosave and native JSON/ZIP bundles without changing its library patch. Match listening levels when comparing effects that change volume.

Both **CSD (MIDI)** and **CSD (SCORE)** omit bypassed effect synthesis, initialization, performance-controller channels and assets used only by those effects. Incoming audio connects directly to downstream consumers; necessary strip and route controls remain. Exported CSDs contain no live bypass crossfade stage. Native exports retain all effect definitions so they can be reactivated after import.

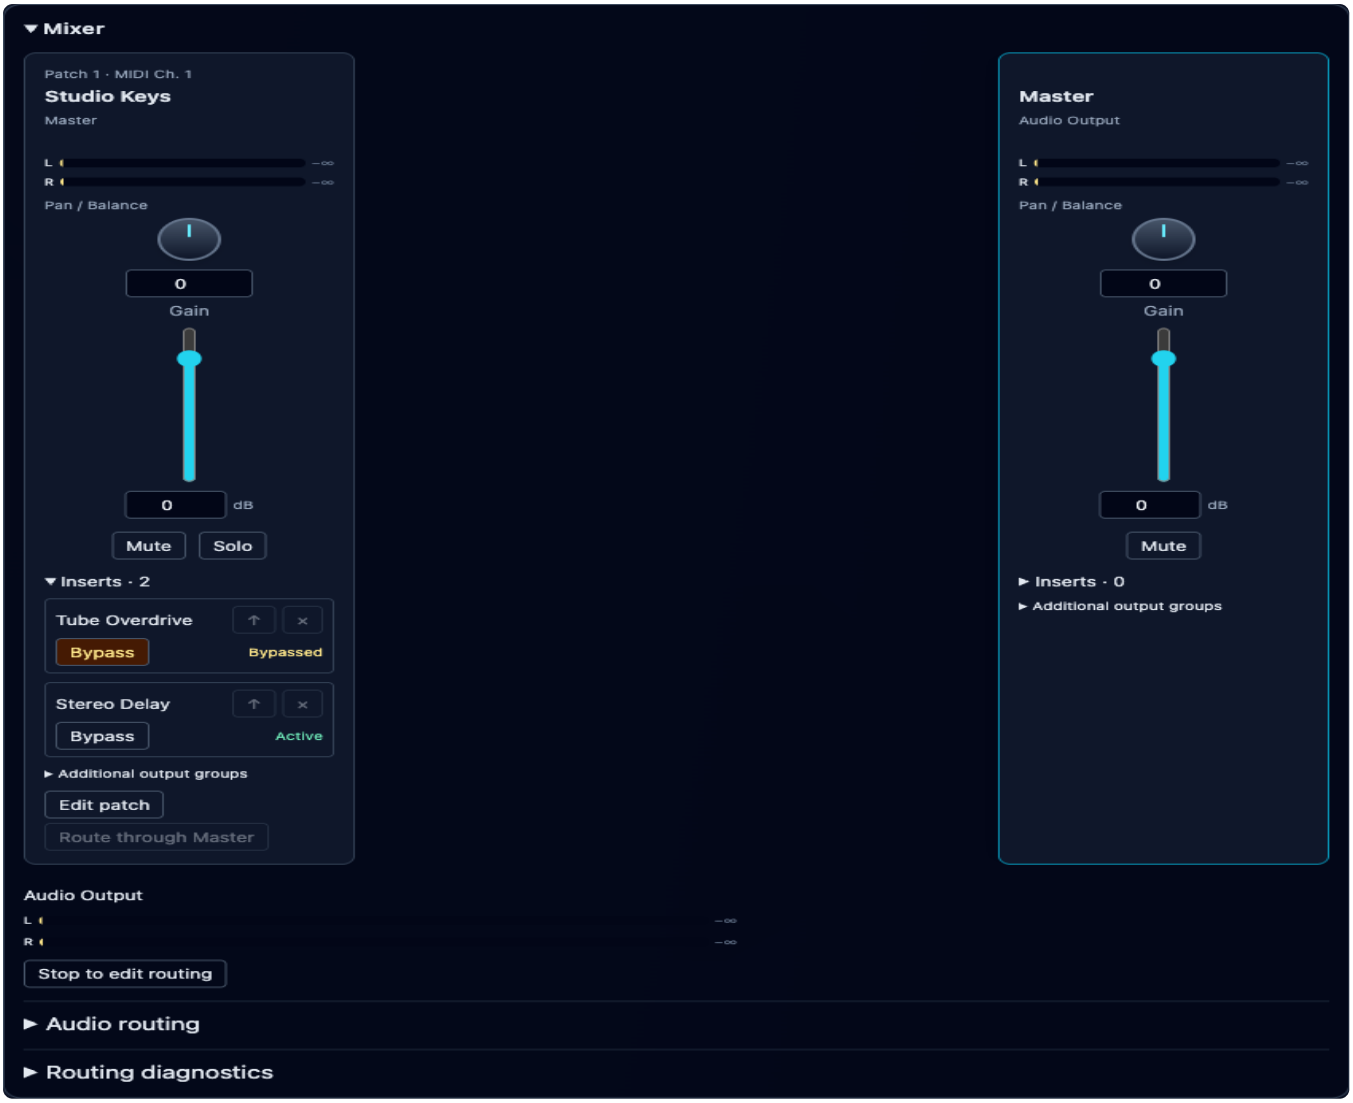

Each insert has its own Bypass control and explicit state. The remaining chain stays active.

Add a stereo send

1. Stop the engine. Add the saved effect as a rack instance and use its **Route through Master** button to connect its main output.
2. In **Audio routing**, choose the instrument under **Source** and the return under **Destination**.
3. Leave **Stereo** checked and verify **Exact channel mapping**, for example `left / right → left / right`. Choose **Send** in the Route selector and click the routing row's **Add**.
4. On the source strip, find **Send → [return name]**. New sends start at silence; enter a value such as -12 dB to hear the return.
5. Choose **Post-fader** when the send should follow the source's fader/balance, or **Pre-fader** for an independent send amount. The source's Mute still silences either kind of send.

Adding a send does not replace the main output route. Avoid adding duplicate connections: parallel paths sum and can increase output level. Up to 1,024 individual channel connections and 64 rack instances (including processors, excluding internal Master) are supported; a stereo pair uses two connections.

Controls

- Faders span –60 to +12 dB, with an exact-silence position below –60. Clearing the numeric value selects silence. Double-click resets to 0 dB.
- Sends span silence to +6 dB. New sends start silent and post-fader. Pre/post can change during playback.
- Rotary controls support vertical dragging, arrow keys, Page Up/Down, Home/End, exact numeric editing and double-

click reset. Shift-drag provides finer adjustment.

- Stereo balance preserves both channels at unity in the center and attenuates the opposite channel toward either side. Explicit mono-to-stereo routing uses equal-power pan. Custom mono mappings stay unchanged.
- Mute blocks all outgoing paths, including pre-fader sends. Processors keep running, allowing downstream tails to decay.
- Multiple solos are allowed. Solo preserves required upstream inputs and downstream processors. Soloing a return excludes its sources' unrelated dry branches. Mute takes precedence. Master has no Solo control.
- Gain, balance, send and mute/solo changes use finite 20 ms ramps. Playback starts at the saved values, without an initial fade.

Meters show stereo peak and RMS. CLIP indicates a peak at or above 0 dBFS; no automatic limiting, normalization or clipping is applied. The final **Audio Output** meter includes Master and direct paths. Connectivity diagnostics describe structure; meters measure actual audio activity.

Routing and repair

Choose a source, destination and explicit channel mapping before adding a main route, send or custom connection. Stereo groups expand into their exact member ports. Uncheck **Stereo** to select one exact source port and destination inlet at a time, including custom or mono mappings. Group display names are not inlet names; legacy label suggestions never rename ports.

Open **Destination matrix** to inspect incoming routes grouped by destination. Expand a route row to change its source port, destination instance or destination port; **Remove** deletes that grouped connection. This is also where broken references can be repaired instead of discarding the performance.

Open **Routing diagram** to follow the same connections as source → destination cards. Expand **Exact channel mapping** on a card to inspect port names and processing stages. Click source/destination names to select routing endpoints; use the strip's **Edit patch** action to open its instrument graph.

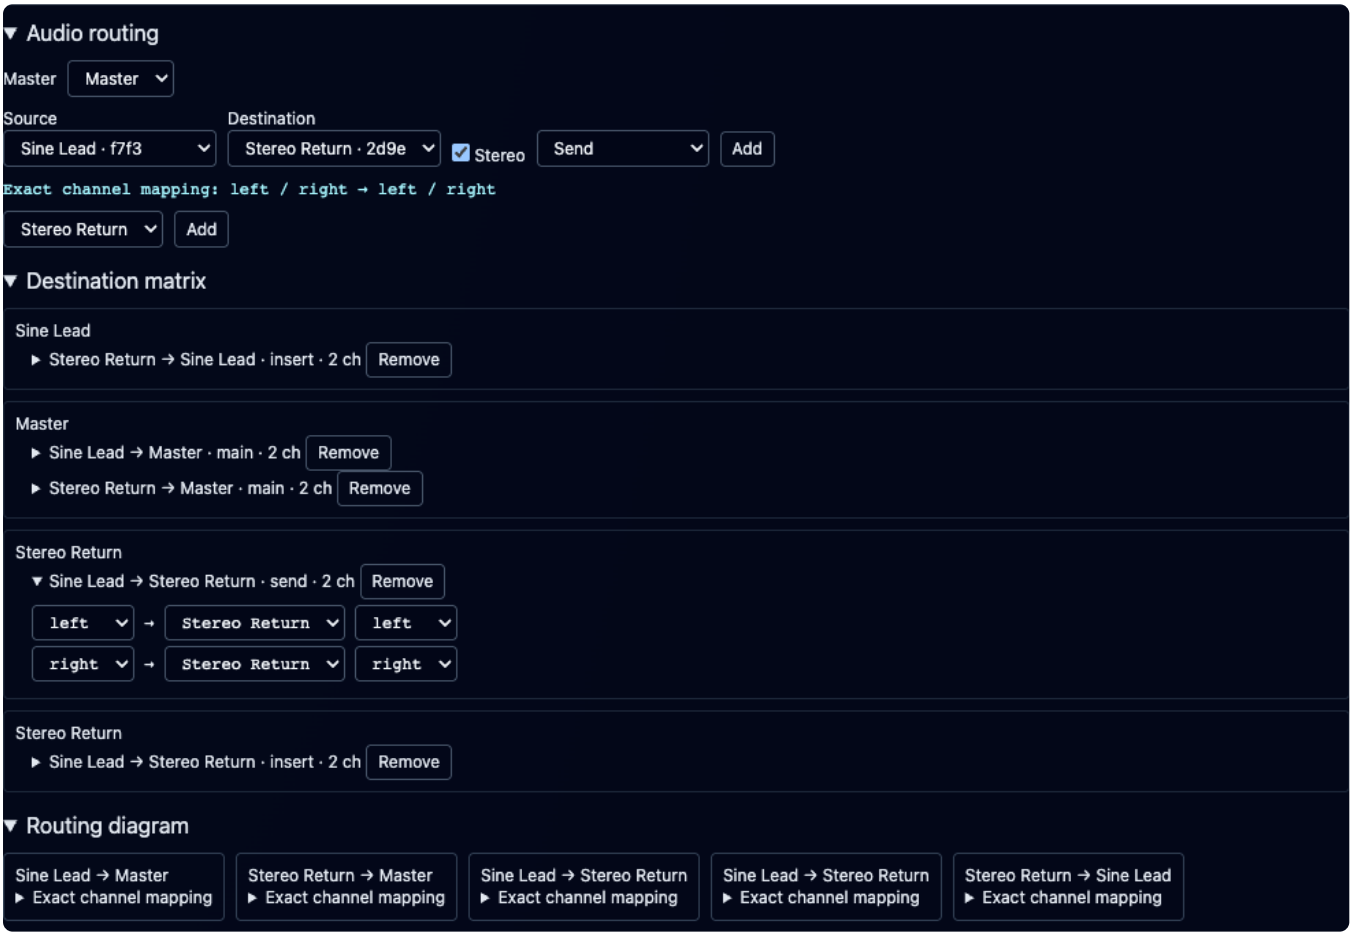

The send's left and right mappings expanded in Destination matrix. Matrix and diagram show the same routes, including the dedicated insert's input and return-to-strip connections.

Route through Master replaces the selected strip's main routes with stereo connections to the designated Master in this performance. Sends and custom routes are retained. Master is a fixed internal stereo output with volume, balance, mute, meters and effect inserts. It is always available, has no editable patch and occupies no library entry or rack slot. Its controls and inserts belong to the performance. **Audio Output** remains available as an editable starter for custom output processors.

New guided instruments connect to Master automatically. Loading an older performance preserves its authored direct/custom paths. Master gain or mute affects only paths routed into Master: a direct strip can still be audible. Check the **Direct Audio Output** badge and the final Audio Output meter when tracking down sound that remains after lowering Master.

Topology changes require stopping the engine. **Stop to edit routing** makes that action explicit. Existing mixer controls remain live, including with **Parallel playback** enabled: faders, balance, mute, solo, sends and insert bypass continue to affect the sound. Invalid references and genuine feedback cycles block playback and CSD export. Broken routes remain available for repair. Unused ports, incomplete stereo pairs and parallel direct/processed paths are warnings.

Troubleshoot routing and silence

Expand **Routing diagnostics** and click **Check routing** after editing connections. Diagnostics describe connectivity, while meters report measured sound. Clicking an instance-related diagnostic can open its patch for inspection. A **Muted or silent path** entry selects the corresponding strip.

Invalid insert assignments appear in red, including an insert whose processor or owning strip is missing, or which refers to itself. These errors are detected even when all visible connections to the deleted instance have already been removed. Backend validation failures appear directly beside **Check routing**. Previous check results are discarded after changes to the rack, patch definitions, routing or mixer settings; check again after editing.

Feedback loops appear above the mixer strips as soon as their connections are present. Each path lists instrument and effect names, instance IDs (to distinguish repeated uses of a patch), MIDI channels, exact audio ports and processing stages. **Before inserts** is the instrument's raw output; **Processed strip** is its output after insert returns; **Input** feeds the effect's processing. Ordinary insert chains returning to their owner's strip are valid.

Choose **Show offending connections** to open the routing diagram and enable **Only feedback connections**. Red borders and **Feedback loop** badges identify the affected strips, inserts, sends and route groups; the Destination matrix highlights the same connections. Expand **Exact channel mapping** to inspect or remove individual channel connections. Turn off the filter to inspect surrounding routes. Highlights update immediately after a routing edit, and all connections become visible again once the loop is gone.

The display shows one closed path for each connected group of feedback loops and highlights every cyclic connection in that group. A group may contain additional overlapping loops. Unrelated incoming routes and downstream output routes are not highlighted. Break every remaining loop, including both stereo channels; muting, bypassing or lowering a send does not remove its processing dependency.

For example, sending a lead's processed strip into its own Reverb insert closes **Lead → Reverb → Stereo Widener → Lead**. Use a separate shared Reverb instance with its own output return if that is the intended effect. The screenshot below uses a temporary copy of Effectomania's earlier routing.

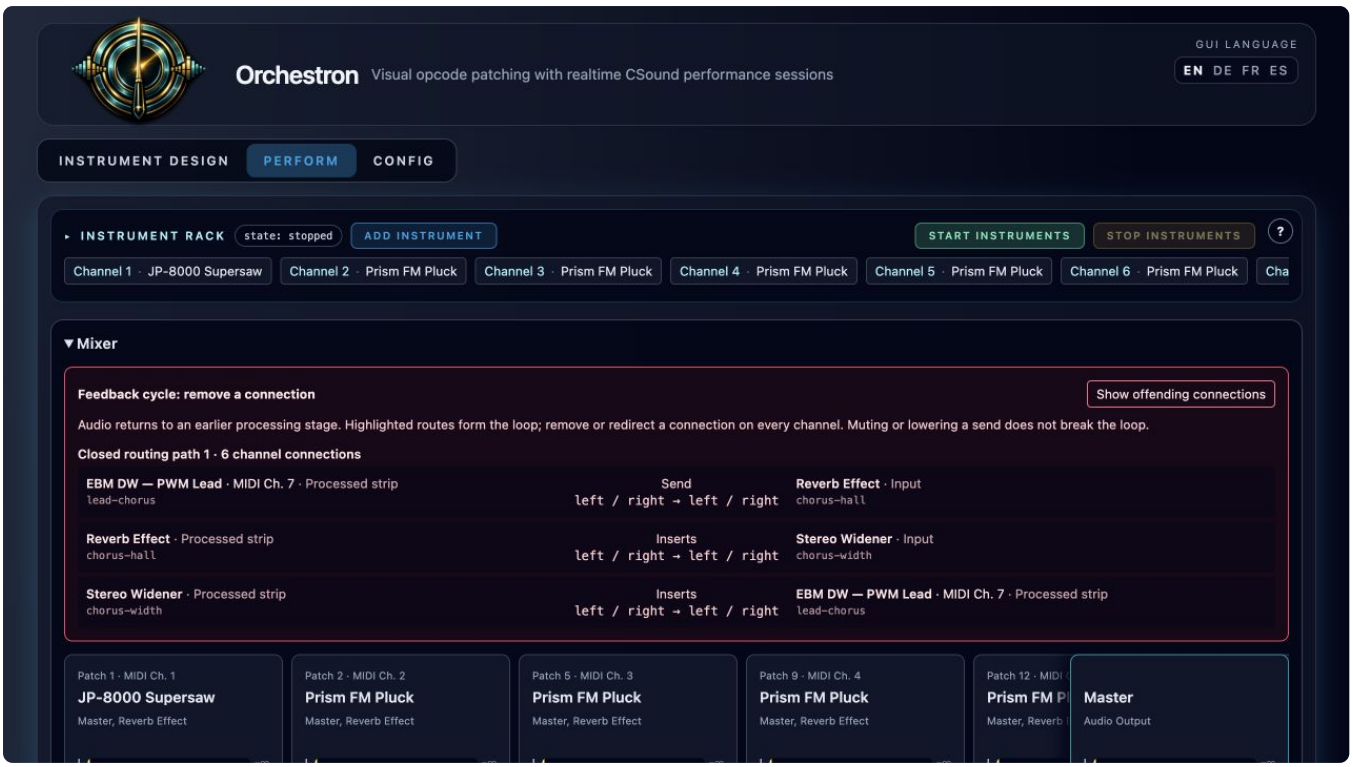

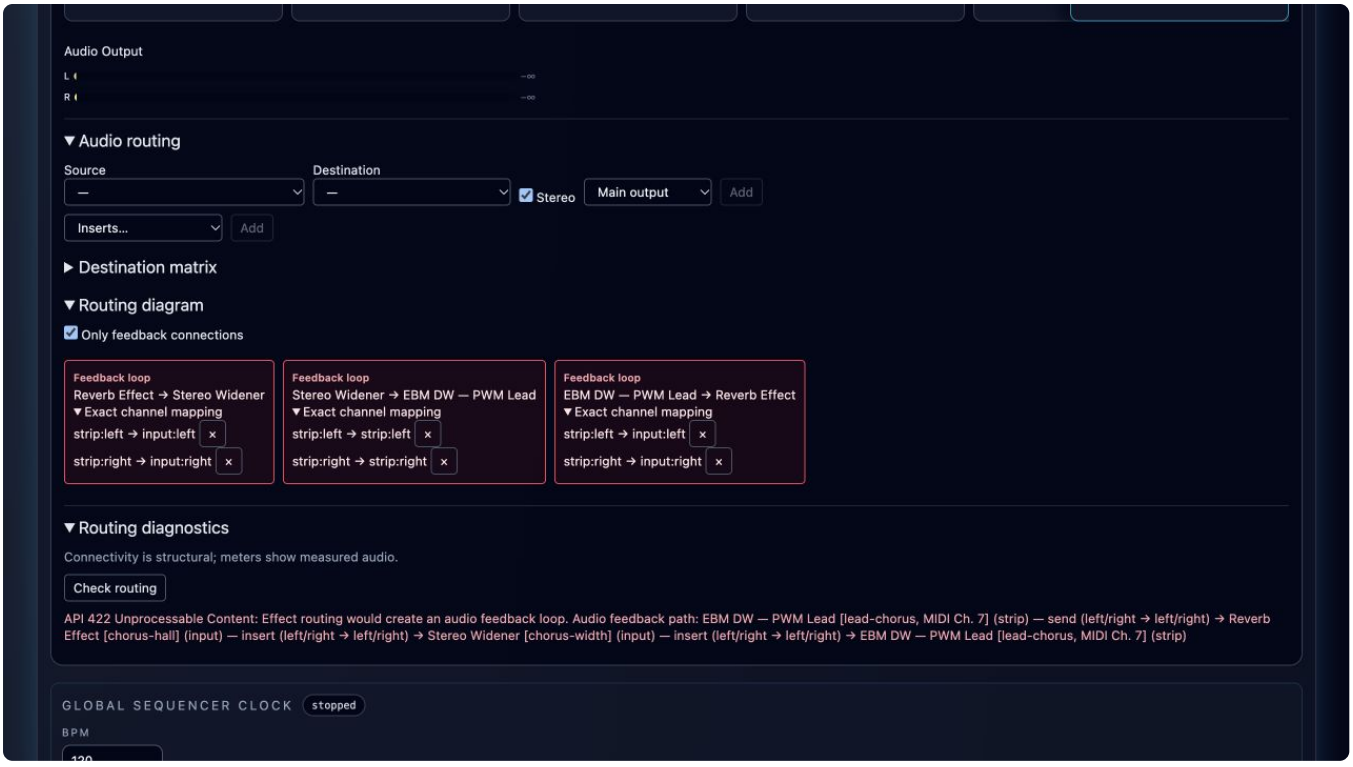

Symptom or diagnostic	What to check
Repair route / Missing reference	Restore the referenced patch or repair the destination and exact port names in Destination matrix.
Invalid insert ownership	Repair the insert assignment that refers to a missing instance or to itself. Removing audio connections alone does not repair an invalid insert assignment.
Feedback cycle	Follow the named path and choose **Show offending connections**. Remove or redirect a connection on every affected channel. Playback and CSD export remain blocked until all cycles are resolved.

Symptom or diagnostic	What to check
Connect an audio output / Select an audio source	Follow the source to an output sink; route audio into effects that require inputs.
Complete stereo mapping / Unused ports	Check both channel mappings and the declared audio groups. These structural warnings do not prove a patch is silent.
Review parallel output paths	Check whether both direct and processed copies intentionally reach the output. Serial insert chains, including Master inserts, do not trigger this warning.
Muted or silent path	Clear Mute or raise a fader/send from silence; also check active solos and whether the source is producing notes/audio.
CLIP	Reduce gain along the contributing paths. No limiter is applied automatically.

The same routing diagnostics appear in CSD exports. A serial insert feeds back into its owner's strip before reaching Audio Output. A parallel-output warning means the other branch can reach Audio Output through a separate exit, such as a shared effect return or an additional output from an insert processor.

A valid route can still be silent: here the Sine Lead strip is muted.

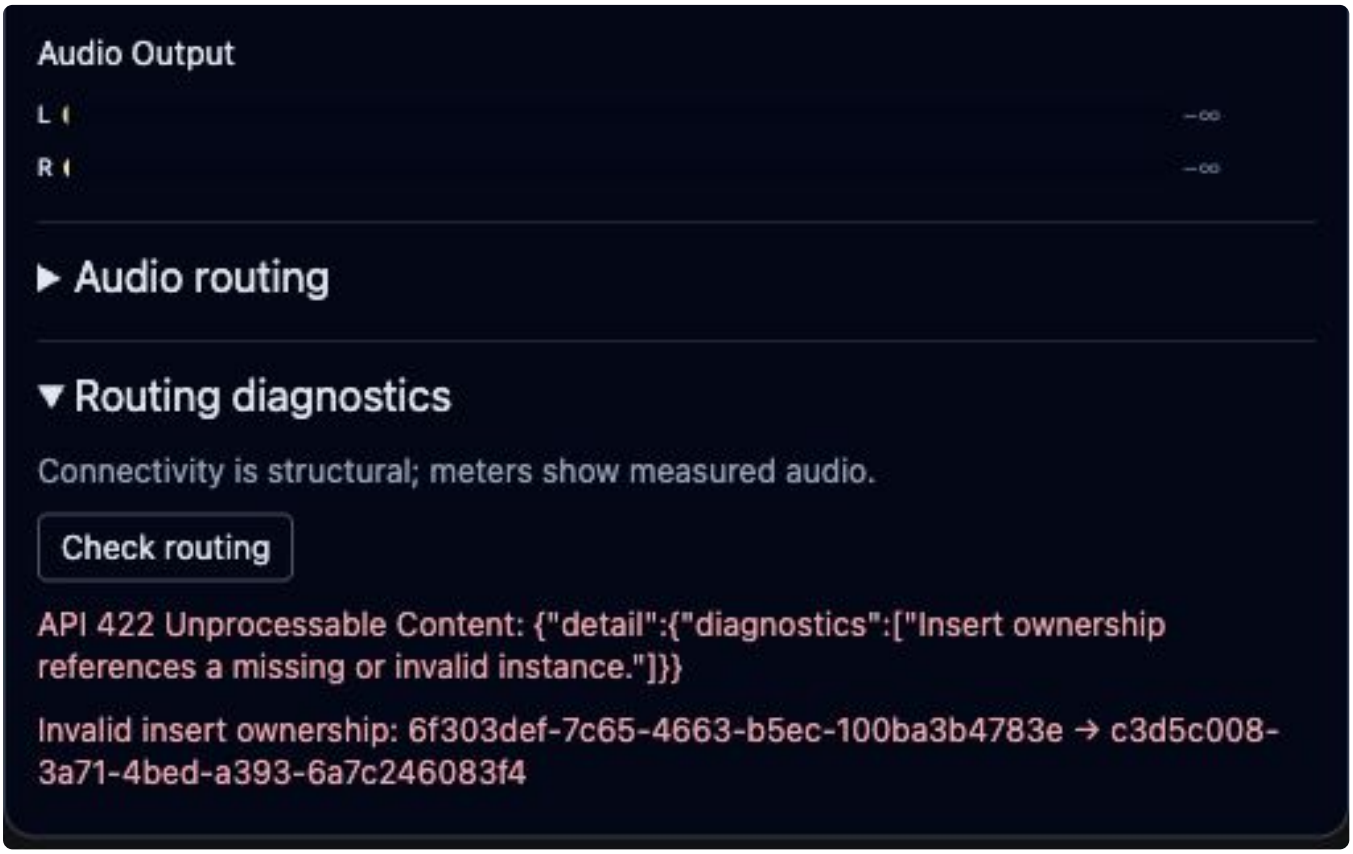

A diagnostic preview of the damaged Sunlit Strings performance: the missing insert assignment is flagged locally and the backend failure appears in the same section. The saved performance was repaired separately.

Persistence and compatibility

Save Performance retains routes, strip gain, balance, mute, bypass, sends, inserts and instance settings. Clone and native exports preserve them too. Mixer automation recording is not included. See Performance Import / Export for portable

backups and offline rendering.

Older performances are upgraded when opened. Older Level controls become audio gain; they no longer scale note velocity, so velocity-sensitive instruments can change timbre. Authored note velocities are retained.

Upgrading existing Masters

The current Master is a fixed output strip. Customized older Master instruments remain as processing instruments routed into it. If a legacy Master is missing, stop playback and use **Replace missing Master with internal output**. Check any unresolved channel mappings in the routing matrix, then save the performance.

Library migration and recovery procedures are in the developer reference.

Navigation: Up | Previous | Next

Using Performance Controls

Navigation: Up | Previous | Next

Use an instrument’s performance controls to adjust envelope times, tone or other sound parameters for your song. Each rack instance owns its settings, even when several instances use the same patch.

Define a Controller

To add or wire a control, follow Defining Performance Controls in Instrument Design. The rest of this guide explains using controls already supplied by an instrument.

Tune a Rack Instrument

Open the instrument's card in the rack. Related controls may appear in labeled groups such as Filter or Envelope. The layout adjusts to browser width and zoom without losing unfinished edits. To change group names or wiring, use Defining Performance Controls.

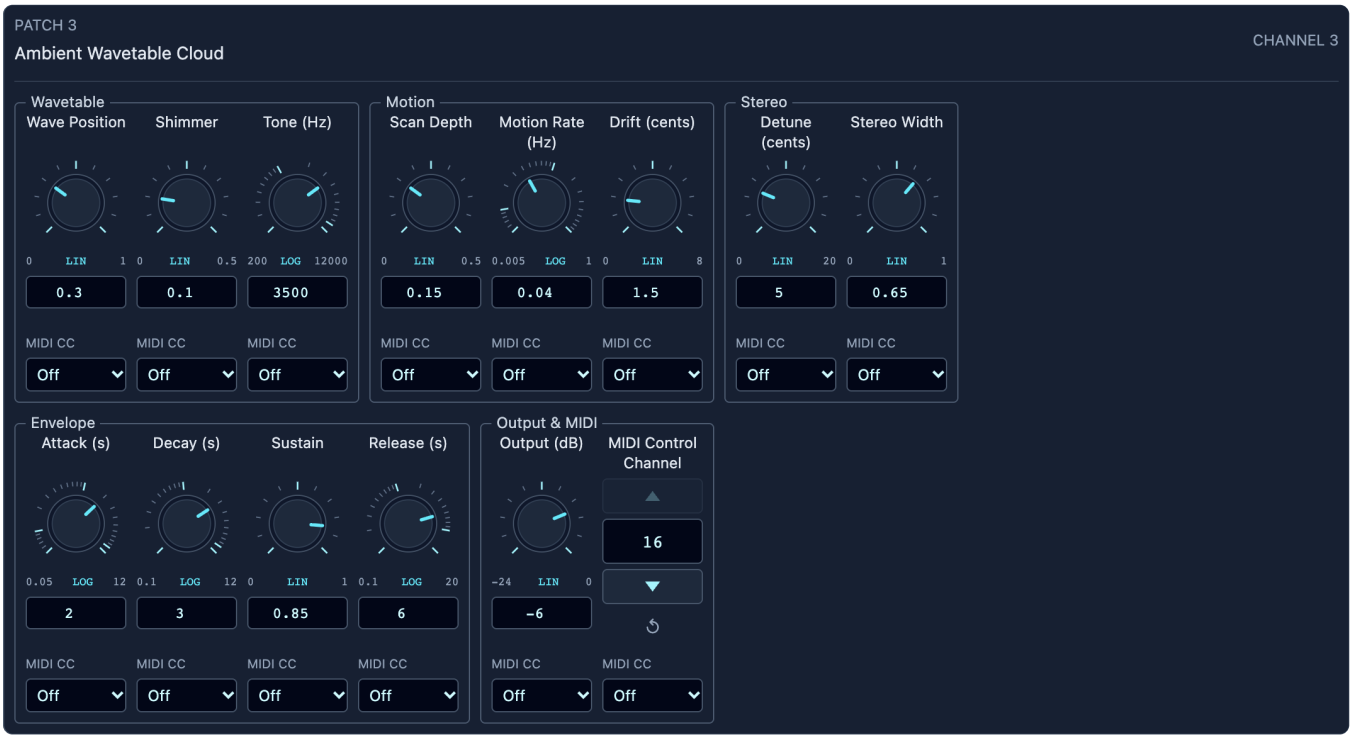

Float retains its knob, caption, scale dashes, endpoints, LIN/LOG indicator and current value. Linear ticks are equally spaced; logarithmic ticks follow ratios and mark decade boundaries.

- Drag vertically, or use arrow keys. Hold Shift for fine adjustment.
- Enter an exact number below the knob; Enter or leaving the field applies it. Escape cancels.
- Home and End choose the endpoints. Double-click, Delete, or Backspace on the focused knob resets to the patch default.

Integer shows an upward triangle above the value and a downward triangle below it. Each click changes the value by exactly one; a button disables at its limit. Type a whole number or use Up/Down, Home and End. Enter or leaving the field applies an entry; Escape cancels it. Fractional entries are rejected, and whole numbers outside the configured range are clamped. Uncommitted drafts survive collapsing and expanding the rack.

Bool shows a toggle with visible On/Off text. Click it, or focus it and use Space or Enter. Off outputs `0` ; On outputs `1` . Integer and Bool each provide a **Reset to patch default** button (↺), which removes the instance override.

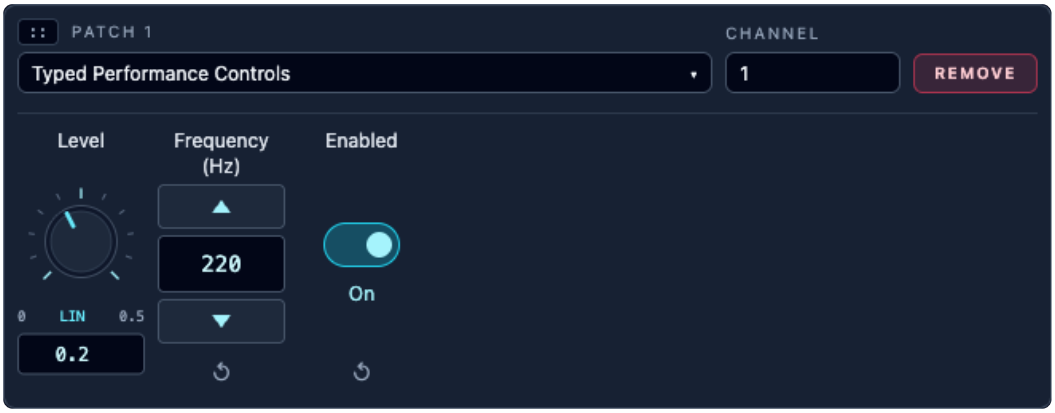

Changes affect newly started notes. Sounding notes keep their initialized value. For continuous instruments, Orchestron detects each controller's behavior from its connections:

- **Live**: every connected input accepts K-rate, so edits apply during realtime playback.
- **Requires restart**: at least one input requires I-rate. If a controller feeds both kinds, all of its destinations keep the initial value until the rack restarts.
- **Not connected**: the controller has no destinations.

Input formulas follow the same connection rules. If/Switch selectors remain initialization settings. Save the instrument after rewiring; changing a controller's type or detected timing while it runs requires restarting the rack. Existing instruments are analyzed automatically when loaded or compiled. The same behavior applies with **Parallel playback**, including fallback to single rendering.

Both **CSD (MIDI)** and **CSD (SCORE)** exports keep continuous-instrument settings fixed at initialization, including controls marked Live. Note and legato-phrase behavior is unchanged.

For an instrument using `midi_legato`, settings are sampled at the start of each **phrase**. Touching and overlapping notes keep the same voice and settings; the next note after a gap reads current values.

For sound-specific settings and recipes, see Supplied Instruments and Effects. The same rack controls work across instruments, even when their parameter names differ.

Additional instrument examples and Stereo Chorus depth are covered in that chapter.

Receive MIDI CC in a Performance

Each controller in the Perform rack has a **MIDI CC** selector: **Off** or **1–50**. The assignment is saved on that rack instance in the performance. It does not change the instrument definition. All controls receive on the instrument's MIDI channel, set once at the top of its rack slot; choose several target channels in a controller sequencer to control several instances together. Multiple controls can share the same CC. Continuous instruments and effects have no MIDI channel, so assignment is unavailable.

The engine must be running, although the arranger may be stopped. External input uses the selected session MIDI input;

manual MIDI controller lanes and controller sequencers also work. Ordinary MIDI delivery continues, including any patch-authored CC response. Muted controller lanes do not change mapped settings.

Type	Incoming CC value 0–127
Float	Scale across Minimum–Maximum using the controller's LIN/LOG scale
Integer	Scale across Minimum–Maximum and round to the nearest integer; ties go toward positive infinity
Bool	0–63 is Off; 64–127 is On

Values 0 and 127 reach the exact endpoints. MIDI has 128 positions, so large integer ranges skip some values. Widgets follow incoming MIDI, and Save Performance stores the resulting values. Mouse and keyboard edits remain available; the next matching CC can change the value again. Repeated CC messages apply even if their raw value is unchanged. Assignments can change during playback without restarting the rack. A new assignment waits for the next delivered CC. Off keeps the current value; Reset to patch default keeps the assignment.

These remain I-rate settings: new notes and new legato phrases use the changed value, while held notes retain their initialized setting. Mapping CC does not add continuous modulation to an already sounding note. Native bundles preserve the assignments. Both CSD export modes preserve authored CC-driven changes; external hardware gestures are not recorded into the arrangement.

Save, Reload, and Export

Untouched controls follow the patch default. After a value changes, an explicit instance override remains until reset, including when it is subsequently set to the current default. App state retains edits immediately through the existing persistence workflow; Save Performance is required to update a saved performance. Native JSON/ZIP bundles and both CSD (MIDI) and CSD (SCORE) exports preserve instance values. The CSD contains its initial settings and runs without a controller client. Mouse edits do not send MIDI CC. Optional assignments receive CC input as described above.

Controls are identified independently of their labels, so duplicate labels can appear. Changing the selected patch clears its overrides. Editing a patch range clamps saved values. Changing to Integer rounds overrides to the nearest integer (ties toward positive infinity) and clamps to the range. Changing to Bool maps values below `0.5` to `0` and values at least `0.5` to `1`. Deleting a controller removes its override. These adjustments show a notice in the rack; untouched defaults remain implicit. Standalone patch export uses defaults. Isolated draft audition starts from defaults and offers temporary compact controls; its values survive an audition restart for the same draft but never change a saved performance.

The Stereo Widener, Bit Crusher, and Harmonizer examples expose four, five, and six performance controls respectively. Check each controller's **Live** or **Requires restart** label to see when edits take effect. Harmonizer aligns its dry signal with its measured 32 ms processing delay at 48 kHz.

Automation from instrument clips

The clip arranger's **Generate automation** action creates one bound lane per distinct assigned CC. It derives routing from this rack instance and retains suspended automation when assignments are switched Off. Curves support song-length durations, ten points and editing directly inside clips. See Instrument automation clips for generation and split/merge rules. I-rate behavior remains unchanged: new notes or phrases adopt the automated value.

Navigation: Up | Previous | Next

MIDI Setup and Inputs

Navigation: Up | Previous | Next

This page covers practical MIDI setup and how external MIDI binding works in Orchestron's unified browser-clock runtime.

External MIDI Input Binding In Orchestron

External MIDI input is bound in the **Instrument Design Runtime panel** (not on the Config page).

Workflow:

1. Open Instrument Design.
2. Use the Runtime panel `MIDI Input` dropdown.
3. Select the desired input:

- `internal:loopback` for the built-in app-only loopback and default fallback binding
- a `host_bridge` input published by `host-midi-helper` for external hardware or DAW MIDI

1. Start the instrument session and play from your hardware/DAW.

The selected MIDI input is reflected in the Perform page footer status line.

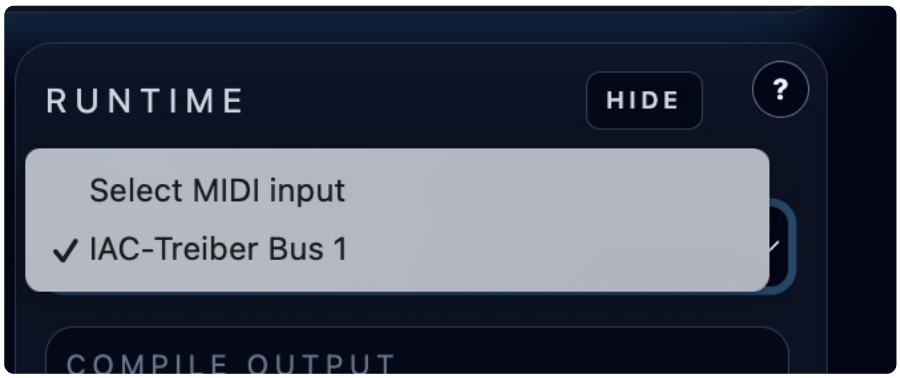

Runtime panel MIDI input dropdown expanded with available MIDI inputs.

Internal Performance Controls Do Not Require External MIDI

These features generate MIDI/control events internally and do not require an external MIDI input device or host bridge:

- melodic sequencers
- controller sequencers
- arpeggiators
- piano rolls (on-screen keyboard)
- manual MIDI controller knob lanes

External MIDI input is only needed when you want to play Orchestron from outside the app. Internal MIDI is always delivered directly into the session engine and ignores the external input binding.

That means:

- switching the session input does not disable sequencers, arpeggiators, piano rolls, or manual controller lanes
- binding a helper-provided external input adds another source into the same engine scheduler

External MIDI Bridge

External MIDI arrives through the optional host helper in `host-midi-helper/`.

General setup pattern:

1. Start the backend with `VISUALCSOUND_HOST_MIDI_TOKEN` set.
2. Start `host-midi-helper` on the same host OS with the same token.
3. Bind the helper-published input in the Runtime panel.

If no helper is running, `internal:loopback` remains available and internal performance tools still work.

macOS Loopback Setup (IAC Driver)

For DAW/app -> Orchestron MIDI routing on macOS:

1. Open **Audio MIDI Setup**.
2. Open **MIDI Studio**.
3. Open **IAC Driver**.
4. Enable **Device is online**.
5. Add/select an IAC bus (for example `IAC Driver Bus 1`).
6. Route your DAW/app MIDI output to that IAC bus.
7. Select that input in Orchestron's Runtime panel.

Troubleshooting MIDI Inputs

No External MIDI Inputs Listed

- Ensure your OS/device/virtual bus is enabled (for example IAC Driver on macOS, ALSA/JACK route on Linux, loopMIDI on Windows).
- Confirm `host-midi-helper` is running and uses the same token as the backend.
- Restart the helper after changing system MIDI configuration.

Perform Page Footer Shows `midi input: internal:loopback`

- Internal app MIDI is working.
- External hardware/DAW MIDI will not reach the session until a helper-provided input is bound.

Sequencer/Piano Roll Works But Hardware MIDI Does Not

- This usually means internal controls are working, but no external helper-provided MIDI input is bound.

Related Pages

- Runtime Panel and Compilation Workflow
- Live Status and Safety Controls

Navigation: Up | Previous | Next

Piano Rolls

Navigation: Up | Previous | Next

Piano Rolls provide manual performance input with scale-aware keyboard highlighting.

Use the pen beside the device name to rename it. See Device Names for editing controls, validation, and import/export behavior.

Collapse the Panel

Collapse the Piano Rolls header to hide the keyboards and their settings. Add and help remain available; adding a keyboard expands the group. Panels start expanded and remember your choice across view switches until browser reload. Collapsing does not stop playback or discard edits.

Purpose

Piano Rolls are for live/manual playing on an on-screen keyboard while the instrument engine is running.

They coexist with melodic sequencers and controller sequencers, which means you can jam manually on top of sequenced playback.

Starting a piano roll is independent from the multitrack arranger transport. If the instrument engine is stopped, `Start` brings up the engine so the keyboard can play without starting arrangement playback.

Add / Remove / Enable

Each piano roll supports:

- `Add Piano Roll`
- `Remove`
- `Start / Stop` (enable/disable that piano roll)

A piano roll only sends notes when:

- that piano roll is enabled
- instruments are running

Piano Roll Controls

Each piano roll has:

- `MIDI Channel (1..16)`
- `Velocity (0..127)`
- `Scale`
- `Mode`
- `Autochord` (off by default)
- `Chord notes (3, 4, or 5, default 3; available when Autochord is on)`

MIDI Channel and Velocity are grouped together, with a small gap before Scale, Mode, Autochord, and Chord notes. All labels appear above their controls; the Autochord checkbox matches the height of the adjacent selectors.

Each MIDI-driven rack instrument requires a unique MIDI channel. Set the piano roll to the channel of the instrument you want to play.

Autochord

Enable **Autochord** to extend a highlighted key into an ascending chord. **Chord notes** counts the pressed note, so in C major

pressing C4 plays:

Chord notes	Sounding notes
3	C4–E4–G4
4	C4–E4–G4–B4
5	C4–E4–G4–B4–D5

The chord starts on the key you press and stacks every other highlighted note, continuing into higher octaves. It follows the same scale/mode as the keyboard. In **mixed** mode, it uses only notes common to the active scales; the resulting intervals may differ from thirds. For example, C major and G major share C, D, E, G, A, B, so pressing C4 with three notes plays C4–E4–A4. Sparse intersections can produce octave-spaced voicings. Keys outside the highlighted set still play single notes.

All chord notes use the keyboard's MIDI channel and velocity, and visible sounding keys light up together. Release ends the chord; notes shared with another held chord stay on until their last held key releases. Changing scale, mode, channel, size, or the toggle affects the next press without altering held chords. Upper voices may extend beyond the visible keyboard, but notes above MIDI 127 are omitted.

Each piano roll retains its own toggle and size through Save/Load, Clone, autosave, and native JSON/ZIP export/import. Turning Autochord off retains the chosen size. Older performances default to off/3. These controls apply to the on-screen keyboard; live gestures are not recorded into sequencers or offline CSD exports.

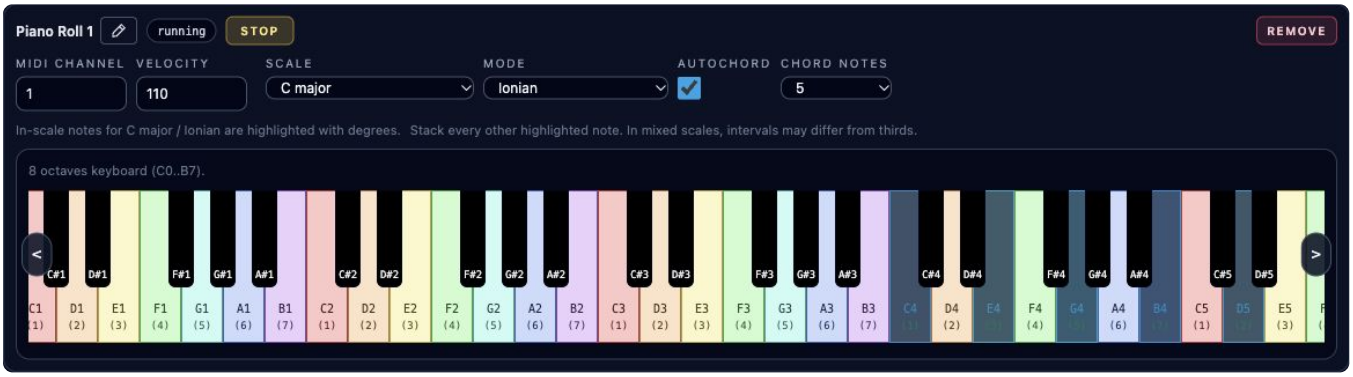

Autochord with five notes selected and the generated chord highlighted.

Scale/Mode Following Behavior During Live Playback

When the instrument engine is running and a piano roll is enabled, the piano roll follows the currently playing melodic and Score Sequencer patterns for harmonic guidance. Clip playback uses the pattern at the audible song position; Play edit uses the active audition pattern. Selecting or editing a different pattern, or queuing an audition, does not change the keyboard's guidance before it plays.

While following, Scale and Mode show the playing theory and are disabled. Key highlighting and Autochord use that same theory. The keyboard's own scale and mode remain saved separately. When no melodic source is playing—including stopped or paused tracks, clip gaps, rests between patterns and song end—the keyboard restores its own scale and mode and both selectors become editable. Starting the keyboard alone preserves these settings.

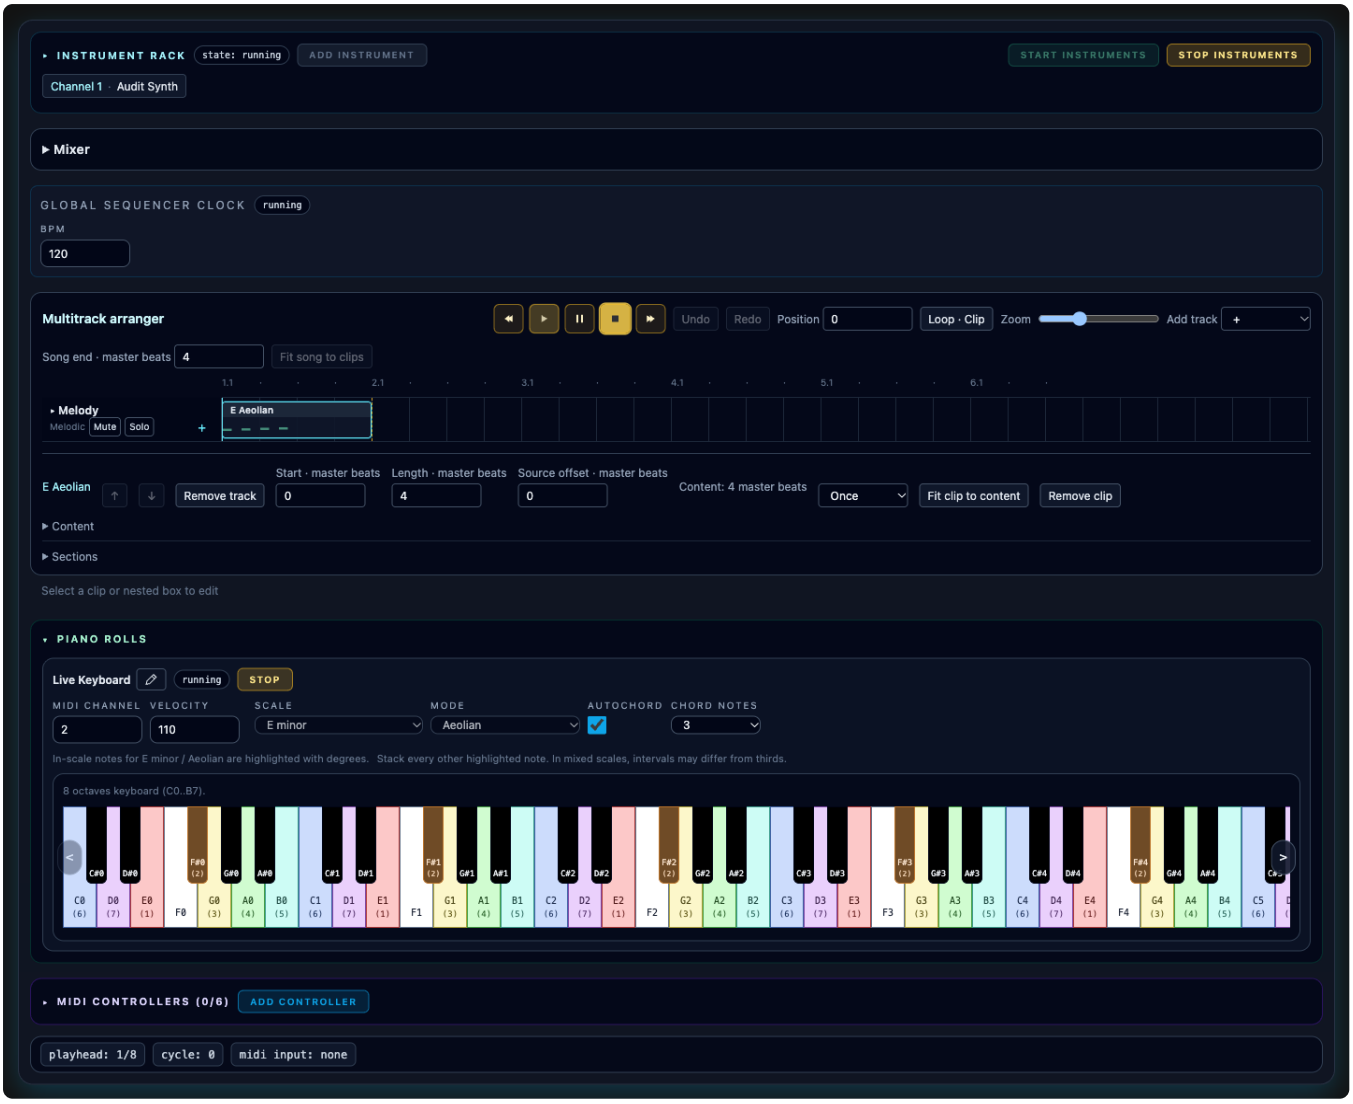

The keyboard follows the playing E Aeolian clip; stopping melodic playback restores its own scale and mode.

Shared Scale/Mode Case

If running melodic sequencers agree on the same scale+mode:

- the piano roll shows that shared scale/mode
- keys are highlighted according to that theory

Mixed Scale/Mode Case

If running melodic sequencers disagree:

- piano roll controls show `mixed`
- highlighting is computed from the intersection of the active theories
- only notes common to the active scales/modes are highlighted

This is especially helpful when multiple tracks create potentially conflicting harmonic contexts.

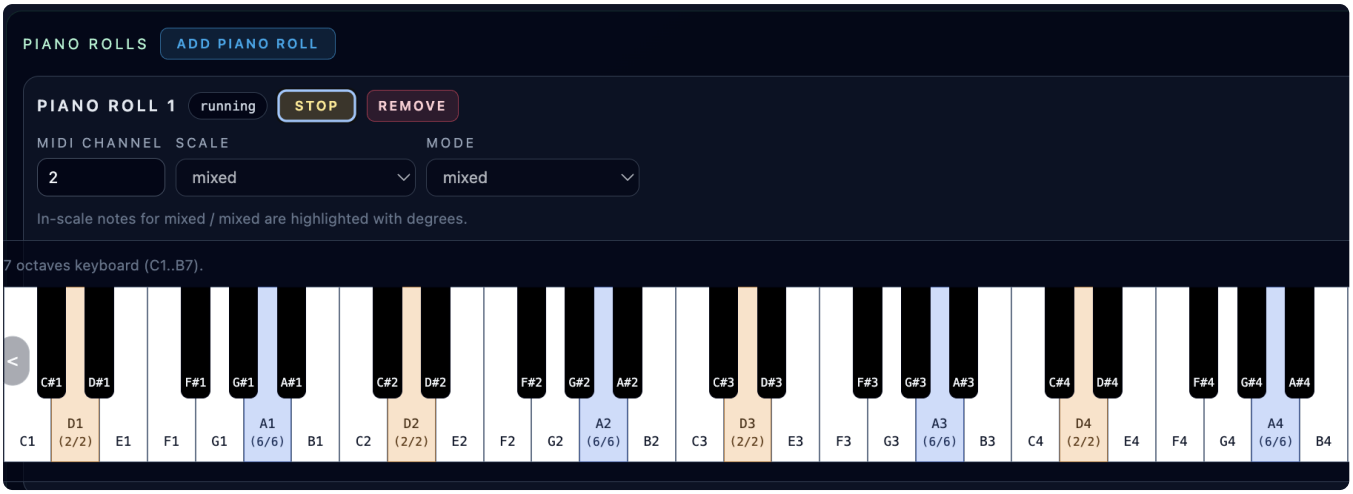

Piano roll with scale-aware key highlighting and mixed-theory display behavior.

Keyboard Range and Layout

The piano keyboard spans:

- `C0 .. B7` (96 notes)

UI behavior:

- Full-width horizontal keyboard section
- Scroll arrows appear when the keyboard overflows the viewport
- Smooth horizontal scrolling controls

Key Highlighting and Degree Labels

Highlighted keys show:

- in-scale status
- scale degree labels (`1..7`)
- rainbow degree coloring for quick visual recognition

In mixed-theory situations, a key can show combined degree labels (for example `2/5`) when it belongs to multiple active track theories.

Playing Notes (Pointer Interaction)

- Press/hold on a key sends `note_on`
- Release sends `note_off`
- Works for both white and black keys

The UI tracks held notes and releases them on pointer cancellation, lost capture, panel collapse, keyboard removal, or when the piano roll or instrument engine stops.

Practical Use Cases

- Live melody or bass improvisation on top of sequencer patterns
- Testing an instrument patch in musical context without external MIDI hardware
- Performing harmonically guided lines while melodic sequencers define the key center

Navigation: Up | Previous | Next

MIDI Controllers

Navigation: Up | Previous | Next

The MIDI Controllers panel provides manual CC control lanes with interactive knobs.

Controllers use compact cards in both arranger layouts: one column on narrow screens, two on small screens, three on wide screens, and four on extra-wide screens. Additional controllers wrap onto the next row.

Use the pen beside the device name to rename it. See Device Names for editing controls, validation, and import/export behavior.

Collapse the Panel

Collapse the MIDI Controllers header to hide the knobs and their settings. The count, Add button and help remain available; adding a controller expands the group. Panels start expanded and remember your choice across view switches until browser reload. Collapsing does not stop playback or discard edits.

Overview

This panel is for manual MIDI Control Change performance (hands-on CC control), separate from controller sequencer automation.

Capacity

- Up to **6** controller lanes can be added

The `Add Controller` button is disabled after 6 lanes are present.

Per-Controller Lane Controls

Each controller lane includes:

- Controller name label (`Controller N` if unnamed)
- Running/stopped state badge
- `Controller #` field (`0..127`)
- **MIDI Channels** checkboxes at the top right: 1–8 above 9–16
- `Start / Stop` enable toggle
- `Remove` button
- Knob control (value `0..127`)
- Numeric value display

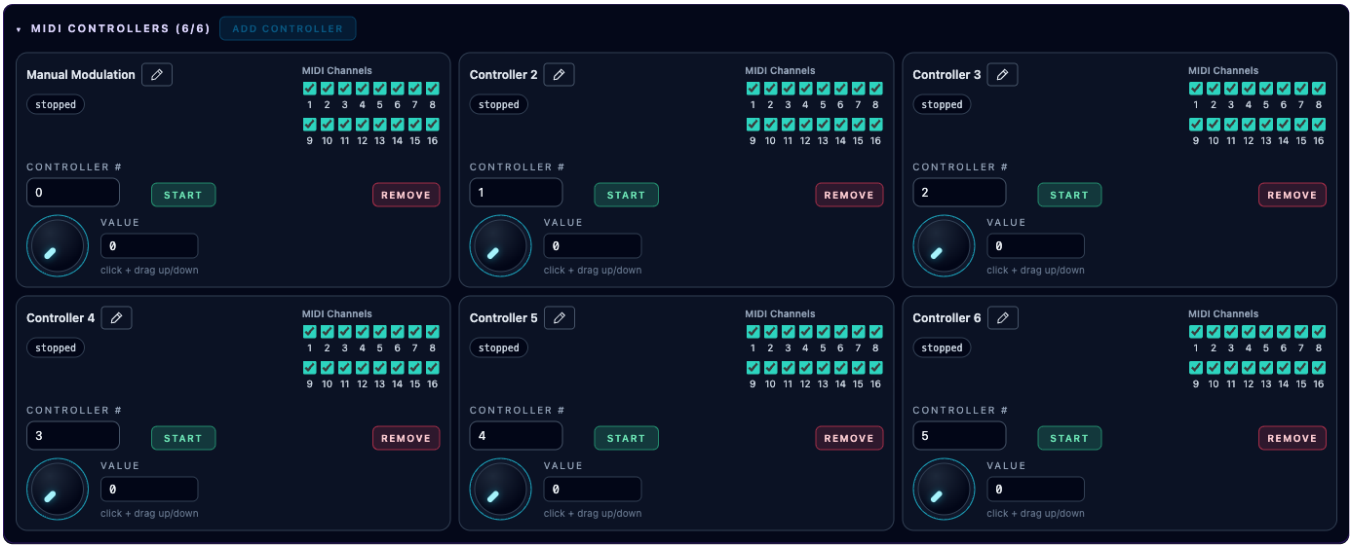

Compact manual MIDI controller cards below the clip arranger, with knobs, CC numbers, channel selections, and values.

MIDI Channels

Check the channels that should receive this lane's CC messages. All 16 are checked by default (OMNI); at least one must remain checked. The last checked box cannot be unchecked until another channel is selected.

When the lane and instrument session are running, changing channels immediately sends the current value to the new selection. Unchecked channels retain their last received value. Starting a lane or session, changing the CC number, and adjusting the knob all use the selected channels.

The selection survives Save/Load Performance, browser reload, native JSON/ZIP export/import, and both CSD exports. Older performances without a selection load with all 16 channels checked. Controllers using the same CC on different channels can have independent values; overlapping mappings receive the latest transmitted value.

Knob Interaction

The knob uses pointer drag interaction:

- Click and drag vertically to change the value
- Up increases the CC value
- Down decreases the CC value

The UI also shows a numeric value readout next to the knob for exact control.

Live Sending Behavior

When a controller lane is enabled and the instrument session is running:

- changing the knob sends MIDI `control_change` messages immediately

This is ideal for expressive live performance adjustments.

Manual Controllers vs Controller Sequencers

Use manual controller lanes when you want:

- direct, improvised control
- tactile-feeling UI knobs
- on-the-fly sweeps and tweaks

Use controller sequencers when you want:

- repeatable automation
- transport-synced movement
- curve-based modulation patterns

Both can exist in the same performance.

Navigation: Up | Previous | Next

Multitrack Arranger

Navigation: Up | Previous | Next

The multitrack arranger brings melodic, Score, drummer, controller and arpeggiator tracks onto one timeline. Clips set where content plays; reusable patterns, groups and supergroups hold the music.

New creates an empty arrangement and one piano keyboard. Choose **Add track** to create a track with its first clip and pattern, then assign its instrument and MIDI channel. Double-click empty track space to add another clip. Use a clip's **+** to add a new pattern or reuse existing content.

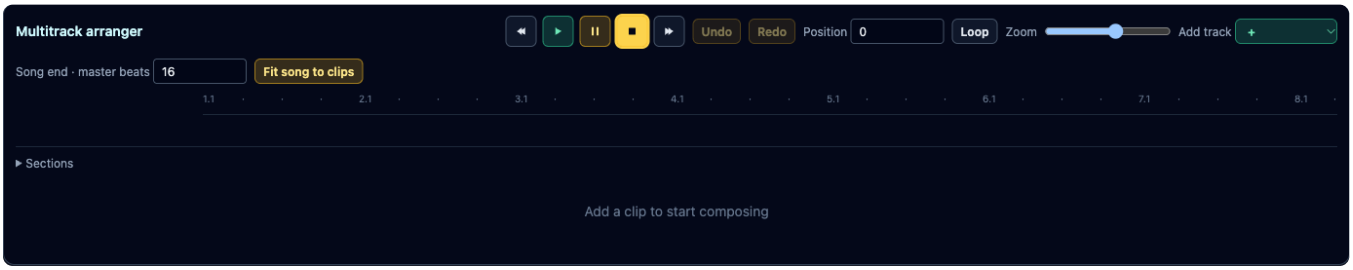

The top-right **?** opens help specifically for this clip workflow in the selected GUI language (English, German, French or Spanish). It covers clip windows, reusable content, musical and structural drafts, editing scope, transport, sections, transposition and instrument automation. Close it with **Close**, Escape or a click outside the dialog.

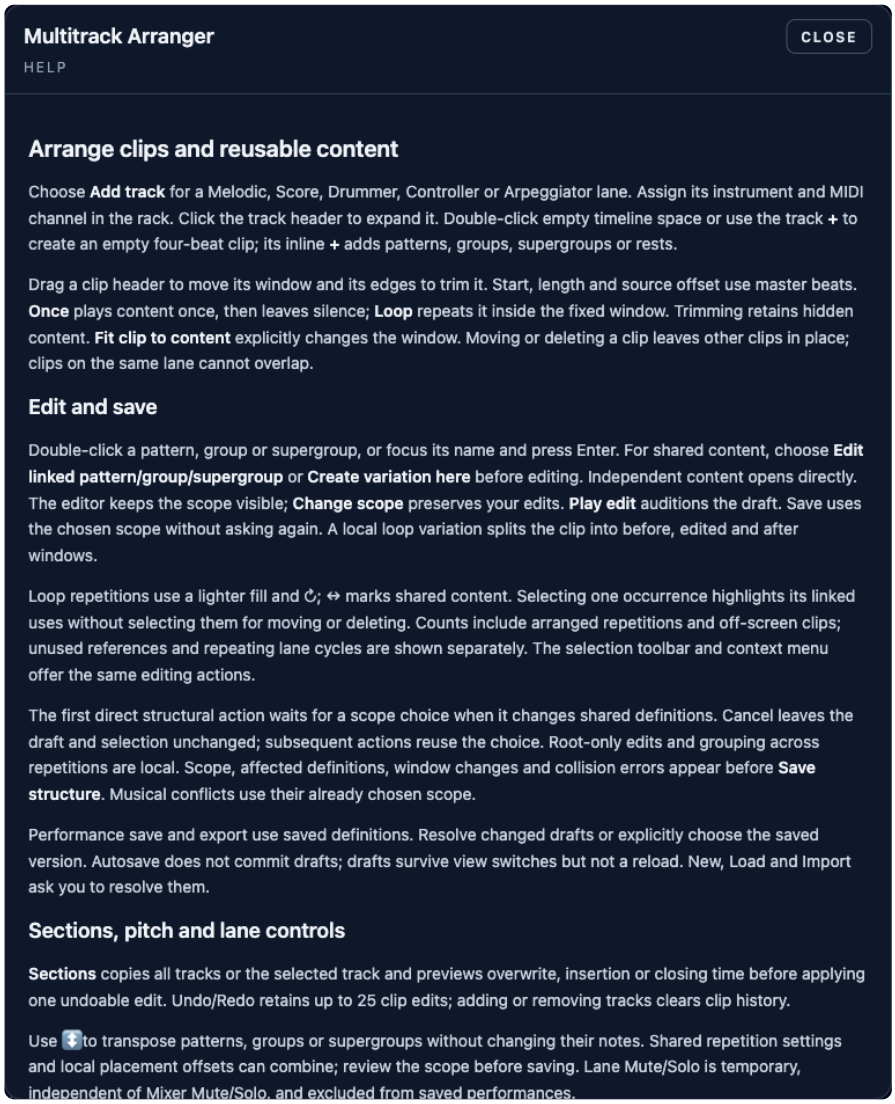

Arrange fixed windows

Unselected tracks stay compact. Click a track header to expand its clips into nested pattern, group and supergroup boxes. Drag a clip header to move it, or its borders to trim it; the inspector also accepts numeric start, length and source offset in master beats. Double-click empty lane space to create an empty four-master-beat Once clip immediately, snapped to quarter beats. The track **+** appends the same empty clip. Its inline **+** opens the content picker; inserted content replaces silence, preserving any remaining rest. Longer content can grow the window when space permits. Moves and deletion leave other clips in place. Overlap on the same lane is rejected.

The clip window and content duration are independent. **Once** plays the content from its source offset and leaves silence after its end. **Loop** repeats content within the window. Trimming preserves the original notes and content; **Fit clip to content** explicitly changes the window. A local edit of one generated repeat splits its original window into before, edited and after parts.

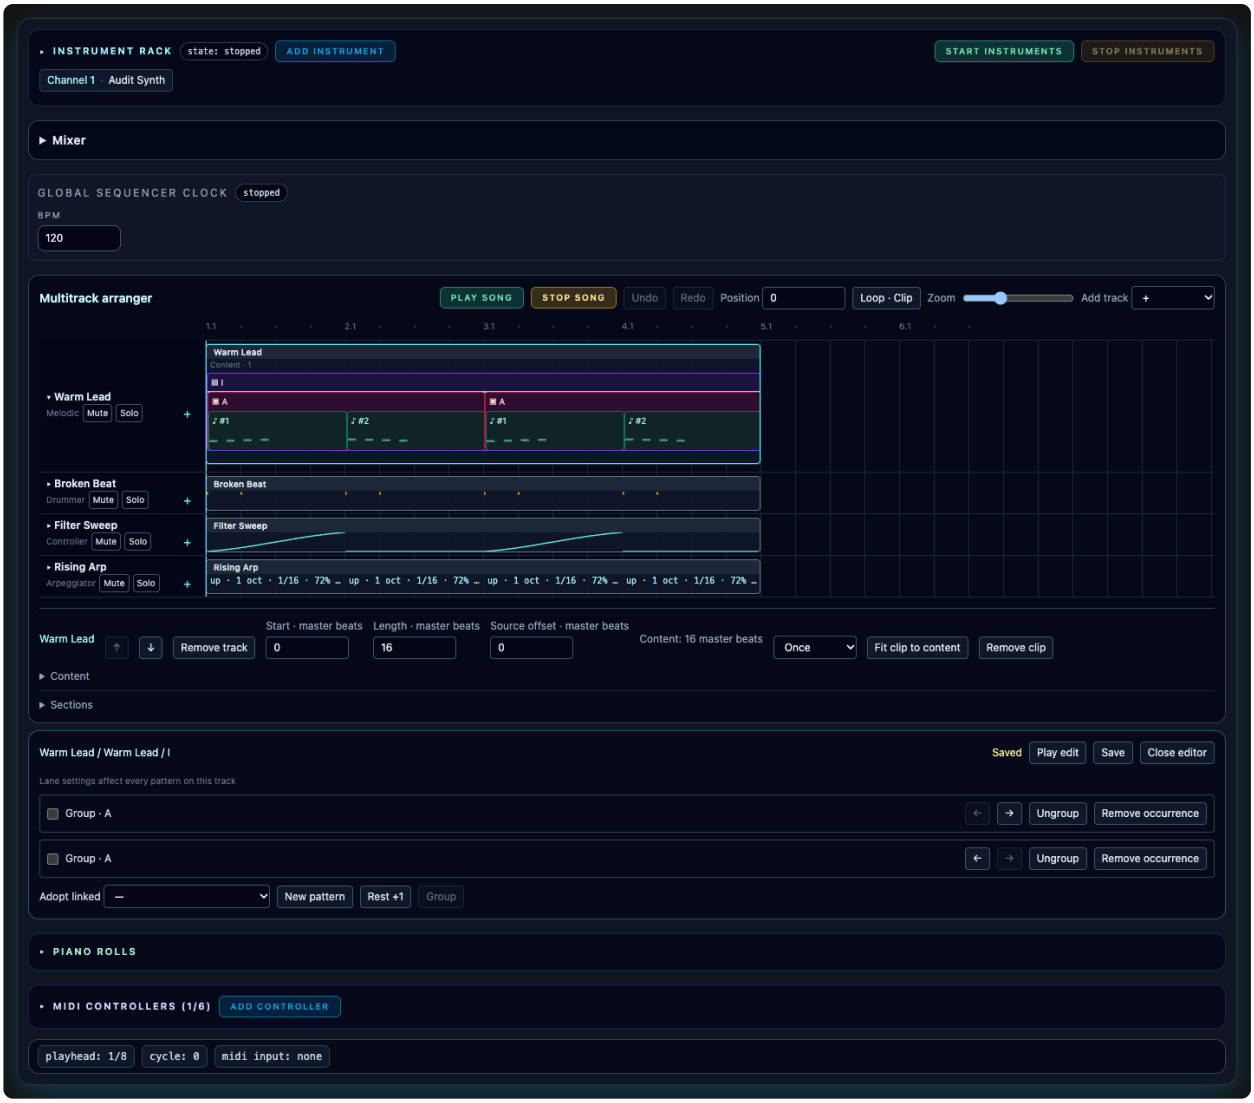

Open **Content** for the selected clip to adopt existing definitions, create a new pattern, add a rest, reorder children, or select adjacent children for Group/Supergroup. Grouping retains reusable definitions. Ungroup opens one level. Removing an occurrence retains its definition. Groups and supergroups have their own contextual child editor.

Sections copies a range of all tracks or the selected track. Paste and overwrite, paste and insert, insert time and close time show a preview before applying one undoable transaction. Track-order arrows change the visible lane order. Removing a track removes its source and drafts; source creation/removal clears clip history because source topology is outside clip-only undo.

Undo/Redo retains up to 25 clip edits through performance save/load, autosave recovery and native bundles. History stores changes compactly so ordinary edits do not repeatedly duplicate every pattern in the saved performance.

Temporary lane mute and solo

Use **Mute** to silence a track or **Solo** to hear it on its own. These controls leave the track's clock and place in the song unchanged. Several tracks can be soloed together; an explicit mute still silences its track. Drum rows are controlled together.

Track Mute/Solo is separate from the mixer, so tracks sharing an instrument can be controlled independently. Existing notes release normally. A muted controller track sends no new values; its target keeps the last value received. Soloing an arpeggiator also keeps its required input tracks active.

These controls are temporary: they are excluded from saved performances and exports and reset when a performance is loaded, imported or replaced.

Instrument automation

Select an instrument clip and use **Generate automation** to create one bound controller lane per assigned MIDI CC. Edit song-length curves directly inside expanded clips, with at most ten points and explicit Save/Discard. MIDI channel and CC follow the rack instrument. See Instrument automation clips for generation, duration, split/merge and suspended-target behavior.

Edit a pattern

For a Score track, add a clip, use its inline **+** / **Add content → New pattern**, then double-click the nested pattern to open the **Score sequencer**. Hold its piano keys to audition pitches or double-click them to append quarter notes. Choose **Edit linked pattern** or **Create variation here** before editing shared content. Save applies that choice directly. See the Score pattern walkthrough.

Double-click a nested box, or focus its name and press Enter, to open its sequencer or arpeggiator. The breadcrumb identifies the pattern you are editing; draft status, **Play edit** and **Save** sit above the musical controls. You can edit notes/chords/holds/strums, sparse score notes, drum velocities/ratchets, both controller datasets, and arpeggiator parameters/presets. Lane-wide timing, routing and instrument settings affect the whole track.

Melodic and Score clips use compact Note maps. Drums use step dots with a fixed velocity gradient. Arpeggiators show parameters, because their pitches depend on incoming chords. Controller Step Mode uses bars; Curve Mode uses the real interpolation. Selecting a curve pattern also opens the production curve editor inside its expanded lane, sharing the same draft as the full editor. Drag, click and double-click retain their usual behavior; the point list offers numeric position and value editing, including independent endpoint values.

Choose **Edit linked pattern** to change every occurrence or **Create variation here** to change only this one. The pattern editing guide owns the scope, save and draft rules. A variation within a loop can split its window into before, edited and after parts while preserving the song's timing.

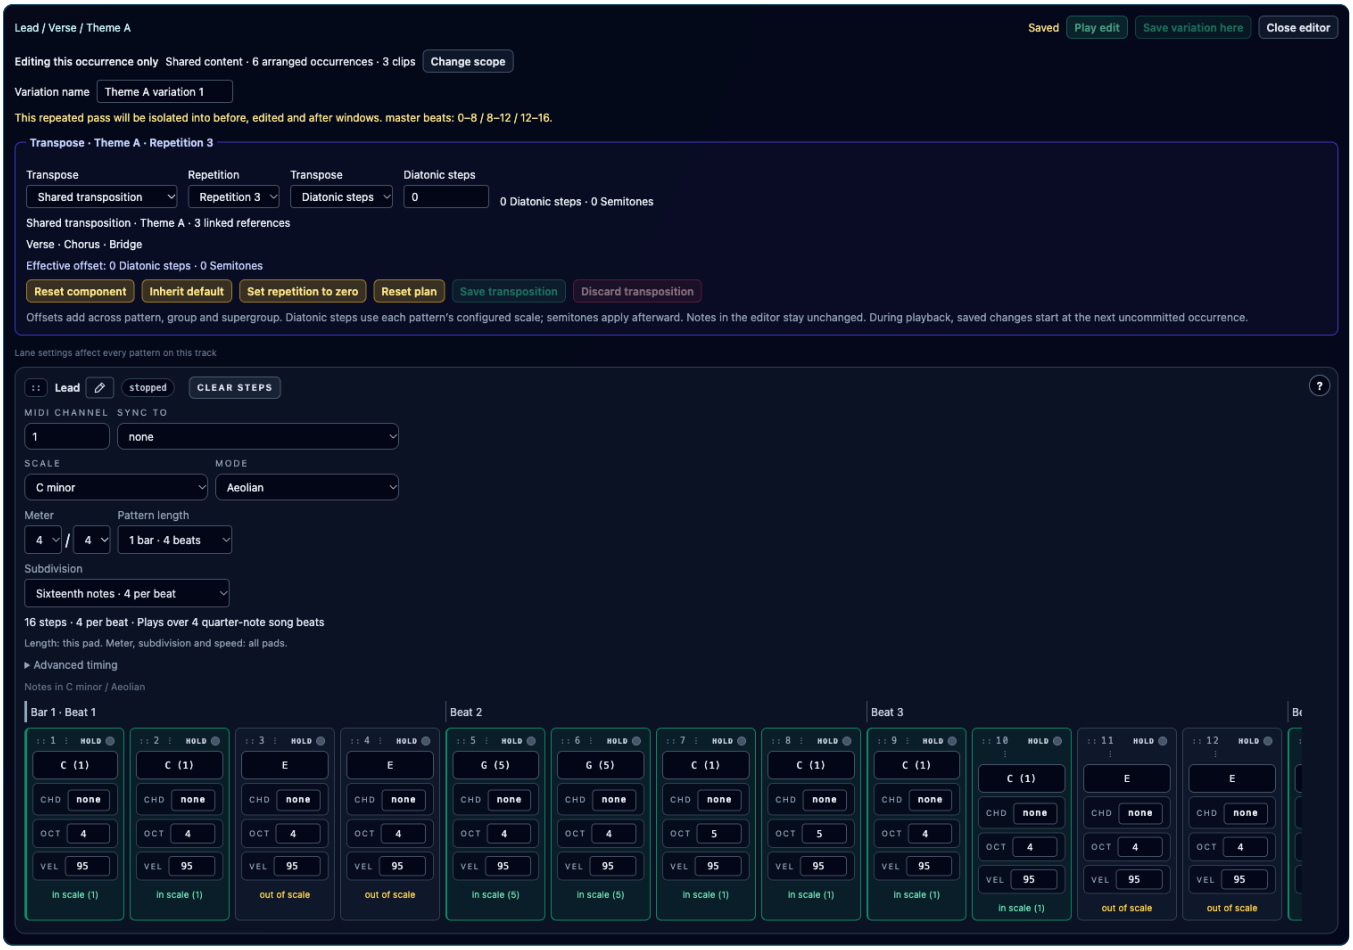

Play and jam

Play (triangle) and **Play edit** start instruments when needed. The tape controls provide rewind/forward by one beat, Play/resume, Pause, and Stop. Pause holds the song cursor; Stop returns to the selected loop start or song start. With the song stopped, Play edit auditions its target. With the song playing, the editor audition temporarily replaces its source lane while other lanes continue. Song start/stop/seek preserves the editor audition; Stop edit returns to the arrangement's current position or silence when the song is stopped. Audition state appears in the compact lane header. The open editor follows the audible pattern position with its existing step highlights, score playhead, controller cursor or arpeggiator status; standalone audition leaves the song cursor still. Closing the editor or leaving Perform ends the audition and retains its draft.

Arpeggiator **Play edit** waits for incoming notes or a chord on its input channel. Once notes arrive, its existing rhythm-step highlight and generated-note previews follow the audition. The source's saved enable switch does not need to be on. Both save scopes preserve playback and its indicators when the saved performance is reloaded.

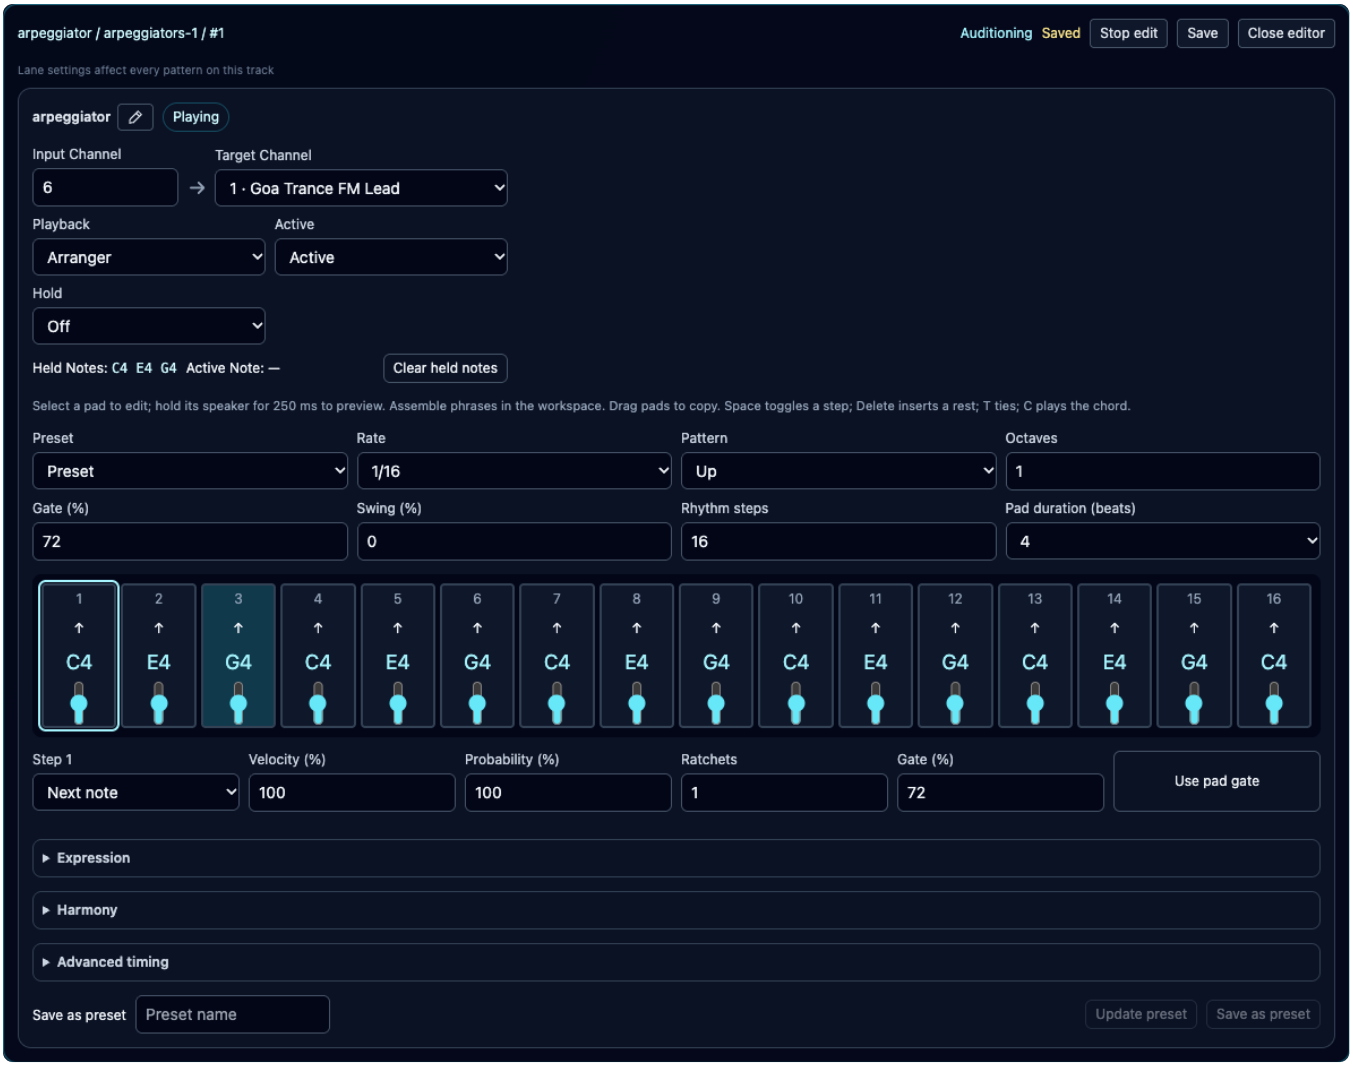

Loop repeats the selected clip's time window across all tracks. With no clip selected, it repeats the whole song from the start through **Song end**, including trailing silence. Click empty timeline space or the ruler outside a clip to deselect it; clicking a clip selects it again. Click **Loop** again to turn it off. This is separate from **Loop** in the clip inspector, which repeats its content only within the fixed clip window. The playhead and Position field follow song position through loops and remain stationary during standalone audition. Newly created clips are selected automatically.

Loading selects a clip even when looping is off: a matching saved loop takes priority, then a clip at the cursor, then the earliest clip. Tracks stay compact and the editor stays closed. Explicitly clearing the selection stays cleared through Play, Stop and Rewind. You can toggle the loop without reopening its editor. See the restored loop selection.

Position accepts a seek destination in master beats. While you edit it, playback continues without replacing your text or moving focus. Press Enter or leave the field to seek once; Escape cancels and resumes displaying the current song position. Empty or negative values are discarded. Loading or importing a performance clears an unfinished position edit.

Song end shows the actual playback boundary in master beats, marked by a dashed amber line. Moving or shortening clips retains trailing silence. To stop at the last clip instead, press **Fit song to clips**, or enter a Song end at or beyond the final clip. This is an undoable arrangement edit, saved with the performance; it does not cut off clips. Position and Song end use master beats, while the ruler uses bar/beat labels.

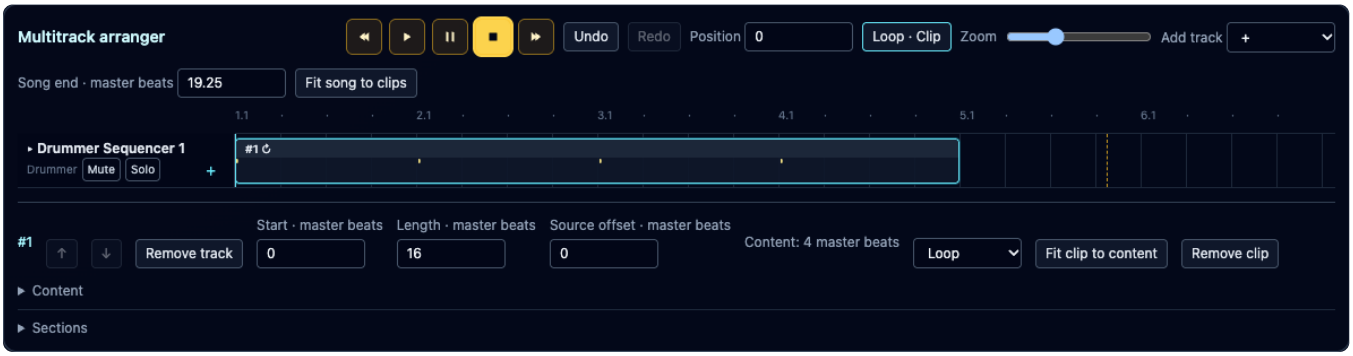

Move **Zoom** all the way left, or focus the slider and press **Home**, to see the entire song through **Song end**, including trailing silence. The overview adjusts when the song length or timeline width changes. Zoom back in for beat markers and detailed clip editing.

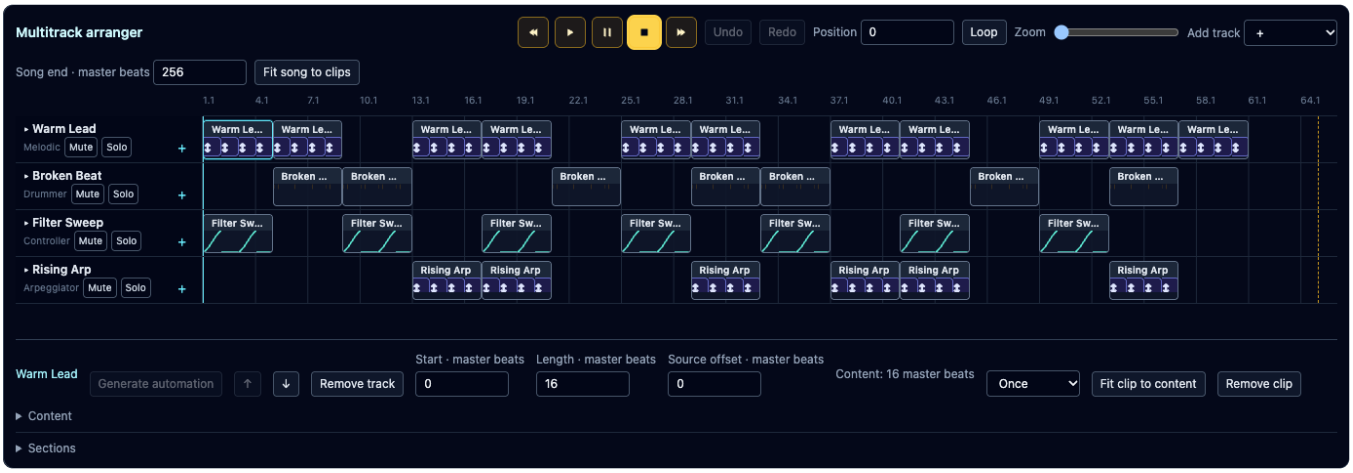

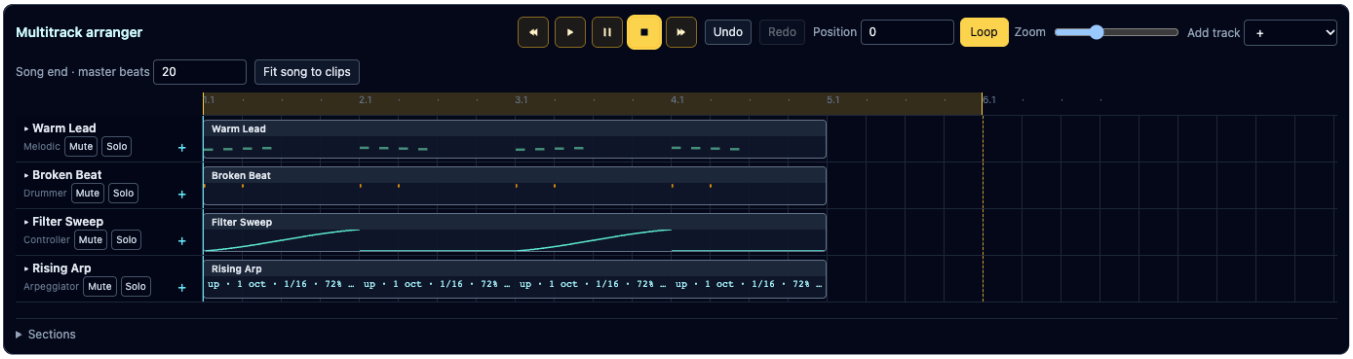

The actual piano keyboard and configurable MIDI controller rows sit below the arranger/editor, each full width. They retain Start/Stop, routing, channel/OMNI, scale/autochord, velocity and CC controls. Their existing Auto/Live/Listening buffer policy still applies. An enabled device does not force low-latency buffering indefinitely. Arpeggiator Arranger/Live clock settings remain separate from clip placement.

Transport buttons wrap onto another row when the viewport or enlarged text leaves too little space. Track headers reserve less than half the viewport, leaving the timeline reachable at enlarged text sizes. Transposition controls wrap inside their dialog. The timeline keeps its own horizontal scrolling.

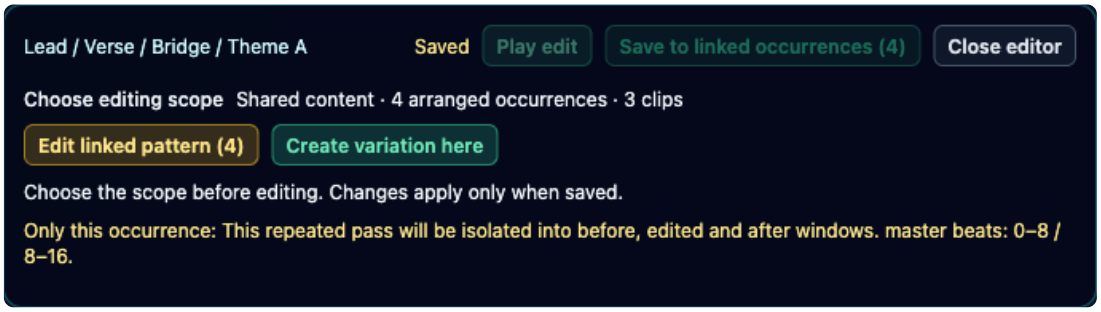

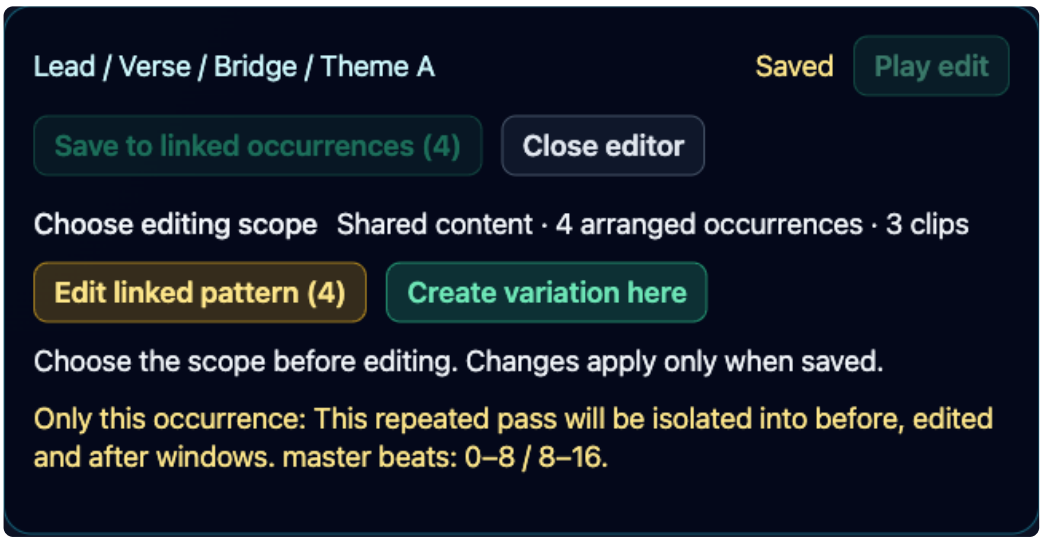

Assemble content directly on the timeline

Expand a track and click a content header to select it. Shift-click another header to select the inclusive chronological range. Shift+Left/Right extends or contracts the range, including across loop repetitions and beyond the visible area. **Group** combines consecutive patterns/rests; **Supergroup** combines consecutive groups. Enter a name in the selection toolbar before grouping. Right-click provides the same grouping actions. Double-click opens the musical/container editor.

Base content uses a stronger fill; generated loop repetitions use a lighter fill and ↻. ↔ marks shared content independently of repetition. Selecting one occurrence gives it a strong outline and its related uses a quiet dashed outline. Related uses are not selected for dragging, deleting or grouping. The summary counts nested occurrences, partial repetitions and off-screen clips. Unused saved references are indicated separately, as are repeating lane cycles. Names and colors do not determine sharing.

Selecting several items disables the single-content editing actions.

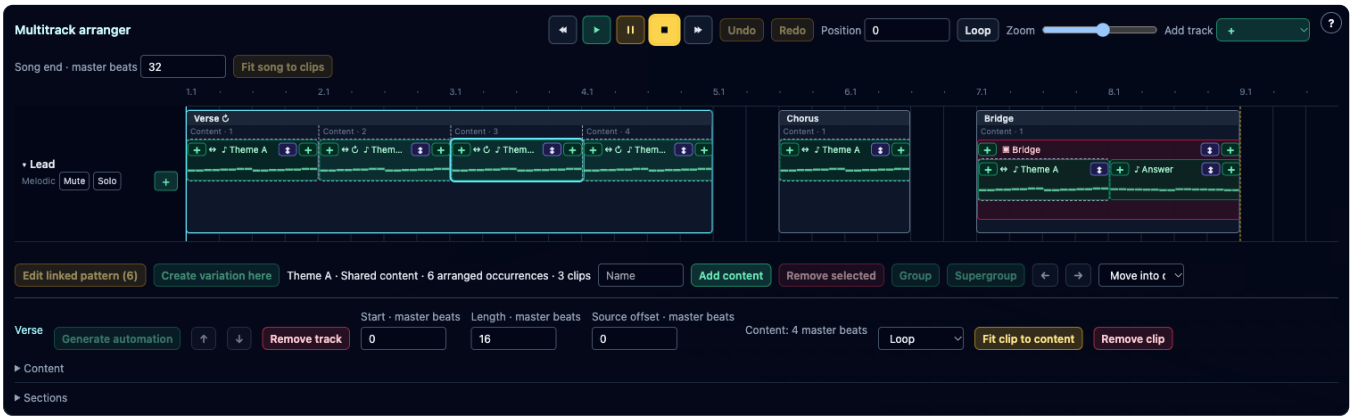

Content · 1, Content · 2, … label loop repetitions. For four repetitions `A A A A`, click the pattern in Content 2, Shift-click the pattern in Content 3, choose **Group**, then **Save structure**. The result is one clip containing `A [A A] A`. The outer clip keeps its name, start and length; the selected occurrences become a group. Its patterns remain linked to their original definitions.

Grouping across repetitions converts the clip to **Once** with explicit ordered content. Complete source content from cropped first/last repetitions is retained, including hidden notes and rests. Timing, neighboring clips and Song end stay unchanged. Repetition-specific transposition remains attached to the original occurrences, including after later edits to shared profiles. This save applies only to this clip; it does not offer an all-clips scope. The conversion and grouping form one draft Undo action and one saved arranger Undo action. Undo restores the compact loop.

Cross-repetition selection applies to content directly inside one clip. Nested selection stays within one container occurrence; selection cannot cross clips or tracks. Command/Control and Option/Alt are not range-selection shortcuts. While a selection spans repetitions, dragging, reordering, removing and ungrouping are disabled. Group it first to obtain ordinary editable content. Existing hierarchy limits still apply. Save/Discard/Cancel resolves an existing structural draft before crossing repetitions; if saving changes the selected occurrences, select the range again. Capacity rejection preserves the selection and existing drafts.

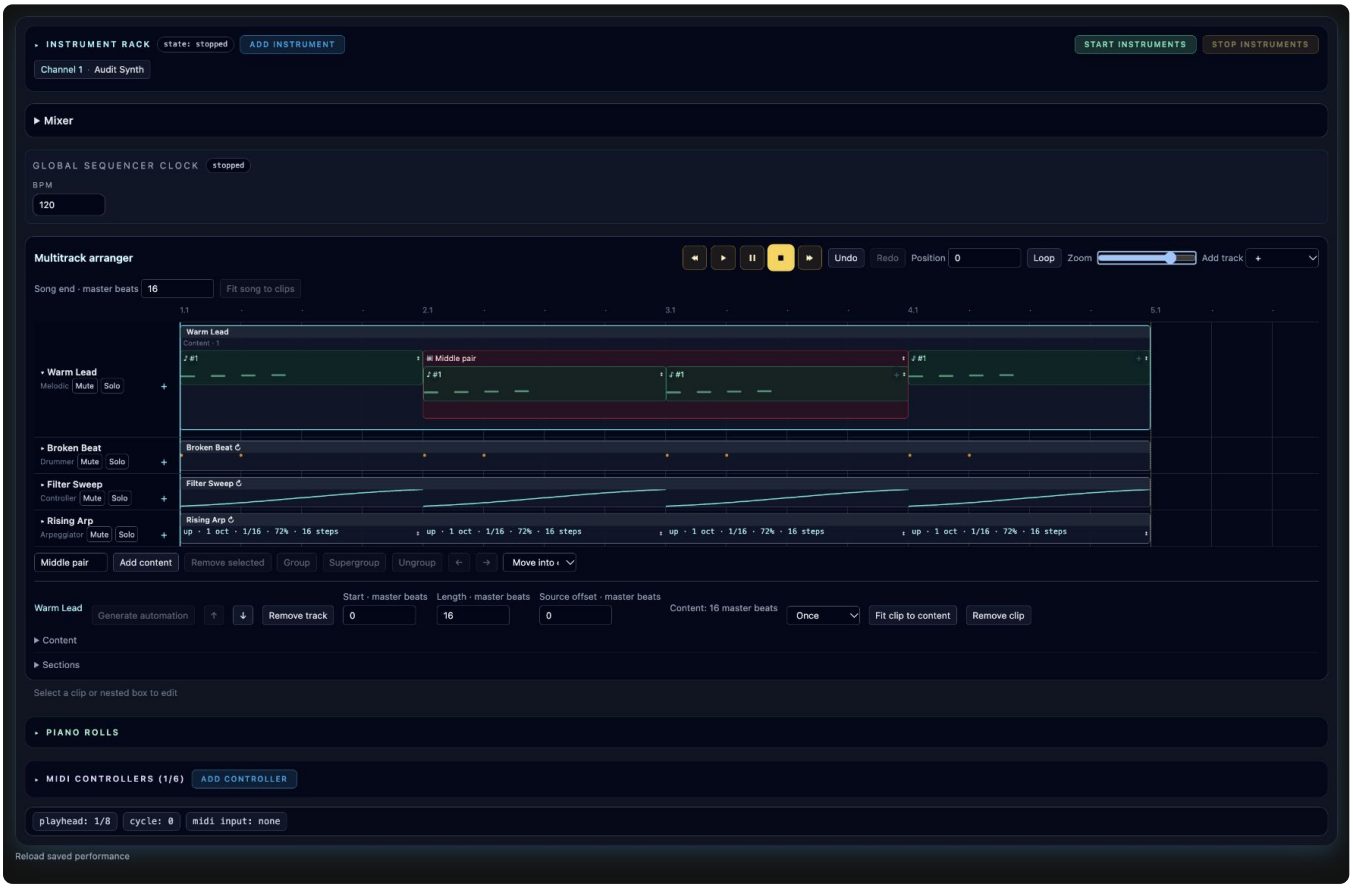

Use **Add content** or the **+** insertion markers before, between or after boxes to search the track's existing definitions. **Adopt linked** retains sharing; **Create variation** gives the inserted definition its own identity. The picker also offers **New pattern** and explicit rests. Unused saved definitions remain available. Double-click an empty area or rest inside a clip to open the same picker at the clicked position, snapped to a quarter master beat. Inside a group or supergroup, the insertion belongs to that container. Existing content stays in place; an insertion that overlaps it is rejected. Content headers still open their editors on double-click, and automation curves retain their point-editing gestures.

Expanded tracks keep the cyan **+** insertion buttons fully visible with 21 × 15 px targets. Collapsed tracks use smaller 14 × 10 px transposition buttons. Content names, indigo ↕ transposition buttons and the final **+** each have their own space in the header. The content boxes retain their dimensions and timing positions.

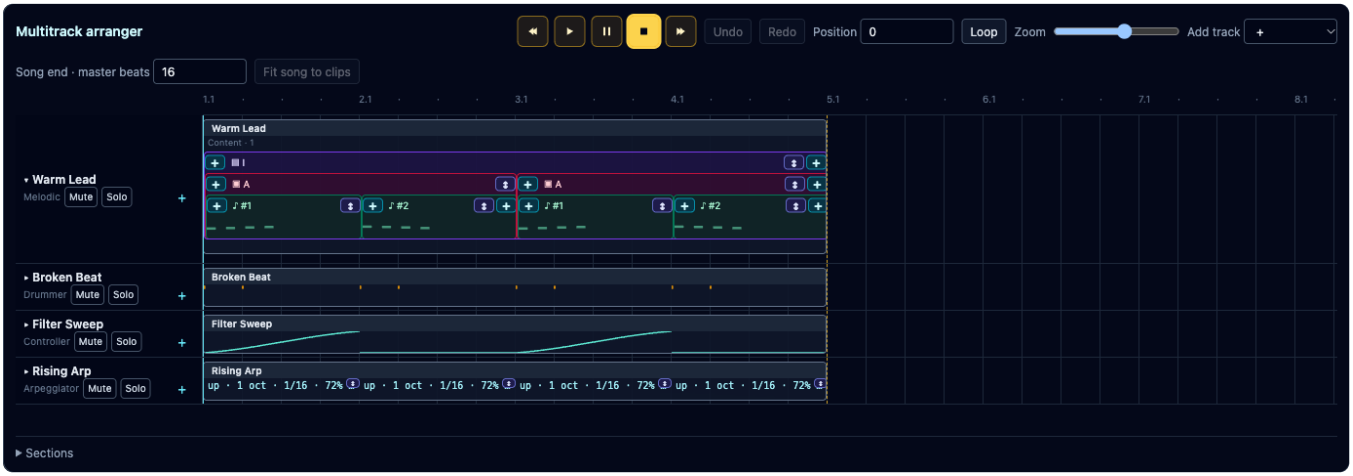

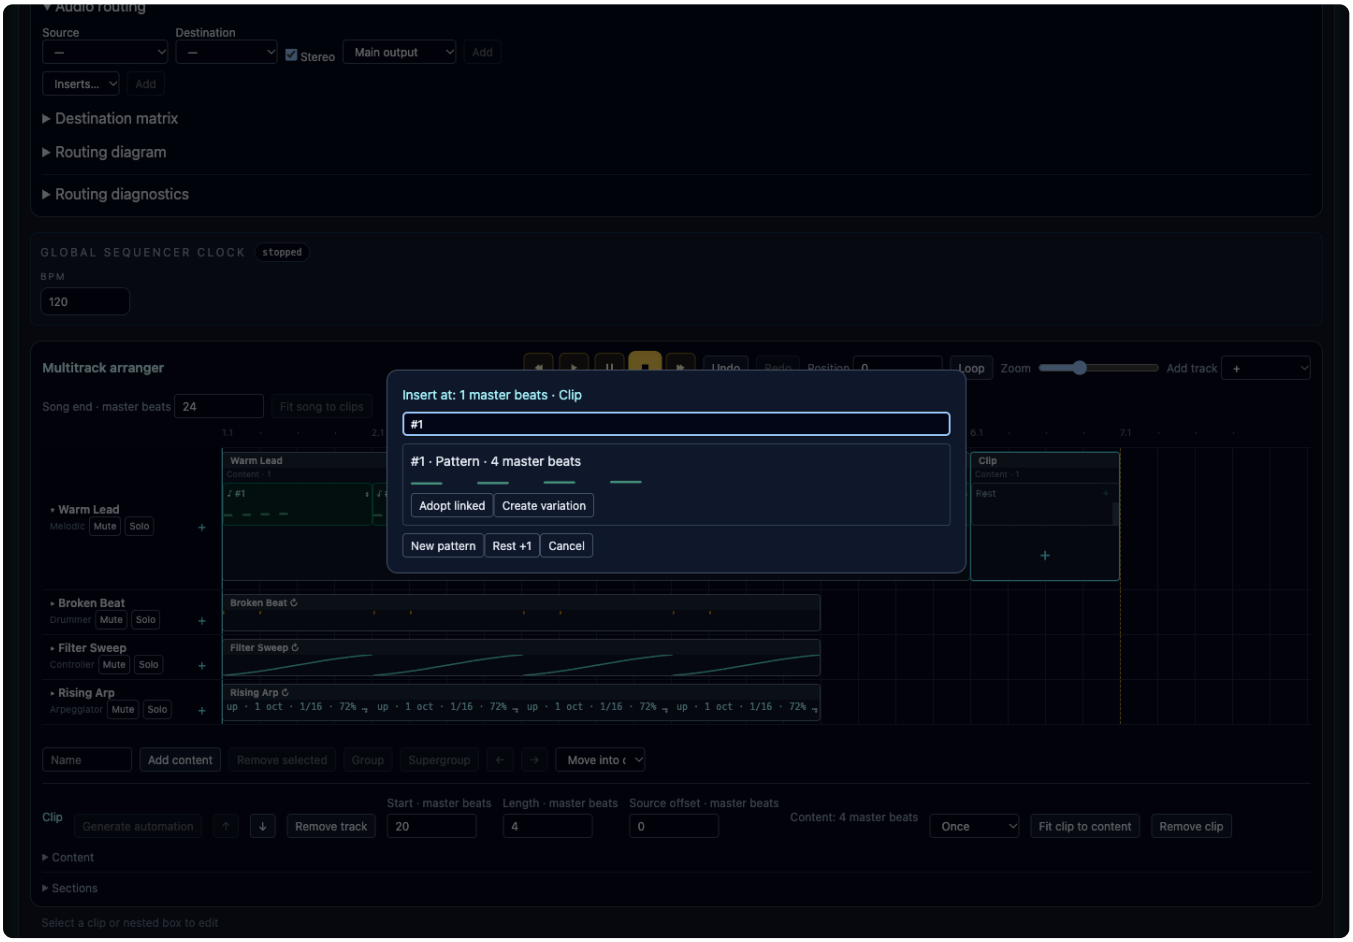

Right-click a pattern, group or supergroup and choose **Duplicate** to append a linked occurrence at the clip root, immediately after its last musical item. Consecutive sibling selections retain their order and internal rests. Trailing silence is used first; unused silence remains. The clip grows only when needed, and neighboring clips stay in place. **Ungroup** appears only when

all selected items are groups or supergroups. Rest-only and cross-repetition selections cannot be duplicated.

For a looped clip, Duplicate appends **once after the final repetition**, converting the original repetitions to explicit content in one Once clip. Complete cropped source content is retained; the duplicate follows the full retained final item. Existing repetition-specific transposition and definition links are preserved. This structural save affects only the selected clip, and one Undo restores the loop. Resolve any pending structure before duplicating a loop, then reselect the content if prompted. Capacity rejection retains the selection; a collision leaves the draft available for correction.

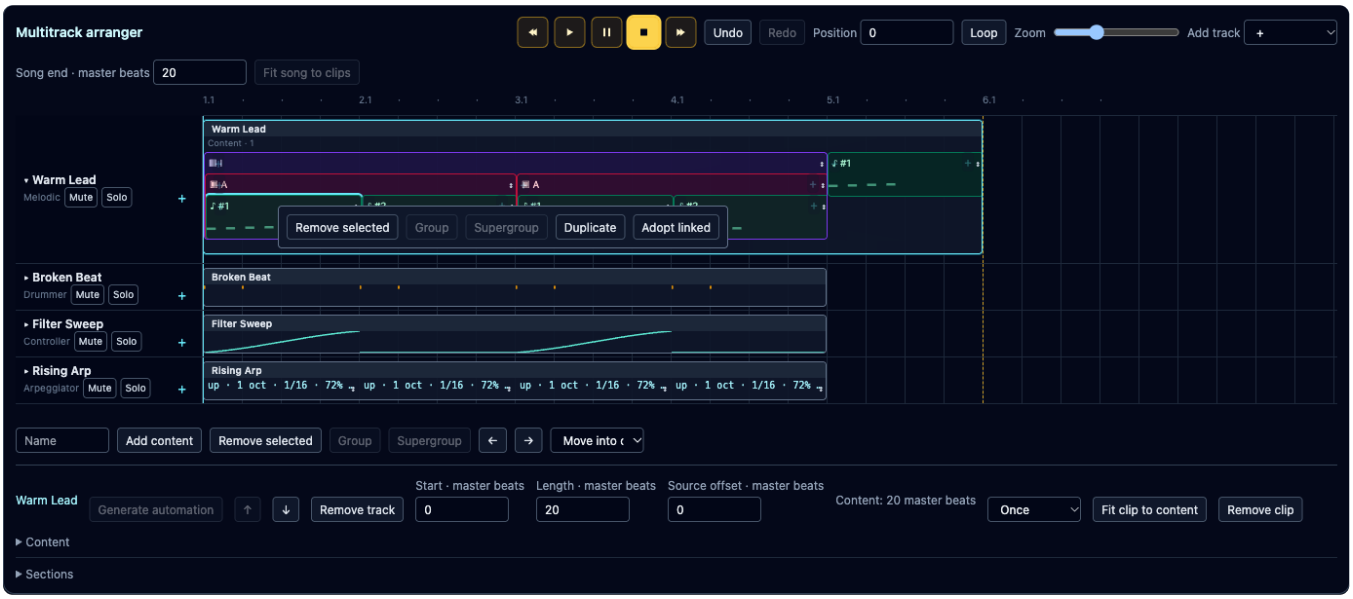

Drag a box or contiguous selection to a quarter-beat position. Moves within a clip/pass leave rests at the source and consume silence at the destination; other content stays in place. Overlapping drops are rejected. Between clips on the same track, dragging copies linked content. Between tracks of the same device type, it copies independent definitions using destination timing. Drum keys must have unambiguous destination rows; controller modes must match. Compact clips accept root drops, and hovering expands the destination track. The destination text shows the position. Self-nesting and moves across loop passes are rejected. Left/Right on a focused content header nudges it by a quarter beat. The toolbar move arrows and **Move into container** retain their explicit reorder behavior. Escape or cancelling a pointer gesture leaves the structure unchanged.

Use **Remove selected**, the matching context-menu action, or Delete/Backspace while a box is focused to remove patterns, rests or containers. The same actions work inside groups and supergroups. Removing content closes its position within the container but does not shrink the clip window; unused pattern definitions remain available. Empty containers disappear. Removing an entire clip’s content leaves a timeline gap; removing one loop pass locally preserves its other passes. The first change to shared structure asks for scope before entering the draft; Save structure then uses that choice.

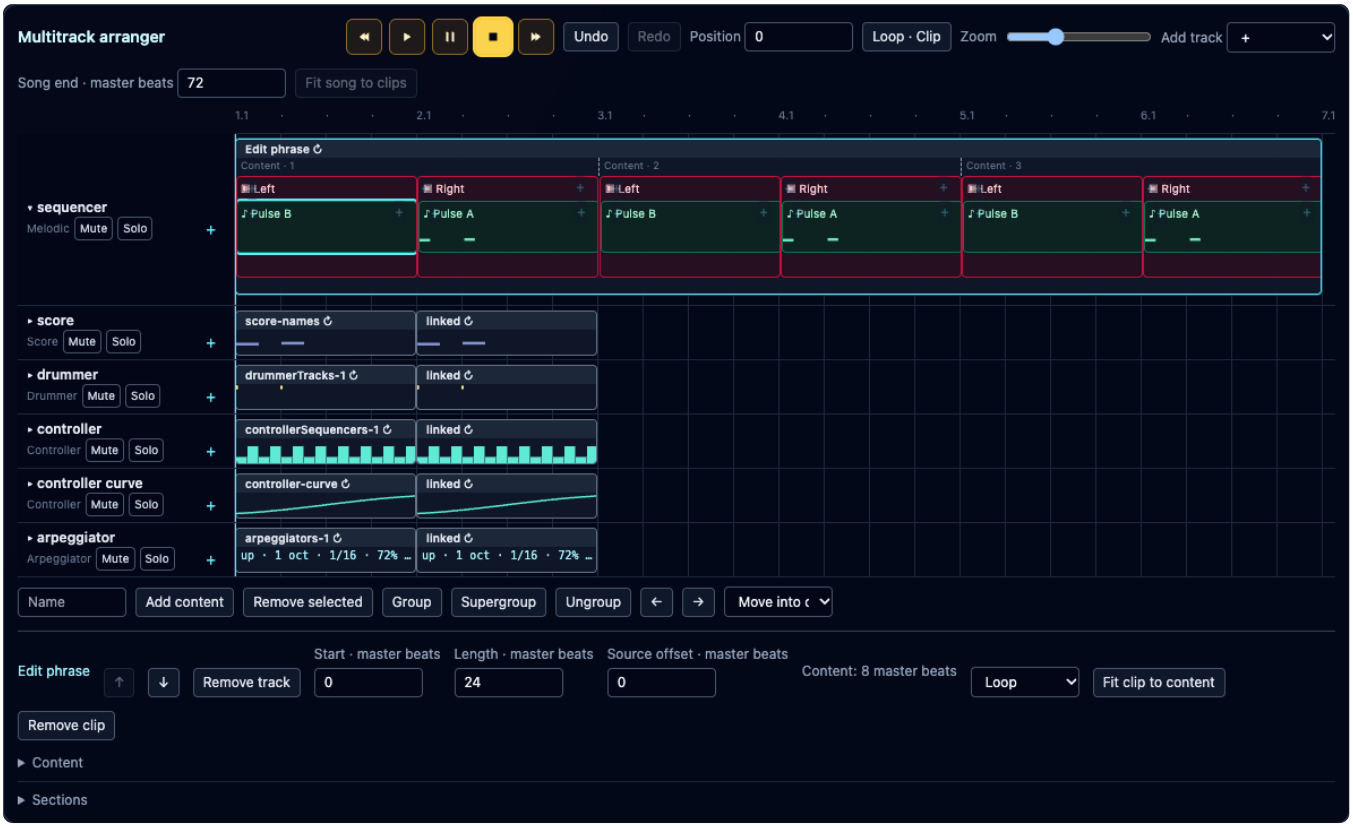

Structural work is a session-only draft: the song continues to play saved content. **Save structure · local** or **Save structure · all linked** applies the established scope atomically; **Discard** abandons it. Draft Undo/Redo is separate from the arranger's saved history. Performance save/export offers the existing saved definitions when unresolved drafts remain; loading another performance requires resolving them.

Longer content grows the affected windows if space permits. Loop windows preserve their full and partial repetition span. Saving all linked occurrences evaluates every affected clip together. A collision rejects the complete save, retains the draft and names the blocking clips; later independent clips are never shifted. Shortening content does not automatically shrink clips. Valid growth extends Song end when necessary.

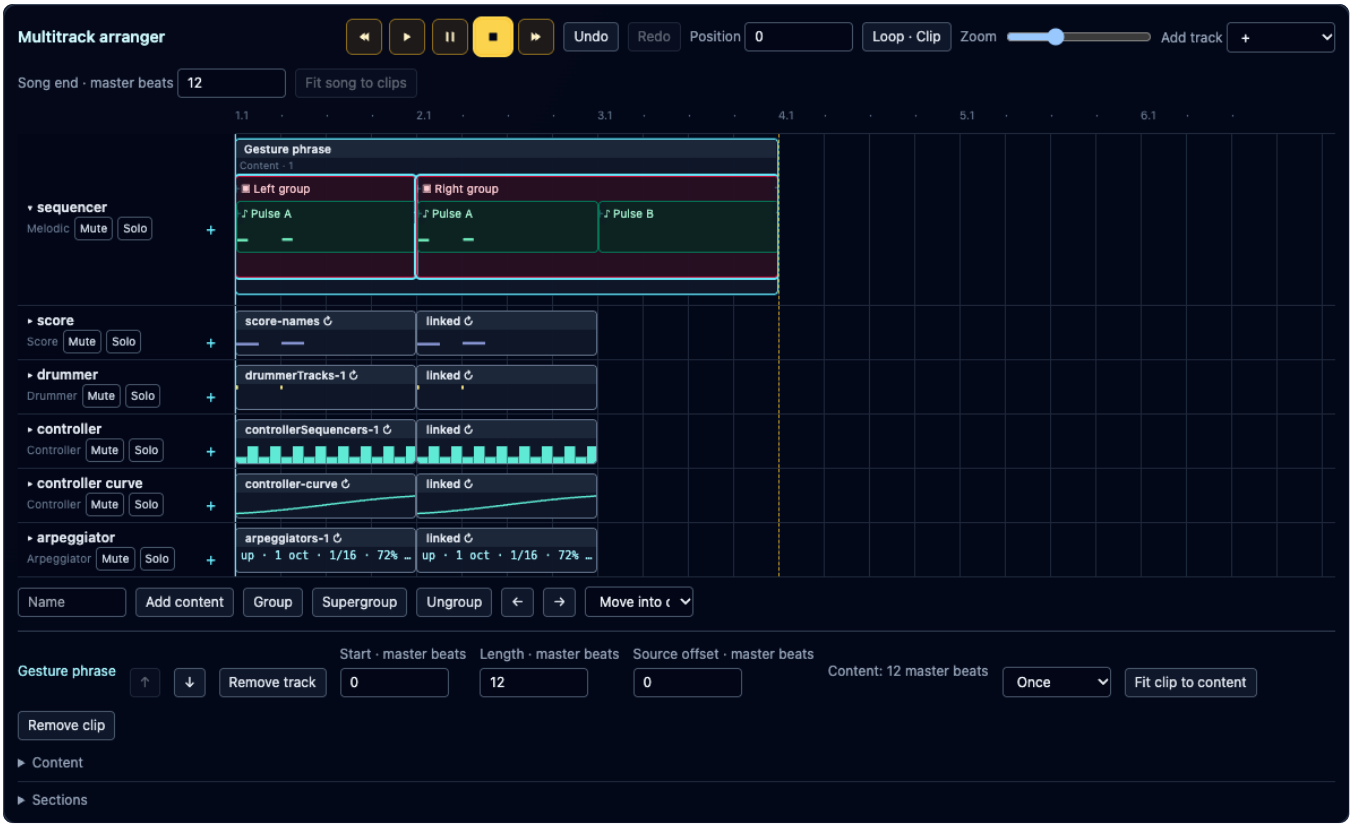

Container inspector and musical draft conflicts

Double-click a group or supergroup to open its child inspector. Adopt linked content, add a new pattern or rest, reorder, group, ungroup, rename or remove children there. These controls edit the **same structural draft** as the expanded timeline. Choose the scope before editing shared containers. The editor shows the affected definitions and clip boundaries, including collision rejection and empty-container removal. Save structure uses the established choice. Draft Undo/Redo and Discard apply to both surfaces.

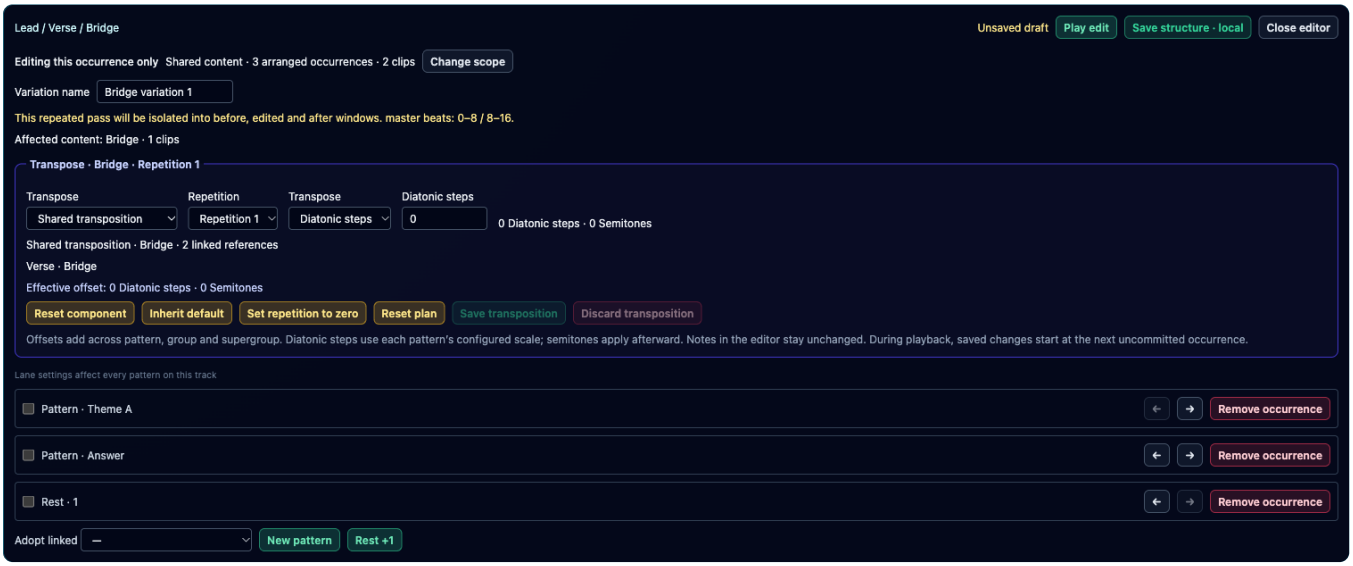

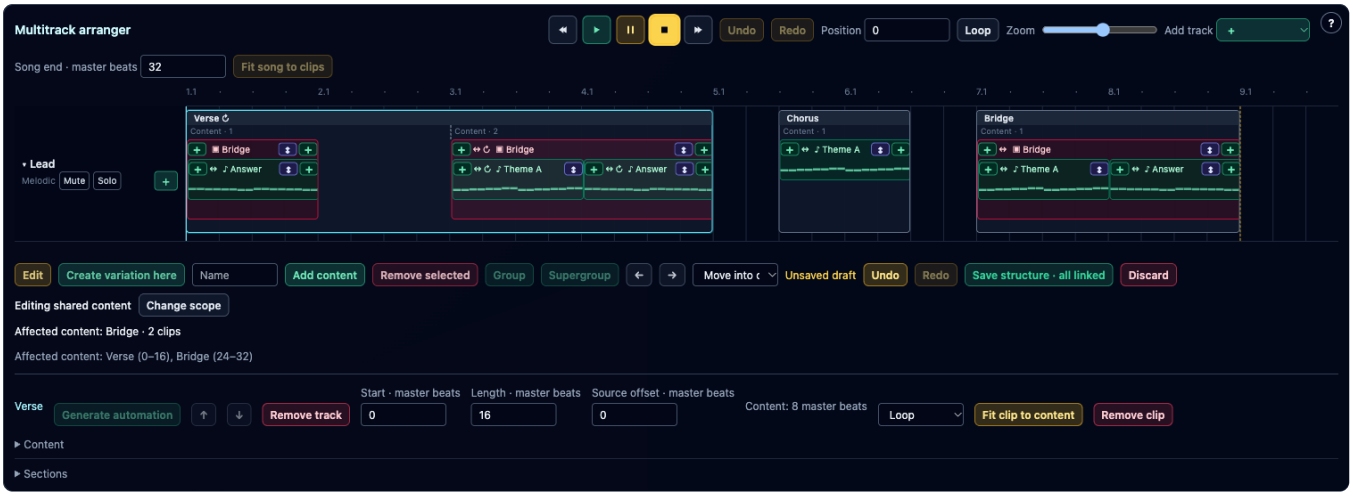

If the pending action intersects an unsaved musical draft, choose Save, Discard or Cancel. Save uses that pattern's established editing scope, then resumes the pending edit. Discard removes only the conflicting draft and resumes; Cancel keeps both drafts and leaves the arrangement unchanged. Other musical drafts remain available. If several drafts conflict, resolve them individually. A failed musical save keeps the draft and displays its error inside the dialog.

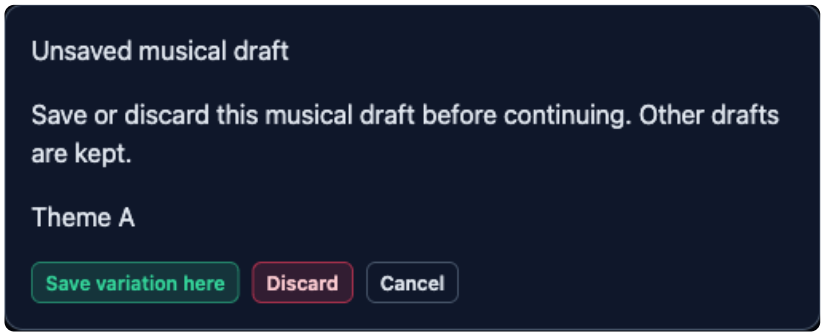

When a musical save forks shared parents or splits a loop pass, the pending structural edit follows the new occurrence paths and retains draft Undo/Redo. Following passes created from that same source during resolution still move when its edited portion grows. Independent clips stay in place. If resolving a conflict makes the pending structural save valid, it resumes using the established scope; otherwise the draft remains available with an error.

Container **Play edit** auditions a snapshot of the current structure. Further structural edits remain visual and do not re-prepare audio; stop and restart Play edit to hear the changed structure. Saving ends the editor audition through the existing lifecycle. Musical sequencer controls retain their existing live-edit behavior.

Transpose clips without changing notes

Click the transposition badge on a pattern, group or supergroup in the expanded arranger lane. This opens**Transpose in arranger** for that occurrence without opening or replacing the musical editor. Compact lanes also offer badges on their pattern previews; expand the lane to edit a group or supergroup directly. The existing **Transpose** controls in the detail editor remain available.

Transposition buttons always show a larger ↕ symbol with an indigo outline and background. Active buttons also show the effective signed offset when space permits: **D** means diatonic steps and **C** means chromatic semitones. Enclosing offsets are included, so each pattern shows its complete total. Hover or focus a button to identify the occurrence and its effective offset. Narrow active buttons retain their brighter indigo highlight. Click any button to add or change a transposition, including on a generated repetition.

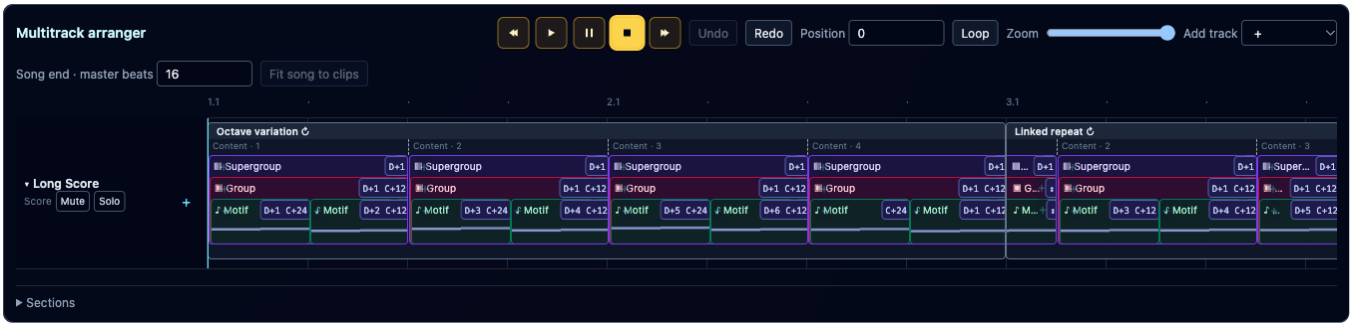

Choose **Diatonic steps** to move through each pattern's configured scale and mode, or **Semitones** for chromatic movement. Negative values move down. Switching the displayed component preserves the other value. All pattern, group and supergroup contributions add; diatonic steps apply first, then semitones. Notes outside MIDI 0–127 after the complete transform are silent. The musical editor continues to show the original notes; timeline Note maps show the sounding pitches.

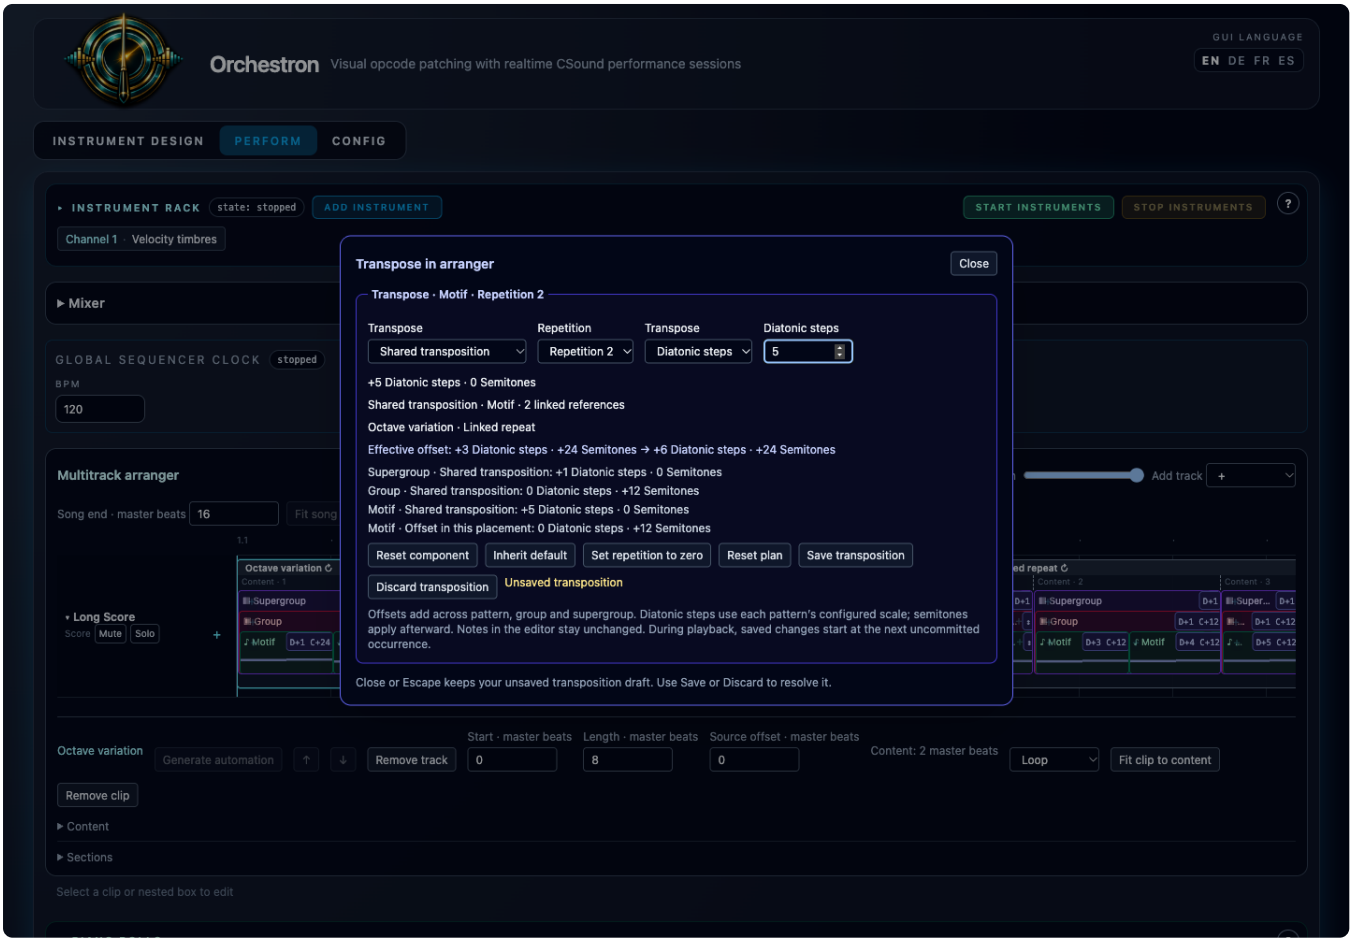

Choose the owner before changing a value:

- **Shared transposition** changes the selected definition everywhere it is used, including other clips. Its repetition settings are shared too.
- **Offset in parent** belongs to this reference inside its parent. Reusing that parent also reuses this offset.
- **Offset in this placement** belongs only to this nested location in this clip. It adds to the shared and parent settings and does not create a new pattern or split the clip.

The contribution list shows the selected item and its enclosing offsets. The inspector lists affected linked clips, compares saved and draft totals, and reports selected-pattern notes that fall outside the MIDI range. Each leaf pattern uses its own scale and mode, including when a group contains patterns with different scales.

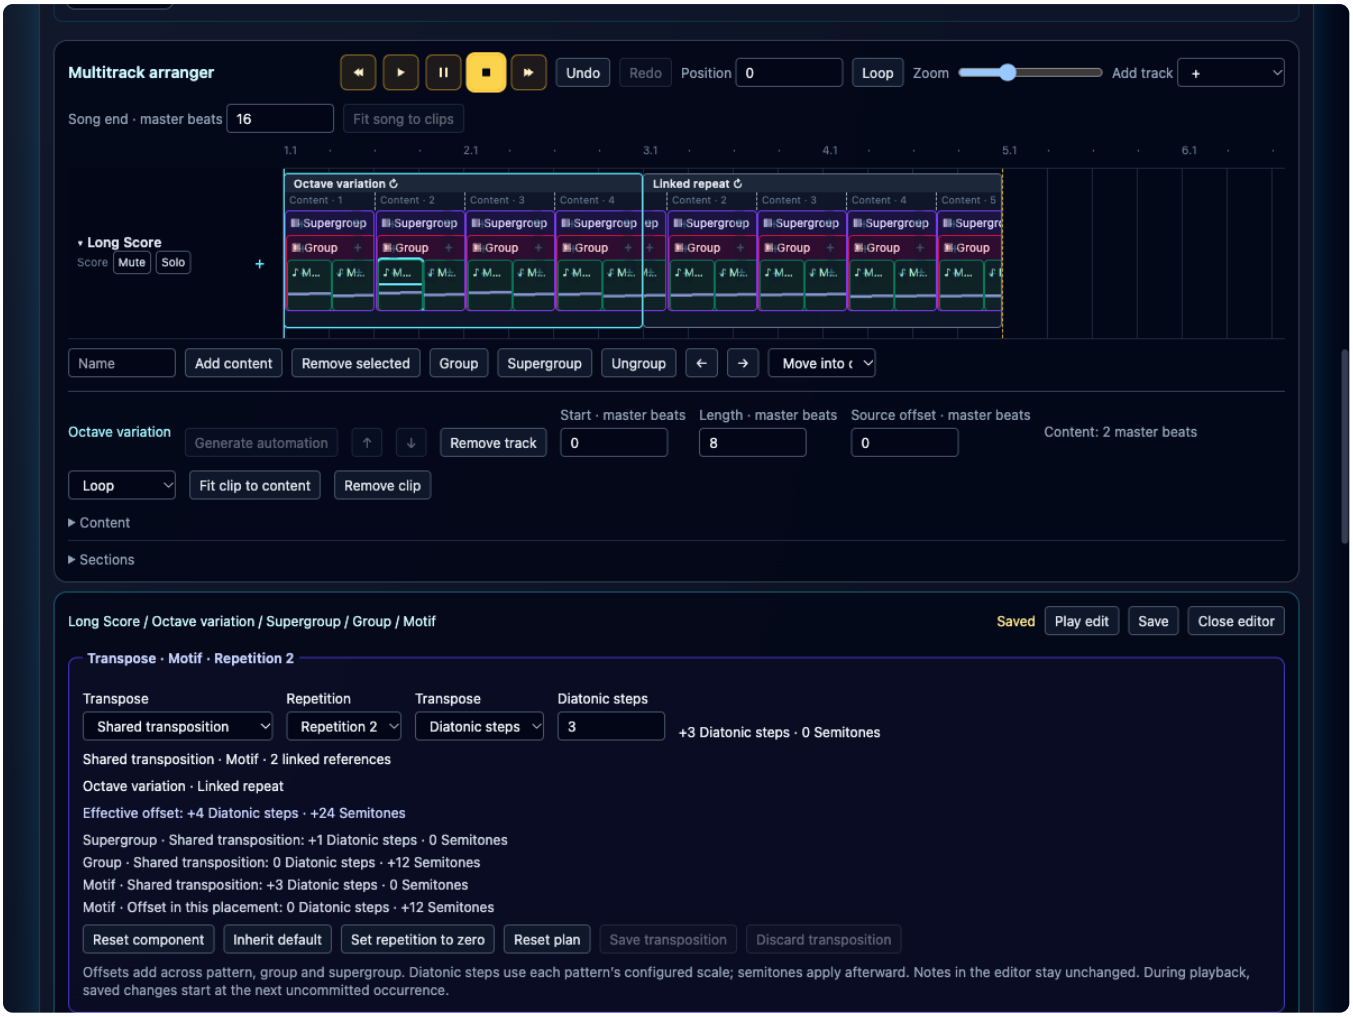

Choose **Default** for the plan's base value, or select a generated repetition in the timeline and edit that **Repetition**. Repetition numbers follow the source content: a cropped clip can begin partway through repetition 2. They do not count how often transport has played. Descendants normally inherit their enclosing reference's repetition number. Grouping across repetitions retains each child's original number: a new group uses the first selected occurrence's number, while its children can retain different numbers. Each enclosing offset is evaluated at its own reference's number.

An explicit repetition replaces its own plan's default; other plans still add. **Set repetition to zero** stores an explicit zero pair. **Inherit default** removes that repetition entry. **Reset component** clears only the displayed component; **Reset plan** clears the default and every repetition entry. Shortening a clip keeps hidden repetition settings for later extension.

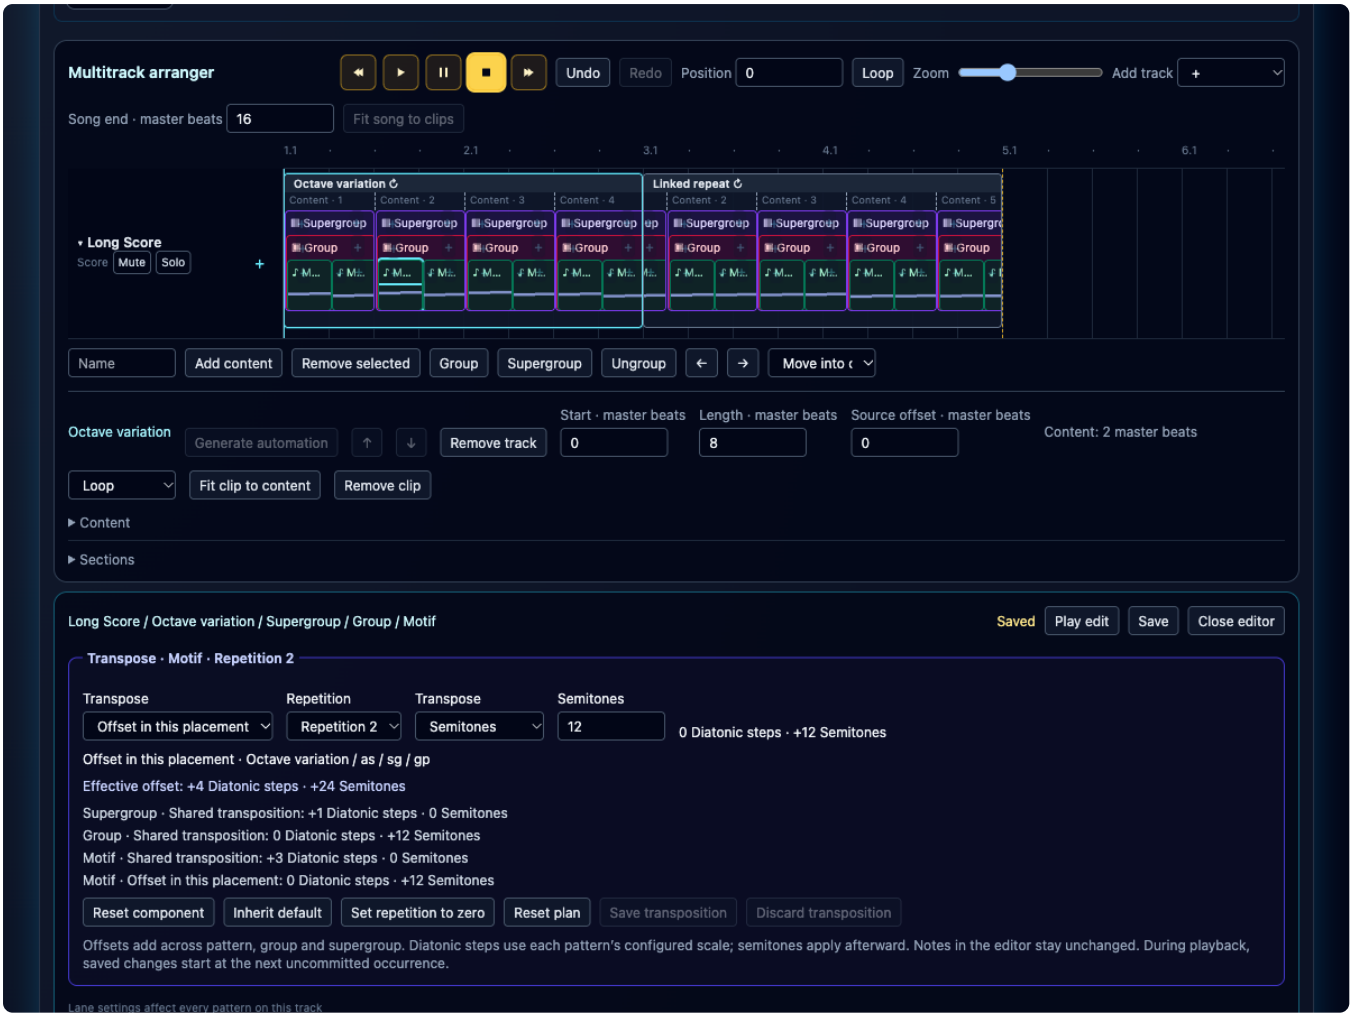

Save transposition commits the draft as one undoable edit. **Discard transposition** leaves saved music unchanged. The pattern editor's **Save** can commit musical and transpose changes together. Shared transposition retains its explicitly selected scope even when musical notes are saved as a local variation; that new variation receives its own copied profile. Conflicting edits report an error instead of overwriting another change.

Saving from the lane closes the dialog and updates affected badges and note previews, including linked clips. **Close**, **Escape**, or clicking the backdrop keeps the unsaved draft; its badge turns amber and gains an asterisk. Reopen it to **Save transposition** or **Discard transposition**. Badge totals continue to show saved settings until Save. Keyboard users can focus a badge and press **Enter** or **Space**; closing the dialog returns focus to that badge. Existing musical drafts are preserved.

Play edit auditions the full selected definition from its own start, using the selected repetition and all enclosing offsets. It retains each reference's repetition number while auditioning; a group assembled across repetitions therefore preserves its children's distinct transpositions. Transposition and structural drafts are snapshots for audition. Saving during arrangement playback changes the next occurrence whose notes have not started; an early attack can already belong to the next repetition. A pending indicator names the target clip, repetition and leaf position. Seek or restart uses the latest saved settings immediately.

Melodic, Score and Arpeggiator lanes support these controls. For arpeggiators, the new mapping follows their existing transpose and quantizer, and uses the pattern's configured scale. Bypass remains unchanged. Drummer, controller and instrument-automation lanes do not offer pitch transposition.

Use a current Orchestron version when sharing performances with transposition. Native bundles and both CSD exports retain saved transposition settings; unsaved drafts and auditions are not exported.

Resize patterns on the track

In an expanded track, drag the lower left or right border of a pattern. The box follows your drag, and**Length** shows the new total duration in master beats. The opposite edge stays fixed; length follows the pattern editor's supported beat increments. Notes remain relative to the pattern start, and shortening retains hidden notes. Growing consumes adjacent silence; shortening inserts rests. Occupied time cannot be overwritten. Focus a pattern border and use Left/Right for keyboard resizing. Escape cancels a gesture. Clip borders continue to trim the outer playback window.

A trailing rest has a visible right-edge handle. Drag it or focus it and press Left/Right to change its duration in quarter-master-beat steps. The handle stays reachable at the visible clip edge when the rest extends beyond it; the next resize starts there. Shortening a nested trailing rest also shortens its enclosing containers when they end there. Other content keeps its position. A rest can be shortened to zero when other content remains; an otherwise empty clip keeps a positive duration.

Trimming a clip changes only its playback window and preserves hidden patterns, notes and rests. To remove a silent tail from the content, shorten or delete the trailing rest and **Save structure**. Then use **Fit clip to content** to shorten the outer window too. Moving or trimming a clip keeps any pending structural edits and their Undo/Redo history. A genuine change to the underlying content still requires resolving the draft. For repeated clips, use **Only this occurrence** when changing the rest duration would move later repetitions.

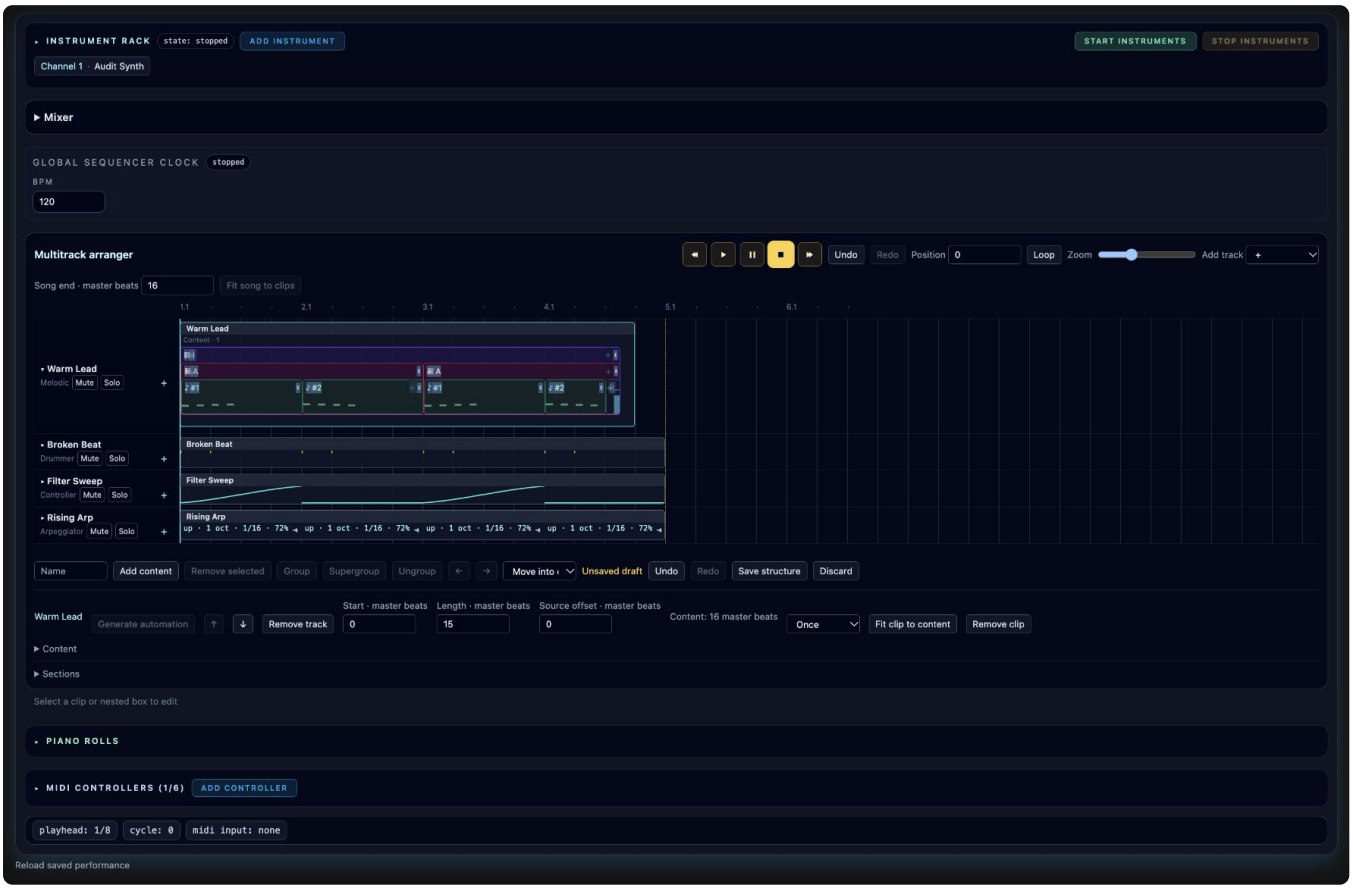

Placement, copying and pattern resizing share the structural draft and Save/Discard workflow. A first direct action that changes shared definitions is held for **Only this occurrence** or **All linked definitions**. Cancel or Escape leaves the existing draft and selection unchanged. After choosing, the action enters the draft once; subsequent actions reuse the scope. **Change scope** changes the whole draft without losing edits. If a later action requires local scope, choose to switch the entire draft or cancel that action. Root-only changes, cross-repetition grouping and loop duplication are local. All linked editing validates every affected use; collisions retain the draft.

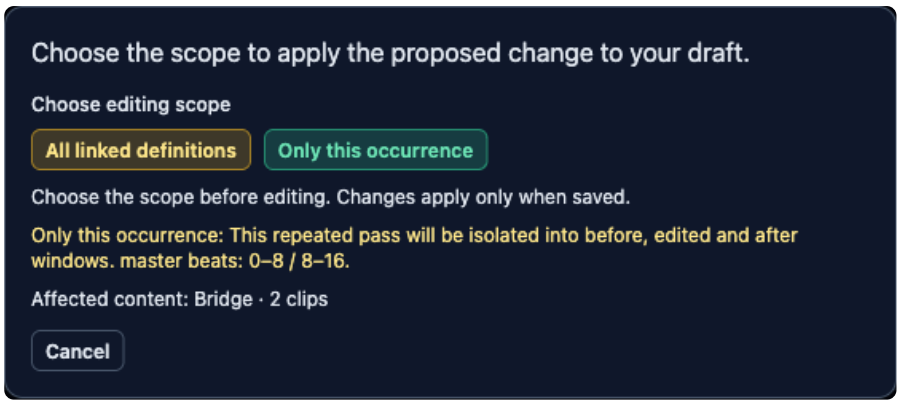

New empty clips commit immediately and support arranger Undo.

A reserved strip below the last track keeps the horizontal scrollbar clear of collapsed clips.

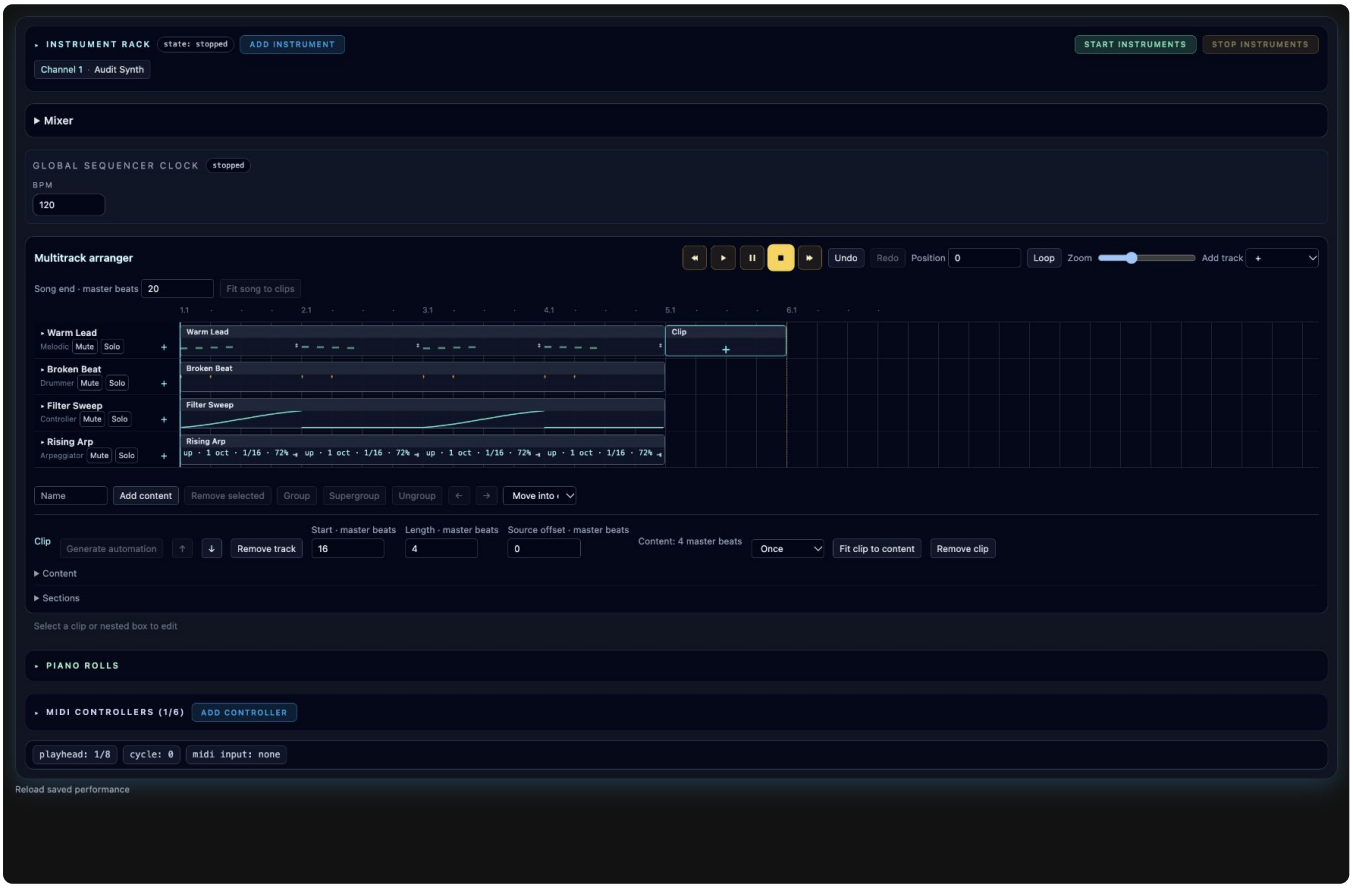

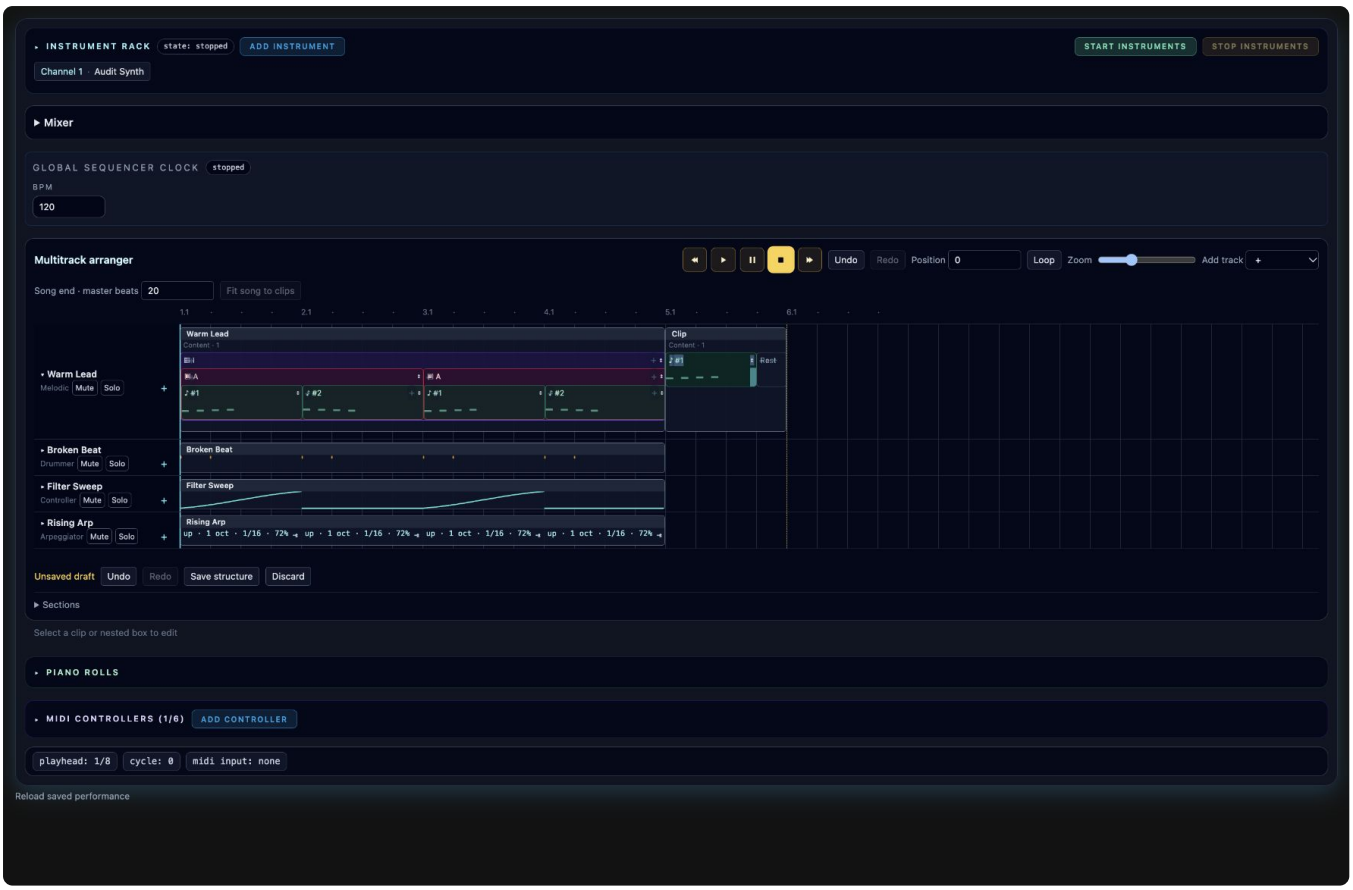
INSTRUMENT RACK
state: stopped
ADD INSTRUMENT
START INSTRUMENTS
STOP INSTRUMENTS
Channel 1 · Audit Synth
Mixer
GLOBAL SEQUENCER CLOCK
stopped
BPM
120
Multitrack arranger
Undo
Redo
Position
0
Loop
Zoom
Add track
Song end · master beats
20
Fit song to clips
Warm Lead
Broken Beat
Filter Sweep
Rising Arp
Clip
Unsaved draft
Undo
Redo
Save structure
Discard
Sections
Select a clip or nested box to edit
PIANO ROLLS
MIDI CONTROLLERS (1/6)
ADD CONTROLLER
playhead: 1/8
cycle: 0
midi input: none
Reload saved performance

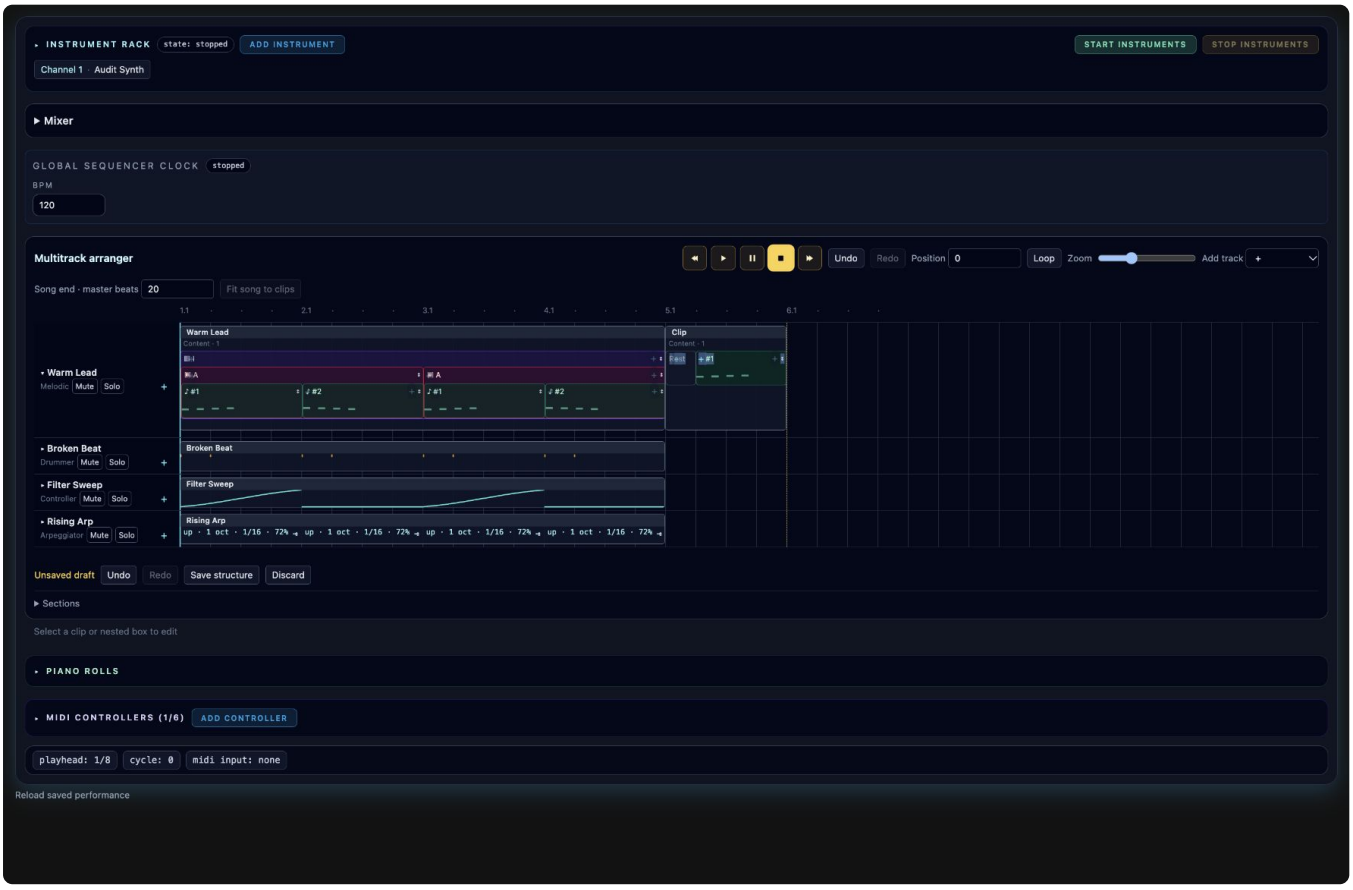
INSTRUMENT RACK
state: stopped
ADD INSTRUMENT
START INSTRUMENTS
STOP INSTRUMENTS
Channel 1 · Audit Synth
Mixer
GLOBAL SEQUENCER CLOCK
stopped
BPM
120
Multitrack arranger
Undo
Redo
Position
0
Loop
Zoom
Add track
Song end · master beats
20
Fit song to clips
Warm Lead
Broken Beat
Filter Sweep
Rising Arp
Clip
Unsaved draft
Undo
Redo
Save structure
Discard
Sections
Select a clip or nested box to edit
PIANO ROLLS
MIDI CONTROLLERS (1/6)
ADD CONTROLLER
playhead: 1/8
cycle: 0
midi input: none
Reload saved performance

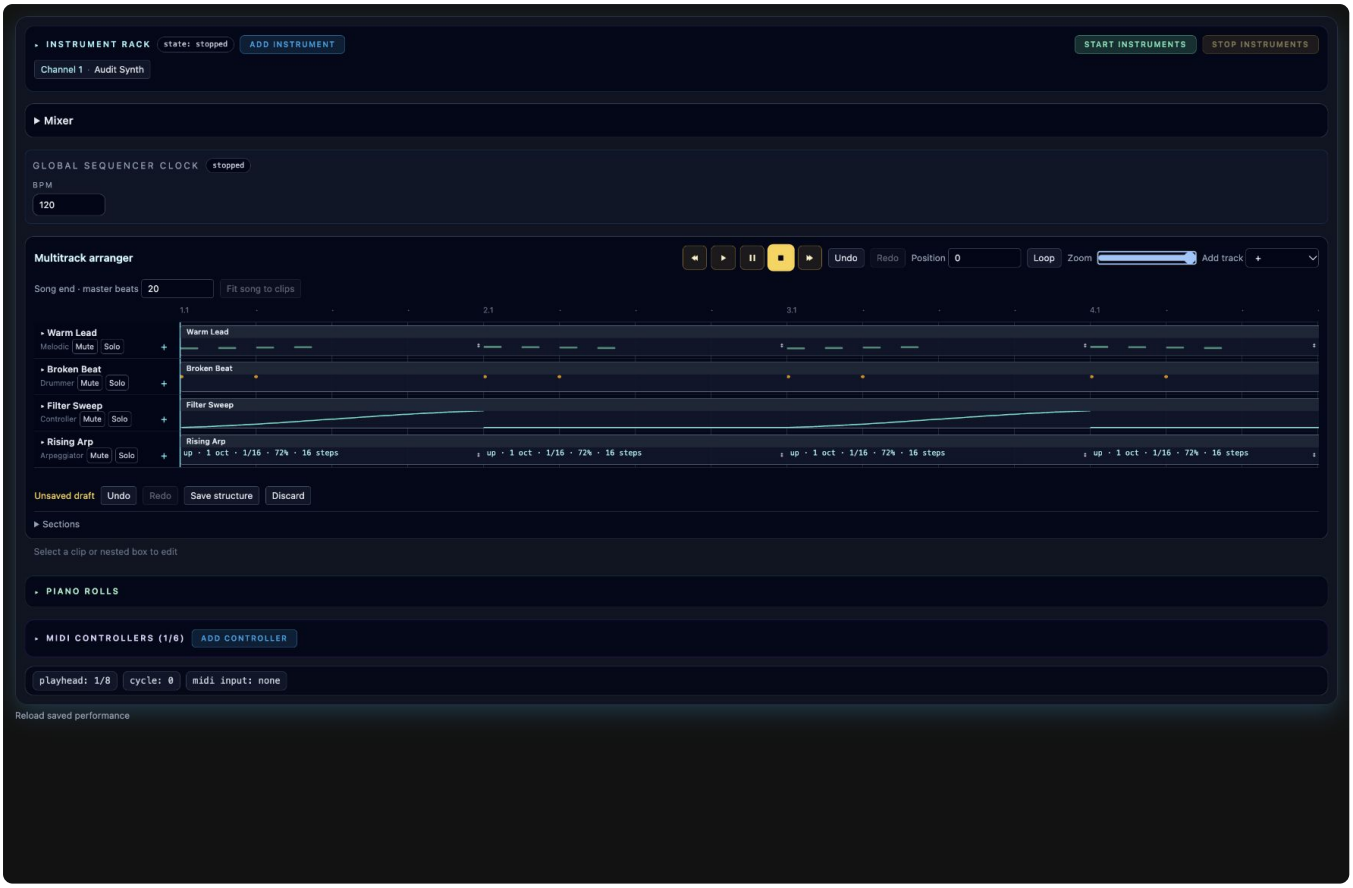

Navigation: Up | Previous | Next

Patterns and Reusable Phrases

Navigation: Up | Previous | Next

A pattern holds the musical content for one track. A clip places that content on the song timeline. The same pattern can appear in several clips, so you can reuse a phrase and choose whether an edit affects one occurrence or all linked occurrences.

Create and reuse patterns

1. Choose **Add track** in the multitrack arranger.
2. Edit the first pattern that opens with the new track.
3. For another clip, double-click empty track space or use the track's **+**.
4. Open the clip's **+** or **Add content** and choose **New pattern**, then double-click the pattern to edit it.

Use **Adopt linked** in the content picker to reuse a saved definition. **Create variation** inserts an independent copy. Removing an occurrence from the timeline keeps its definition available for reuse.

Melodic patterns hold notes, chords, holds, strums and timing offsets. Drummer patterns hold hits, velocities, ratchets and ramps for each row. Score patterns hold notes with independent starts and durations. Controller patterns retain curves and step values. Arpeggiator patterns hold rhythm, expression, harmony and random-seed settings.

Pattern duration and clip duration are independent. Trimming a clip hides content without deleting it; **Fit clip to content** explicitly matches the window to the content. Meter, subdivision and playback speed belong to the track, while tempo is shared across the performance.

Groups and supergroups

Expand a track to see the patterns and rests inside its clips. Click one content header, then Shift-click another or use Shift+Left/Right to select adjacent content. **Group** combines patterns and rests; **Supergroup** combines groups. Name the phrase before grouping.

Groups and supergroups are reusable definitions. **Ungroup** opens one level and retains the definition. Use the content picker to insert saved phrases into another clip on the track. See Assemble content directly on the timeline for selection across repetitions, rests, movement and duplication.

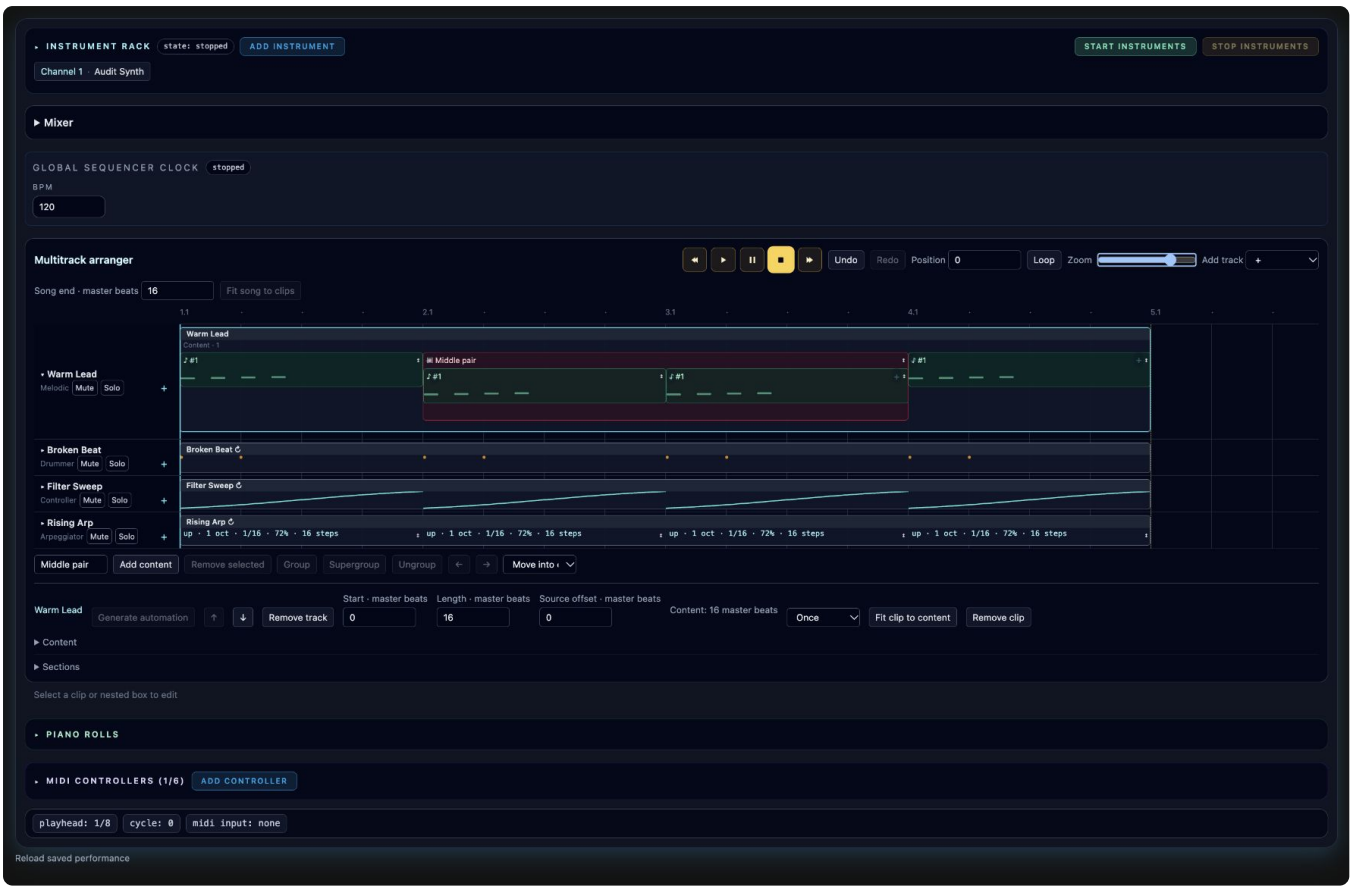

Edit and save

Double-click a pattern, group or supergroup to edit it. The breadcrumb identifies the selected occurrence. Musical edits and structural edits have separate drafts.

Edit linked pattern/group/supergroup and **Create variation here** choose scope before editing shared content. Independent content opens directly. The scope remains visible and can be changed while retaining edits; Save applies it without another question. The first structural action that changes shared definitions also waits for a scope choice. Cancel keeps the draft and selection. Save/Discard/Cancel resolves conflicting drafts before another operation proceeds.

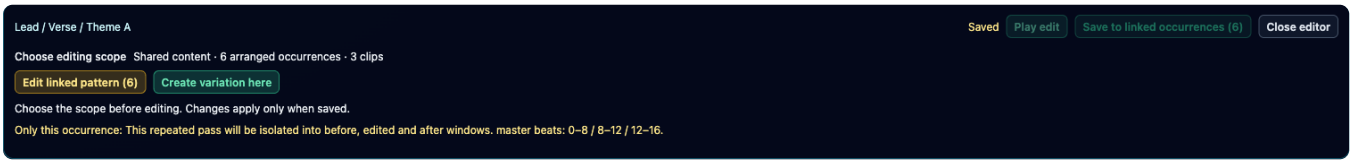

Performance save, autosave and exports use saved definitions. Resolve changed drafts before saving the performance, or explicitly choose the saved definitions. Drafts survive view switches, but a browser reload does not restore them.

Audition a phrase

Use **Play edit** to hear a draft and **Stop edit** to return to the song. Audition does not save your work. SeePlay and jam for how previews interact with song playback, looping and arpeggiator input.

Transpose reusable content

Use the arranger's transposition controls to change pitch without rewriting notes. Shared repetition settings and local placement offsets can combine. Review the affected occurrences before saving. See Transpose clips without changing notes.

Navigation: Up | Previous | Next

Sequencer Timing

Navigation: Up | Previous | Next

Use this guide when tracks need different meters, triplet grids or playback speeds while sharing one song tempo.

Meter, pattern length and subdivision

For a one-bar triplet bass alongside ordinary drums, choose **Meter 4/4**, **Pattern length 1 bar · 4 beats**, and **Subdivision Eighth-note triplets · 3 per beat**. The bass shows **12 steps in four labeled beat groups**. Leave Playback speed at 1×. At 120 BPM both one-bar patterns last two seconds.

Pattern length applies to the selected pad: any integer from 1–16 local beats for melodic/drum pads, or 1–32 for controller pads. Whole bars show both bars and beats; partial bars show beats. Meter, subdivision and playback speed apply to all pads of that sequencer. Subdivision offers 1, 2, 3, 4, 6 or 8 steps per meter beat. Musical names follow the denominator: three steps per beat means eighth-note triplets in /4 and sixteenth-note triplets in /8.

Global BPM counts quarter notes. A local beat is a quarter note in /4 and an eighth note in /8: one bar of 6/8 lasts three song beats (1.5 seconds at 120 BPM). Changing the meter numerator changes grouping without resizing pads. Changing the denominator retains the beat count and steps, but changes elapsed duration. Compound-meter accent grouping is not applied.

Changing subdivision keeps numbered step positions, notes, holds, velocities and timing offsets. It changes the number of visible steps within the same pattern duration; it does not redistribute the music. Hidden steps remain stored and return when you expand the pattern. Controller keypoints keep their normalized positions. Each pad has a 128-step limit; choices that would exceed it on any affected pad are disabled with an explanation.

Advanced timing → Playback speed contains the existing ratios, with multipliers: 1:1 = 1×, 3:2 = 1.5×, 2:1 = 2×. A speed other than 1× stays visible when Advanced timing is closed. The summary shows step count, subdivision and elapsed quarter-note song beats, separate from the pad's local length. Bar boundaries are stronger than beat boundaries, including incomplete final bars; no silent cells are added. Controller guides use bar.beat labels.

Timing edits prepare while playback continues. If preparation fails, the current music keeps playing and the draft remains available. Older /8 patterns are converted to eighth-note beats when opened, preserving their duration and step positions; the displayed beat count can therefore change.

Navigation: Up | Previous | Next

Melodic Sequencers and Step Editing

Navigation: Up | Previous | Next

Melodic sequencers are step-based pattern sequencers used for note playback.

This page covers the **melodic sequencer** editor. Drum-machine style programming is documented in Drummer Sequencers.

Use the pen beside the device name to rename it. See Device Names for editing controls, validation, and import/export behavior.

See Sequencer timing for triplets, beat grouping and editing behavior.

Open and close the editor

Double-click a pattern in the arranger to open its editor. **Close editor** returns to the timeline and retains its draft. Closing the editor ends **Play edit** auditioning; the song's transport remains independent.

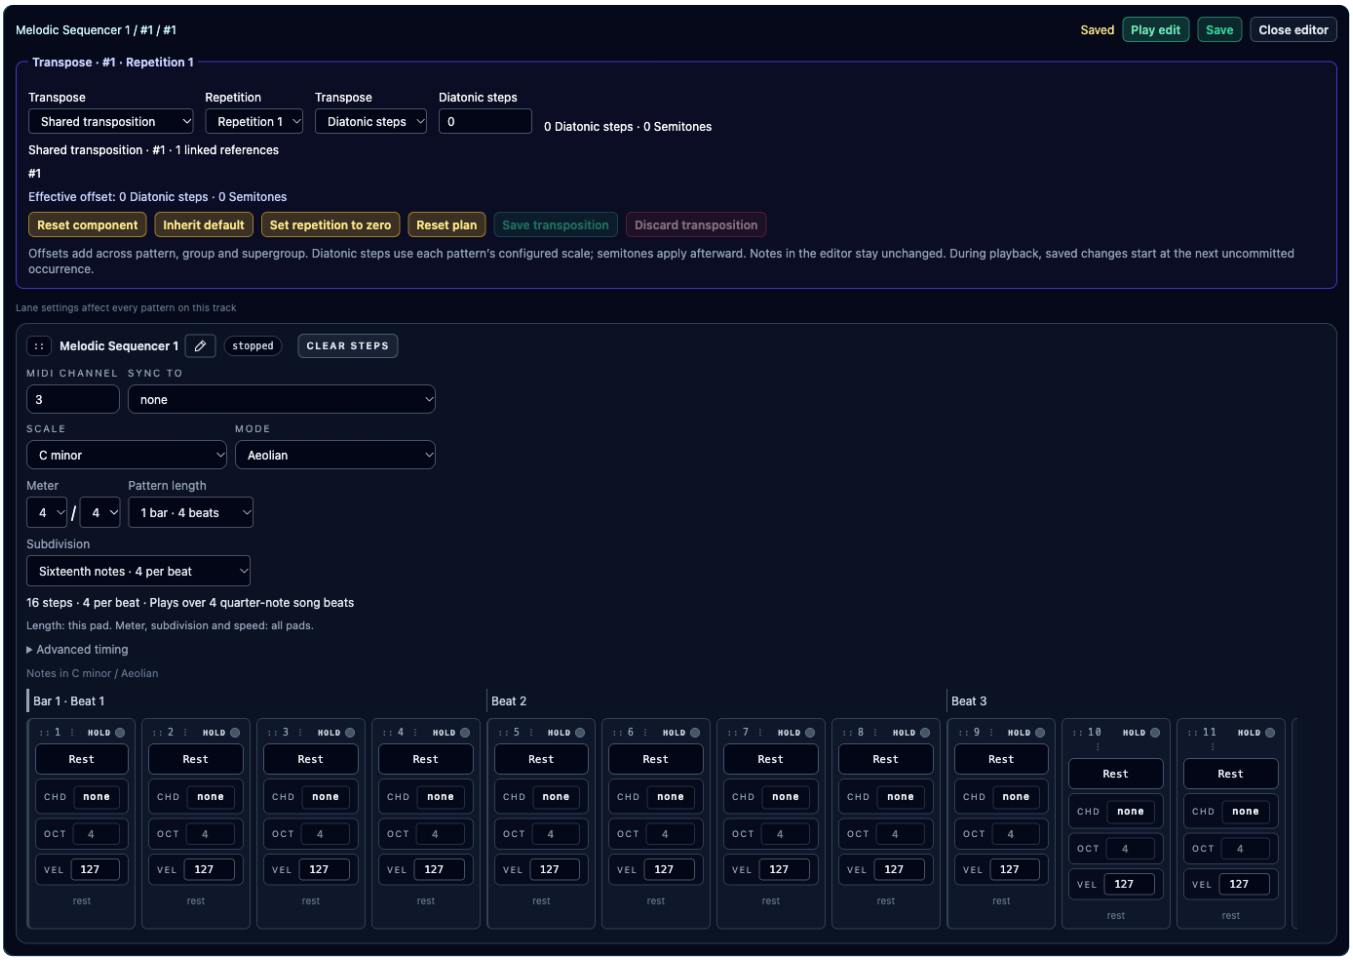

Global Sequencer Clock

The sequencer section contains a global clock with:

- Running/stopped state badge
- `BPM` field (range `30..300`)

The backend runs the step timing (native runtime clock), which reduces browser timing jitter during playback.

Adding / Removing Melodic Sequencers

Choose **Add track → Melodic** in the arranger, then add a clip and a pattern. Select a clip on the track to access **Remove track**. A performance can contain up to 16 melodic sequencers.

Per-Sequencer Controls

The melodic pattern editor provides:

- sequencer state badge (running/stopped or queued start/stop state)
- `Clear Steps`
- `MIDI Channel (1..16)`
- `Scale` (root + scale type)
- `Mode` (Ionian, Dorian, Phrygian, Lydian, Mixolydian, Aeolian, Locrian)
- `Meter (2..7 over 4 or 8)`
- `Subdivision` (1, 2, 3, 4, 6 or 8 steps per local meter beat)
- `Advanced timing → Playback speed (1:1, 2:1, 3:2, 4:3, 3:4, 5:4, 4:5, 7:4)`
- `Pattern length` (every integer from 1–16 local beats; whole bars show both units)

The step editor width is derived from the sequencer's own timing and length:

- steps per beat = the sequencer's selected `Subdivision`
- pad steps = `beats * steps per beat`
- default timing (`4/4`, grid `4`) gives `16` steps for a `4`-beat pad
- odd meters can still use matching one-bar pad lengths such as `3` beats in `3/4` or `5` beats in `5/4`

`Advanced timing → Playback speed` changes how quickly the sequencer advances relative to the shared transport beat without changing its stored pad length, meter, or grid:

- `1:1` keeps normal speed
- ratios above `1:1` make the melodic sequencer run faster
- ratios below `1:1` make the melodic sequencer run slower

This is separate from `Meter` and `Subdivision`, so you can combine polymeter and true polyrhythm on the same performance page.

Queued Start/Stop State Labels

When changes are queued to take effect at the next cycle boundary, the sequencer state label can show queued states such as:

- starting at step 1
- stopping at step 1

Scale, Scale Type, and Mode

Orchestron separates:

- Scale root (for example `C`, `Db`, `F#`)
- Scale type (`major`, `neutral`, `minor`)
- Mode (7 diatonic modes)

This supports guided note entry and better readability in the sequencer step editor and piano roll highlighting.

Edit patterns in the arranger

Double-click a pattern inside an expanded clip to open this editor. The breadcrumb identifies the occurrence being edited. **Play edit** auditions its draft; during song playback, other tracks continue. **Save** lets you update **Only this occurrence** or **All linked occurrences**. Cancel keeps the draft.

Assemble patterns, groups, supergroups and rests directly inside clips. See Patterns and Reusable Phrases for linked content, variations and structural editing.

Step Editor (Per-Step Note Programming)

Each step cell supports:

- Rest or note selection
- Chord selection, including the standard `5` power chord (root + perfect fifth)
- Hold toggle (`HOLD`)
- Octave field (separate from pitch class)
- In-scale/out-of-scale status display
- Playhead highlighting during playback

Step Note Entry Workflow

Each step has a pitch-class dropdown grouped into:

- `Rest`
- In-scale notes (for the current sequencer scale/mode)
- Out-of-scale notes

The octave is edited separately via the `OCT` field.

This split note/octave design makes pattern entry faster than selecting full MIDI note numbers.

Theory Aids In Steps

- In-scale notes are highlighted
- In-scale notes show degree labels (`1..7`)
- Out-of-scale notes remain selectable (for chromatic writing)
- Enharmonic note labels are available where relevant (for example sharps/flats)
- Chord choices are grouped as none, diatonic, or chromatic for the selected step root and mode. A `5` power chord is diatonic only when both the root and perfect fifth are in the current mode.

Hold Toggle

Each step includes a `HOLD` toggle.

Use `HOLD` to sustain behavior across steps according to the sequencer runtime logic and the note/gate pattern.

Step Cell Visual States

Step cells visually indicate:

- Active playhead step (when the melodic sequencer + transport are running)
- In-scale programmed note
- Out-of-scale programmed note
- Rest / hold combinations

Early and Late Notes

Timing moves an individual attack by **–50% to +50% of one local step**, in 1% increments. Zero means **On grid**. The millisecond readout follows the current tempo, grid and beat ratio: at 120 BPM with 4/4, Subdivision 4 and speed 1×, –20% is 25 ms early. Timing belongs to each pad and does not change its length, meter or playback speed.

Moving a note also moves its release, preserving its length and any HOLD extension. A following attack can shorten an ordinary preceding note to avoid overlap. Timing moves the entire chord, including its strum. Strummed notes retain their individual durations and can overlap later steps. For coincident ordinary attacks, the later logical step wins; for strummed

notes this applies only to identical pitches.

An early first step plays before the boundary when the same pad repeats. On a fresh start or a different-pad launch it plays at the boundary instead. Queued stops, pauses and finite arrangement ends suppress repeat anticipation. A command received after an anticipated attack has already sounded cannot undo that attack.

Copying steps or pads preserves timing. Clear Steps resets it. Inactive drum cells retain their timing for reactivation. Old performances load on grid; Save/Load, browser restoration, native bundles and both CSD export modes preserve offsets. Timing edits during playback use the normal coalesced live-edit workflow.

Right-click a melodic note or chord, press **Shift+F10** while focusing a step, or click its ⋮ button to open **Step properties** above the grid. Use the **Timing** slider or signed percentage field; the reset arrow restores zero. Timing remains available for rests. Nonzero timing appears as a clickable badge in the step. The slider supports arrow keys; the note/chord copy handle also copies timing.

Strummed Chords

In **Step properties**, select **Strum → Low → High** or **High → Low**. Notes play in pitch order, including the chord's octave and voicing. **Off** plays the chord together. **Spread** sets the time from the first to the last note as an integer **0–100% of one local grid step**; the slider, percentage field and reset share the same value. The millisecond readout follows tempo, meter, subdivision and playback speed. Spread uses the next grid boundary, even if that step is empty.

For a three-note chord at 120 BPM, 4/4, Subdivision 4 and speed 1×, a step lasts 125 ms:

Spread	First note	Second note	Last note
0%	0 ms	0 ms	0 ms
40%	0 ms	25 ms	50 ms
100%	0 ms	62.5 ms	125 ms

Timing shifts all these attacks together. At 100%, the last note lands at the next step relative to the shifted start. Each note keeps the original chord duration, including HOLD extensions, so later notes may overlap following steps. A repeated pitch retriggers: the newer note replaces that pitch's old release. Different pitches sounding together at a step boundary are retained. This is one strum per chord occurrence; melodic steps have no ratchet controls.

Select a note and chord to enable Strum. Removing the chord or changing the note to a rest disables strumming but retains the settings for reuse. **Off** retains Spread; **0%** plays simultaneously. A compact **↑ 40%** or **↓ 40%** badge opens the same panel. Step/pad copies, hidden steps, saves, browser restoration and native bundles preserve the settings; **Clear Steps** resets them. MIDI and SCORE exports use the same attacks and durations.

The panel edits the displayed pad and closes on pad changes, collapse, removal, New/Load/Import, or Escape. Direction, spread and timing edits to a strum already in progress apply on its next occurrence; live HOLD edits still update its release. Same-pad repeats can carry delayed notes across their boundary. Stops, seeks, a different pad, arrangement rests and finite endings cancel unfinished strums and release their notes.

Clear Steps

`Clear Steps` resets the current melodic sequencer step contents (for the active pad pattern context) so you can quickly reprogram it.

Related Features

- Linked patterns, local variations, transposition and grouping are documented in Patterns and Reusable Phrases.

Navigation: Up | Previous | Next

Drummer Sequencers

Navigation: Up | Previous | Next

Drummer sequencers are drum-machine style step sequencers for fixed MIDI drum keys.

They are optimized for per-step drum hit programming instead of melodic note/chord entry.

Use the pen beside the device name to rename it. See Device Names for editing controls, validation, and import/export behavior.

See Sequencer timing for triplets, beat grouping and editing behavior.

Open and close the editor

Double-click a pattern in the arranger to open its editor. **Close editor** returns to the timeline and retains its draft. Closing the editor ends **Play edit** auditioning; the song's transport remains independent.

What Makes Them Different

Compared to melodic sequencers, drummer sequencers:

- use MIDI key numbers (`0..127`) per row (drum instrument selection)
- do not use scale / mode / chord controls
- let you program multiple drum rows per step
- store velocity per active drum hit (per row + step)

Adding / Removing Drummer Sequencers

Choose **Add track → Drummer** in the arranger, then add a clip and a pattern. Select a clip on the track to access **Remove track**. A performance can contain up to 16 drummer sequencers.

Per-Drummer-Sequencer Controls

The drummer pattern editor provides:

- sequencer state badge (running/stopped or queued start/stop state)
- `Clear Steps`
- `+ Key` (add another drum row)
- `MIDI Channel (1..16)`
- `Meter (2..7 over 4 or 8)`

- `Subdivision` (1, 2, 3, 4, 6 or 8 steps per local meter beat)
- `Advanced timing → Playback speed` (`1:1`, `2:1`, `3:2`, `4:3`, `3:4`, `5:4`, `4:5`, `7:4`)
- `Pattern length` (every integer from 1–16 local beats; whole bars show both units)

Drum Rows (`Keys`)

The left side of the drummer grid contains vertically stacked drum rows.

Each row has:

- row index label
- MIDI key input (`0..127`)
- row remove button (`x`)

Auditioning Drum Keys While Editing

When the instrument engine session is running, changing a row key sends a short one-shot MIDI note preview on the drummer sequencer's MIDI channel.

This helps identify which drum sound is mapped to a given MIDI key number.

Drum Step Grid (LED Matrix)

The grid is row-based:

- rows = selected drum keys
- columns = `beats * steps per beat`

This keeps the `Keys` column horizontally aligned with the LED rows.

LED States

- inactive step: dark LED with visible border
- active step: red LED
- current playing column: highlighted across every drum row
- active step in the current playing column: flashing green LED

Velocity (Per Hit)

Velocity is shown by LED color saturation (stronger color = higher velocity).

The LED border color stays visible even at low velocities, so low-velocity active hits remain distinguishable from inactive steps.

While dragging vertically on a hit, the UI also shows a live numeric `velocity: xx` readout next to the edited step.

Editing Workflow

- Click an inactive LED to activate a hit
- Click an active LED (without dragging) to deactivate it
- Click and drag vertically on an LED to change velocity (`0..127`)
- While dragging, watch the live `velocity` readout for the exact value
- Keyboard:
- `Enter` / `Space` toggles the hit
- `Arrow Up` / `Arrow Down` adjusts velocity

Early and Late Notes

Timing moves an individual attack by **–50% to +50% of one local step**, in 1% increments. Zero means **On grid**. The

millisecond readout follows the current tempo, grid and beat ratio: at 120 BPM with 4/4, Subdivision 4 and speed 1×, –20% is 25 ms early. Timing belongs to each pad and does not change its length, meter or playback speed.

Moving a hit also moves its release. A following attack in the same drum row can shorten the preceding hit to avoid overlap. If neighboring attacks in that row land at exactly the same instant, the later logical step wins. For repeated hits within a step, see Ratchets and Velocity Ramps.

An early first step plays before the boundary when the same pad repeats. On a fresh start or a different-pad launch it plays at the boundary instead. Queued stops, pauses and finite arrangement ends suppress repeat anticipation. A command received after an anticipated attack has already sounded cannot undo that attack.

Copying steps or pads preserves timing. Clear Steps resets it. Inactive drum cells retain their timing for reactivation. Old performances load on grid; Save/Load, browser restoration, native bundles and both CSD export modes preserve offsets. Timing edits during playback use the normal coalesced live-edit workflow.

Drag a hit **left/right** to adjust timing, or **up/down** to adjust velocity. The first drag direction locks the property until release; clicking without dragging still toggles the hit. Changed hits show a signed percentage and a displaced indicator.

Use **Left/Right** arrow keys for timing and **Up/Down** for velocity. Right-click a hit, or press **Shift+F10**, to open **Step properties**, which includes Timing controls with a slider, numeric field and reset. The controls affect that hit in the displayed pad. Closing or collapsing the editor ends the gesture.

Edit patterns in the arranger

Double-click a pattern inside an expanded clip to open this editor. The breadcrumb identifies the occurrence being edited. **Play edit** auditions its draft; during song playback, other tracks continue. **Save** lets you update **Only this occurrence** or **All linked occurrences**. Cancel keeps the draft.

Assemble patterns, groups, supergroups and rests directly inside clips. See Patterns and Reusable Phrases for linked content, variations and structural editing.

Pattern length uses local beats, while meter, subdivision and playback speed belong to the track. All drum rows share the pattern duration. Choose a duration that fits the meter to create one-bar patterns in odd meters.

Typical Uses

- Kick / snare / hat step programming
- Percussion layers with multiple rows
- Switching between groove variations using pattern pads
- Automating drum pattern changes with the arranger while performing on piano roll/controllers

Ratchets and Velocity Ramps

Right-click a drum cell, or focus it and press **Shift+F10**, to open **Step properties**. Set **Ratchets** to 1–8: 1 is the normal single hit, and higher counts divide one local step into evenly spaced hits. An **×N** badge identifies the roll in the grid. At 120 BPM with Subdivision 4 in 4/4, four hits occur at 0, 31.25, 62.5 and 93.75 ms within a 125 ms step.

Enable **Velocity ramp** and choose **Final velocity** (0–127). The first hit uses the cell's velocity and the remaining hits interpolate to the final value. For four hits, 100 → 40 produces 100, 80, 60, 40. Zero-velocity hits are silent. Disable the ramp for constant velocity. At one hit the ramp controls are disabled, but their values are retained.

Timing shifts the whole roll while preserving its spacing. The next active cell in the same drum row cuts off unfinished repeats; an equal-time collision belongs to the later step. On a fresh start or a different-pad launch an early first roll starts at the boundary. On a confirmed same-pad repeat it can begin early, and a late final roll can finish across that boundary. Stopping or seeking that device, launching a different pad, entering an arrangement rest or reaching its finite end cancels unfinished rolls. A roll already in progress keeps its count, ramp and timing when those settings are edited; the next occurrence uses the prepared changes.

Editing these properties does not activate an inactive cell. Disabling and re-enabling a cell preserves them; copying pads and saving/exporting preserves them, including hidden steps. Clear Steps restores one hit without a ramp. The panel edits the

displayed pad and closes when changing pads, collapsing the drummer section or removing the target. Escape also closes it. Existing click toggles, horizontal timing drags, vertical velocity drags and arrow-key editing remain available.

Navigation: Up | Previous | Next

Score Sequencers

Navigation: Up | Previous | Next

Score Sequencers hold long, polyphonic passages with an independent start, duration and velocity for every note. Choose **Add track → Score** in the arranger. Its first clip and pattern open automatically; double-click an existing pattern to reopen its Score editor. Match its MIDI channel to a rack instrument. A new pattern starts at four local bars.

The **?** button at the top right opens the standard help modal in English, German, French or Spanish, following the GUI language. Use it from the Score pattern editor.

Select a Score clip and use **Remove track** to remove its track.

Draw and edit notes

The vertical keyboard stays on the left while the score scrolls. All 128 MIDI pitches remain available. Coloured rows and keyboard markers show scale degrees; notes outside the selected mode remain editable. Each pad remembers its own scale and mode. Changing these controls updates guidance without moving notes.

- **Play a key:** press and hold a vertical piano key to audition its pitch at velocity 100 through this device's MIDI channel. Release to stop; focused keys also respond to Enter and Space. The instrument engine starts automatically if needed, without starting the score or arranger, and remains running until stopped. Releasing during startup cancels that pending note. Collapsing, leaving the editor or losing focus releases held keys.
- **Collect a melody:** double-click a key to append a quarter note at velocity 100 to the edited pad. Start at zero if empty, otherwise exactly after the latest note end, including hidden notes. Subdivision is ignored: a quarter occupies

one local beat in /4 or two in /8. Pattern length stays fixed; increase it first if the whole note will not fit. Each insertion is one Undo action and selects and reveals the new note. Insertion works without audio. Collect pitches, then arrange their rhythm and duration.

- **Draw:** drag empty space to set a note's start and duration. Drag a note to move it; drag either edge to resize it. Creation, movement and resizing snap to the selected subdivision. Changing subdivision leaves existing notes in place.
- **Select:** drag a rectangle around notes. Cmd/Ctrl-click adds or removes a note from the selection; dragging selected notes moves them together.
- **Copy / Paste:** copy selected notes and paste them at the last clicked score position. The clipboard is shared between this device's pads.
- **Note properties:** right-click a note, or select it and press **Shift+F10**. Set velocity or add a chord from **In scale** or **Out of scale**. Added tones inherit the root's start, duration and velocity, and remain independently editable. Identical pitch/start/duration combinations are not added again.
- **Time zoom**, **Pitch zoom**, and **Fit score** control the view. Scroll to reach other pitches or later passages. The white playhead follows audible time on the sounding pad, independently of editing another pad.
- **– ♪ / + ♪** transpose notes by a scale degree using the same calculation as melodic sequencers. Hold Shift for an octave. **Copy pad to** copies the whole score, length, scale and mode to another pad.

Keyboard commands apply within the focused editor: Delete/Backspace removes selected notes; Cmd/Ctrl+A selects all; Cmd/Ctrl+C/V copies/pastes; Cmd/Ctrl+Z undoes and Cmd/Ctrl+Shift+Z redoes. Escape or a cancelled pointer gesture discards the note preview. A completed drag is one undo action.

The score has a separate 25-action undo/redo history per pad. Selection, clipboard, history, zoom and scroll survive collapse and view switches during the session. New, Load and Import reset this temporary editor state. These settings are not stored in a performance or its arranger history.

Duration and hidden notes

The grid ends exactly at Pattern length. Drawing a note, dragging its right edge, or moving it can place its end exactly on that boundary. Resizing its left edge keeps the end fixed. Blank space after the boundary is outside the pattern and does not create notes.

Pattern length accepts 1–8,192 whole local beats and displays bars, beats and elapsed time at the current tempo and playback speed. Meter, subdivision and speed apply to all pads in the device; length applies only to the selected pad. Global tempo remains quarter-note based. See Sequencer timing.

Set length to arrangement uses the other arranged lanes and rounds up to a whole local beat. An arrangement exceeding the score capacity is rejected.

Shortening keeps the music at its original positions. Notes starting beyond the end are silent, and releases crossing the end are clipped during playback. The **Hidden notes** count includes both cases. Expand the length to recover them. **Trim hidden notes** removes later notes and cuts crossing notes to the current end; Undo restores the original material.

Each pad supports 10,000 stored notes, including hidden notes. A performance supports 16 Score Sequencers and 50,000 score notes in total. Operations exceeding these limits are rejected as a whole. General save and export size limits also apply.

Transport and arrangement

Use the arranger's **Play**, **Pause** and **Stop** controls for the song. **Play edit** auditions the open Score pattern while retaining the current song cursor. During song playback it replaces that track while the other tracks continue. **Stop edit** restores the arrangement at its current position.

Use clips, groups and supergroups to arrange Score patterns on the timeline. Edit partial passages inside the Score editor. Pattern duration and clip duration are independent; trimming a clip retains the notes beyond its window.

Different pitches sustain independently. A repeated pitch retriggers it and replaces its previous release. Starting, seeking or returning from an audition inside a sustained note retriggers only its remaining duration.

Live edits apply at a render boundary. Unchanged sounding notes remain held; changed durations update their releases; deleted or moved notes release. New attacks behind the playhead wait for the next occurrence. A failed preparation keeps the playing version and the edited draft. Shortening a finite arrangement past its cursor stops it without rewinding.

Save, export and authoring

Save/Load, clone, autosave and native JSON/ZIP bundles preserve saved Score patterns, note identities, hidden material, timing and arrangement. Save the pattern draft before saving or exporting the performance.

Both MIDI and SCORE exports retain the musical timing, including the remaining duration of notes already sounding at the export range's start. Score notes can also feed arpeggiators through their input MIDI channel.

Optional agent-assisted score authoring is covered in Performance Creator.

Clip transposition

The clip arranger supports shared repetition transpositions and independent nested offsets. Saved settings are preserved by native bundles and both MIDI and SCORE CSD exports. Use a current Orchestron version when sharing transposed performances.

Create a Score pattern in the clip arranger

Score patterns are edited directly from their clips in the multitrack arranger.

1. Choose **Add track → Score** and match its MIDI channel to a rack instrument.
2. Use the track's first pattern, or add another clip with the track's **+** or by double-clicking empty lane space.
3. For an empty clip, use its inline **+** / **Add content**, then **New pattern** to insert a Score pattern. Existing definitions can also be adopted as linked patterns.
4. Double-click the nested pattern box, or focus its name and press **Enter**. The Score sequencer opens in context with the clip breadcrumb and draft controls.
5. Hold piano keys to try pitches and double-click to collect quarter notes; then move and resize notes. The **?** opens the same multilingual help.
6. Use **Play edit** to audition the pattern. Choose **Edit linked pattern** or **Create variation here** before editing shared content. The chosen scope remains visible, and Save applies it without another scope question. A local save creates a variation only when content changed. Save the pattern before saving or exporting the performance to include these edits.

Pattern length and the clip window are independent. Saving musical edits keeps the clip boundaries; use **Fit clip to content** when the clip should match its content.

Navigation: Up | Previous | Next

Arpeggiators

Navigation: Up | Previous | Next

An arpeggiator processes MIDI notes and drives an existing rack instrument. Send chords from a melodic sequencer, an on-screen keyboard, or external MIDI to its **Input Channel**. Choose the named instrument in **Target Channel**. Both controls are always visible in the top row, with **Input Channel → Target Channel** showing the MIDI direction. Inputs must be unique; an arpeggiator cannot target another arpeggiator input.

Start and Synchronize

Arranger follows song playback, seeks and loop selections. Incoming chords join the next subdivision of the shared beat; chord changes keep the phrase position. Seeks and loop wraps reconstruct the phrase at the destination, and playback stops at Song end. The arpeggiator needs input notes or a held chord to produce sound.

Live runs independently while the instrument engine is active. The first incoming note establishes the Live clock; subsequent chords retain its pulse. Both modes use global BPM, including tempo changes.

Restart controls phrase position without moving the subdivision clock:

- **Continue:** keep advancing through chord changes (default).
- **First held note:** restart the phrase on a fresh key gesture.
- **Every beat / Every bar:** restart at the next subdivision of each master beat/bar.

Straight, triplet and dotted rates use musical beat positions. Swing alternates long and short intervals with an unchanged total duration for each pair.

Hold and Processing

Hold Off is the default. Releasing all input notes silences the output, preserving rests in a source sequence. **Hold and replace** remembers the last chord after release; the next fresh gesture replaces that chord. **Add/remove notes** toggles pitches explicitly. **Clear held notes** releases the arpeggiator's notes and clears pending input and launches.

Active generates the pattern; **Bypass** forwards incoming notes to the target; **Mute** consumes notes silently. Stop, removal, routing changes and processing/playback-mode changes release notes owned by this arpeggiator.

Musical patterns

Choose **Add track → Arpeggiator** in the arranger, add a clip and a pattern, then double-click the pattern to edit it. Each pattern stores its order, rate, octave range, transpose, expression, harmony, rhythm, rotation and random seed.

Rhythm length and pattern duration are independent. A rhythm has 1–32 steps; pattern duration is 1–8 or 16 master beats. A five-step rhythm can continue through repeated four-beat occurrences of the same pattern. Changing patterns restarts the musical cursors. Arrangement rests silence output while incoming pitches continue to be tracked.

Use linked patterns to repeat a phrase, or **Create variation** for an independent version. Groups and supergroups collect reusable content inside clips.

Audition and save

Use **Play edit** to audition the open pattern with incoming notes or a chord. During song playback it temporarily replaces its track while the other tracks continue. **Stop edit** restores the track at the current song position. Hold and Active/Bypass/Mute remain effective during audition.

Save offers **Only this occurrence** or **All linked occurrences**. It preserves the rhythm, expression, harmony and random seed. See Patterns and Reusable Phrases.

Edit the Rhythm Grid

Select a step, then choose its action:

Action	Result
Next note	Play the next pitch in the selected order
Chord-note position	Play a relative position in the ascending, octave-expanded input pool; oversized positions wrap
Rest	Advance time silently
Tie	Extend the preceding note/chord without retriggering; silent without a preceding note
Chord	Play the entire expanded input pool

Velocity bars scale incoming accents in percent. Each step also has probability, a gate override and **1–4 ratchets**. Gate supports **5–200%**; values above 100% overlap different pitches. **Use pad gate** clears the step override. Rests and probability misses keep the automatic note cursor stationary unless **Advance through rests** is enabled. **Note repeats** repeats an automatic pitch before moving on. **Rotation** changes the rhythm's starting step.

Keyboard editing: Left/Right selects adjacent steps; Space toggles a rest; Delete/Backspace inserts a rest; T inserts a tie; C plays the chord; N restores Next note. Double-click also toggles a rest. Controls have accessible text labels.

Note previews use large 18 px labels. Narrow grids scroll horizontally; long chord labels stay on one line with an ellipsis. Hover over a step for the complete note list and scale degrees, also available to screen readers.

With **Follow source** or **Custom** scale handling, step borders use the piano keyboard’s degree colours: **1 red, 2 orange, 3 yellow, 4 green, 5 teal, 6 blue, 7 violet**. Degrees describe the generated pitches after transposition and quantization. Follow source uses each note’s incoming scale context, with the pad’s scale as fallback. Octave repetitions keep the same colour. Chords show equal, distinct border segments for their degrees, ordered 1–7. **Scale Off**, missing previews, rests and ties keep neutral borders. Previews and colours appear only when editing the playing pad. A cyan outer outline marks the selected step independently of its degree colour; the playback background and keyboard focus remain separate.

Expression and Harmony

Input velocity preserves accents; Fixed and Random velocity remain available. Timing and velocity humanization change individual strikes. Timing variation stays inside each strike's scheduling window and never moves the next clock position.

Repeat variation repeats seeded random choices; **Evolve** varies them by cycle. **New variation** changes the saved seed. Each pad has its own seed, independent of other devices; identical input and transport positions produce reproducible events in playback and exports.

Octaves expand the input pool before applying note order. Down descends through the complete range. Up/Down reverses without duplicating endpoints. **Octave by octave** retains an alternative traversal in Advanced timing.

Scale handling defaults to **Off**. **Follow source** uses the source's root/mode and falls back to the pad's configured scale when absent. **Custom** always uses the pad's scale. The effective scale and resolved pitch preview help explain the output. Playback highlights follow audible PCM status.

Presets, Live Edits and Collapse

Presets apply only to the pattern being edited. **Modified** marks changes from its preset. **Update preset** overwrites a user preset; **Save as preset** creates a new one. Built-in presets remain available. Routing, playback settings and other patterns are preserved.

Live edits coalesce for 80 ms and apply at audio render boundaries, preserving transport and pending launches. Invalid edits retain the working configuration. Status updates never resubmit musical settings. Collapse hides visual work while preserving playback, pattern selection, details and drafts.

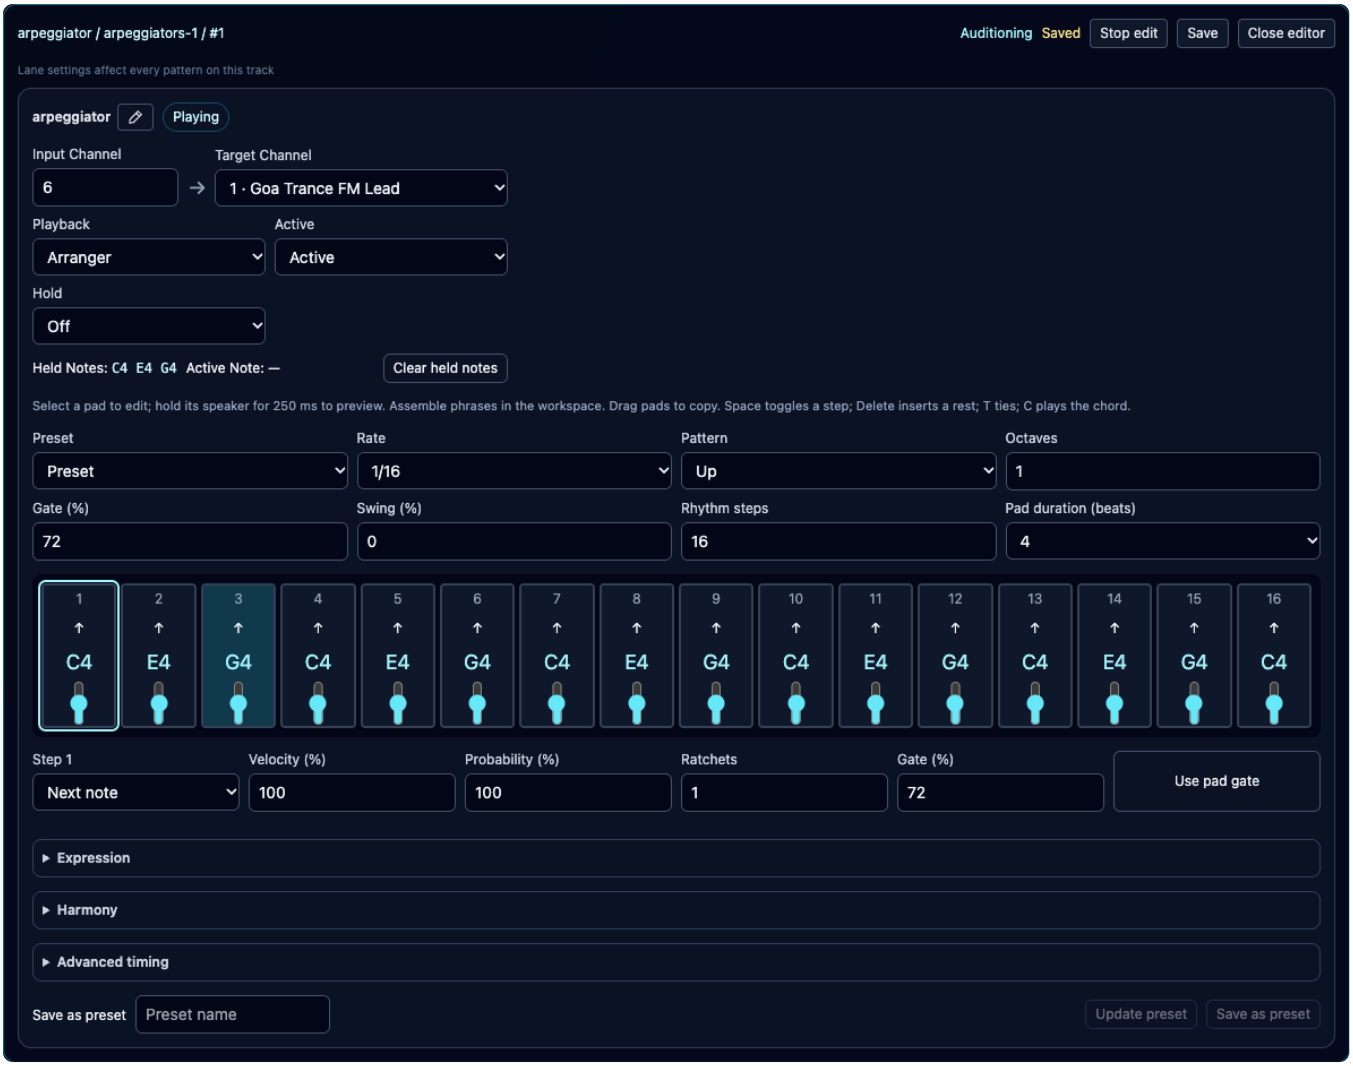

Clip transposition

The clip arranger supports shared repetition transpositions and independent nested offsets. Saved settings are preserved by native bundles and both MIDI and SCORE CSD exports. Use a current Orchestron version when sharing transposed performances.

Navigation: Up | Previous | Next

Controller Sequencers

Navigation: Up | Previous | Next

Controller Sequencers automate MIDI CC values using editable curves or discrete steps.

Use the pen beside the device name to rename it. See Device Names for editing controls, validation, and import/export behavior.

See Sequencer timing for triplets, beat grouping and editing behavior.

Open and close the editor

Double-click a pattern in the arranger to open its editor. **Close editor** returns to the timeline and retains its draft. Closing the editor ends **Play edit** auditioning; the song's transport remains independent.

What A Controller Sequencer Does

A controller sequencer sends timed MIDI Control Change messages on its selected channels. **Curve Mode** (the default) samples an editable curve. **Step Mode** sends an integer at each populated grid cell. The mode applies to every pattern on the track; switching modes preserves both datasets.

Typical uses:

- filter sweeps
- morph controls
- modulation depth animation
- macro movement synchronized to transport

Adding / Removing Controller Sequencers

Choose **Add track → Controller** in the arranger, then add a clip and a pattern. Select a clip on the track to access **Remove track**. A performance can contain up to 16 controller sequencers.

Per-Controller-Sequencer Controls

Each controller sequencer provides:

- Running/stopped state badge
- `Controller # (0..127)`
- **MIDI Channels** checkboxes, 1–16 in one row beside `Clear Steps`
- `Meter (2..7 over 4 or 8)`
- `Subdivision` (1, 2, 3, 4, 6 or 8 steps per local meter beat)
- `Advanced timing → Playback speed (1:1, 2:1, 3:2, 4:3, 3:4, 5:4, 4:5, 7:4)`
- `Pattern length` (every integer from 1–32 local beats, subject to 128 steps)
- CC label preview (`CC N`)
- **Curve Mode / Step Mode** toggle
- Curve editor or step grid
- **Reset Value** (`0..127`, default `0`) in Step Mode

MIDI Channels

Check the channels that receive the sequencer's CC messages. All 16 are checked by default (OMNI); at least one must remain checked. The row scrolls horizontally on narrow screens. This is a track setting shared by all its patterns.

Channel edits use the same 80 ms live-edit synchronization as curve edits and take effect at the next engine block without restarting transport. In Curve Mode, the current curve value is sent to the new selection; removed channels retain their last

received value. In Step Mode, departing CC/channel pairs receive Reset Value, newly targeted pairs receive Reset Value, and continuing pairs hold their state until the next populated step.

Selections survive Save/Load Performance, browser reload, native JSON/ZIP export/import, and both CSD export modes. Older performances without the setting load with all 16 channels checked.

Pattern Length (1–32 local beats)

This defines the pattern repeat length relative to the controller sequencer's own timing. Step Mode shows `length × subdivision` cells, up to 128; shortening hides trailing cells without deleting them. Keypoints stay normalized across the full pad duration, so the same curve shape stretches automatically when you choose a longer beat length.

`Advanced timing → Playback speed` changes how quickly the controller sequencer moves through that curve relative to the shared transport. Faster ratios create repeating automation polyrhythms without changing the stored keypoint positions.

Instrument-bound automation

In the clip arranger, **Generate automation** creates controller lanes from the instrument rack's CC assignments. These lanes lock routing to the rack instance and support song-length curves with at most ten points. Curve duration uses master beats independently of the retained Step Mode grid. See Instrument automation clips for inline editing, Save/Discard and mapping changes.

Edit patterns in the arranger

Double-click a pattern inside an expanded clip to open this editor. The breadcrumb identifies the occurrence being edited. **Play edit** auditions its draft; during song playback, other tracks continue. **Save** lets you update **Only this occurrence** or **All linked occurrences**. Cancel keeps the draft.

Assemble patterns, groups, supergroups and rests directly inside clips. See Patterns and Reusable Phrases for linked content, variations and structural editing.

Step Editor and Reset Value

Enter a whole number from **0 to 127** in a cell. Press Enter or leave the cell to apply; Escape cancels. Delete the text to leave a cell empty. Empty cells send nothing and the receiver holds its previous value. Zero sends a real CC message, and repeated identical values each send again at their own step. Beat/bar groups and the highlighted playing step follow the device's own meter, subdivision and playback speed.

Live playback and CSD (MIDI) retain those repeated messages. CSD (SCORE) can omit writes that leave the receiving channel/CC unchanged, preserving playback and all authored step data. Startup writes remain; exports with note-triggered `initc7` patches retain repeats. See offline export.

Reset Value is sent for the chosen CC on the selected channels on start and stop, including finite endings. It is sent before the first step and notes at the start position. Starting in an empty part of the pattern waits for the next populated step. Choose a value appropriate to the controlled parameter: zero is the default, not a universal neutral setting for every CC.

A seek during playback sends Reset Value, then continues from the new position without replaying earlier cells. A stopped seek sends nothing. Pattern changes, loops and arrangement rests hold the last received value without resetting. **Play edit** initializes the controller on entry and resets it when restoring playback.

Entering Step Mode while running initializes Reset Value. Leaving it resets the old step output before the curve takes over. Editing Reset Value takes effect on the next reset. Reset messages obey lane mute/solo, and only address the selected CC/channel pairs; no automatic "Reset All Controllers" message is sent. Sources sharing a CC/channel still use last-message-wins behavior.

See Granular Capture for a complete four-lane stutter, freeze, capture and granular step pattern.

Curve Editor Interactions

The curve editor is an interactive graph view (spline-based curve display and sampling).

When all keypoints have the same value, the curve appears as a horizontal line across the editor. This also applies to the minimum (`0`) and maximum (`127`) CC values.

Supported interactions:

- Click background to add a point (interior points only)
- Drag a point to change position/value
- Double-click an interior point to remove it
- Endpoints are boundary anchors (positions remain at start/end of the curve)

Start and end values are independent. While dragging either endpoint, it snaps to the opposite endpoint's value when the difference is **3 CC units or less**. An amber horizontal guide, endpoint rings and an **Endpoint snap** value label show the snap. Hold **Shift** to disable snapping; pressing or releasing Shift also works without moving the pointer during a drag. Only the dragged endpoint changes. Numeric entry stays exact and never snaps. Interior points do not snap to the endpoints.

To build continuous automation across pads, match each pad's end value to the next pad's start: for example, **20 → 50**, then **50 → 90**. Match a pad's own endpoints when it should repeat without a value jump. Repeating a pad with different endpoint values deliberately returns to its start value. Matching endpoint values does not guarantee matching curve slopes.

Save/load, native bundles and MIDI/SCORE exports retain independent endpoint values. Existing closed curves keep their shape; an omitted end point still defaults to the start value. Curve playback continues to send only changed CC values.

Visual Playback Indicators (When Running)

When the main transport is running and the controller sequencer is enabled, the editor shows:

- a vertical playback position line
- a marker showing the currently sampled CC value on the curve

This makes it easy to understand exactly what value is being sent at each transport moment.

Live Use Notes

- Controller sequencers run alongside melodic sequencers on the same backend transport clock.
- Curve and step edits are combined after an 80 ms editing pause and prepared while the previous configuration keeps playing. They take effect at the next engine block without resetting transport. The editor modifies the displayed pad even when the arranger selected it automatically.
- See Editing During Playback for failure handling and transport behavior.
- Open the pattern you want to edit, then use **Play edit** to audition its draft.
- You can combine automated controller sequencers with manual MIDI controller knob lanes on the same performance page.
- If you automate the same CC number from multiple sources, the last-sent value wins at the MIDI receiver side (plan mappings accordingly).

Inline automation editing

Controller curves can also be edited directly inside an expanded clip. See Instrument automation for generating curves tied to rack controls.

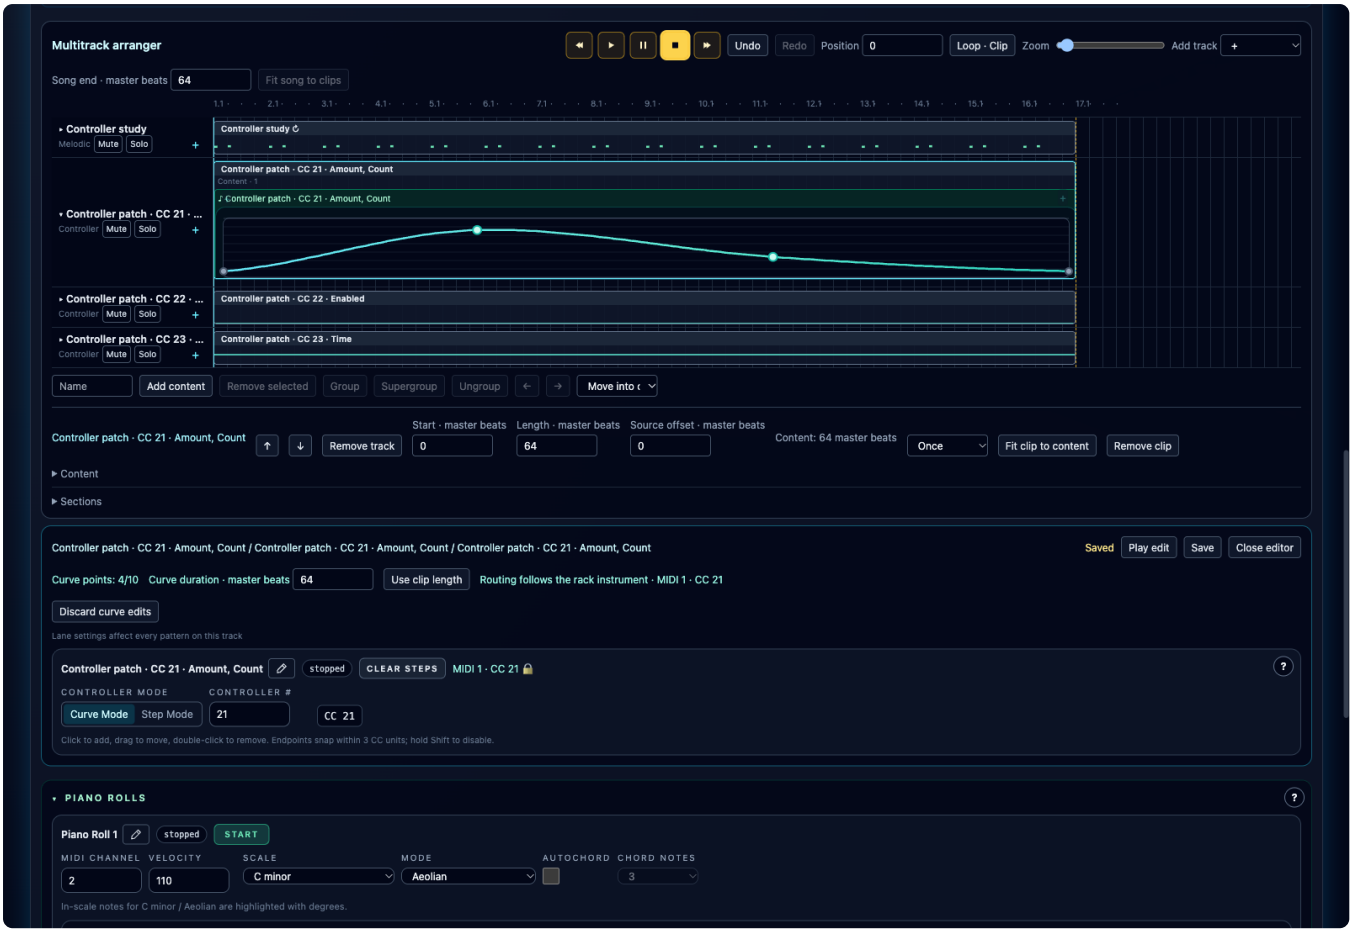

Navigation: Up | Previous | Next

Instrument Automation Clips

Navigation: Up | Previous | Next

In the clip arranger, **Generate automation** creates controller clips owned by a rack instrument. Assign MIDI CCs to its performance controllers in the rack first. Select a Melodic, Drummer, Score or Arpeggiator clip, then choose Generate automation in its inspector. If several rack instances receive that source's output channel, choose the instance in **Automation instrument**. Continuous instruments and invalid or unassigned controls are excluded.

One lane is created per distinct CC, with the instrument and mapped control names in its label. Controls sharing a CC share one lane. Existing lanes are reused. Each new clip starts and ends with the selected source clip and contains one equally long curve, with zero source offset and Once playback. Its initial two points use the first mapped control in patch order and the current rack value, converted through its linear, logarithmic or Boolean mapping.

Generating again preserves existing automation with the same start and length. Other overlaps reject the entire operation. Automation clips can subsequently move, trim, loop or disappear independently. Deleting their source clip or source lane leaves automation intact. Bound automation can exceed 16 controller lanes, within the arranger's 72-lane total; generic controller sequencers retain their 16-device limit.

Edit a slow curve inside its clip

Expand the automation lane and select its curve preview, or double-click the pattern header. Click the curve background to add a point, drag to move it, and double-click an interior point to remove it. Clip headers still move clips, borders trim them, and pattern headers select or move structure. Zoom, scrolling, nested groups, repeated occurrences and source offsets preserve the curve's original coordinates.

The counter allows **10 points including both endpoints**. Endpoint positions stay fixed; their values are independent. Endpoint dragging retains the existing three-unit snapping and Shift override. Use shorter, looped patterns for fast movements.

Curve duration · master beats stretches the normalized curve without changing its points or the clip window. **Use clip length** copies the selected window's duration. Trimming crops the existing content. Extending Song end leaves existing curves at their saved durations. Curves can span up to the arrangement limit of 1,000,000 master beats.

Shared curves follow the same editing scope and saving rules as musical patterns, including when editing their inline points. **Discard curve edits** removes the selected draft. Play edit auditions the draft through the usual workspace controls. Performance save, autosave and native/CSD export retain committed definitions; unsaved curves remain session-only. Committed edits use arranger Undo/Redo. Adding, splitting or merging source lanes clears clip history, following the existing topology rules.

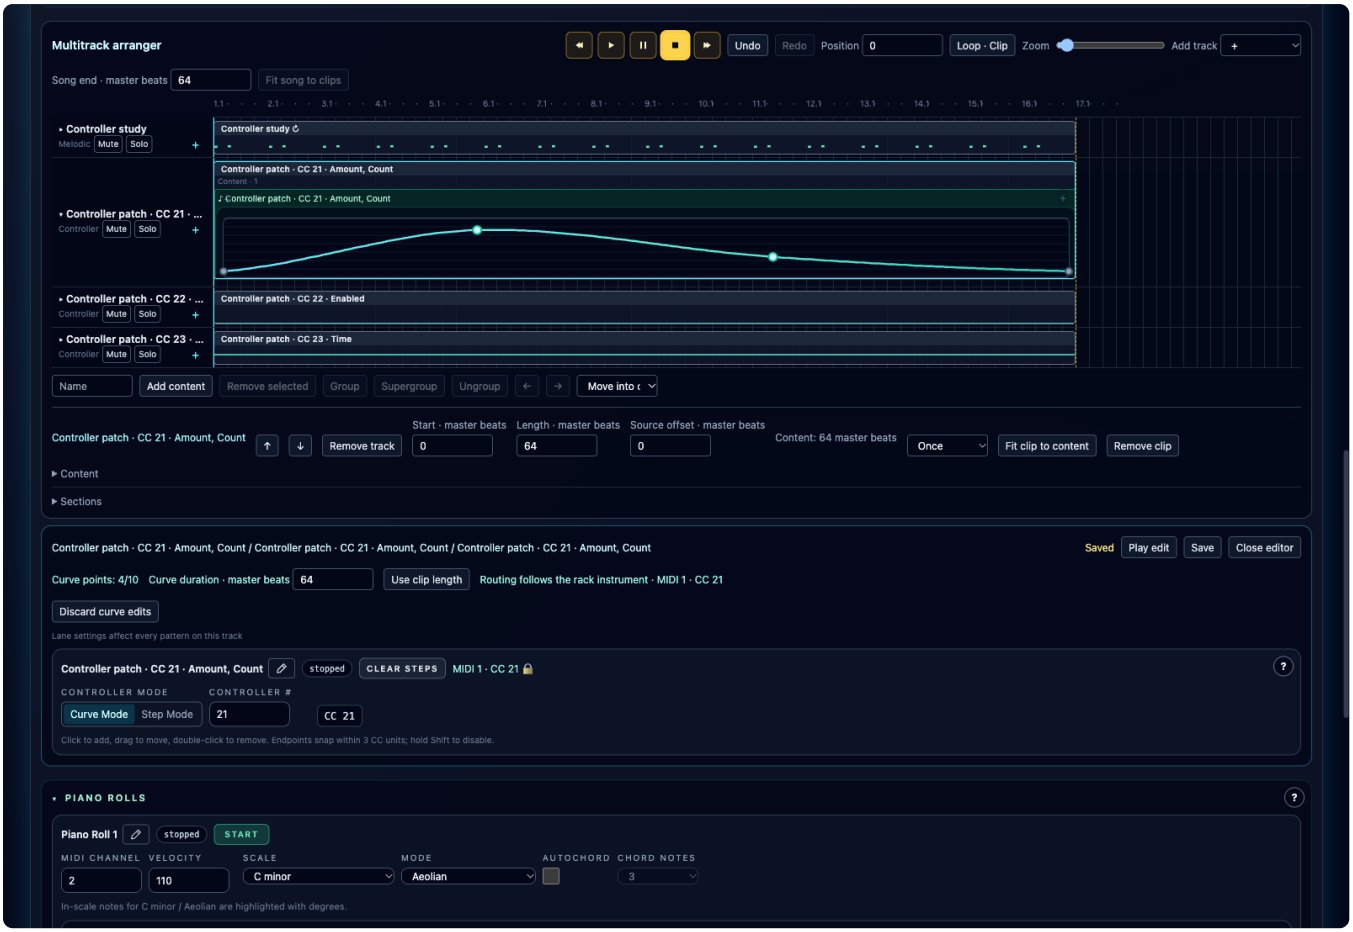

Routing and reassignment

The owning rack instance determines MIDI channel and CC. Both are read-only in the automation editor. Changing its rack channel follows automatically. Changing one mapped CC retargets the lane without changing its music. If controls sharing a lane move to different CCs, their complete automation is copied into independent lanes.

When mappings merge, identical automation combines automatically. If candidates differ, choose which complete lane survives; **Cancel** preserves both the old mapping and all automation. Save or discard affected automation drafts before changing assignments. Live changes prepare first. If preparation fails, the previous playing configuration remains in place and the error is shown.

Off mappings, removed controls and missing instruments retain their automation with a suspended warning. Suspended lanes send no MIDI. Reassignment uses the same split/merge rules. Ownership stays with the rack instance, even if the original note source disappears.

Step mode and compatibility

Step Mode keeps its local timing controls, 128 stored cells, hidden-cell retention and Reset Value behavior. Mode applies to the lane. Switching modes preserves both datasets and their separate durations: the step grid does not resize a long curve, and the curve duration does not resize the grid.

To generate another curve on an existing Step Mode automation lane, switch that lane back to Curve Mode first. Generation preserves its existing music and does not change the lane mode for you.

Use a current Orchestron version when sharing performances with instrument automation. Older performances retain their existing generic controller behavior.

For note instruments, mapped performance settings affect newly started notes or phrases. Continuous controls marked **Live** can change during playback; **Requires restart** controls retain their values until restart. See Using Performance Controls.

Navigation: Up | Previous | Next

Live Status and Troubleshooting

Navigation: Up | Previous | Next

The Perform page includes several status for live operation.

Browser Audio Health

The Runtime panel shows queued PCM and pending-render milliseconds, underrun and overrun counters, and the latest backend render-time/audio-time percentage while a browser-clock session is active.

Audio refills and PCM writes run in a Dedicated Worker. Sequencer step and pad updates are carried with the rendered PCM and applied only when their target frame becomes audible. If the Perform page is hidden, visual animation pauses and resumes from the current audible snapshot instead of replaying missed steps.

Performances are limited to 128 flattened backend note tracks. Each drummer row becomes one backend note track, so the total includes melodic sequencers plus all drummer rows.

Footer Status Bar

At the bottom of the Perform page, Orchestron shows live status values such as:

- `playhead` position (`current step / step count`)
- `cycle` count
- active `midi input` name (for example `internal:loopback` or a helper-provided device)

This footer is useful during performance because it gives quick feedback without opening the Instrument Design runtime panel.

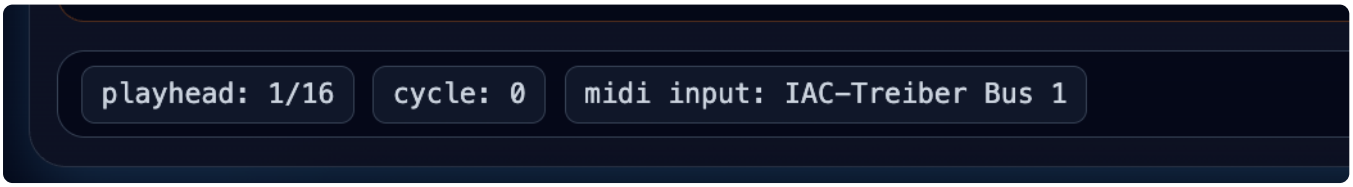

Footer status bar with playhead/cycle/MIDI input.

Engine Stats and Parallel Playback

A narrow **Engine stats** line stays fixed at the bottom of the Perform window. It shows whether the session uses **Single** or **Parallel** rendering. It also contains the optional **Parallel playback** checkbox.

To opt in, stop the instruments, enable **Parallel playback**, then start them again. The engine analyses the connections between whole instruments, inserts, sends and shared effects. It keeps dependent work and shared mutable resources together and distributes independent groups across available CPUs. Instruments are never split internally. No catalog or per-song registration is required, and browser latency settings do not control eligibility.

The number of processes depends on the graph and available container capacity; one process handles the shared downstream effects and output. Built-in read-only sine tables are present in each process. Shared mutable tables, MIDI initialization and shared random state keep their users together. Unsupported resource behavior, insufficient independent work or insufficient CPU/memory capacity selects **Single fallback**. Splitting a graph does not guarantee a speedup: small instruments and shared effects can still dominate the work.

The choice is off by default, resets when the page reloads, and is not saved in a performance or an export. It cannot be changed while instruments are running. Turning it off while stopped returns the next session to single rendering. Audio timing, buffer sizes and CSD exports are unchanged. Sequencer controllers retain their MIDI channels and reach every relevant instrument across the parallel processes.

On Linux, parallel rendering shows the total load on the logical CPUs assigned to its workers, for example `cpu0: 60% cpu1: 45% cpu3: 77%`. These CPU readings include other work using those CPUs. With single rendering, **Server process** shows the

combined CPU usage of the server process: 100% represents one CPU, and the value can exceed 100%. It is not a per-instrument meter. Unsupported platforms or unavailable readings show **CPU usage unavailable**, not zero. Each CPU percentage has a reserved, right-aligned space, so neighbouring readings stay in place as the load changes between one, two and three digits.

- **Green:** 0–50%.
- **Orange:** above 50% through 85%, including the 50–60% interval.
- **Red:** above 85%.

Readings update about once a second while Perform is visible. Switching views or hiding the tab suspends observation; audio keeps playing. Load percentages do not replace the audio underrun and queue indicators. Live-input latency is a separate diagnostic and may still vary.

If a parallel worker fails during playback, the engine stops and reports the error. It does not silently replace the running synths or replay audio. Stop and explicitly restart; automatic fallback can use single rendering on that restart.

Track State Labels (Queued Changes)

A track's state badge can reflect queued actions that apply at the next cycle boundary, for example:

- starting @ step 1
- stopping @ step 1

This helps you understand why a track button press did not take effect immediately while transport is running.

Error Banner (Transport / Runtime Errors)

The Perform page displays an error banner for problems such as:

- no active instrument session when trying to start sequencers
- trying to play piano roll notes before the instrument engine is ready
- pad queue request failures
- controller send failures
- sequencer config sync failures

Editing During Playback

Notes, controller curves, tempo and arrangement edits keep the transport at its current position. Rapid edits are combined after an 80 ms pause; the previous configuration keeps playing while the update is prepared. The completed edit takes effect at the next engine block. Editing a displayed pad changes that pad, including when the arranger has selected it automatically.

Automatic pad changes, playhead movement and status updates do not resubmit the performance configuration. Queued pad launches still wait for their usual pad boundary. Stop cancels pending configuration application. Shortening a finite arrangement past the current position stops playback without rewinding.

If preparation fails, the error banner reports it and the previous configuration continues playing. Your edit remains in the editor. Make another edit or explicitly restart to try again; the app does not continuously retry a failed update. Saving and exporting retain the existing file formats.

MIDI Input Reference

The footer's MIDI input display reflects the session's bound external MIDI input (selected in the Instrument Design Runtime panel). If it shows `internal:loopback`, internal app MIDI is still active, but external hardware/DAW MIDI will not reach the session until a helper-provided input is bound.

See Runtime Panel and Compilation Workflow and MIDI Setup and Inputs.

Live Performance Safety Checklist

1. Confirm instrument session is running.
2. Confirm expected MIDI input is bound.
3. Check track/controller/piano-roll enable states.
4. Save the performance after major routing changes.

Definition audition

A preview is temporary and does not save a pattern or change its authored arrangement. If it fails to prepare, current playback continues and your draft remains available. See Play and jam for the current clip-editor transport controls.

Temporary lane output controls

Lane **Mute** and **Solo** affect future output while the track keeps its place in the song. A muted controller leaves its target at the last value received. See lane controls for the complete behavior and how they differ from mixer controls.

Navigation: Up | Previous | Next

Performance Import / Export

Navigation: Up | Previous | Next

Use **Export** for an editable backup or to share a complete performance. Use**Export CSD (MIDI)** or **Export CSD (SCORE)** to render a song outside Orchestron with Csound.

Choose an export

Action	Use it for	Result
Export	Backup, sharing or later editing in Orchestron	A native `.orch.json` or `.orch.zip` bundle
Export CSD (MIDI)	Offline rendering with a separate MIDI score	A ZIP containing CSD, MIDI, assets and rendering instructions
Export CSD (SCORE)	Offline rendering with notes and controllers in the Csound score	A ZIP containing CSD, assets, rendering instructions and any warnings

Save open musical and structural drafts before exporting. Temporary auditions are not authored content. Exporting does not replace **Save Performance** in your library.

Native performance export

Choose **Export** on the Perform page. The bundle preserves the whole rack, instrument definitions, routes, Master and inserts, mixer settings, device names and channels, saved patterns, arrangement, transposition, instrument settings and MIDI mappings. Arranger Undo/Redo history is also retained when compatible.

When instruments use uploaded audio samples, SoundFonts or analysis files, Orchestron creates a ZIP containing the required assets. Otherwise it can use JSON. Keep the downloaded bundle intact for importing on another machine.

Import a performance

1. Choose **Import** and select an Orchestron performance bundle.
2. In the import options, choose whether to import the performance, its patch definitions, or both. Import both for a self-contained transfer.
3. Resolve any existing names in the conflict dialog.
4. Check the rack and routing, then play and save the imported performance.

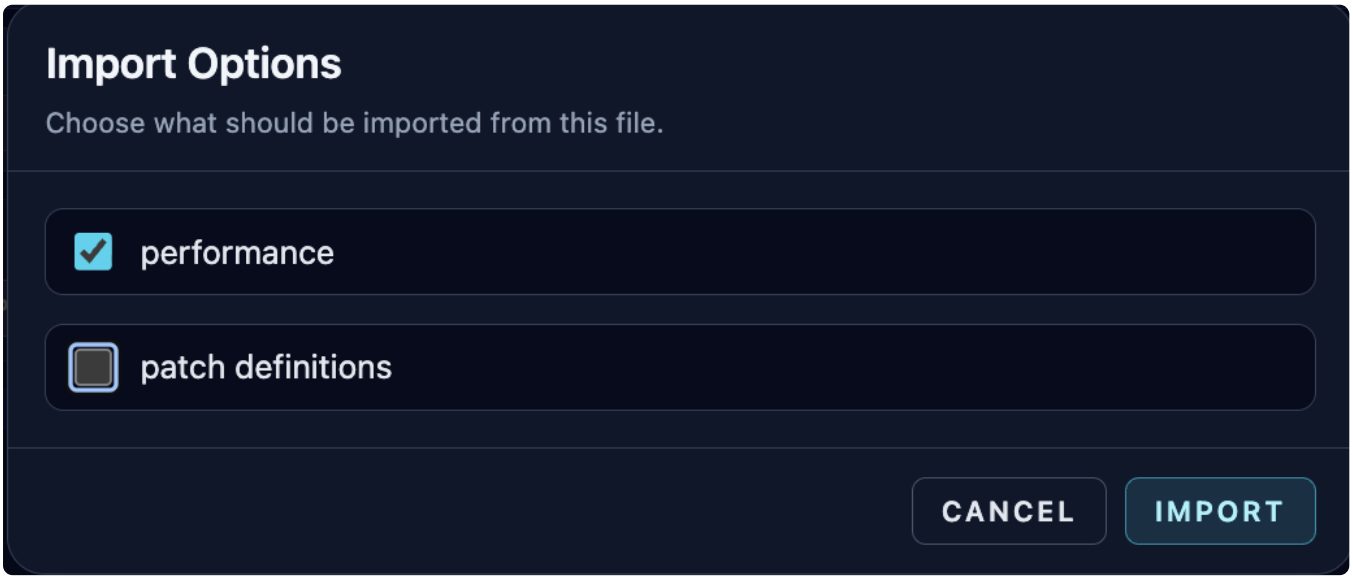

For patch-name conflicts, choose overwrite, rename or skip. For the performance itself, choose overwrite or rename; to omit the performance, change the import options instead. Skipping patch definitions requires suitable instruments to exist in the

destination library. Import validation reports unresolved instrument assignments.

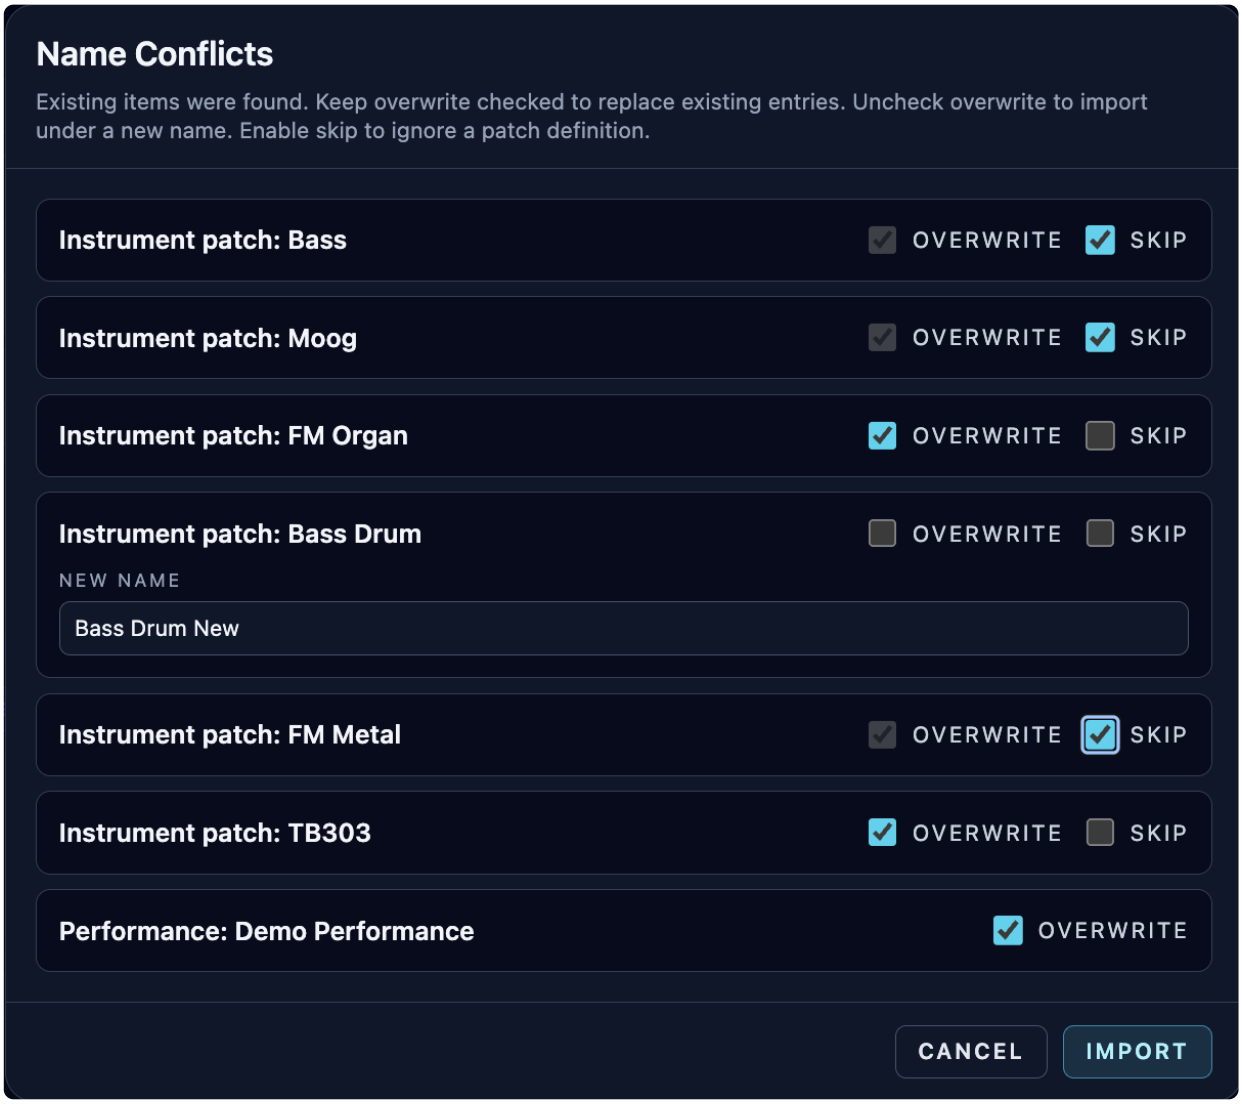

Render with Csound

1. Finish the arrangement and set a finite song end. Disable looping for the offline export.
2. Choose **Export CSD (MIDI)** or **Export CSD (SCORE)**.
3. If SCORE export asks for controller accuracy, choose a maximum deviation and continue.
4. Extract the ZIP. Open the included `README.txt` and run its Csound command from the extracted directory.
5. Read `WARNINGS.txt` when present before relying on the rendered result.

The packages render at **48 kHz**, `ksmps=1`, with **32-bit float WAV** output and effect release tails. Keep those settings when you need the documented timing and controller accuracy. Floating-point output retains headroom; check levels before converting to an integer audio format.

Both modes include instrument instance settings and seed enabled manual MIDI controller values on their selected channels. Live keyboard gestures are not recordings and are not added to the exported arrangement.

Controller accuracy

SCORE export preserves eligible original controller curves. Where a curve needs approximation, **Maximum deviation** accepts **1–10 CC units**, default **1**, on the 0–127 MIDI scale. A larger allowance can reduce data at the cost of accuracy. Changing the allowance may have no effect when original curves can be used directly. Step patterns and abrupt resets retain their events.

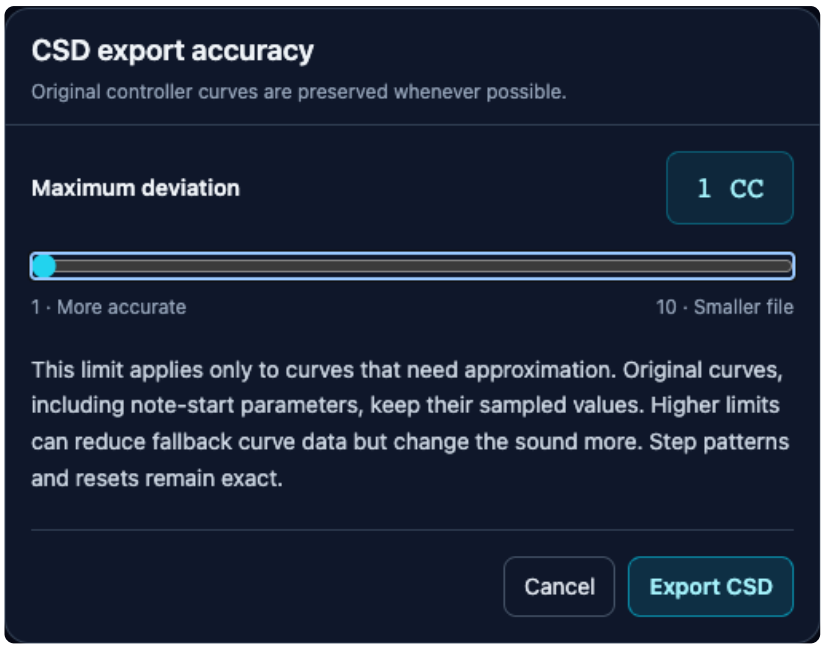

Any actual curve approximation is described in the exported warnings. MIDI export retains the captured MIDI events.

Which instruments are rendered

Native bundles retain the complete rack. Offline CSD exports select note instruments assigned to devices and continuous instruments used by routes, Master or inserts. Assigned devices count even when stopped or empty. Continuous generators with direct output are included; unused rack instances and their exclusive assets are omitted.

Bypassed inserts pass their inputs onward without processing. Their exclusive assets are omitted from the CSD package, while native bundles preserve them for later use. Direct output paths retain their documented Master bypass behavior.

Assets and export limits

Upload or import samples, SoundFonts and analysis files through Orchestron. Raw backend filesystem paths cannot be used as substitutes for bundled assets.

Offline exports must be finite. The range and its absolute end are limited to 65,536 transport steps, with at most 16 notes on one step and an event budget of 200,000. If an export exceeds a limit, shorten or simplify the exported arrangement. Large bundles may also reach the server's configured size limit.

Compatibility and recovery

Use a current Orchestron version to open performances using newer features. Older supported performances retain their music when imported; missing newer settings receive defaults. Review routing and the internal Master when opening older files. Incompatible arranger history may be discarded with a notice while the performance remains usable.

A raw sequencer snapshot is also accepted for legacy workflows, but it is not a self-contained instrument bundle. Prefer native exports for backups and transfer.

See saving and workspace recovery, instrument import/export, or the optional authoring tools. Developers can consult the export technical reference.

Navigation: Up | Previous | Next

Supplied Instruments and Effects

Navigation: Up | Previous | Next

Use these guides to explore ready-made sounds and processors. Each guide links to an instrument file and explains its controls. Import the file in **Instrument Design**, save it, then select it in Perform. The graph remains editable.

Choose a sound or processor

- Supplied Instruments
- Standard Effects

Add an instrument to a performance

A playable instrument creates sound from MIDI notes. Add it to the rack, assign an unused MIDI channel, and send notes from a piano roll or track on that channel. Route its Stereo Output to Master or through your chosen effects.

A continuous effect processes incoming audio while the instrument engine runs. It needs an audio source, not a note track. Use an insert to process one strip, or a send and return to share an effect across several sources. The Standard Effects guide explains both.

Some continuous instruments generate their own sound without an input. Read the individual guide before choosing its routing. Keep initial levels low enough to leave headroom when combining sounds.

Change and keep a sound

The rack's performance controls belong to that instance. Save Performance to retain those settings. To change the underlying synthesis or create a new sound, edit and save the graph in Instrument Design. See using performance controls.

Navigation: Up | Previous | Next

Supplied Instruments

Navigation: Up | Previous | Next

These instruments respond to MIDI notes. Import and save one, add it to a performance, then use a piano roll or sequencer on its assigned channel. Each guide gives the playing behavior and controls specific to that sound.

Instrument guides

- Physical Steel-String Guitar
- Ambient Techno Instruments
- HSB Mono
- Enough and Logan
- Goa Trance FM Lead

For processors that receive audio from another instrument, see Standard Effects.

More instrument examples

The Analog Drumkit has shared **BD snappiness** for notes 35/36, shared **Hi-hat tuning** for 42/46, and **Snare tuning** for 38. Each defaults to **1**, preserving the original sound. Side stick 37, clap 39, electric snare 40, toms 41/43/45/47/48/50, crash 49, ride 51 and tambourine 54 each have one individual characteristic control. New notes adopt changes.

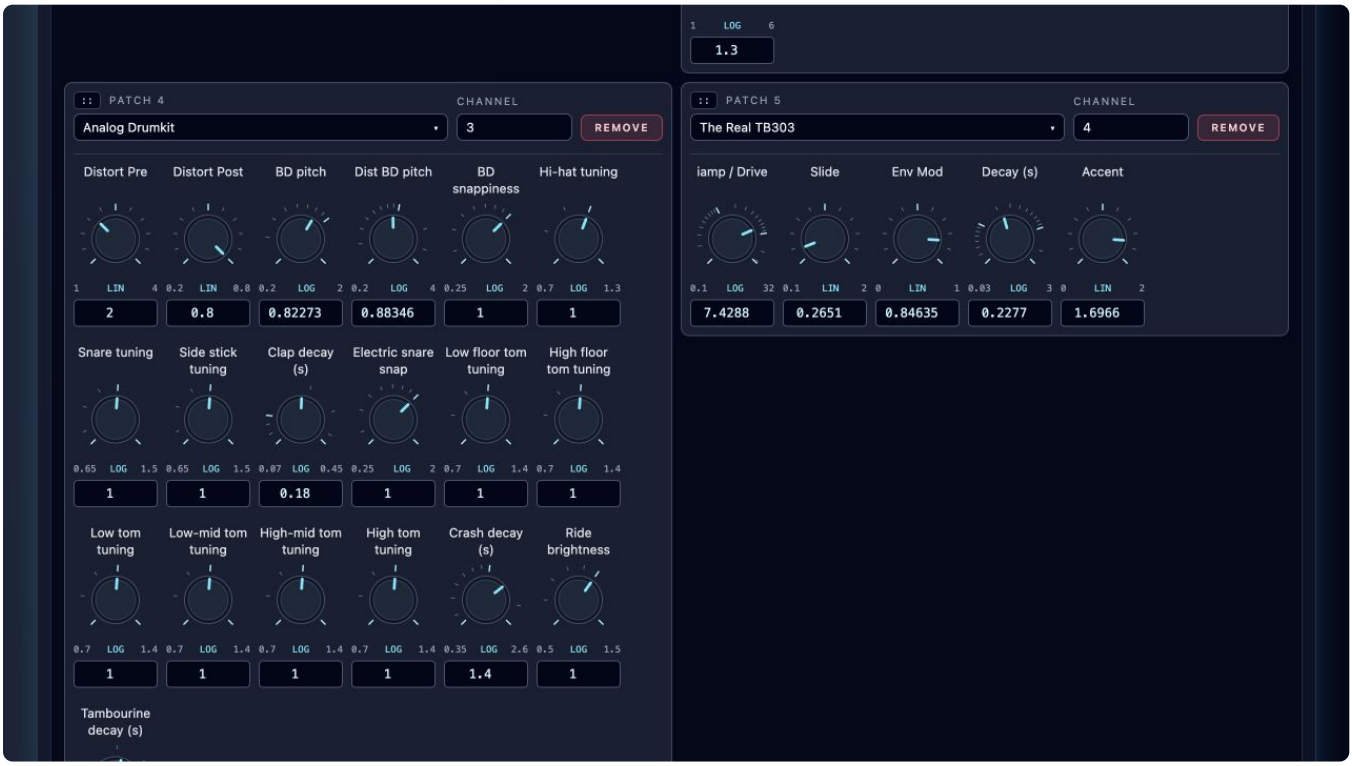

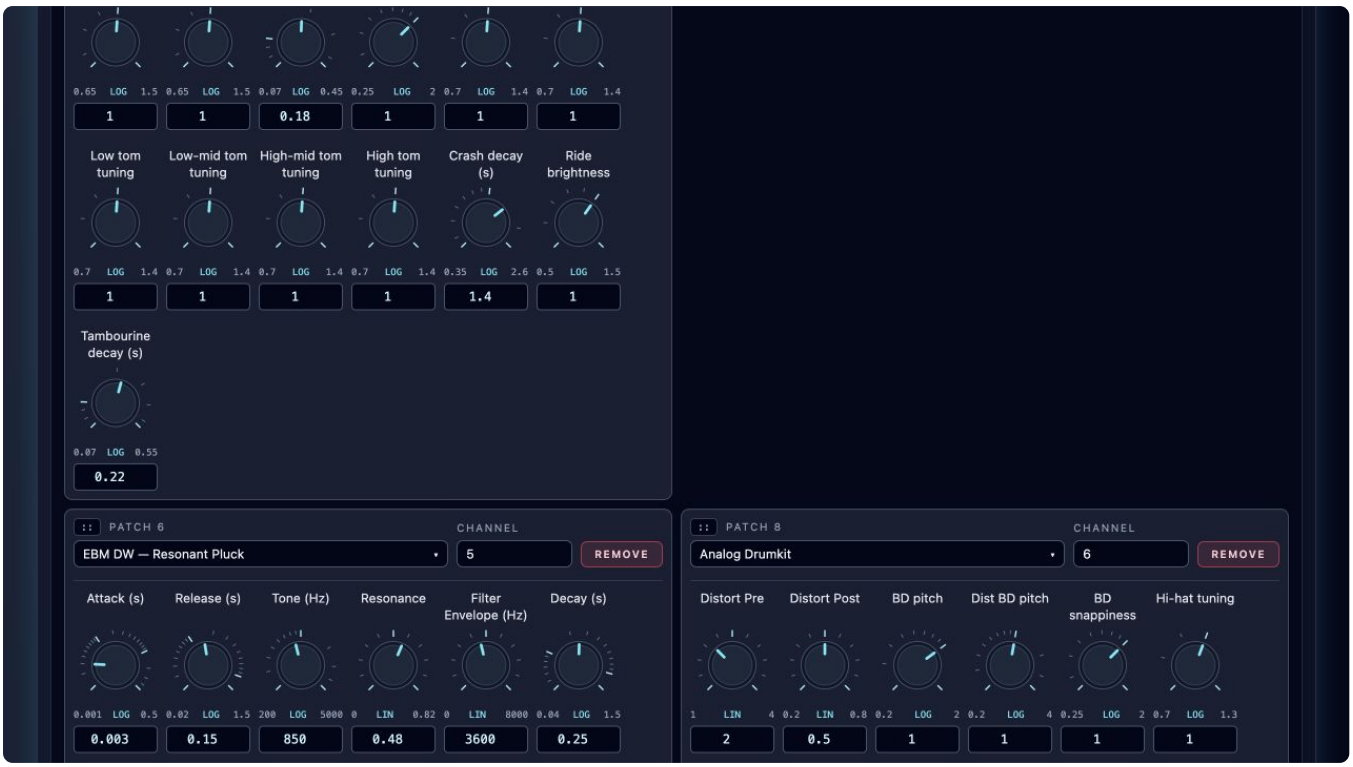

Granular Synthesis

Download Granular Synthesis. This instrument exposes 30 settings for its two `grain3` streams: harmonic source shaping, grain windows, density and duration, pitch and phase scatter, random distributions, internal motion, filters, ADSR and stereo width. Pan and level use the rack mixer. The controls are grouped into Source, Grains, Pitch, Grain Position, Motion, Stereo, Filter and Envelope. Settings apply to new notes; held notes continue their initialized internal modulation. Import its native bundle in Instrument Design, then add it to a performance to explore the grouped controls.

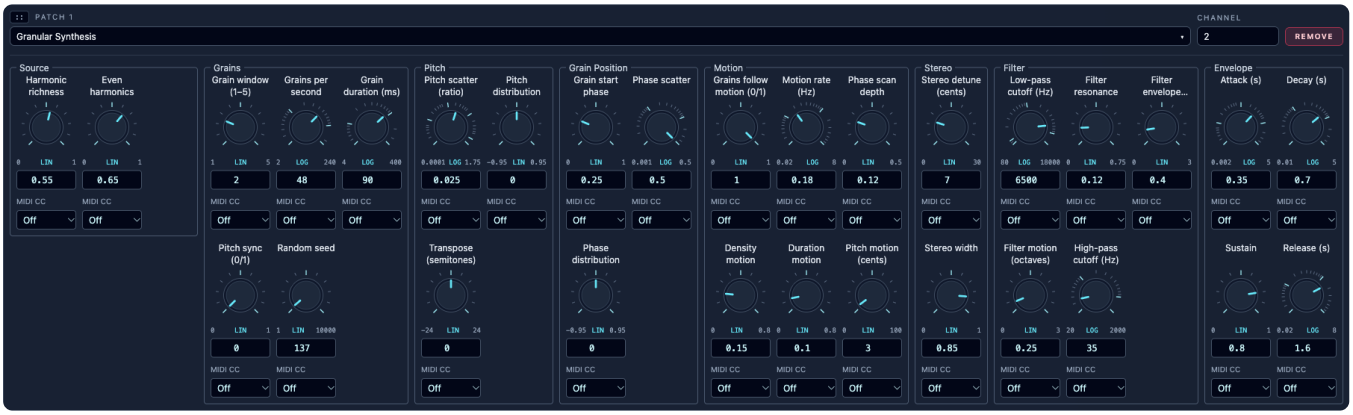

Minimoog Monophonic also groups its 40 controls by oscillator, modulation, noise, filter, envelopes and performance settings. Ambient Wavetable Cloud groups its 14 controls into Wavetable, Motion, Stereo, Envelope and Output & MIDI. These group names only organize the rack; sounds, controller IDs and saved settings remain unchanged.

Navigation: Up | Previous | Next

Physical Steel-String Guitar

Navigation: Up | Previous | Next

Steel String Guitar — Physical Model is a polyphonic Melody instrument. Each note has its own pluck and string state. Overlapping notes and manually timed strums ring independently. The two existing steel-string guitars remain available.

Import `Steel_String_Guitar_Physical_Model.orch.instrument.json` from the native instrument download, add it to Perform, and route its Stereo Output to Master. The output is centered and dry. Play E2–E6 for the intended register; use note lengths to control damping. There is no legato, sample playback, convolution, reverb, or pitch-transposition stage.

Playing controls

These are per-instance performance settings. Changes apply to new notes; held notes retain their settings.

Control	Range	Default	Effect
Finger → Plectrum	0–1	0.35	Sharper contact and more attack brightness toward Plectrum.
Brightness	0–1	0.55	Contact softness and high-frequency retention inside the string.
Body resonance	0–1	0.45	Amount of fixed wooden-body resonances.
String sustain	1–8×, logarithmic	3	Nominal string decay multiplier.
Release	0.05–0.8 s, logarithmic	0.18 s	Damping and final body fade after note-off.
Pluck position	0.08–0.35	0.18	Contact location as a fraction of string length.
Palm mute	0–1	0	Shorter, darker decay.
Natural variation	0–1	0.2	Small changes to contact, position and energy.

Velocity changes excitation energy, contact duration and brightness. Natural variation does not move note timing or wobble pitch. Chords and strums use the individual note timing you author.

Editable synthesis

The editable graph separates the pluck, string model, body resonance and output shaping. Clone the instrument before changing these blocks to keep the supplied sound available. The dry model leaves room for your own reverb or other mixer effects.

The string model technical reference describes its inputs, physical model, stability and measurements.

Graph and controls

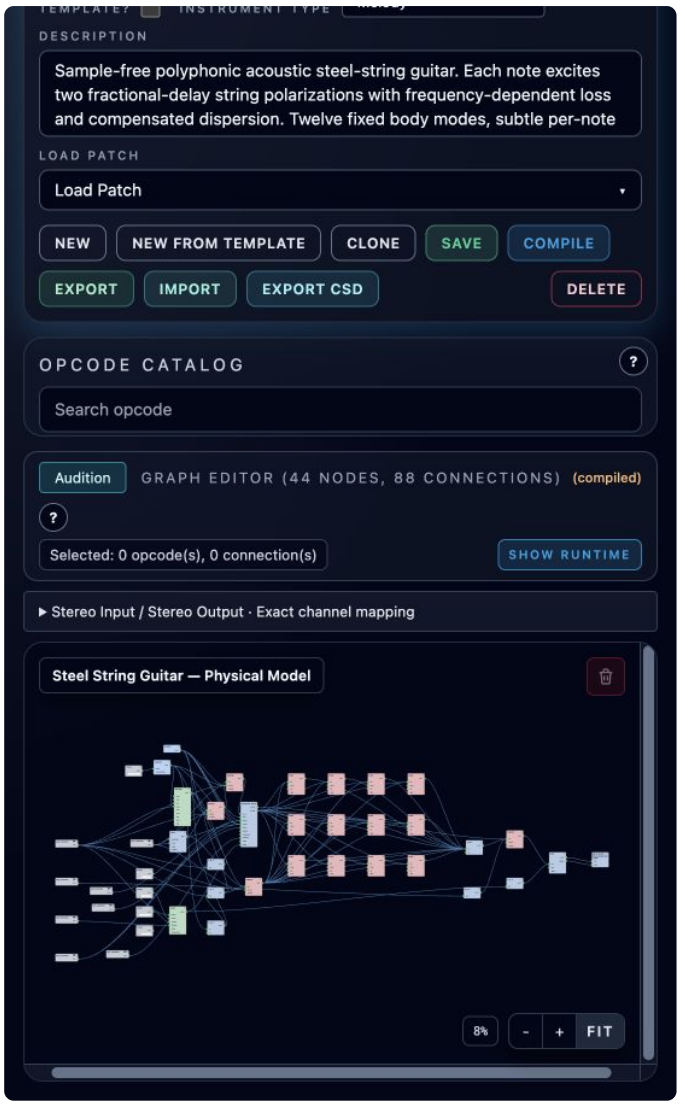
DESCRIPTION
Sample-free polyphonic acoustic steel-string guitar. Each note excites two fractional-delay string polarizations with frequency-dependent loss and compensated dispersion. Twelve fixed body modes, subtle per-note
LOAD PATCH
Load Patch
NEW
NEW FROM TEMPLATE
CLONE
SAVE
COMPILE
EXPORT
IMPORT
EXPORT CSD
DELETE
OPCODE CATALOG
Search opcode
Audition
GRAPH EDITOR (44 NODES, 88 CONNECTIONS)
(compiled)
Selected: 0 opcode(s), 0 connection(s)
SHOW RUNTIME
▸ Stereo Input / Stereo Output · Exact channel mapping
Steel String Guitar — Physical Model
8%
-
+
FIT

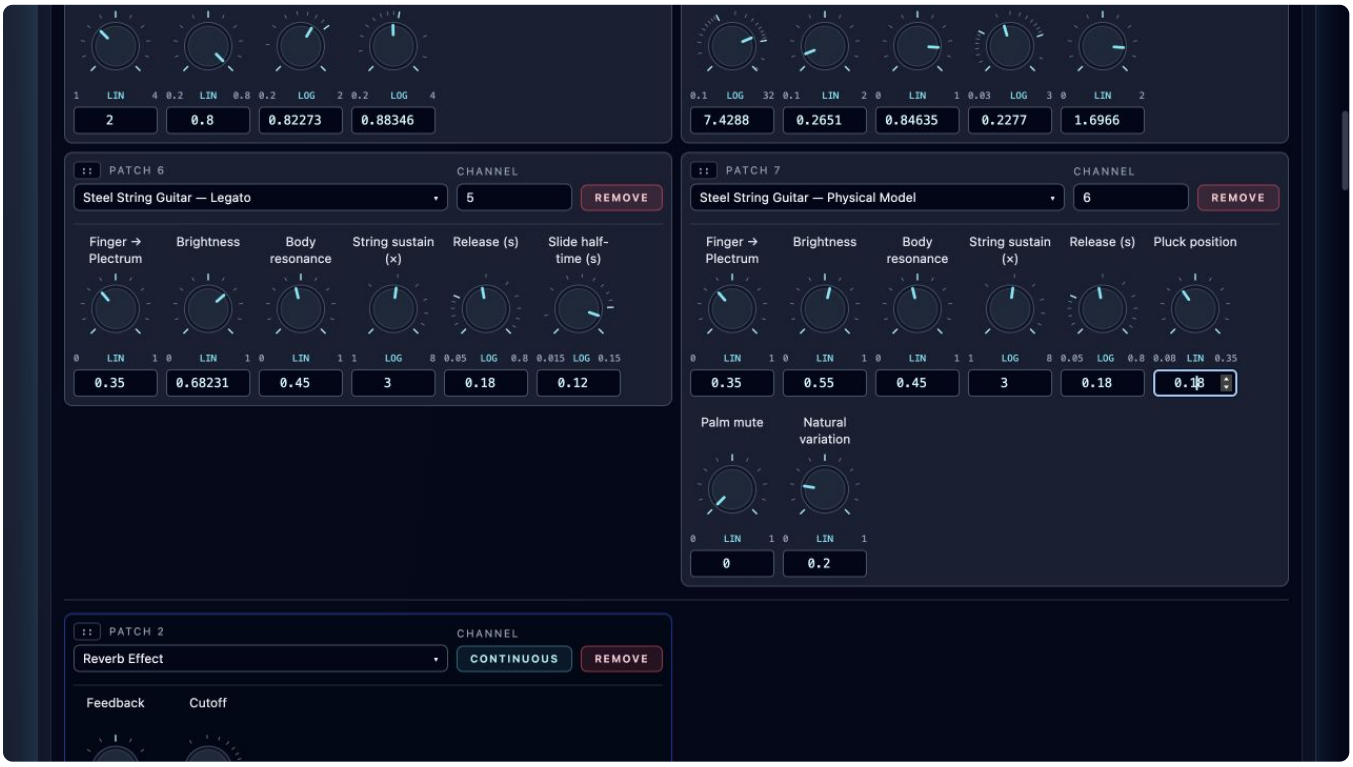
2
0.8
0.82273
0.88346
7.4288
0.2651
0.84635
0.2277
1.6966
PATCH 6
CHANNEL
Steel String Guitar — Legato
5
REMOVE
Finger → Plectrum
Brightness
Body resonance
String sustain (×)
Release (s)
Slide half-time (s)
0.35
0.68231
0.45
3
0.18
0.12
PATCH 7
CHANNEL
Steel String Guitar — Physical Model
6
REMOVE
Finger → Plectrum
Brightness
Body resonance
String sustain (×)
Release (s)
Pluck position
0.35
0.55
0.45
3
0.18
0.13
Palm mute
Natural variation
0
0.2
PATCH 2
CHANNEL
Reverb Effect
CONTINUOUS
REMOVE
Feedback
Cutoff

Verification and previews

Judge the sound in your own arrangement: dry instrument previews do not include a recorded guitar's microphone, room or overlapping strings. The technical reference above describes the measurements and comparisons.

References

- Csound mode: body resonators and Q.
- Csound deltap3: fractional delay taps.
- Csound setksmps: local sample processing.

Navigation: Up | Previous | Next

Ambient Techno Instruments

Navigation: Up | Previous | Next

Ambient Rounded Sub and **Ambient Wavetable Cloud** are dry, velocity-sensitive, polyphonic instruments. Import their native JSON files, assign note channels in Perform, and route **Stereo Output** through the mixer to **Master**.

- Ambient Rounded Sub: sine/triangle bass, gentle saturation, low-pass tone and optional amplitude pulsing. Suggested MIDI notes **24–55**. Its two output channels are identical.
- Ambient Wavetable Cloud: four morphing waveforms, three detuned voices, slow scanning and pitch drift, stereo width and an octave shimmer layer. Suggested MIDI notes **36–84**. Add shared rack delay/reverb for spacious ambience.

The cloud's shimmer is an octave oscillator layer. The instrument's sustained textures come from its oscillators and modulation. Waveforms range from a sine through warm and hollow spectra to bright upper harmonics. At high pitches the tables reduce upper harmonics; the octave layer fades away near its upper limit. The bass has stable pitch and no detuning. Neither instrument needs sample assets.

The authoring guide includes rebuild, library-save and audition commands. Native instrument exports, saved instance settings, and MIDI/SCORE CSD exports retain the controls.

Performance and MIDI controls

Performance settings apply when a new note starts. Existing notes retain their base settings; MIDI CCs reshape held notes and release tails with a **30 ms half-time**. Envelope and Output controls are performance-only.

MIDI Control Channel defaults to **16**, independently of the rack's note channel. The control-channel knob is rounded to the nearest whole channel. Set MIDI controller tracks to that channel and use **Reset Value 0**. A CC value of **0** preserves the performance setting. Higher values add intensity (this refers to the value of each assigned CC, not MIDI controller number 0). New notes never reset shared CC values. Instances using the same control channel share automation; choose a different channel to control copies independently.

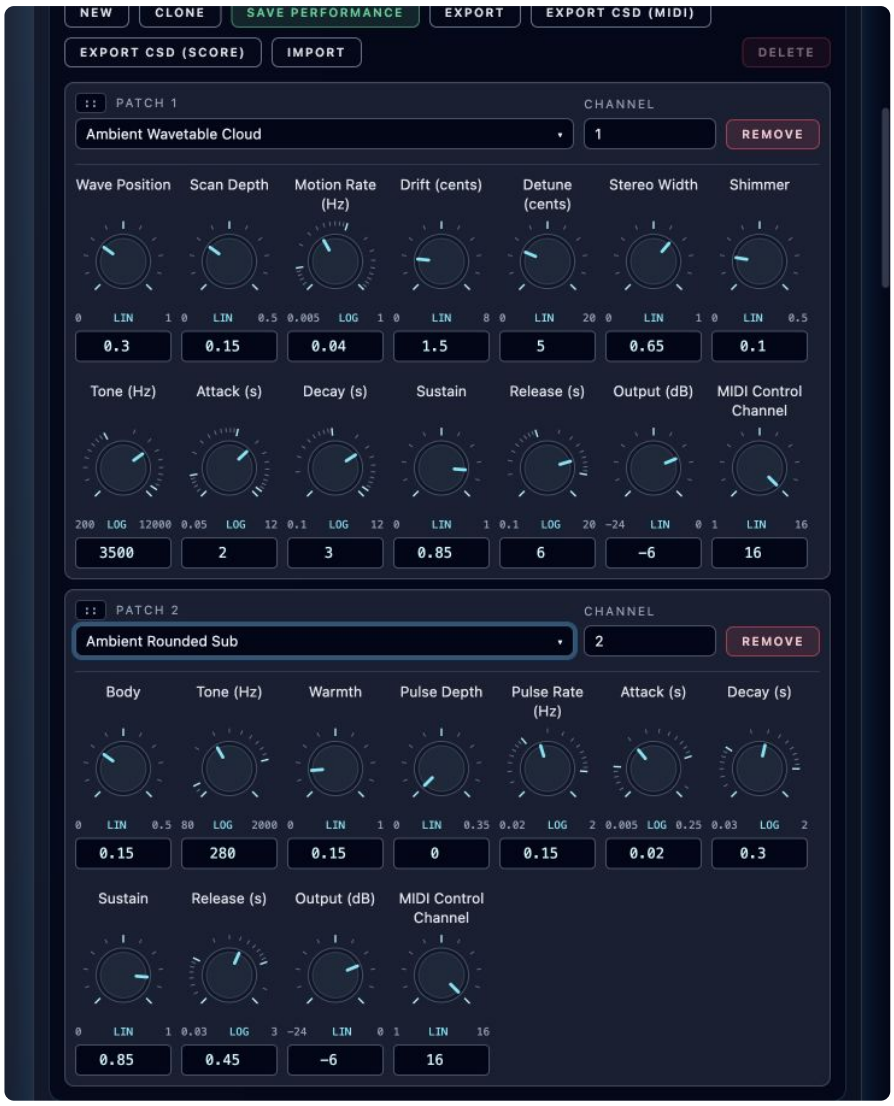

The rack above uses note channels 1 and 2 while both instruments listen for their distinct timbre CCs on control channel 16.

The mappings below leave existing Granular Capture and Stereo Moog Filter CC20–23 available.

Rounded Sub

Control	Performance range	Default	MIDI CC	At CC127
Body	0–0.5	0.15	24	Triangle blend 0.5
Tone	80–2,000 Hz	280 Hz	25	Base ×4, maximum 4,000 Hz
Warmth	0–1	0.15	26	Saturation amount 1
Pulse Depth	0–0.35	0	27	Amplitude modulation depth 0.35
Pulse Rate	0.02–2 Hz	0.15 Hz	28	Base ×4, maximum 4 Hz

Warmth blends in a normalized soft saturation curve; zero bypasses saturation. Body adds a small triangle component to the sine. Tone filters their combined sound. Pulse Depth modulates amplitude without changing pitch; Pulse Rate is in Hz and does not follow tempo.

Wavetable Cloud

The 14 performance controls use five rack groups: **Wavetable** (position, shimmer, tone), **Motion** (scan depth, rate, drift), **Stereo** (detune, width), **Envelope** (ADSR), and **Output & MIDI** (level, control channel).

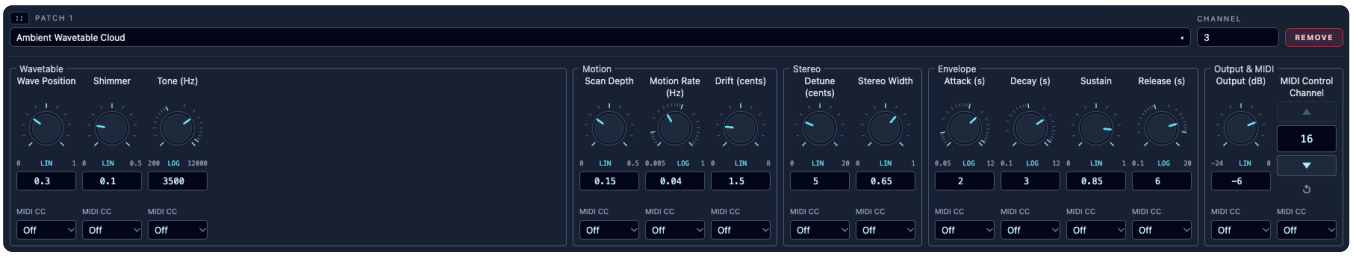

Control	Performance range	Default	MIDI CC	At CC127
Wave Position	0–1	0.3	102	Position 1
Scan Depth	0–0.5	0.15	103	Excursion 0.5
Motion Rate	0.005–1 Hz	0.04 Hz	104	Base ×4, maximum 2 Hz
Drift	0–8 cents	1.5 cents	105	Pitch drift 8 cents
Detune	0–20 cents	5 cents	106	Side-voice detuning ±20 cents
Stereo Width	0–1	0.65	107	Full width
Shimmer	0–0.5	0.1	108	Octave-layer mix 0.5
Tone	200–12,000 Hz	3,500 Hz	109	Base ×4, maximum 16,000 Hz

Wave Position selects the center of the waveform scan. Scan Depth moves around that center, clamped to 0–1. Separate slow oscillators vary each voice's scan and pitch, with different rate ratios to avoid a short repeating movement. Motion Rate controls their overall speed in Hz. Drift changes pitch over time; Detune sets the fixed spacing between voices. Stereo Width 0 centers all voices. Shimmer balances the fundamental voices with two voices one octave above.

For normalized MIDI value `u = CC / 127`, amount controls use `base + u × (maximum − base)`. Tone and rate use `base × 4^u`, capped as shown. CC64 therefore approximately doubles a frequency/rate. An amount already set to its maximum has no further upward travel.

Envelope and output

Control	Sub range; default	Cloud range; default
Attack	0.005–0.25 s; 0.02 s	0.05–12 s; 2 s
Decay	0.03–2 s; 0.3 s	0.1–12 s; 3 s
Sustain	0–1; 0.85	0–1; 0.85
Release	0.03–3 s; 0.45 s	0.1–20 s; 6 s
Output	–24–0 dB; –6 dB	–24–0 dB; –6 dB
MIDI Control Channel	1–16; 16	1–16; 16

Long cloud releases overlap subsequent chords. Set the mixer level for the number of simultaneous notes and any external effects.

Starting sounds

Start from defaults and keep all assigned CCs at zero for these recipes:

Sound	Settings
Pure sub	Body 0, Warmth 0, Tone 2,000 Hz
Warm pulse	Body 0.35, Warmth 0.65, Pulse Depth 0.25, Pulse Rate 1.2 Hz
Dark drone	Wave Position 0.08, Tone 700 Hz, Shimmer 0, Motion Rate 0.015 Hz
Shimmer	Wave Position 0.75, Shimmer 0.4, Tone 10,000 Hz, Stereo Width 0.9
Drifting cloud	Scan Depth 0.4, Motion Rate 0.12 Hz, Drift 6 cents, Detune 12 cents, Stereo Width 1

Navigation: Up | Previous | Next

HSB Mono

Navigation: Up | Previous | Next

Import HSB Mono, add it to the performance rack, and route **Stereo Output → Master**. Play MIDI notes 36–84 for the intended bass/lead range. The sound blends five sine-wave octaves.

Performance parameter	Range	Default	MIDI CC
Tonality (octave cycle)	0–1	0	20
Brightness (octaves)	–2 to +2	0	74
Legato half-time (s)	0.001–0.5, logarithmic	0.060	5
MIDI control channel	1–16	16	—

Send CCs on the selected control channel, independently of the note channel. For each mapped controller, **value 0 recalls the rack setting**; values 1–127 span its full range. Use controller-sequencer **Reset Value 0**. CC changes have 10 ms smoothing and affect held notes; rack settings refresh at phrase start. Instances sharing a control channel share MIDI automation.

Touching or overlapping notes share one voice and envelope; the newest held note wins. Gaps up to approximately **20 ms** retain pitch glide, while longer gaps start directly at the new pitch. A positive gap restarts the amplitude envelope. Larger half-times produce slower slides. The 20 ms window is fixed in this patch and quantized to the control block. Existing monophonic support and The Real TB303 are unchanged.

Tonality is cyclic: 0 and 1 give the same pitch class. Brightness moves the spectral weighting across octaves. The amplitude envelope is 6 ms / 150 ms / 80% / 80 ms and responds to note velocity.

Authoring details and reproducible previews.

Navigation: Up | Previous | Next

Enough and Logan

Navigation: Up | Previous | Next

Enough and Logan are dry, polyphonic Melody instruments adapted from joel / ebmtranceboy's altamax Csound project. Import Enough or Logan, add it to the Perform rack, and route its Stereo Output to Master. Both graphs remain editable.

Enough

Five sine oscillators form an evolving FM network. The 20× operator modulates the 4× operator; that signal and a separate 9× operator modulate a 3× operator, which modulates the final 1× carrier. Modulation depths recede during a held note, leaving a pure sine after two seconds at the default Evolution setting.

This graph uses ordinary `oscil3` nodes. It intentionally omits the source's tiny 0.0001 self-feedback term. Cubic-interpolated oscillators also differ from the source's sample-by-sample phase advancement and rounded table lookup. This is an editable adaptation of the original algorithm, not a sample-identical recreation.

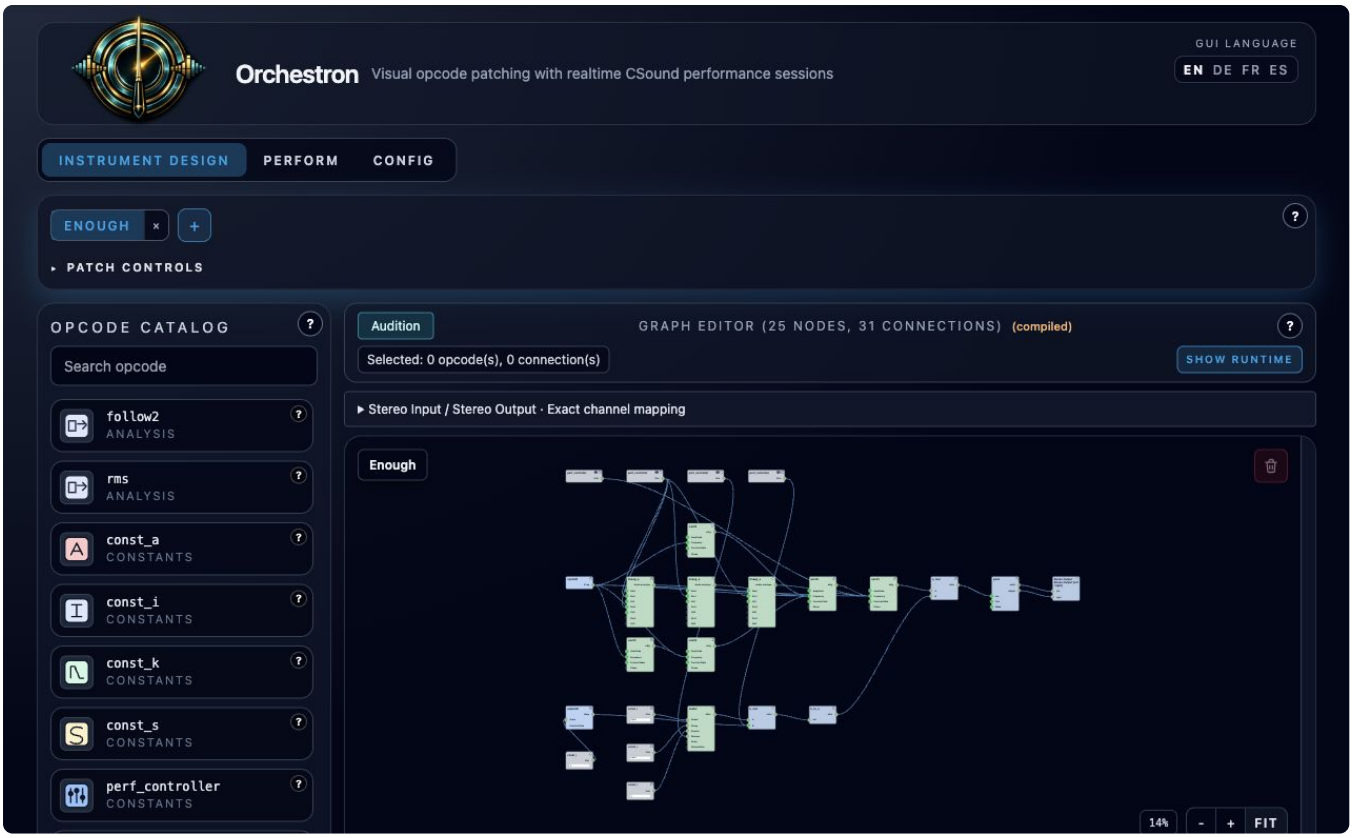

Logan

Two audio-rate phase ramps drive mirrored and limited waveform shapes. Their combination alternates between two cubic-shaped waveforms at the fundamental frequency. The shape moves from 0.93 to 0.03 over five seconds by default. The phase-shaping network retains the source equations.

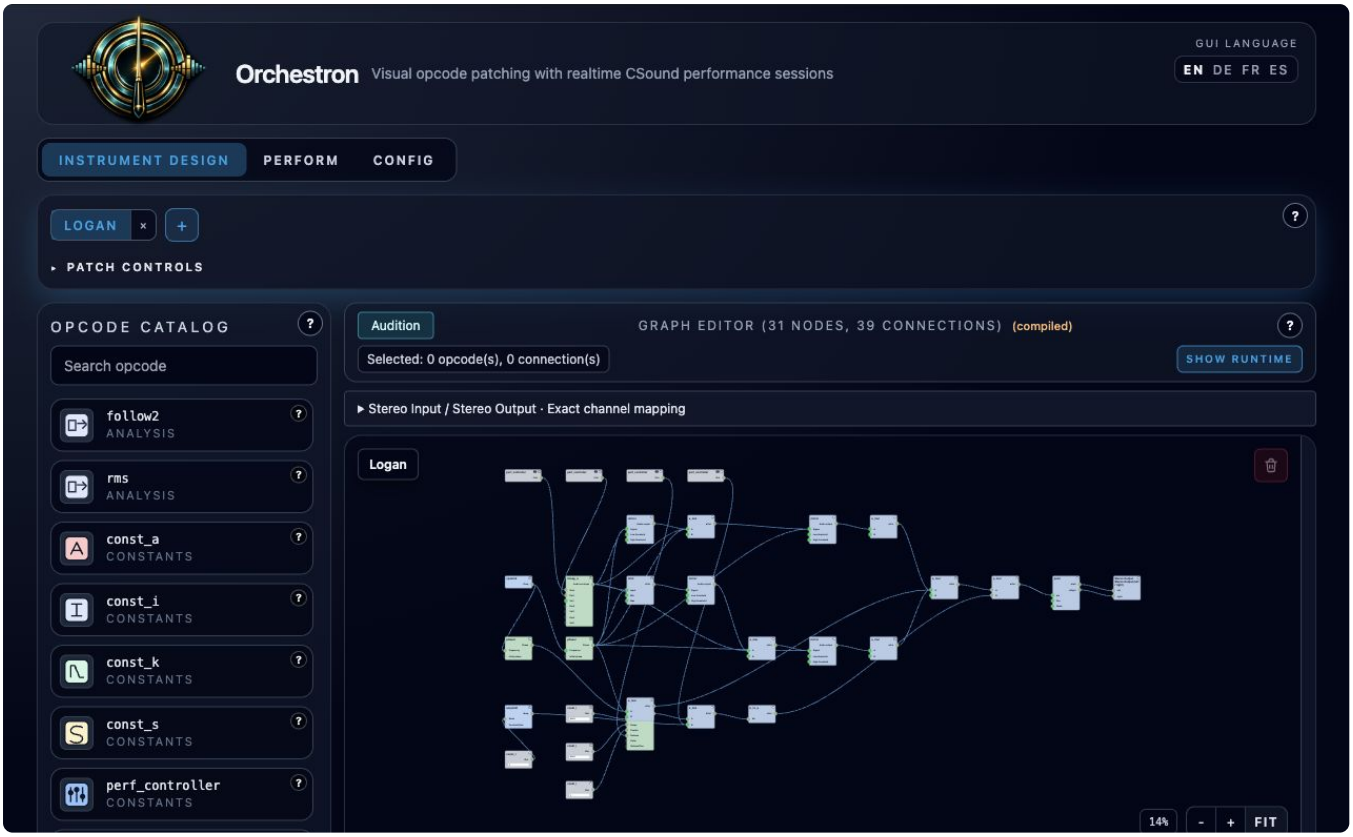

Performance controls

Control	Enough	Logan
Timbre	Final FM-depth multiplier, 0–2; default 1	Initial shape, 0.05–0.97; default 0.93
Evolution	Multiplies modulation-envelope times, 0.25–4; default 1	Shape sweep duration, 0.1–20 seconds; default 5
Release	0.01–3 seconds; default 0.2	0.01–3 seconds; default 0.5
Level	Per-channel gain, 0–1; default 0.5	Per-channel gain, 0–0.25; default 0.06

Settings are read when each note starts. Held notes retain their settings. Velocity changes final loudness without changing modulation depth. Both instruments add a 1 ms MIDI attack and decay with full sustain, and use a release-sensitive envelope. Centered stereo output preserves the source's per-channel gain; chord levels add together, so reduce Level or the mixer strip for dense chords. The source project's delay and reverb are not included.

Reusable building blocks

The catalog also provides `phasor` (an audio-rate phase ramp), `mirror` (reflection between thresholds), and `linseg_a` (an audio-rate linear envelope). Existing `linseg` nodes retain their control-rate output. `linseg_a` emits native Csound `linseg`; its final value holds after the last segment. Omit optional duration and endpoint pairs together.

Navigation: Up | Previous | Next

Goa Trance FM Lead

Navigation: Up | Previous | Next

Import the Goa Trance FM Lead instrument, save it in Instrument Design, then add it to Perform. Assign a MIDI channel and route its output through the mixer to Master. Use a piano roll or melodic track on that channel to play it.

Play and shape the sound

The instrument combines cross-FM synthesis with two resonant high-pass filters. Start with its default performance settings and a modest mixer level. Try single notes before adding chords, then adjust the supplied controls while comparing the sound. Save Performance to keep the settings for that rack instance.

The filters track pitch and velocity differently in the two channels. Resonance is limited automatically at settings that would otherwise make a filter unstable. The requested control value is retained, but effective resonance can stop increasing as you move higher in the register. The guard is evaluated for each new note.

Resonant gain can still exceed full scale. Leave mixer headroom, especially for high notes or chords. The guard prevents the known filter runaway; it does not act as an output limiter.

Edit the instrument

The graph remains editable. If you change the filter structure or oversampling, review the stability protection as part of that change. The filter stability technical reference preserves the derivation, guard formulas and reproduction procedure.

Navigation: Up | Previous | Next

Standard Effects

Navigation: Up | Previous | Next

These supplied stereo effects can process individual instrument strips or the combined mix as Master inserts. They can also run as separate rack instruments receiving effect sends. Import and save an effect in Instrument Design to make it available in Perform.

Effect Guides

- Tube Overdrive
- Stereo Delay
- Stereo Flanger
- Autopanner
- Parametric Equalizer
- Stereo Widener, Bit Crusher, and Harmonizer
- Granular Capture: Stutter, Freeze and Grains
- Stereo Moog Filter
- Stereo Auto Wah Wah
- Trance Gate

Inserts and Sends

An **insert** processes a strip's entire signal. Select the instrument strip or Master, choose an effect under **Inserts…**, and click **Add**. Use the effect's own Mix control for dry/wet balance.

An **effect send** adds a parallel path to a separate effect instrument, called a **return**. Several instruments can share it, each with its own send level. See Audio Mixer and Routing for the routing controls.

Add a Shared Send Effect

1. Stop the instruments before changing the rack or routing. In **Perform → Add Instrument**, select the saved effect. It appears as a continuous instrument with its own mixer strip.
2. Connect the effect's output to **Master**. Keep each source instrument's main output connected to Master as well, so its original signal remains audible.
3. In **Mixer → Audio routing**, select the instrument under **Source** and the effect under **Destination**. Leave **Stereo** checked, verify the channel mapping, choose **Send** in the Route selector, and click **Add**.
4. On the source strip, adjust **Send → [effect name]** to control how much signal reaches the effect. New sends start at silence; try **–12 dB** initially. **Post-fader** follows the source's volume and balance; **Pre-fader** is independent of those controls. Repeat for other instruments that should share the effect.
5. Start the instruments, balance the send levels and the return's fader while listening, then **Save Performance**. Existing send levels remain adjustable during playback.

For a parallel return, use **wet-only** output where supported, so the original signal is not added twice. For example, set Stereo Flanger's **Mix to 1**. The supplied Stereo Delay includes dry audio, and Granular Capture passes dry audio while inactive; these are generally better used as inserts. Each guide below explains the effect's routing and mix behavior.

Stereo Chorus Depth

Download Stereo Chorus and add it as a mixer insert. **Amount** ranges from `0.001` to `0.1` seconds (default `0.02`); **LFO Frequency Multiplier** ranges from `0.1` to `10` (default `1`). Higher Amount increases modulation depth and wet delay. Check each control's **Live** or **Requires restart** label.

Update older copies before using large Amount values, then restart the rack. The technical notes explain the corrected delay range.

Navigation: Up | Previous | Next

Tube Overdrive

Navigation: Up | Previous | Next

Download Tube Overdrive and import it using **Instrument Design → Import**. This separate continuous instrument adds warm stereo saturation with standard Csound opcodes. The original **Overdrive** example remains available for comparison.

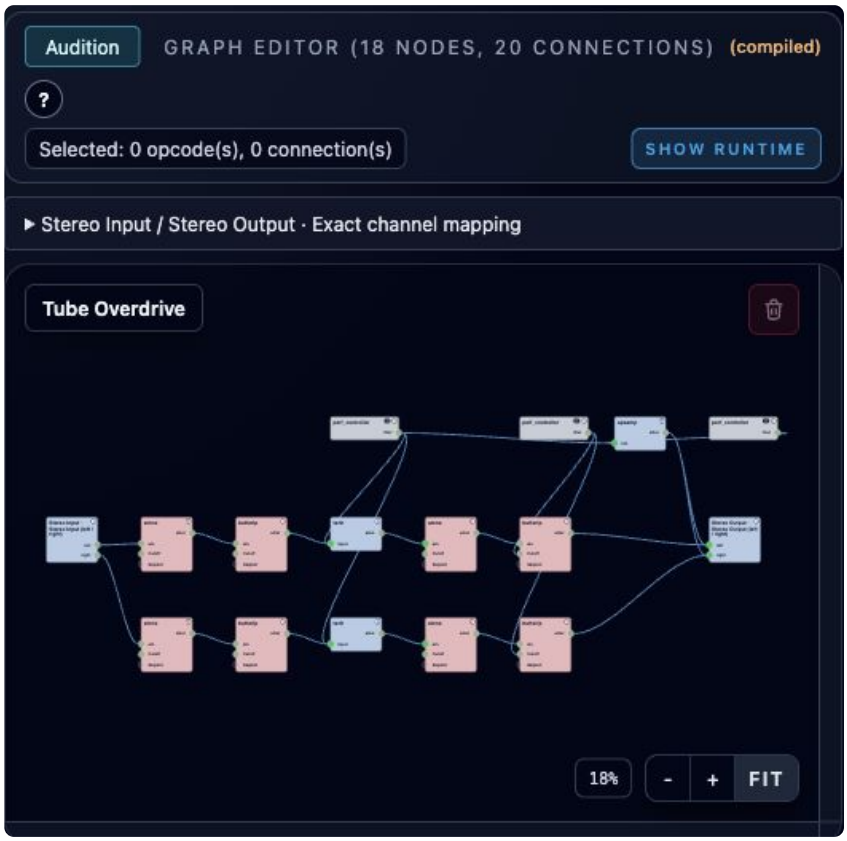

Use as an insert

1. Import the instrument and open **Perform** with the instruments stopped.
2. Select the source strip in **Mixer**, choose **Tube Overdrive** in **Inserts…**, and click **Add**. It can also process the combined signal as a Master insert.
3. Set its performance controls, then start the instruments. Controls marked **Live** respond during playback; restart the rack for controls marked **Requires restart**.
4. Adjust Output to match the untreated level when judging the sound. Save Performance to retain that instance's settings.

The stereo input and output use the exact channel names `left` and `right`. The output is fully processed, with independent left/right signal paths. If used as a return instead, mix the dry and processed signals using the routing and send levels; the patch has no internal dry/wet control.

Controls

Controller	Minimum	Maximum	Default	Scale
Drive (dB)	0	24	12	Linear
Tone (Hz)	2,000	12,000	6,500	Logarithmic
Output (dB)	–18	+6	–3	Linear

Drive increases gain before saturation. Start around 3–6 dB for subtle color, 12 dB for moderate breakup, and 18–24 dB for stronger overdrive. Input level also determines how hard the effect is driven. **0 dB Drive is not bypass.**

Tone sets the output low-pass cutoff: lower values soften the highs; higher values retain more bite. **Output** trims the processed level. Partial gain compensation reduces level changes as Drive increases, but does not equalize perceived loudness. Positive output gain can exceed full scale; lower Output or the downstream mixer gain if the meter clips.

These are per-instance performance settings, not MIDI CC automation. Their stable IDs are `tube_drive_db`, `tube_tone_hz`, and `tube_output_db`. Native performance bundles and both CSD export modes retain the selected settings.

How the graph works

Both channels follow the same chain:

```
Stereo Input → atone (20 Hz) → butterlp (12 kHz)
             → biased tanh → atone (10 Hz) → butterlp (Tone)
             → compensated Output gain → Stereo Output
```

The input filters remove subsonic energy and reduce very bright material entering the nonlinear stage. The shaping formula is:

```
aShape = tanh(aPre * ampdb(iDriveDb) + 0.15)
aAC atone aShape - 0.148885033623318, 10
aTone butterlp aAC, iToneHz
aOutput = aTone * ampdb(iOutputDb - iDriveDb / 2)
```

The small bias makes the curve asymmetric, adding even harmonics alongside odd harmonics. Subtracting `tanh(0.15)` keeps silence at zero; the following high-pass removes DC generated by the signal. One shared `upsamp` node converts the computed gain to audio rate for the output formulas. All processing uses existing opcodes and input formulas, so the graph remains editable.

This is a tube-style coloration effect, not a model of a particular amplifier. It has no cabinet simulation, power-supply sag, or oversampling. Filtering does not remove aliasing already generated by the nonlinear stage. High notes and strong Drive can expose this limitation; use the bright-signal comparisons in the example and validation guide to judge the sound for your source.

References: Csound tanh, tube-distortion techniques.

Navigation: Up | Previous | Next

Stereo Delay

Navigation: Up | Previous | Next

Download Stereo Delay and import it using **Instrument Design → Import**. This continuous stereo effect adds tempo-synced echoes while returning the original left and right channels at unity gain.

Use as an insert

1. With instruments stopped, open **Perform → Mixer**.
2. On a source strip or Master, open **Inserts…**, select **Stereo Delay**, and click **Add**.
3. Set Tempo to your performance's BPM and Beats to the desired echo interval.
4. Start the instruments. Controls marked **Live** respond during playback; restart the rack for settings marked **Requires restart**.
5. Save Performance to retain that instance's settings.

The named stereo input and output channels are `left` and `right`. The effect already includes dry audio; using it as a parallel return also adds that dry signal to the mix. It has no dry/wet control or limiter. Overlapping echoes can exceed full scale, especially with high Feedback; allow headroom in the source or downstream mixer.

Controls

Controller	Minimum	Maximum	Default	Scale
Tempo (BPM)	40	180	120	Linear
Beats	1	16	4	Linear
Feedback	0	0.9	0.5	Linear
Stereo	0	1	0.5	Linear

Tempo and Beats set the interval to `60 × Beats / Tempo` seconds, rounded to the nearest audio sample. At 120 BPM, Beats 1 means 0.5 seconds and Beats 4 means 2 seconds. Fractional Beats are supported. The full range is approximately 0.333–24 seconds. Beats counts quarter-note beats, independent of the sequencer's meter. Enter Tempo manually; this instrument does not read sequencer tempo.

Feedback scales every echo. At 0.5, the first echo is multiplied by 0.5, the second by 0.25, the third by 0.125, and so on. Zero produces only the unchanged dry signal. This gain is applied before the Stereo blend.

Stereo controls echoes only. At zero, every echo preserves its original left/right channels. At one, each echo uses the average `(left + right) / 2`: the first goes entirely left, the second entirely right, then left again. Intermediate values linearly blend these behaviors. Averaging keeps identical left/right input from doubling in level; a one-sided signal becomes half as large in the fully panned branch, and opposite-phase channels can cancel there.

Native bundles and both CSD export modes retain the selected instance values.

Edit the sound

Open the saved instrument in Instrument Design to explore its delay paths and performance controls. Clone it before changing the graph if you want to keep the supplied version.

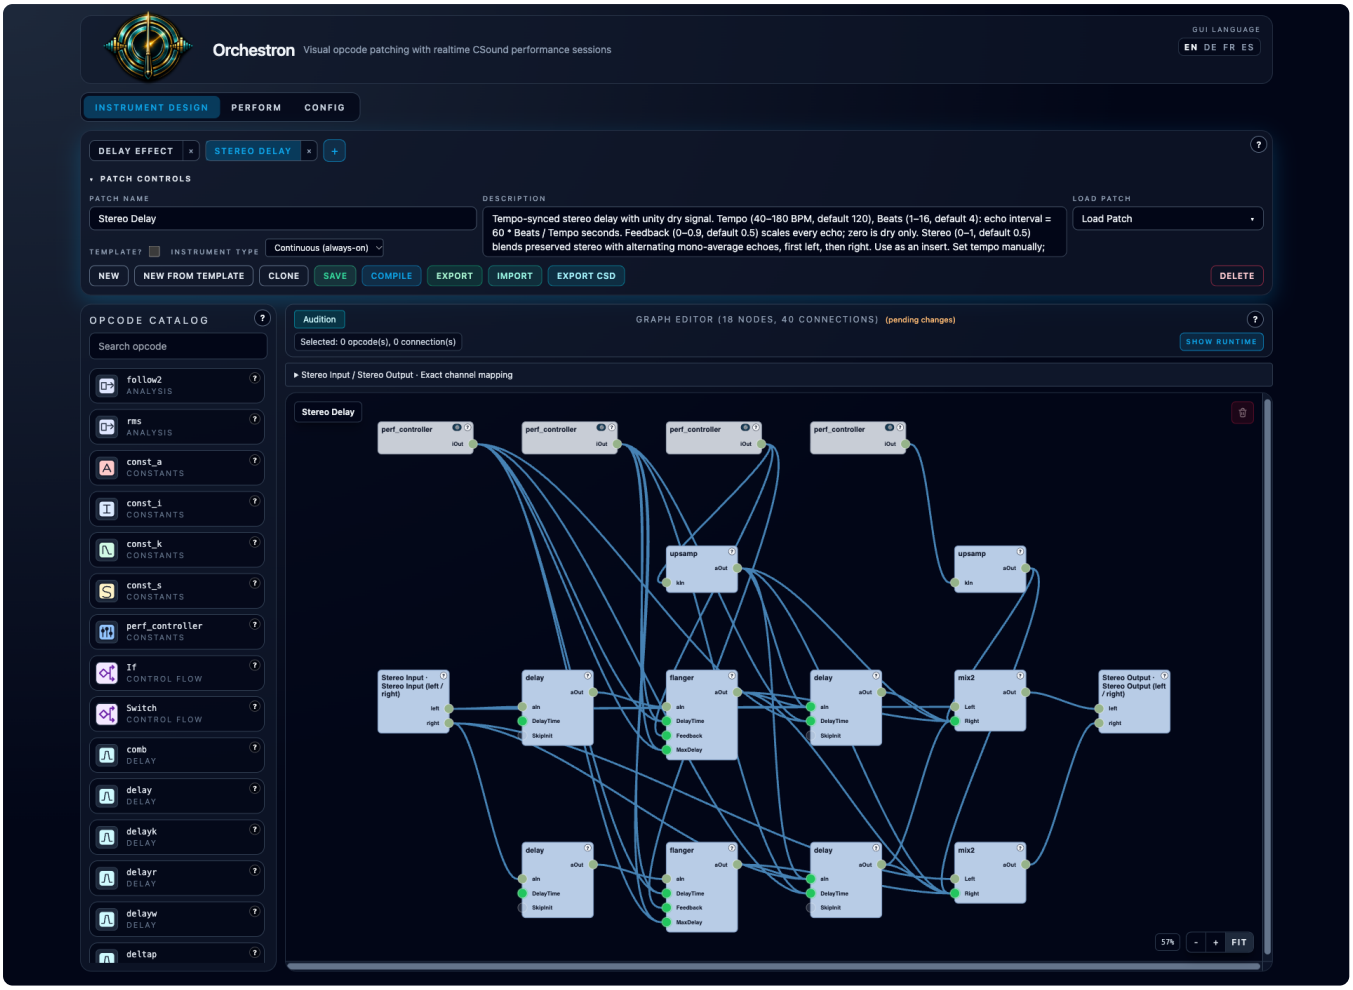

The graph uses existing `delay`, `flanger`, `upsamp`, `mix2`, `inleta`, `outleta` and `perf_controller` nodes. See the example and validation guide for graph details. Long intervals with high Feedback need long render tails; the instrument does not extend the export duration automatically.

References: Csound flanger, Csound delay.

Navigation: Up | Previous | Next

Stereo Flanger

Navigation: Up | Previous | Next

Download Stereo Flanger and import it with **Instrument Design → Import**. This continuous effect mixes each channel with a short, moving delay to produce sweeping comb-filter notches. Left and right retain separate audio paths; Stereo Phase offsets their sweeps.

Use in a performance

1. With instruments stopped, open **Perform → Mixer**.
2. On an instrument strip or Master, open **Inserts…**, select **Stereo Flanger**, and click **Add**.
3. Start with the defaults, then adjust the six performance controllers.
4. Controls marked **Live** respond during playback. Restart the rack for settings marked **Requires restart**.
5. Save Performance to retain that instance's settings.

For a shared send effect, set Mix to **1** so the return contains only wet audio. Stereo Input and Stereo Output both expose `left` and `right` channels. The effect needs routed audio and does not use MIDI notes.

Controls

Controller	Range	Default	Scale
Rate (Hz)	0.02–5	0.25	Logarithmic
Delay (ms)	0.2–5	1	Linear
Depth (ms)	0–5	2	Linear
Feedback	–0.9–0.9	0.5	Linear
Stereo Phase (deg)	0–180	90	Linear
Mix	0–1	0.5	Linear

Rate sets the sweep speed in cycles per second. The default completes one cycle every four seconds; tempo synchronization is manual.

Delay is the shortest delay. **Depth** is the sine wave's amplitude, with the same amount added as a positive offset. The delay therefore sweeps from `Delay` to `Delay + 2 × Depth`: **1–5 ms** at the defaults. Depth zero produces a static comb filter. Across all settings, delays stay between **0.2 and 15 ms** inside a 20 ms buffer.

Feedback sharpens the resonances. Negative values invert successive repeats, changing the comb's character. Zero removes recirculation but keeps the moving delay. Wet gain is multiplied by `1 − abs(Feedback)` to compensate for feedback buildup; increasing feedback can reduce average loudness while emphasizing the resonant frequencies.

Stereo Phase sets the right sweep's phase relative to the left. Zero gives matching sweeps; 90 degrees gives a quarter-cycle offset; 180 degrees makes them sweep in opposite directions. It does not cross-feed audio between channels.

Mix crossfades linearly from untouched dry audio at zero to compensated wet audio at one. Intermediate settings produce the familiar flanging interference; wet-only settings emphasize pitch modulation and feedback coloration.

Edit the sound

The graph uses two audio-rate `poscil3` sine oscillators, two `flanger` delay lines, and ordinary stereo input/output blocks. Stable controller IDs are `sf_rate`, `sf_delay`, `sf_depth`, `sf_feedback`, `sf_phase`, and `sf_mix`. Native bundles and both CSD modes preserve the graph and instance settings.

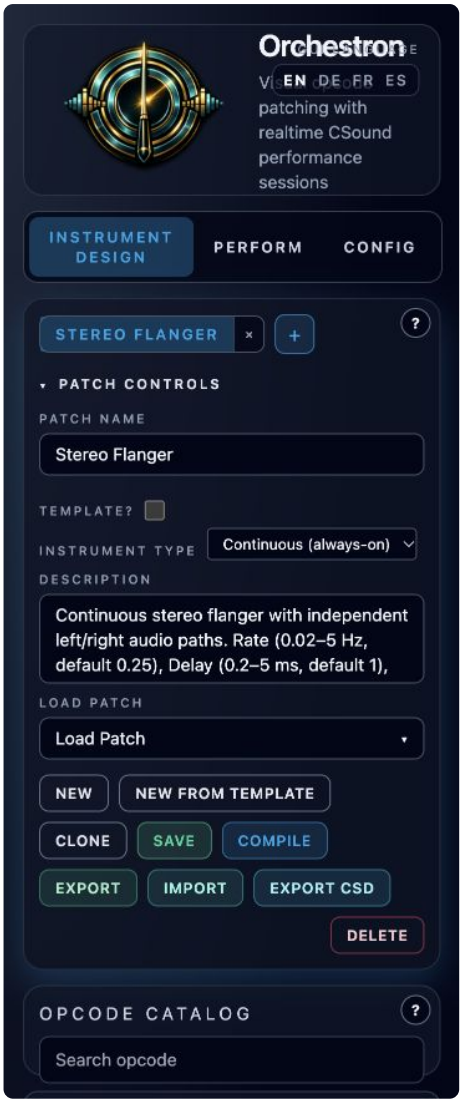

See the example guide for the editable patch, reproducible previews, and regression checks.

References: Csound flanger, Csound poscil3.

Navigation: Up | Previous | Next

Autopanner

Navigation: Up | Previous | Next

Download Autopanner and import it with **Instrument Design → Import**. This Continuous stereo effect uses a smooth sine wave to move the balance between left and right. Each input stays in its original channel; the effect never sums or cross-feeds them.

Use as an insert

1. With instruments stopped, open **Perform → Mixer**.
2. On a source strip or Master, open **Inserts…**, select **Autopanner**, and click **Add**.
3. Enter your performance's tempo and the number of quarter-note beats per full cycle.
4. Start the instruments. Controls marked **Live** respond during playback. Restart the rack for settings marked **Requires restart**.
5. Save Performance to retain that instance's values.

Use an insert for the full sweep. A parallel send leaves the original signal audible alongside the panned return and reduces the apparent movement. The grouped Stereo Input and Stereo Output both expose `left` and `right`. The instrument needs routed audio, not MIDI notes.

Controls

Controller	Range	Default	Scale
Tempo (BPM)	40–180	120	Linear
Beats per cycle	0.125–16	4	Linear
Depth	0–1	1	Linear

One complete cycle lasts `60 × Beats per cycle / Tempo` seconds. At 120 BPM, four beats take two seconds; one beat takes half a second. Fractional beats are supported. Beats count quarter notes regardless of the sequencer's meter. The full range is approximately 0.0417–24 seconds per cycle.

Tempo is manual. The effect does not read performance tempo, follow arranger position, or align to its bars. Starting the effect resets its phase to the center, moving toward the right first. Arranger seeks do not reset this cycle.

Depth 0 passes both channels through unchanged. **Depth 1** alternately attenuates each channel to silence. At the center of a full-depth sweep, each channel has gain 0.5 (about –6 dB). Each channel's gain stays between zero and one; the effect adds no gain, latency, feedback or tail. A signal present only on the left stays on the left and fades in and out; it is not transferred to the right. For identical left/right input, the summed amplitude stays constant throughout the sweep at a fixed Depth, although perceived loudness can vary.

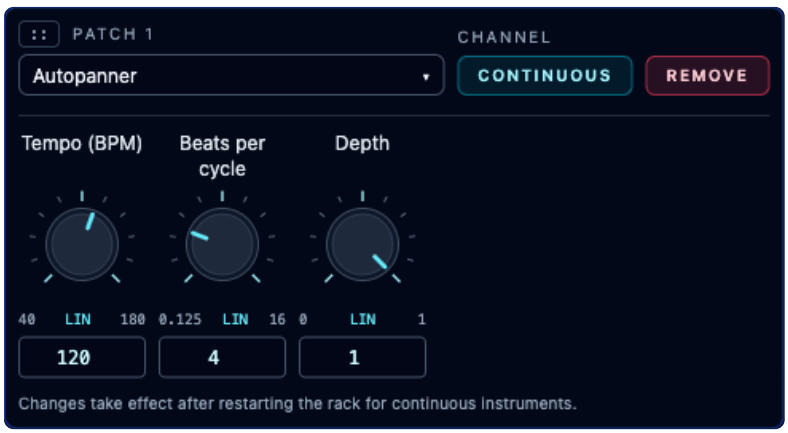

Editable graph and exports

With `sine = sin(2π × time × Tempo / (60 × Beats))` , the graph computes:

- `outputLeft = inputLeft × [1 − Depth × (1 + sine) / 2]`
- `outputRight = inputRight × [1 − Depth × (1 − sine) / 2]`

One audio-rate `poscil3` oscillator feeds separate audio multipliers. An `upsamp` node promotes the fixed Depth setting for audio formulas. The graph uses existing opcodes and ordinary grouped inlets/outlets. Stable controller IDs are `ap_tempo` , `ap_beats` and `ap_depth` . Native bundles and both CSD export modes preserve the graph and each instance's settings.

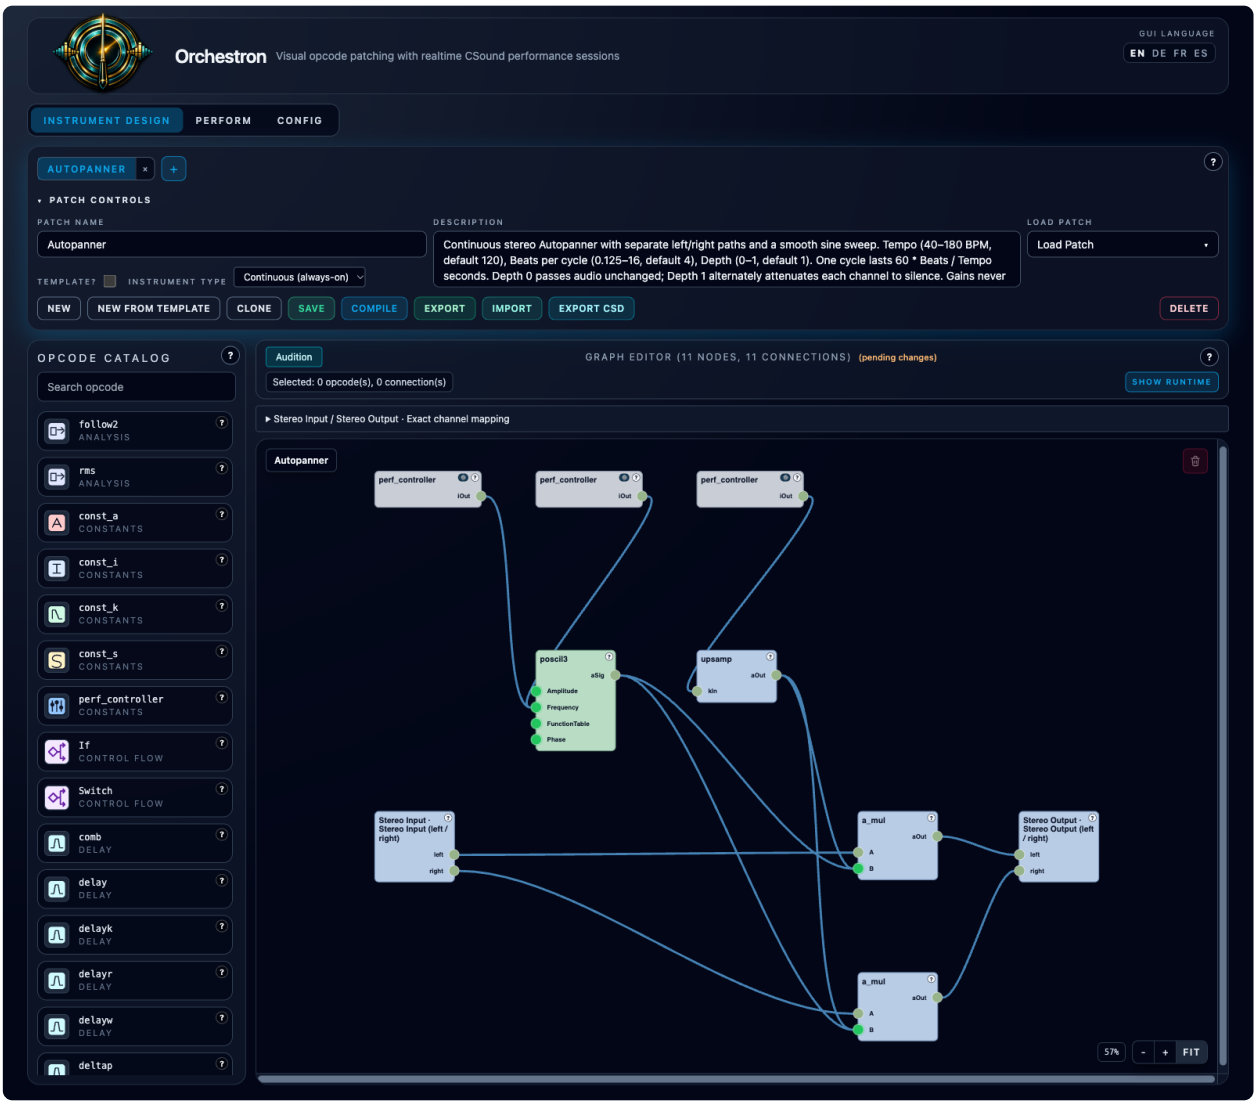

See the example guide for the builder, matched previews and regression checks.

Navigation: Up | Previous | Next

Parametric Equalizer

Navigation: Up | Previous | Next

Download Parametric Equalizer and import it using **Instrument Design → Import**. This continuous stereo effect has seven peaking bands, with independent left/right signal paths and linked gain settings. Add it to a source strip or Master using **Perform → Mixer → Inserts…**.

Controls

Center	Typical use	Gain range	Default
63 Hz	Sub/bass weight	–20 to +20 dB	0 dB
160 Hz	Bass warmth	–20 to +20 dB	0 dB
400 Hz	Low-mid body	–20 to +20 dB	0 dB
1 kHz	Midrange	–20 to +20 dB	0 dB
2.5 kHz	Presence	–20 to +20 dB	0 dB
6.3 kHz	Definition	–20 to +20 dB	0 dB
16 kHz	Air	–20 to +20 dB	0 dB

All seven controls use a **linear dB scale**. Zero is neutral; –20 dB gives amplitude gain 0.1, and +20 dB gives 10. The graph converts each setting with `ampdb(db)` before feeding `pareq.kv`, as in the Csound EQ example.

Controls marked **Live** respond during playback. Use Save Performance to retain the instance settings. Allow headroom when boosting: overlapping bands add gain, and this EQ has no limiter. Its named stereo inputs and outputs are `left` / `right`. Use it as an insert; a parallel dry return changes the resulting EQ response.

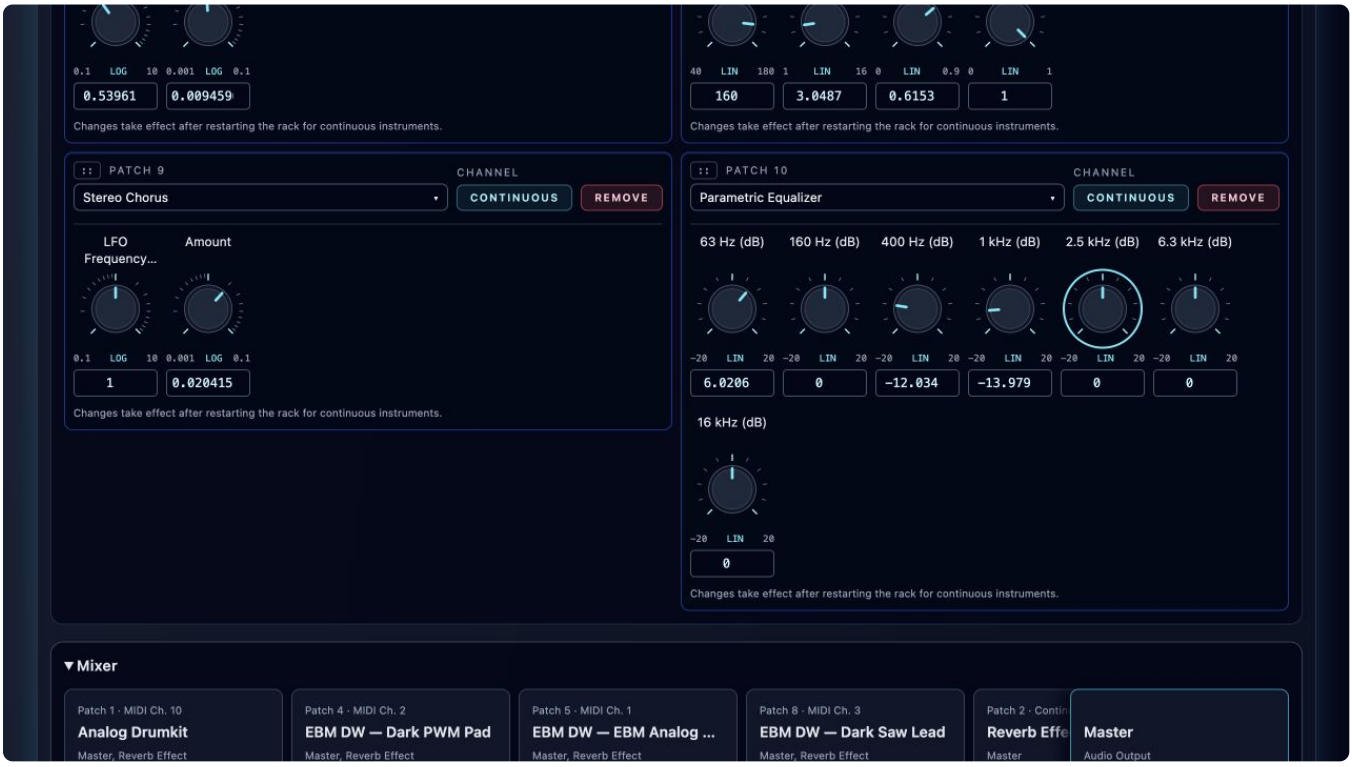

Frequency spacing and Q

The rounded centers are approximately 1⅓ octaves apart. Nominal **Q = 1.05** gives broad tonal shaping with useful overlap, approximately 1⅓-octave half-gain bandwidth at low frequencies. These are general-purpose defaults; edit the frequency constants and paired Q formulas in Instrument Design for narrower corrections. Digital frequency warping narrows the

uppermost bell near Nyquist. The supplied graph runs at 48 kHz, keeping 16 kHz below Nyquist.

For symmetric boosts/cuts, each `pareq.kq` input uses `1.05 * ampdb(db / 2)`. This adjusts Csound's denominator Q by the square root of linear gain to match the Audio EQ Cookbook's peaking convention. Nominal musical Q stays 1.05; the raw opcode Q varies with gain. The two channels use identical settings. Seven filters run **in series per channel**, giving unity gain when all controls are zero.

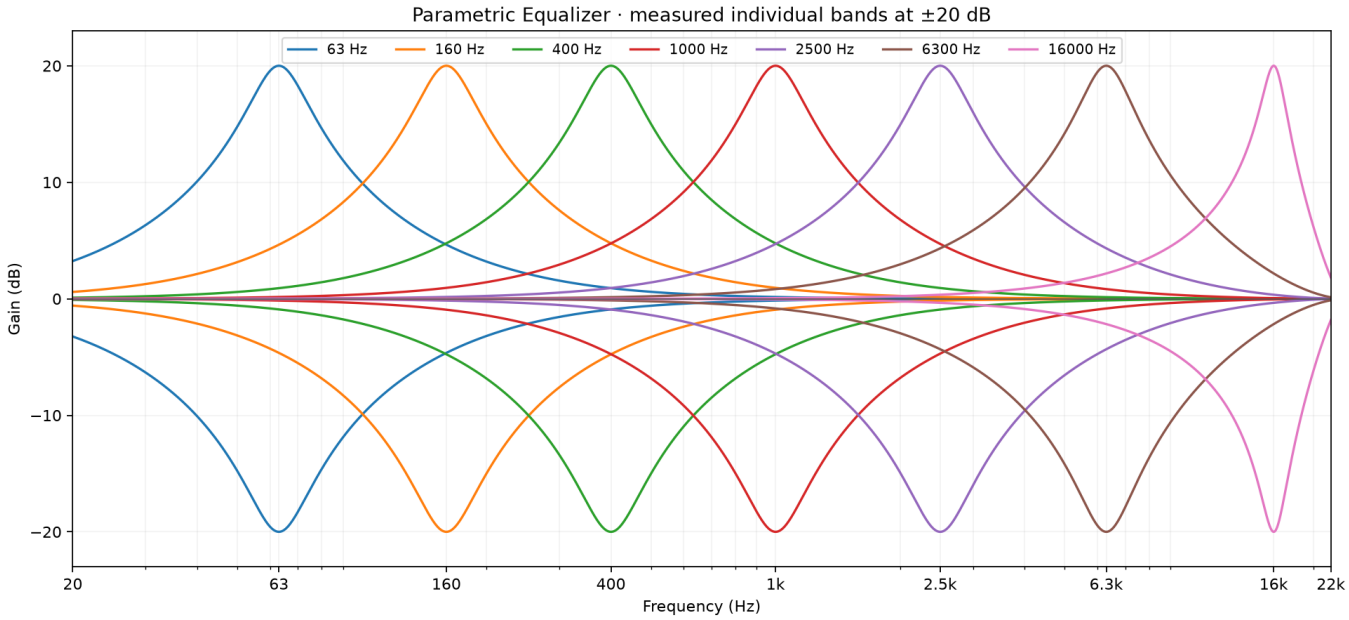

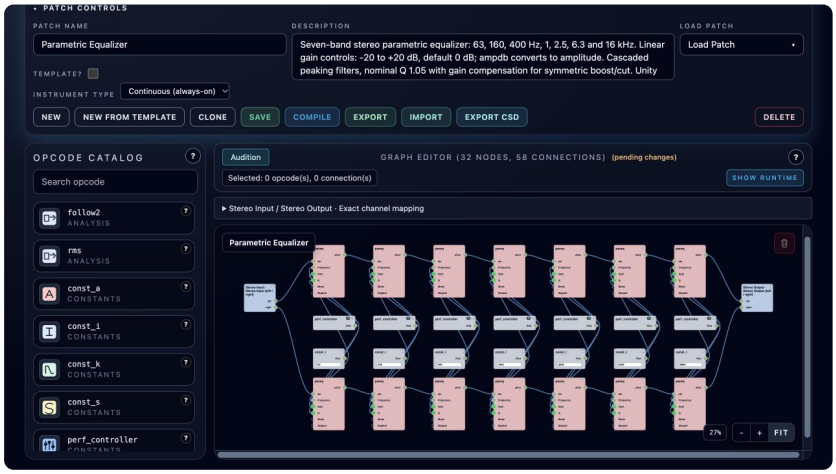

Updating the original three-band instrument

The earlier graph summed three full-range EQ outputs in parallel. At neutral settings this produced amplitude gain 3, or **+9.54 dB**. The new cascade corrects that behavior, so existing mixes can sound quieter and their EQ curves change.

The original three controller IDs are retained, with centers changing from 100 → 63 Hz, 600 → 400 Hz and 1.2 → 1 kHz. Four new bands default to zero. When updating an older library manually, convert any stored amplitude override with `20 * log10(old value)`; for example, 2 becomes +6.0206 dB and 0.2 becomes –13.9794 dB. Range clamping alone does not convert units. Keep a backup of old performances and reload open browser sessions after migration.

See the example and validation guide for the preserved original graph, rebuild commands and audio checks.

Navigation: Up | Previous | Next

Stereo Widener, Bit Crusher, and Harmonizer

Navigation: Up | Previous | Next

These three continuous instruments process incoming stereo audio without MIDI notes. Add an instrument to the rack and use its grouped Stereo Input and Stereo Output as a mixer insert. Each instance has its own performance settings; controls marked **Live** update during playback, while **Requires restart** identifies I-rate or mixed connections. The examples include editable graphs and native instrument files.

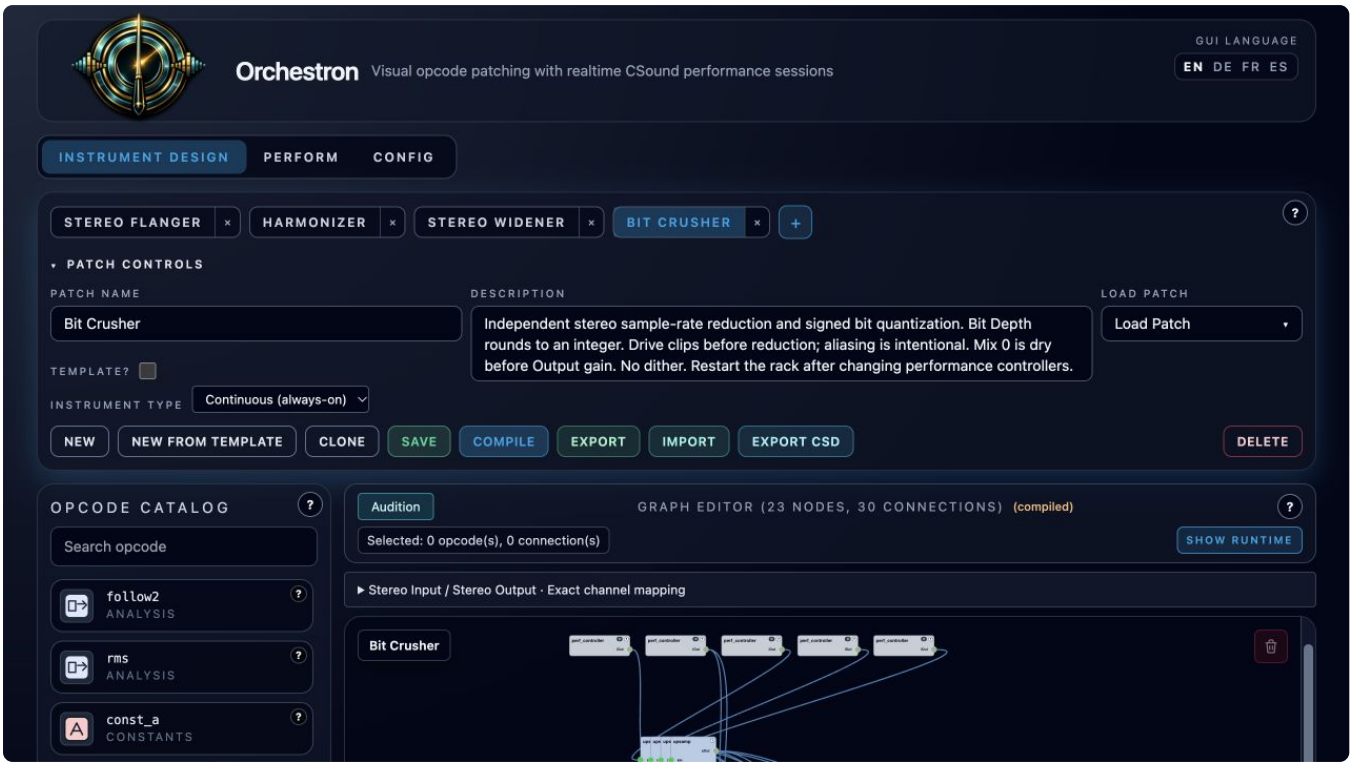

Stereo Widener

Controller	Range	Default
Width	0–2	1.5
Mono Widening	0–1	0
Delay (ms)	1–30	12
Output (dB)	–24–0	–6

Width scales the side signal while preserving the center. Width 0 makes mono; Width 1 with Mono Widening 0 preserves the original stereo image before Output gain. Width above 1 exaggerates the stereo difference. Mono input stays mono until Mono Widening is raised.

Mono Widening adds the difference between the center signal and a delayed copy to the side signal. Try Mono Widening 0.5–1 with Width 1–2 on a mono synth or pad. Delay controls the coloration and apparent space. The added side cancels when summed to mono; the mono sum remains the original center signal multiplied by Output gain. High width can increase individual channel peaks, so leave headroom.

Bit Crusher

Controller	Range	Default
Bit Depth	2–16	8
Downsample Factor	1–64, logarithmic	4

Controller	Range	Default
Drive (dB)	0–24	0
Mix	0–1	1
Output (dB)	–24–0	–6

Each channel is driven and hard-clipped, reduced in sample rate, and quantized. Bit Depth rounds to the nearest integer. Downsample Factor 1 preserves the sample rate; 4 produces an effective 12 kHz rate at a 48 kHz engine rate. Fractional factors are supported. Aliasing and stepped waveforms are intentional parts of the sound.

At 2 bits, quiet inputs below ±0.25 quantize to silence; raise Drive to bring them above that threshold. Drive affects the processed path only. Mix 0 gives the original input, and Mix 1 gives the processed signal; Output gain follows the blend. For a send return use Mix 1.

Harmonizer

Controller	Range	Default
Voice 1 Interval (semitones)	–24–24	+7
Voice 2 Interval (semitones)	–24–24	+12
Voice 1 Level	0–1	0.7
Voice 2 Level	0–1	0.5
Mix	0–1	0.5
Output (dB)	–24–0	–6

Two spectral pitch shifters add fixed harmony voices. Intervals round to whole semitones. The defaults add a fifth and an octave. Set either voice level to zero to mute it. If the level sum exceeds 1, both levels are divided by their sum before mixing with the dry signal.

Left and right channels are processed separately. The effect does not detect a key or scale and does not preserve vocal formants. Sustained notes and chords suit this algorithm; sharp transients can smear or ring. There is no feedback tail.

Measured processing latency is 1536 samples (32 ms at 48 kHz) on the validated Csound 6.18 runtime. The dry path uses the same delay, including at Mix 0. Use Mix 1 for a send return; a separate undelayed parallel dry route will still differ in timing.

Csound pvscale reference · Csound fold reference

Download the effects

- Stereo Widener
- Bit Crusher
- Harmonizer

Import the desired file in Instrument Design. Graph formulas and measurement procedures are in the technical reference.

See Granular Capture for live MIDI-controlled stutter, freeze and granular playback.

Navigation: Up | Previous | Next

Granular Capture: Stutter, Freeze and Grains

Navigation: Up | Previous | Next

Granular Capture is one continuous stereo effect with Stutter and Granular playback and a shared Freeze control. Add it as a mixer insert, feed its grouped stereo input, and route its grouped output onward. It records incoming audio continuously, even while playing a frozen snapshot. The two channels keep separate samples and share capture timing and grain positions.

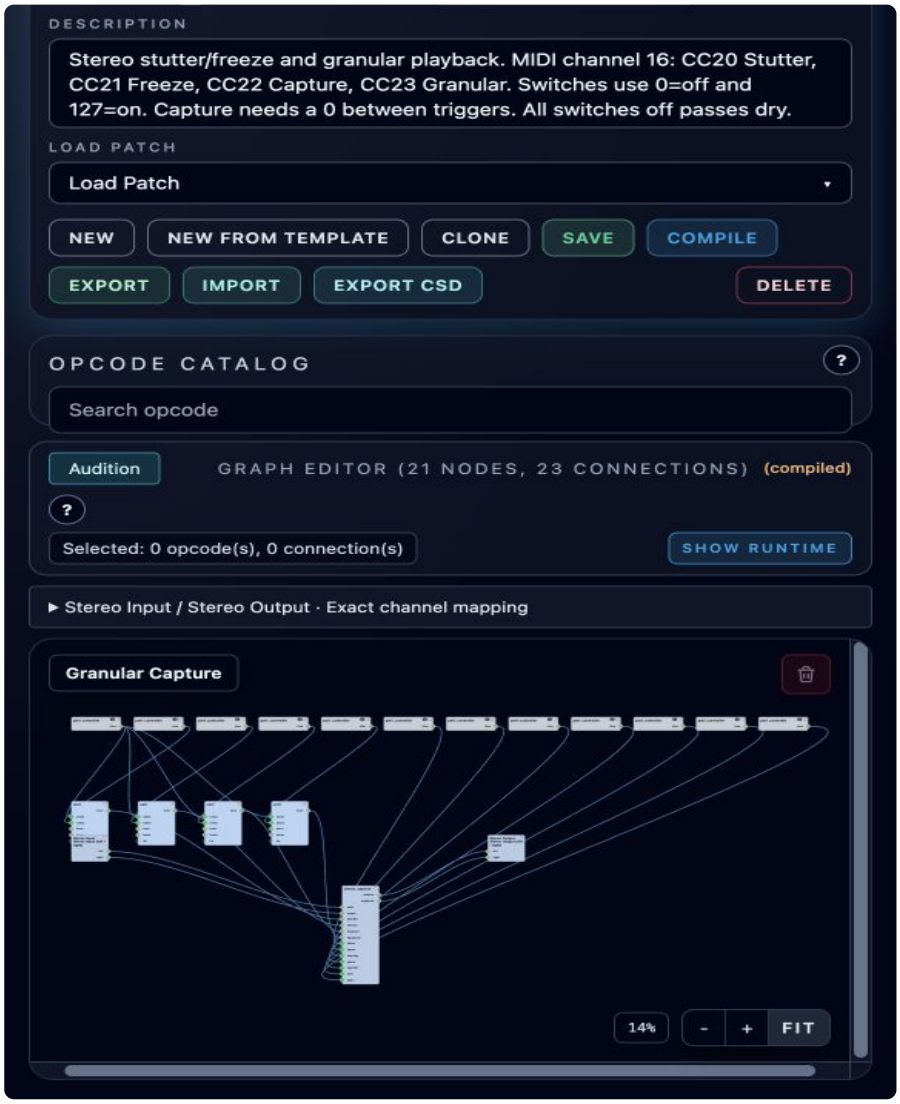

MIDI switches

The default **MIDI Control Channel is 16**. This setting is independent of MIDI note activation: the instrument remains Continuous and needs no note channel. Each switch has its own configurable CC number. Choose distinct CCs and avoid collisions with other receivers. Mapping knobs stop at 95 because Csound handles RPN/NRPN and some higher CCs specially. Send the initial Reset Value 0 before recording; standard CCs such as volume can start above zero. Instances sharing a channel and CC respond together; select another channel or CC set for independent instances.

Switch	Default CC	0–63	64–127
Stutter	20	Release stutter	Capture and repeat the recent slice; an existing Freeze snapshot is retained
Freeze	21	Release frozen source	Capture and hold the recent slice

Switch	Default CC	0–63	64–127
Capture	22	Arm the next trigger	Rising edge replaces the snapshot, including while frozen
Granular	23	Loop playback when Stutter/Freeze is on; otherwise dry	Play windowed grains from live history or the held snapshot

Use **0 and 127** in step cells. Capture is a value-edge trigger: put **0 between successive 127 values**. Repeating 127 sends another MIDI message but does not produce another rising edge. Capture alone updates the snapshot without enabling wet playback. Stutter or Freeze alone produces a loop; Granular on produces grains. Freeze off resumes live grains unless Stutter is still holding the snapshot. All three playback/hold switches off returns to dry audio with Output gain.

Configure a step pattern

1. Add a Controller Sequencer and select **Step Mode**.
2. Select only MIDI channel **16** and set **Controller # 20**.
3. Set **Reset Value 0**. Enter `127` where stutter should begin and `0` where it should end. Empty cells hold the current state.
4. Add lanes for CC21, CC22 and CC23 when needed, each with Reset Value 0. Enable Freeze before switching to Granular to turn a repeat into a sustained texture.
5. For recapture while frozen, send CC22 `127`, then `0`, then another `127` at the next capture point.

A four-beat pattern at 120 BPM with four steps per beat can use:

Lane	Populated steps (one-based)
Stutter, CC20	3: 127; 7: 0
Freeze, CC21	9: 127; 15: 0
Capture, CC22	13: 127; 14: 0
Granular, CC23	11: 127; 16: 0

The reproducible example includes a matching sequencer example. Sequencer start/stop and finite endings send Reset Value. Pad changes and loops hold the current value. A controller stopped with a nonzero Reset Value can leave an effect active. Muted controller lanes do not send their reset; release the corresponding CC manually or restart the rack.

Performance settings

These settings belong to each rack instance. Controls marked **Live** respond during playback; **Requires restart** identifies initialization settings. MIDI switches act during playback.

Setting	Range	Default
MIDI Control Channel	1–16, rounded	16
Stutter / Freeze / Capture / Granular CC	0–95, rounded	20 / 21 / 22 / 23
Slice	20–1000 ms	125 ms
Grain Size	5–200 ms	50 ms
Grains / Second	1–80	40
Grain Pitch	–24–24 semitones, fractional allowed	0
Position Spread	0–1	0.5
Mix	0–1	1
Output	–24–0 dB	–6 dB

Slice is the **stutter repeat period**, with an extra 5 ms of history captured for its overlapping seam. At 120 BPM, 125 ms is a

sixteenth note. This is a millisecond setting, not automatic tempo synchronization: use `60000 / BPM / subdivisions per beat` and restart after adjusting. Grain Pitch affects granular playback only. Position Spread zero starts each grain at the beginning of the slice; higher values scatter starts within it. Sparse short grains sound choppy; overlapping long grains produce a smoother sustained texture.

Mix zero preserves dry audio, with Output gain and no processing delay. Inactive playback also passes dry. Use an insert: a parallel send would duplicate dry audio while inactive. Freeze can intentionally sustain indefinitely, so send a release before the end of an offline render. Capture before enough history exists includes silence. Engine stop/restart clears both history and snapshot buffers.

Playback limits

Each instance has its own capture history. Loop seams and state changes use short crossfades. High Grain Pitch values can introduce aliasing; reduce the interval if the result becomes too harsh. This is a time-domain freeze, so it retains a slice of the waveform rather than holding a spectral analysis.

Use normal sequencer step lengths for the MIDI switches. Extremely short on/off pulses can pass between audio updates and fail to trigger a capture. Stop/restart clears captured audio.

Download and further details

- Download Granular Capture
- Processor design and technical checks

Navigation: Up | Previous | Next

Stereo Moog Filter

Navigation: Up | Previous | Next

Stereo Moog Filter is a Continuous stereo resonant low-pass effect. Import and save the native instrument, then add it as an instrument-strip or Master insert. Its grouped Stereo Input and Stereo Output preserve separate left and right audio; both channels share the same cutoff and resonance settings. No MIDI notes are required.

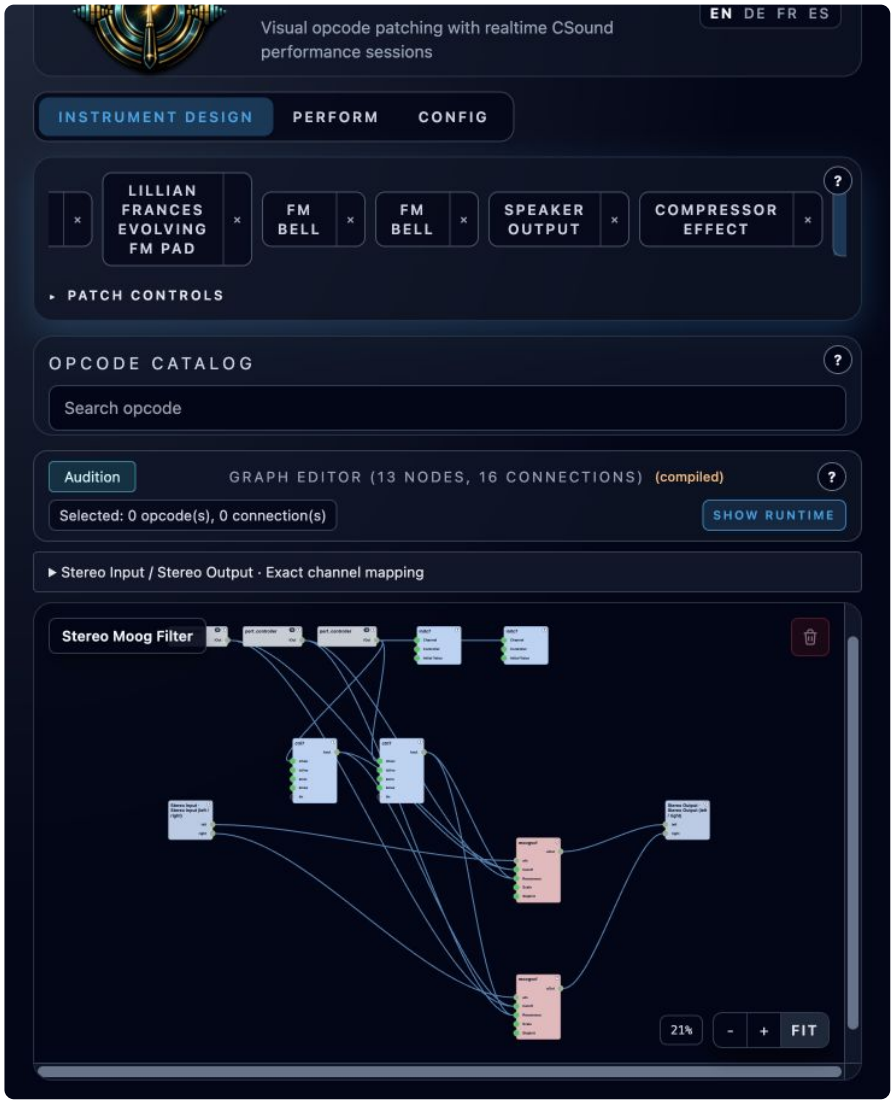

Performance settings

Setting	Range	Default
Cutoff (Hz)	20–20,000, logarithmic	1,000
Resonance	0–1, linear	0.3
MIDI Control Channel	1–16, rounded	16

These settings belong to each rack instance. Controls marked **Live** respond during playback; restart the rack for settings marked **Requires restart**. The effect is fully processed, with no additional output gain, dry/wet blend or smoothing. Use the mixer strip for output level. Resonance near 1 can self-oscillate, so leave headroom.

MIDI control

Controller	Effective value
CC20	Cutoff × CC20 / 127
CC21	Resonance × CC21 / 127

The factors are bounded to 0–1. A value of 127 uses the full performance setting, 64 uses approximately 50.39%, and 0 produces an exact zero factor. The cutoff knob's logarithmic scale does not change the linear MIDI multiplier. For example, Cutoff 4,000 Hz and CC20 64 produce about 2,015.75 Hz.

MIDI changes take effect at the next engine control block for both channels together. A cutoff of zero blocks new input, while the filter's existing state continues to decay nonlinearly; it is not an instant output mute. Raising CC20 reopens the filter.

Starting the instrument initializes CC20 and CC21 to maximum, so the performance settings work immediately without a MIDI controller. Later MIDI values replace that initialization. In Controller Sequencers select the effect's MIDI channel, use CC20 or CC21, and set **Reset Value 127** to return to the selected performance setting when the lane stops. Empty steps hold the preceding value.

Shared channel behavior: MIDI state is shared by channel and CC, including startup initialization. Instances on the same channel respond together; starting one resets these two CCs for other receivers. Channel 16/CC20–21 also overlap Granular Capture's default Stutter and Freeze switches. Assign different MIDI Control Channels when using the effects independently.

Graph and exports

Two native moogvcf nodes process the audio channels. Two ctrl7 nodes supply normalized live factors, and two initc7 nodes initialize them once. The graph bounds the factors because Csound's initializer internally stores `127 × value + 0.5`, slightly above full scale when the normalized initial value is 1.

Native instrument and performance exports preserve the graph and instance settings. CSD MIDI and SCORE exports preserve initialization and automation; SCORE uses the explicit-channel controller array. Their event timing can differ by one control block. Controller state itself is runtime state.

Open the imported instrument in Instrument Design to inspect its two filter paths.

Navigation: Up | Previous | Next

Stereo Auto Wah Wah

Navigation: Up | Previous | Next

Stereo Auto Wah Wah is a Continuous stereo filter effect that opens with the input's loudness and closes as the sound fades. Import and save the native instrument, then add it as an instrument-strip or Master insert. It starts without MIDI notes. The output is fully processed; use the mixer for level and insert bypass.

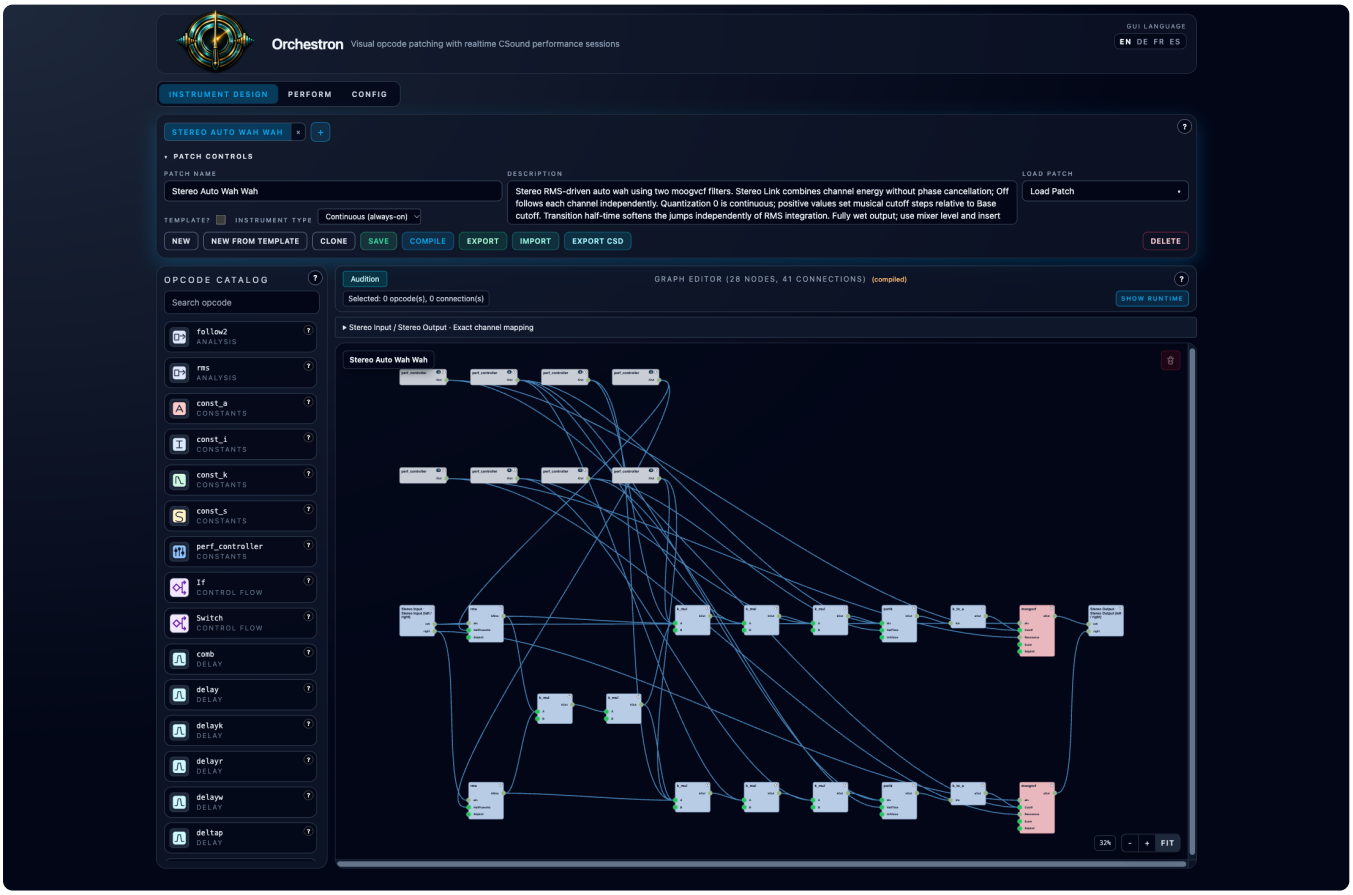

Performance settings

Setting	Range	Default	Use
Base cutoff (Hz)	50–2,000, logarithmic	200	Lowest cutoff, reached as the input falls silent.
Sweep range (octaves)	0–6	4	Distance the envelope can open the filter. Zero keeps a fixed cutoff.
Sensitivity	0.25–32, logarithmic	8	How strongly input loudness opens the filter. Increase for quieter sources.
RMS integration time (ms)	1–1,000, logarithmic	30	Averaging time of the loudness detector. Larger values respond more slowly.
Resonance	0–0.95	0.65	Emphasis around the cutoff frequency. Higher values give a more pronounced wah.
Quantization step (semitones)	0–24	0	Zero is continuous; positive values set the interval between cutoff steps.

Setting	Range	Default	Use
Transition half-time (ms)	1–100, logarithmic	5	Time to move halfway toward each new cutoff target.
Stereo Link	Off / On	On	Use one combined envelope, or let each channel move independently.

Settings belong to each rack instance. Controls marked **Live** respond during playback; restart the rack for controls marked **Requires restart**. These are performance settings, not live MIDI CC controls.

Continuous sweeps and musical steps

Start with Quantization at **0** for a smooth envelope-following wah. Set Sensitivity so quieter notes leave the filter partly closed and strong attacks open it noticeably. If it stays wide open, reduce Sensitivity. With the default Base cutoff and Sweep range, the sweep runs from 200 Hz to approximately 3,200 Hz.

Try **1** for fine semitone steps, **7** for fifths, **12** for octave jumps or **24** for two-octave jumps. Fractional steps are also supported. The steps are relative to Base cutoff: at 200 Hz with octave steps, the targets are 200, 400, 800, 1,600 and 3,200 Hz. This is frequency spacing, not detection of the notes being played. The highest target is limited to Sweep range, so a final step may be shorter when the interval does not divide that range evenly.

RMS integration and Transition half-time serve different purposes. Integration averages the input's loudness **before** selecting a target; Transition softens movement **after** the target is quantized. Keep Transition short for distinct jumps with rounded edges, or raise it for a gliding effect. Long transitions can blur rapid steps. The cutoff remains between 20 Hz and the lower of 18 kHz or 40% of the sample rate.

Stereo response

Stereo Link On combines the channels' RMS energy so both filters move together. Opposite-phase signals do not cancel the detector. With equal signals on both sides, the detected level is the same as either channel alone; with only one side active, its combined detected level is about 3 dB lower.

Stereo Link Off lets each channel's own loudness drive its filter. This can create different left/right movement from staggered notes or stereo percussion. The audio paths always stay separate: linking the detector does not mix sound between channels.

Saving and exports

Save Performance retains each instance's settings. Native bundles preserve the editable graph, Continuous type and grouped stereo ports. Both MIDI and SCORE CSD exports retain the envelope following and softened cutoff steps.

The filter uses moogvcf, with rms detection and portk transition smoothing.

Navigation: Up | Previous | Next

Trance Gate

Navigation: Up | Previous | Next

Trance Gate is a continuous stereo insert with direct MIDI gating and a swept triangle LFO. Import the native instrument, save it, and insert it on an instrument strip or Master. Both channels share the same gain envelope while keeping their audio separate. Use a serial insert: a parallel dry route continues passing ungated audio.

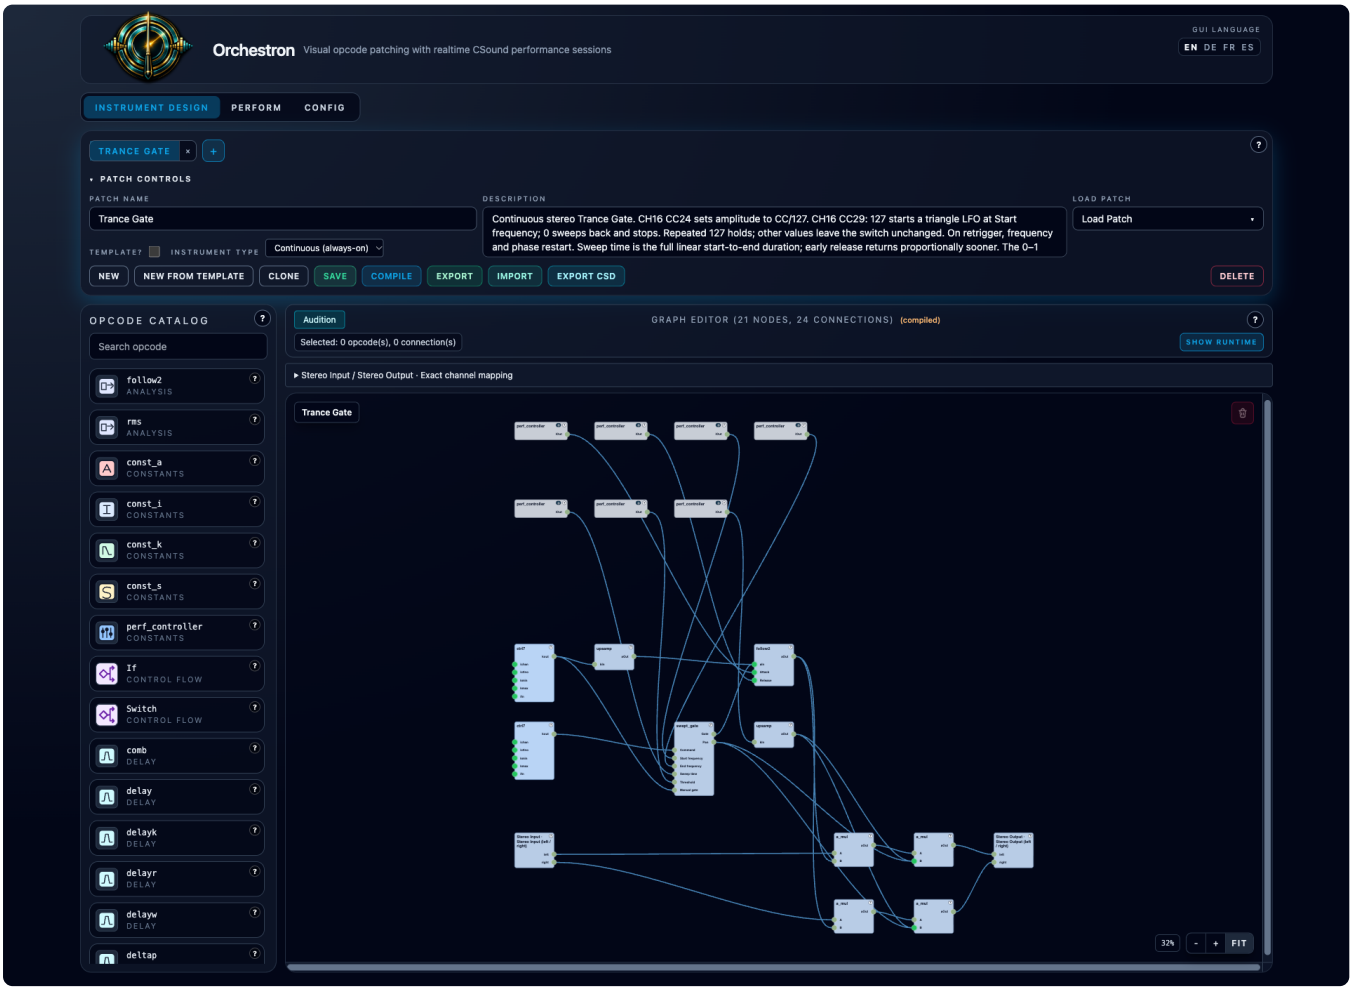

MIDI controls

Both controllers use **MIDI channel 16**, independently of the source instrument's note channel.

Controller	Behavior
CC24	Sets direct gate amplitude to CC/127: 127 is unity, 0 is closed, and 64 is approximately half amplitude.
CC29 = 127	Switches on the LFO at Start frequency with its triangle at zero, rising. Repeating 127 keeps it running.
CC29 = 0	Sweeps back from the current frequency to Start, then stops the LFO and closes its gate.
CC29 = 1–126	Leaves the LFO switch state unchanged.

While on, frequency moves linearly from Start to End in the selected **Sweep time**, then stays at End. On release it returns at the same speed: releasing halfway through a four-second sweep returns in two seconds. Descending sweeps work too. If Start equals End, the LFO runs at that frequency and stops immediately on 0. Switching on during a return sweep resets both frequency and phase.

The triangle runs from **0 to 1** and opens the gate only when strictly above **Threshold**. Threshold 0.5 gives approximately 50% open time at constant frequency; higher thresholds shorten openings. Threshold 1 stays closed; threshold 0 opens everywhere except the triangle's minimum. Frequency is in Hz and is independent of the performance tempo.

The requested gain is the **maximum of CC24/127 and the binary LFO gate**, smoothed by the shared Attack/Release envelope. Keep CC24 at **0** to hear only the LFO. CC24 at 127 keeps audio fully open even after the LFO stops; intermediate CC24 values set the level between LFO openings.

Step sequencer setup

1. Add a Controller Sequencer in **Step Mode**, channel **16**, Controller **29**, **Reset Value 0**.
2. Put **127** where the sweep should start and **0** where it should return. Empty cells and repeated 127 values continue the current sweep.
3. Leave CC24 at 0 for LFO-only gating, or add a separate CC24 lane for direct amplitude steps.
4. Allow time after the last 0 for the return sweep and release envelope. Stopping the audio engine itself stops all processing immediately.

Starting an instance does not initialize or reset either shared controller. A fresh runtime starts silent; an existing CC29 value of 127 starts the LFO when the instance starts. Instances share CH16/CC24 and CH16/CC29 commands but have independent phase, sweep and envelope settings. Avoid using those same channel/CC pairs for unrelated automation. Muted controller lanes hold their last received value; send 0 explicitly when needed.

Performance settings

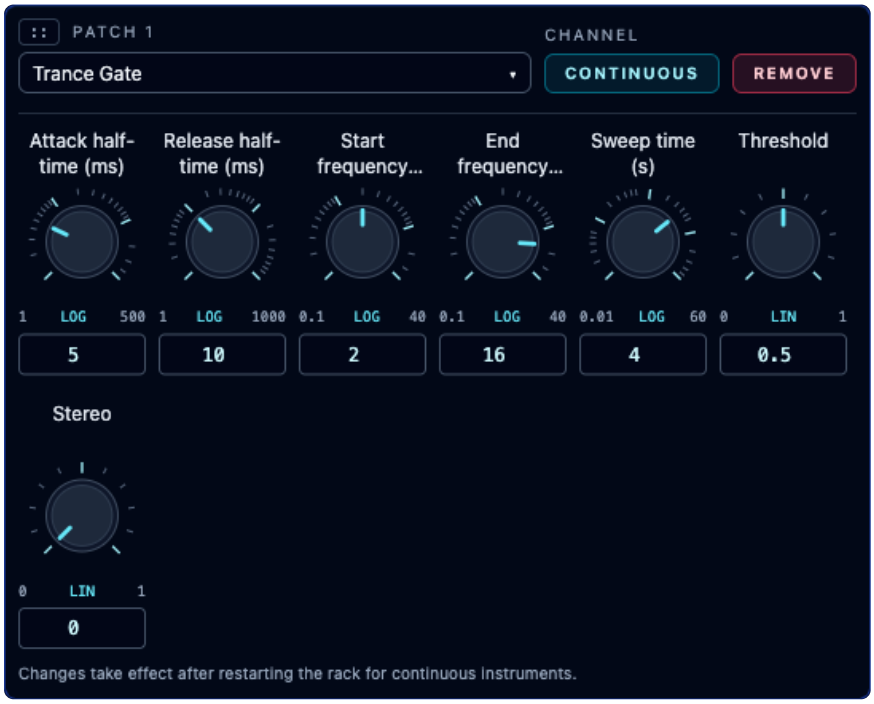

Control	Range	Default
Attack half-time (ms)	1–500, logarithmic	5
Release half-time (ms)	1–1000, logarithmic	10
Start frequency (Hz)	0.1–40, logarithmic	2
End frequency (Hz)	0.1–40, logarithmic	16
Sweep time (s)	0.01–60, logarithmic	4
Threshold	0–1, linear	0.5
Stereo	0–1, linear	0

Attack, **Release** and **Stereo** update live during realtime playback. Restart the rack after changing **Start frequency**, **End frequency**, **Sweep time** or **Threshold**. Both CSD export modes keep all these settings fixed at initialization. MIDI commands act during playback. A half-time halves the remaining distance to the requested gain; ten half-times cover approximately 99.9% of a transition. High LFO rates or long half-times soften the gate and may prevent it from fully opening or closing between pulses.

Stereo movement

Stereo 0 preserves the existing stereo output exactly. Raising Stereo adds channel balance movement, following the same convention as Autopanner: each input stays in its own channel. At **Stereo 1**, pan 0 passes only the left channel and pan 1 only the right; the center gives each channel gain 0.5. A signal present on only one input fades in that channel rather than moving to the other.

- **CC29 / LFO:** pan follows the current 0–1 triangle, including its return sweep. The gate threshold still applies, so only the open portions of that sweep are audible. Retriggering during return resets phase and pan together.
- **CC24 / manual:** each transition from 0 to any positive value samples a fresh noise value and holds that random pan position. Repeated positive values, or changes between positive values, do not resample. The position holds through the release envelope.
- **Both sources:** the running LFO takes priority for pan. CC24 openings still sample a position in the background; if CC24 remains positive after the LFO stops, its held random position takes over. With neither source active, the last pan position holds through release. A new instance starts centered.

Stereo movement does not add gain or crossfeed. Random manual pan positions belong to each instance and are not saved automation, so separate renders can choose different positions.

Graph and exports

Native exports preserve the graph and seven performance controls. MIDI and SCORE CSD exports include the processor and the same explicit-channel automation. The finite export must include enough time for the return and release: send 0 early enough, leave trailing rests, or extend the final `f 0` duration as described in the export's rendering instructions.

See the processor design and reproduction notes for technical details.

Navigation: Up | Previous | Next

Configuration

Navigation: Up | Previous | Next

Use Config to adjust synthesis and browser playback, and consult this chapter when you want to change language, improve responsiveness or understand what survives a reload. Start with the defaults unless you have a specific problem to solve.

Chapter Contents

- GUI Language and Integrated Help
- Audio Engine Settings (Config Page)
- Browser-Clock Latency
- Saving, Workspace Recovery and Defaults

Choose the right setting

- **Instrument engine settings** control synthesis for the patch, including sample rate and block size.
- **Browser playback settings** control buffering and the balance between responsive playing and resilient listening.
- **Workspace preferences** include language and settings restored with your working context.

If you want to connect a keyboard or another MIDI application, follow MIDI Setup and Inputs in Performance. Internal piano rolls and sequencers work without an external MIDI device.

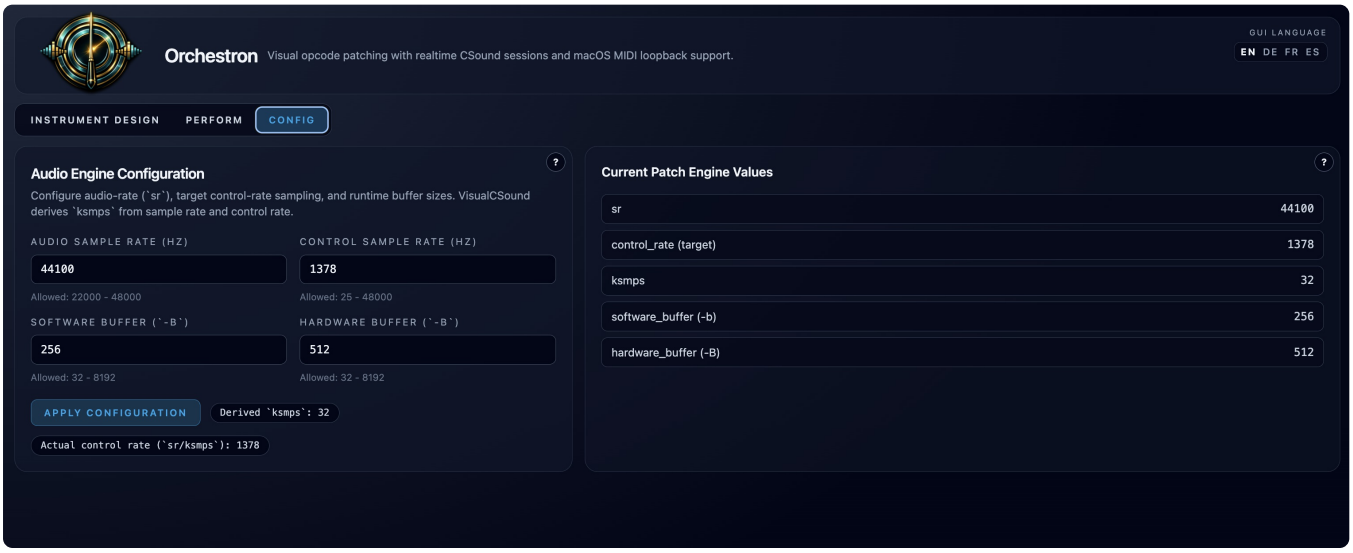

Navigation: Up | Previous | Next

GUI Language and Integrated Help

Navigation: Up | Previous | Next

Orchestron includes multilingual UI labels and integrated documentation so you can stay inside the app while learning and working.

GUI Language Selector

The language selector is in the top-right area of the main app header.

Supported languages:

- English
- German
- French
- Spanish

The selector is presented as quick language links (short labels) for fast switching.

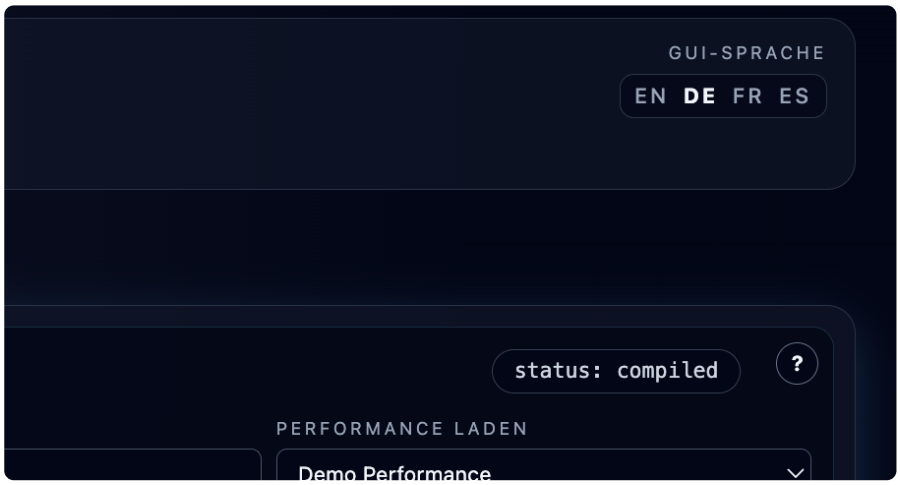

Top-right header language selector links.

What Changes When You Switch Language

Language switching affects user-facing text such as:

- Page labels (`Instrument Design` , `Perform` , `Config`)
- Toolbar/button labels
- Context help pages (integrated help modals)
- Opcode documentation modal content (localized generated markdown)
- Many labels inside advanced editors (including the GEN editor)

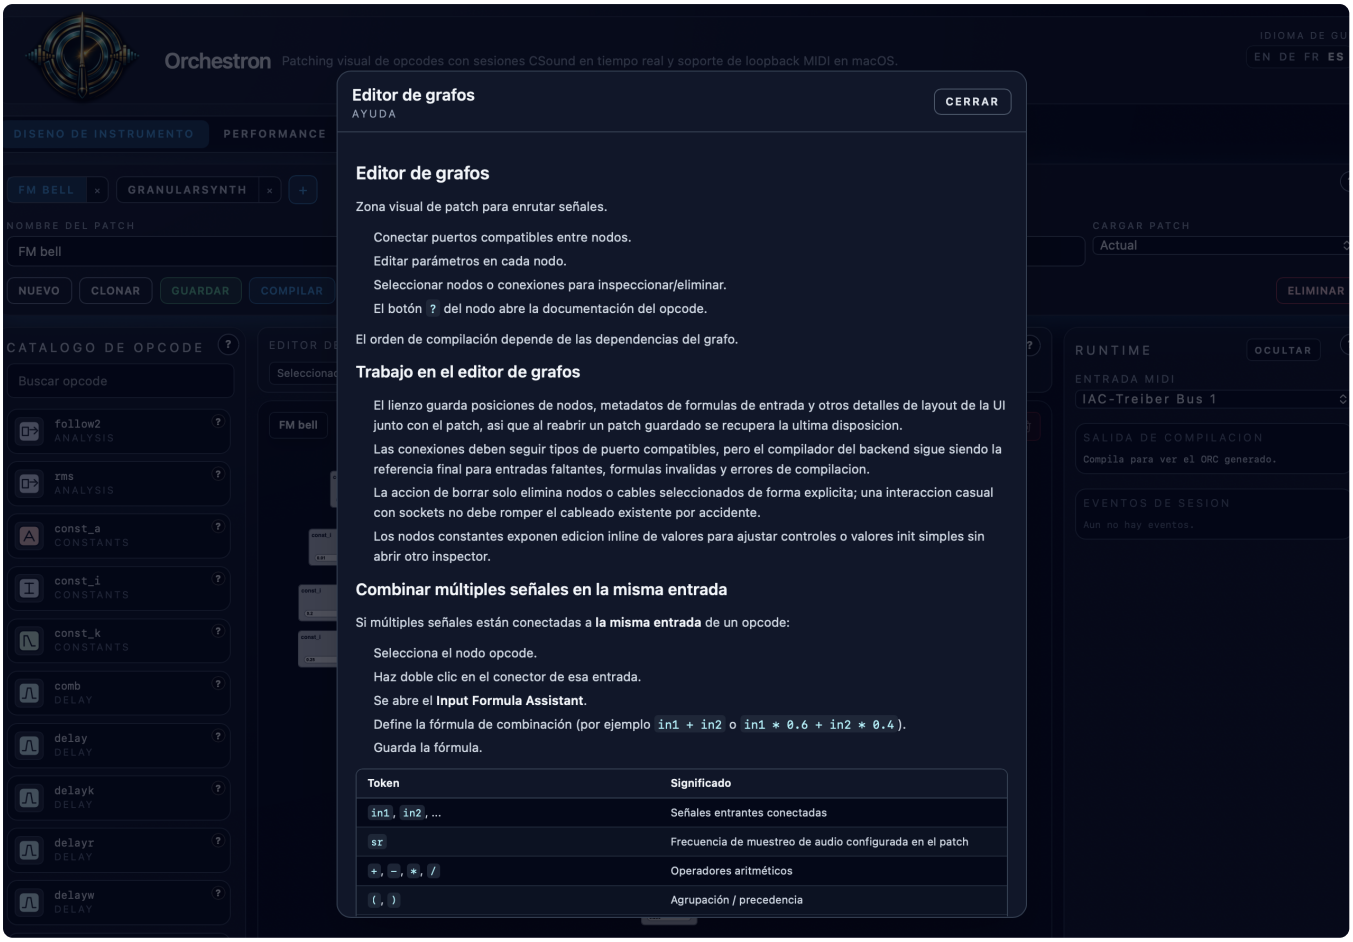

Integrated help and UI localization in multiple supported languages.

Integrated Help (? Buttons)

Orchestron provides context help buttons (?) across the UI.

Examples include:

- Instrument patch toolbar
- Opcode catalog
- Graph editor header/area
- Runtime panel
- Sequencer sections (rack, Mixer, tracks, Score Sequencers, both multitrack arrangers, piano rolls, MIDI controllers)
- Config page sections

Clicking a context ? opens a markdown help modal tailored to that UI area.

Mixer help covers levels, routing, sends, inserts, bypass and feedback diagnostics. Its button remains available while collapsed. The new Multitrack Arranger opens clip-specific guidance for fixed windows, reusable content, drafts, save scopes and playback; pad-based performances retain their compatibility-arranger help. Both use the same localized dialog as the rack, live keyboard and live MIDI controllers.

The Score Sequencer uses this same top-right button in its panel and clip pattern editor, including while collapsed. Its help explains keyboard auditioning, double-click quarter-note insertion, drawing, editing and playback in all four GUI languages.

On the Instrument Design, Perform, and Config pages, these help modals focus on the current tool, device, transport section, or engine setting panel instead of repeating generic UI conventions.

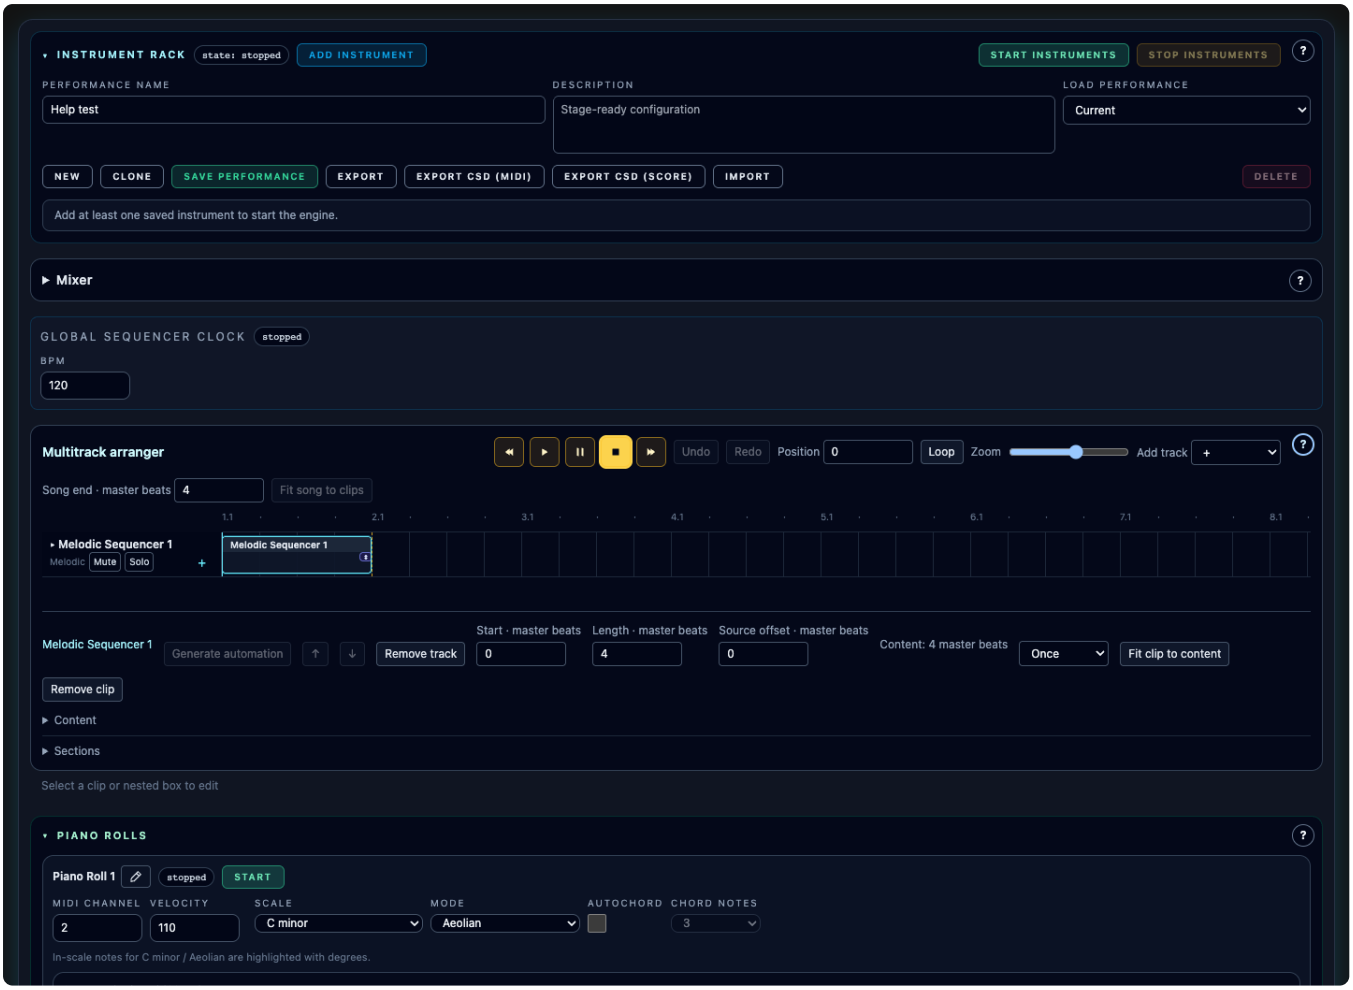
INSTRUMENT RACK
state: stopped
ADD INSTRUMENT
START INSTRUMENTS
STOP INSTRUMENTS
PERFORMANCE NAME
Help test
DESCRIPTION
Stage-ready configuration
LOAD PERFORMANCE
Current
NEW
CLONE
SAVE PERFORMANCE
EXPORT
EXPORT CSD (MIDI)
EXPORT CSD (SCORE)
IMPORT
DELETE
Add at least one saved instrument to start the engine.
Mixer
GLOBAL SEQUENCER CLOCK
stopped
BPM
120
Multitrack arranger
Undo
Redo
Position
0
Loop
Zoom
Add track
Song end · master beats
4
Fit song to clips
1.1
2.1
3.1
4.1
5.1
6.1
7.1
8.1
Melodic Sequencer 1
Melodic
Mute
Solo
Melodic Sequencer 1
Melodic Sequencer 1
Generate automation
Remove track
Start · master beats
0
Length · master beats
4
Source offset · master beats
0
Content: 4 master beats
Once
Fit clip to content
Remove clip
Content
Sections
Select a clip or nested box to edit
PIANO ROLLS
Piano Roll 1
stopped
START
MIDI CHANNEL
VELOCITY
SCALE
MODE
AUTOCHORD
CHORD NOTES
2
110
C minor
Aeolian
3
In-scale notes for C minor / Aeolian are highlighted with degrees.

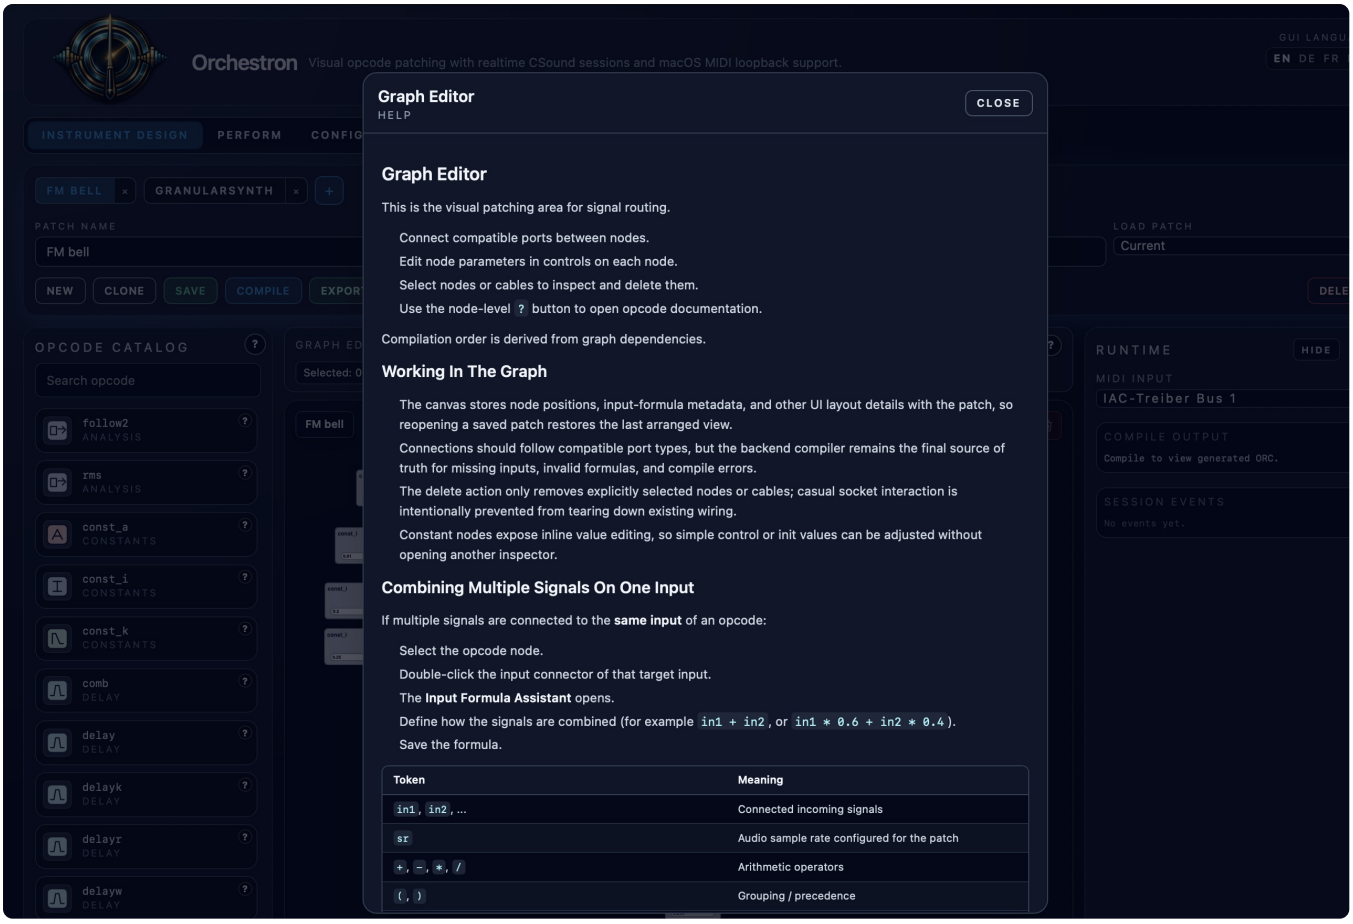

Context help modal integrated into the workflow.

Opcode Documentation (`?` On Opcode Node)

Placed opcode nodes provide a node-level `?` button that opens opcode documentation.

The modal includes:

- Localized description
- Inputs and outputs (with detailed port descriptions)
- Syntax/template summary
- Tags
- `Open Csound Reference` link to the official Csound documentation page for that opcode

Modal Behavior

Both help modals and opcode documentation modals support:

- Close button
- Click outside modal to close
- `Esc` key to close

Best Practice

Use context help for workflow guidance and opcode docs for technical/operator details.

This combination is faster than switching repeatedly between the app and external documentation while designing patches.

Navigation: Up | Previous | Next

Audio Engine Settings (Config Page)

Navigation: Up | Previous | Next

The Config page controls patch-level Csound engine settings, parallel playback batch size and browser-clock latency settings for the current workspace.

Where These Settings Are Stored

The Audio Engine Configuration section is stored in the current patch's `graph.engine_config`.

That means:

- different patches can use different engine settings
- switching patches can change the visible Config page values

Browser-clock latency settings are different:

- they apply to the active browser-clock session runtime
- they are stored in app state, not in the patch
- applying them updates the active browser-clock controller without changing patch data

Parallel Playback

The optional **Parallel playback** control is in the fixed Engine stats line at the bottom of Perform. Enable it while instruments are stopped. Single rendering remains the default and fallback. The engine analyses instrument/effect dependencies and shared resources, then uses available CPU and memory capacity to choose a parallel layout. No per-song catalog is required. The choice does not change engine timing, browser buffering or CSD exports. See Engine Stats and Parallel Playback.

Parallel playback normally exchanges **256 samples** between workers at a time. This reduces communication overhead while each instrument keeps its existing synthesis block size and scheduled note/controller timing. A live action received after a batch has begun takes effect at the next available batch; 256 samples span about **5.3 ms at 48 kHz**. This is separate from the browser's playback queue and does not change exported CSDs.

In **Config → Parallel Playback**, enter a **Batch size (samples)** from **32 to 1024** and click **Apply Batch Size**. The default is **256**. Stop playback before applying; the choice takes effect on the next start with **Parallel playback** enabled in Perform. It is saved in app state across reloads, separately from the parallel playback checkbox, which still resets to off. It is not stored in instruments, performances or CSD exports.

The engine rounds down to whole synthesis blocks, with at least one block per batch. For example, with a 32-sample synthesis block, entering 255 uses 224 samples. Try 128 for shorter batches or 32 to compare with one handoff per synthesis block. Larger batches reduce handoffs but can delay newly arriving live actions; 1024 samples span about 21.3 ms at 48 kHz.

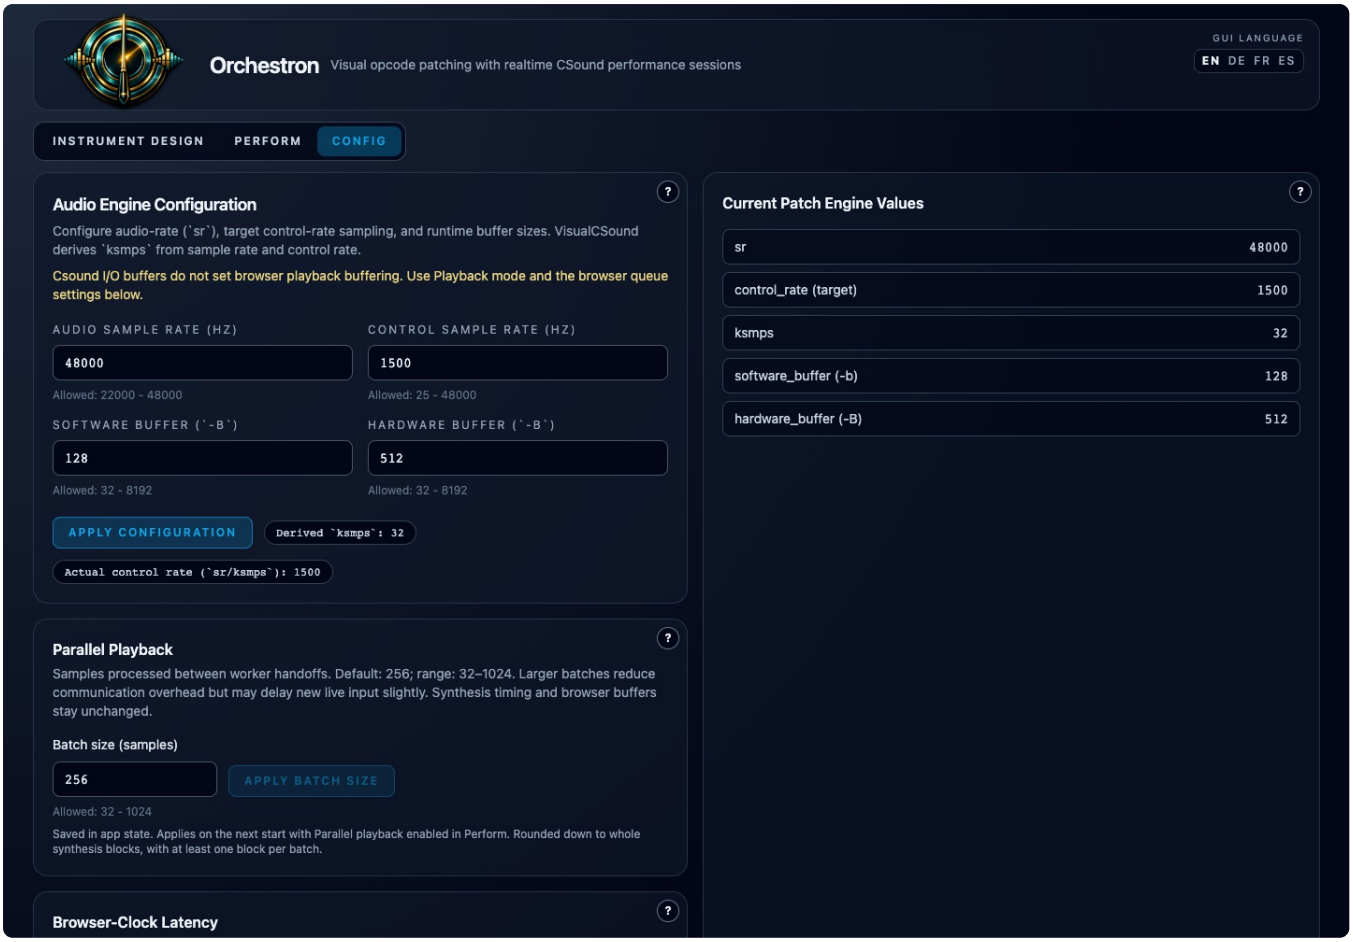

Editable Settings (Config Page)

Patch Engine Section

Audio Sample Rate (`sr`)

- Field label: `Audio Sample Rate (Hz)`
- Allowed range: **22000 .. 48000**
- Integer only

Control Sample Rate (`control_rate` target)

- Field label: `Control Sample Rate (Hz)`
- Allowed range: **25 .. 48000**
- Integer only

This is the target control-rate value used to derive `ksmps` .

Software Buffer (`-b`)

- Field label: `Software Buffer (-b)`
- Allowed range: **32 .. 8192**
- Integer only

Hardware Buffer (`-B`)

- Field label: `Hardware Buffer (-B)`
- Allowed range: **32 .. 8192**
- Integer only

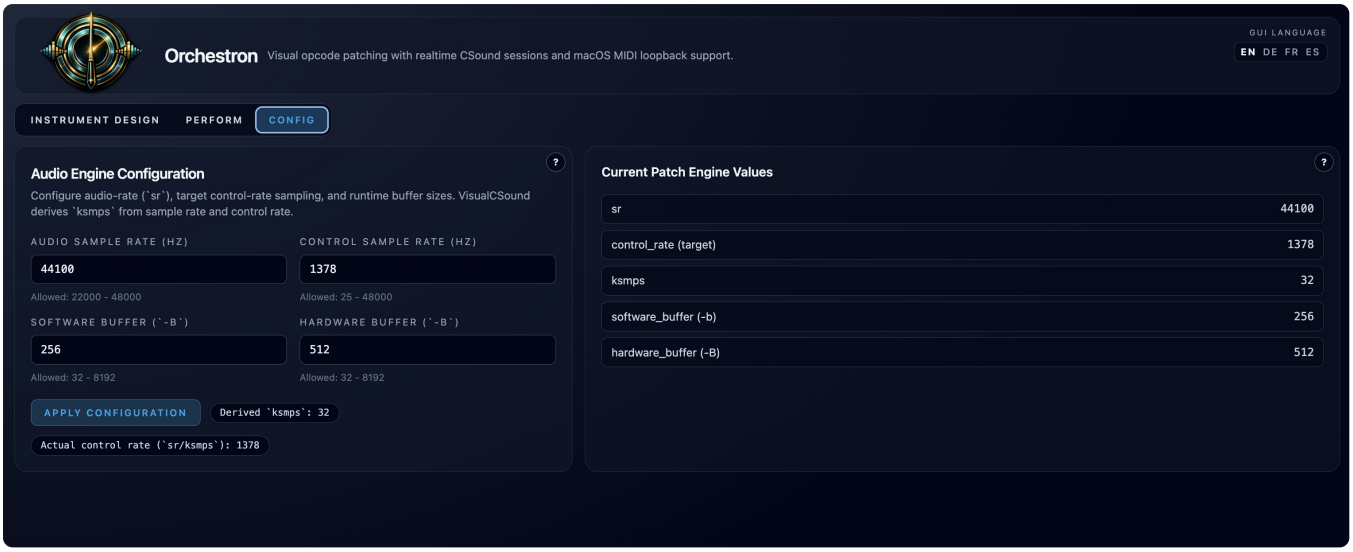

Config page audio engine settings in a valid state.

Derived Values Shown In The UI

The Config page previews derived values before applying:

- Derived `ksmps`
- Actual control rate (`sr / ksmps`)

This helps you see the real runtime effect of your target control-rate choice.

Current Patch Engine Values (Read-Only Panel)

The right-side panel shows the normalized engine values stored in the patch:

- `sr`
- `control_rate`
- `ksmps`
- `software_buffer`
- `hardware_buffer`

These are the current patch's values. A multi-instrument performance takes its shared engine configuration from its first compiled instrument target; this panel is not a global performance playback configuration.

Browser-Clock Latency Section

The Config page shows an additional section for browser-owned audio latency tuning whenever a browser-clock session is active.

See Browser-Clock Latency for the unified browser-clock workflow explanation and the full UI screenshot of this section.

Playback mode and inactivity timeout are shown above the expandable **Live buffering and recovery** fields:

- steady queue low water / high water (`ms`)
- startup queue low water / high water (`ms`)
- underrun recovery boost / maximum underrun boost (`ms`)
- maximum blocks per render request
- steady / startup / recovery parallel request limits
- immediate note render blocks
- immediate note render cooldown (`ms`)

These values control the browser PCM queue depth and render request behavior, which are the dominant latency/stability

knobs on the unified browser-clock path.

The patch engine sample rate and `ksmps` still determine synthesis timing. The Software/Hardware Buffer fields do not control browser playback buffering. Changing them does not enlarge the browser queue; use Browser-Clock Latency for that.

Choose **Auto / Live / Listening** in the Browser-Clock Latency section. Auto is the default, with a configurable 10-second inactivity timeout (1–300 seconds). It detects manual piano/controller and selected external MIDI input, while ignoring sequencer automation and controller initialization. Held notes, sustain and controller gestures prevent the timeout. Queued audio is preserved on every switch, so the first note after Listening may still be delayed. The current state appears in Config and Runtime.

The advanced **Live buffering and recovery** controls retain your existing tuning. Listening adds minimum steady targets of 750/1500 ms and startup targets of 1000/1800 ms, retaining larger custom targets. See Playback Mode for the complete behavior.

The dedicated Browser-Clock Latency chapter shows the browser-clock tuning section in detail.

Field-By-Field Meaning

Steady Low Water / Steady High Water

- These define the normal browser PCM queue target after startup is complete.
- Lower values reduce live-play latency.
- Higher values make clicks and underruns less likely.

Startup Low Water / Startup High Water

- These define the larger queue target used while the stream is still priming after connect or after an underrun recovery phase.
- They are usually kept above the steady values so the stream can stabilize before dropping into lower-latency playback.

Underrun Recovery Boost / Max Underrun Boost

- When the browser detects underruns, VisualCSound temporarily adds extra queue headroom.
- `Underrun Recovery Boost` controls how much extra buffer is added per underrun event.
- `Max Underrun Boost` limits how far that temporary recovery buffer can grow.

Max Blocks Per Request

- Caps the maximum backend render chunk size requested in one browser-clock render request.
- Smaller values reduce head-of-line delay for live notes.
- Larger values reduce request overhead but can make live interaction feel less immediate.

Steady / Startup / Recovery Parallel Requests

- These limits control how many render requests may remain in flight at once.
- More parallelism can keep the queue full more aggressively.
- Too much parallelism can also create larger bursts and less predictable latency.

Immediate Note Render Blocks / Immediate Note Render Cooldown

- A live `note_on` sends a small urgent render request in addition to the normal refill logic.
- `Immediate Note Render Blocks` sets the size of that urgent render burst.
- `Immediate Note Render Cooldown` throttles how often these urgent requests can be sent.
- These controls help manual piano playing feel more responsive, but values that are too aggressive can reintroduce clicks.

Practical Adjustment Order

1. Choose Auto, Live or Listening; Auto defaults to a ten-second inactivity timeout.
2. If clicks remain, raise the steady or startup watermarks slightly.
3. If latency still feels too high, reduce `max blocks per request` carefully.
4. Only then tune immediate note render and parallel request limits.

Validation Behavior

The page validates:

- integer format
- allowed range
- high-water fields must be greater than their matching low-water fields

`Apply Configuration` is disabled until all visible values are valid.

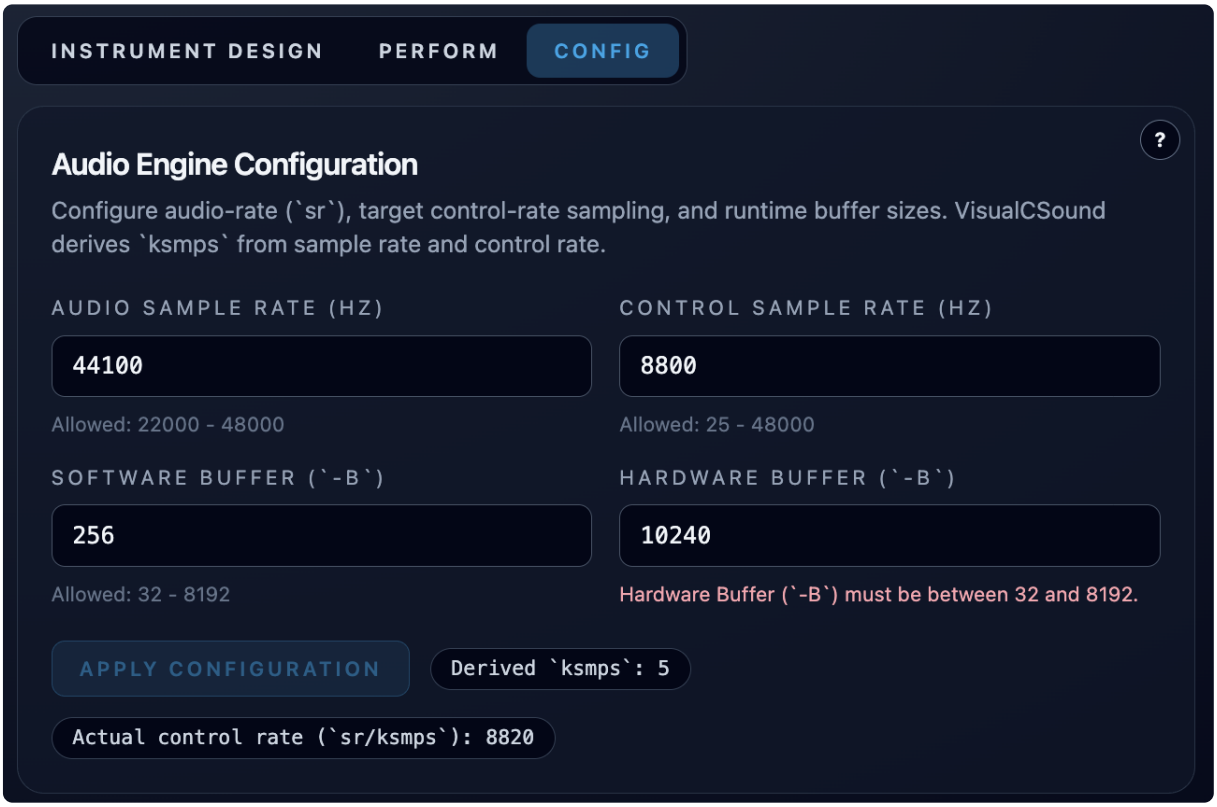

Config page validation messages with invalid values entered.

Apply Configuration

`Apply Configuration` writes the new engine settings into the current patch state.

Practical effect:

- future compile/start operations for the patch use the updated engine config
- the patch must still be saved if you want the change persisted in the patch library

`Apply Browser-Clock Settings` writes the browser-clock latency settings into app state and refreshes the active browser-clock controller if one is connected.

Other Engine Fields Present But Not Edited In UI

The patch engine config also contains values such as:

- `nchnls` (default `2`)
- `0dbfs` (default `1`)

These are part of the patch model, but the current Config page focuses on the user-facing timing and buffer controls listed above.

Tuning Tips

- Lower buffers reduce latency but increase glitch risk.
- Higher buffers improve stability but increase latency.
- Use the Runtime panel and live testing to find a stable setting for your machine and patch complexity.
- On the browser-clock path, start by adjusting browser-clock low/high water, render blocks, and parallel request limits before changing patch `-b/-B`.
- For live MIDI sessions, prefer `ksmps <= 32`. The runtime warns when a live session starts with a larger `ksmps` because MIDI execution is still quantized to the k-period.

Navigation: Up | Previous | Next

Browser-Clock Latency

Navigation: Up | Previous | Next

Orchestron plays sound through your browser on every supported platform. **Config → Browser-Clock Latency** controls the balance between responsive playing and enough buffered audio to withstand brief interruptions.

Playback Mode

Choose **Auto**, **Live**, or **Listening**, then select **Apply Browser-Clock Settings**. Auto is the default, including when restoring a workspace saved before these modes existed. The mode and **Auto inactivity timeout (seconds)** are workspace preferences, stored in app state. The timeout defaults to **10 seconds** and accepts integer values from **1 to 300 seconds**. The current state appears in Config and in the Runtime panel; it is not saved with performances.

Mode	Behavior
Auto	Uses Listening while idle. Starting a piano keyboard or manual MIDI controller, playing notes, moving a manual MIDI control, or receiving external notes/CCs selects Live.
Live	Always uses the configured Live queue targets and render limits, regardless of input activity. Select this before playing if the first note must be responsive.
Listening	Keeps a larger queue for listening to a performance, even when manual input arrives.

Auto returns to Listening after the configured inactivity timeout. Held notes, sustain and active controller drags keep Live selected. Stopping/removing a manual source removes its activity; when no other source is active, Auto can return immediately. Merely leaving a keyboard/controller enabled or an external device connected does not keep Live selected forever. The timeout starts again at note/pedal release or the end of a controller gesture. External input is limited to the session's selected MIDI input. Disconnecting that device, rebinding MIDI input or Panic clears its held state.

Sequencers, controller sequencers and the initial values sent automatically by manual MIDI controls do not trigger Live. Only the original external/manual input triggers it when routed into an arpeggiator.

Switching preserves all queued audio, pending renders and transport position. The browser lets the existing queue drain when changing to Live and replenishes it when changing to Listening. The indicator shows **Switching to Live** while the larger queue is draining. This means the first unexpected note after Listening may still have the previous latency. No samples are skipped to make the switch appear instantaneous.

Open **Live buffering and recovery** to tune the existing advanced controls. Listening uses at least **750/1500 ms** steady low/high water and **1000/1800 ms** startup low/high water, with up to **512 Csound blocks per request**. Larger custom queue targets are retained; other recovery and parallelism settings remain in effect. These queue targets are independent of Csound's synthesis block size. At 48 kHz, 32,768 frames represent about 683 ms and 65,536 frames about 1,365 ms. Queue settings already accept up to 2,000 ms, with a six-second ring buffer providing capacity for recovery headroom.

Render requests remain capped at 512 blocks: that is about 341 ms at `ksmps=32`, but only 10.7 ms at `ksmps=1` and 48 kHz. Listening changes buffering, not synthesis/control rate. Larger queues absorb short stalls; they cannot compensate indefinitely if rendering is consistently slower than playback. Use Audio Queue, Underruns and Render Ratio in the Runtime panel to distinguish these cases.

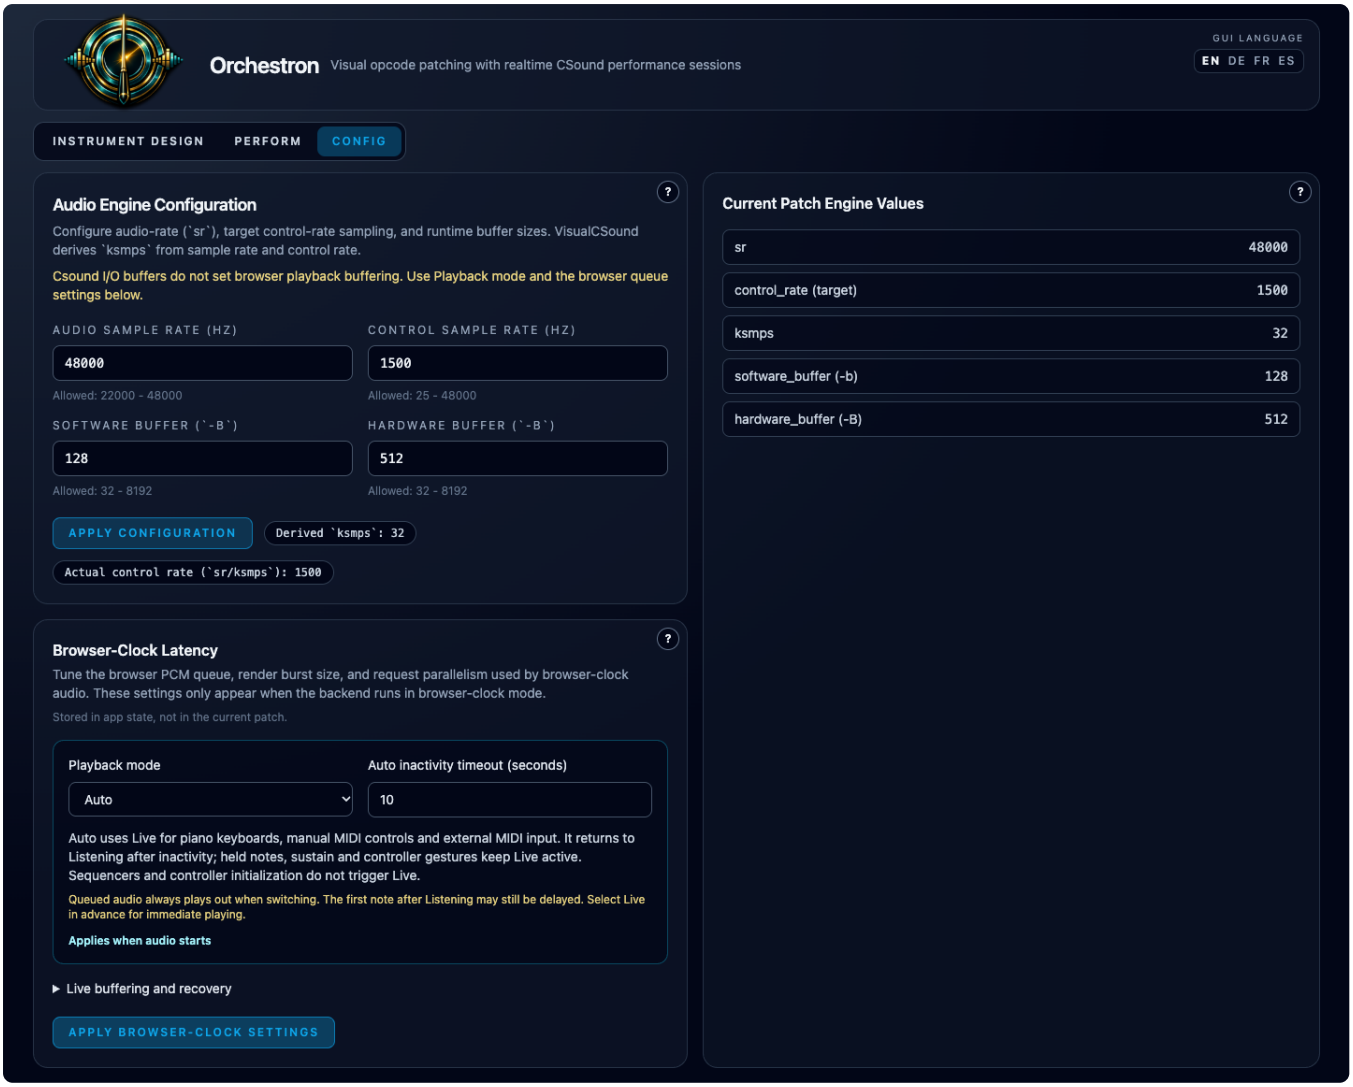

Choose a playback mode in Browser-Clock Latency; expand Live buffering and recovery for advanced tuning.

Why These Settings Matter

The browser keeps a queue of audio ready for playback. A larger queue absorbs brief rendering or network delays but increases the time between pressing a key and hearing a note. Playback begins when the startup queue target is reached. The Runtime panel reports queued audio, pending audio, underruns and Render Ratio so you can judge whether the settings suit your performance.

Lower values usually feel more immediate, but they also leave less headroom for:

- brief rendering delays
- network delays between browser and server
- temporary browser interruptions
- heavier patches that take longer to render

Higher values add safety margin, but manual playing and note feedback feel less direct.

Main Control Groups

Queue Targets

- `Steady Low Water` / `Steady High Water` define the normal browser PCM queue window after startup has stabilized.
- `Startup Low Water` / `Startup High Water` define the larger queue window used while the stream is priming or

recovering.

If you hear underruns after connect or after transport restarts, start by increasing the steady or startup queue targets slightly.

Recovery Headroom

- `Underrun Recovery Boost` adds extra queue headroom after an underrun is detected.
- `Max Underrun Boost` limits how large that temporary recovery buffer may grow.

These controls help the browser recover cleanly from brief dropouts without forcing a permanently larger steady-state latency.

Render Request Size And Parallelism

- `Max Blocks Per Request` limits how much audio the browser asks for in one render request.
- `Steady Parallel Requests`, `Startup Parallel Requests`, and `Recovery Parallel Requests` control how many render requests may stay in flight in each state.

Smaller requests can improve responsiveness, while larger or more parallel requests can make it easier to keep the queue full on slower systems.

Live Note Responsiveness

- `Immediate Note Render Blocks` sends a small urgent render burst after a live note-on event.
- `Immediate Note Render Cooldown` limits how often those urgent note-triggered bursts may happen.

These are the most direct controls for making manual keyboard or MIDI playing feel snappier in browser-clock mode.

Practical Tuning Order

1. Start with **Auto** (default), or choose **Listening** when live interaction is unnecessary.
2. If the stream crackles or underruns, check the Runtime diagnostics. Raise both low and high queue watermarks so refilling also starts earlier.
3. If latency still feels too high, lower the steady queue targets cautiously.
4. If live note response feels sluggish, reduce `Max Blocks Per Request` or adjust the immediate-note controls in small steps.
5. Apply the settings and test again with a realistic patch load.

Keep in mind that these settings are stored in app state for the current workspace. They are not saved into the patch itself.

They also work together with the engine settings from the patch:

- browser-clock settings tune queueing and request behavior
- patch engine `sr` and `ksmps` determine synthesis timing and the audio duration of a render request
- Csound I/O buffers `-b` and `-B` do not control browser queue depth or the browser device buffer on this headless path

Navigation: Up | Previous | Next

Saving, Workspace Recovery and Defaults

Navigation: Up | Previous | Next

Save library entries to keep and share your work. Workspace recovery helps you resume editing after a reload, but it does not replace saving an instrument or performance.

Two kinds of saved state

Action	What it keeps
Save in Instrument Design	The named instrument graph and its settings in the library
Save Performance	The rack, routes, mixer, instance settings and authored musical content
Automatic workspace recovery	Working context such as open instrument tabs, selected view, language and browser-clock preferences
Native Export	A portable editable instrument or performance bundle

Before leaving a session

1. Save instrument changes you want to keep in the library.
2. Save or discard open pattern and structure drafts, with the intended linked/local editing scope.
3. Save the performance.
4. Export important work as a native bundle for a separate backup.

Recovered instrument tabs can contain unsaved graph edits. Musical and structural drafts are session-only: resolve them before reloading.

What stays temporary

Selections, zoom, scroll and drafts survive view switches and panel collapse. New, Load and Import reset this temporary context. Auditions and arranger lane Mute/Solo are not saved music.

Compile and save

Compile checks a graph without saving it. Normal Save compiles first; templates may remain incomplete and are excluded from Perform. See Patch Toolbar and Instrument Tabs.

New instrument defaults

New opens the starter chooser; Empty patch supplies a blank graph. Defaults are 48 kHz audio, a target control rate of 1,500 Hz and a 32-sample synthesis block. See Audio Engine Settings before changing them.

Internal MIDI works without a keyboard. Select external MIDI explicitly when needed.

Older performances and large saves

Supported older files receive defaults for settings they did not contain. Review the mixer and Master routes after import, then save a fresh native backup. Keep the original file until you have checked playback.

If a save or import reports a size limit, retain your current draft, reduce unnecessary material or assets, and try again. Server limits protect storage; a failed save is not confirmation that your library has been updated.

For technical storage and migration details, see the persistence reference.

Navigation: Up | Previous | Next

Appendix: Optional Authoring Tools

Navigation: Up | Previous | Next

These optional tools let a coding agent create instruments and performances from a musical brief. You can use Orchestron entirely through its graphical interface; the tools are included here for readers who want to automate authoring or inspect the underlying workflow.

Detailed skill links open the repository sources so they are usable from the PDF as well. The example prompts refer to your local checkout.

Chapter Contents

- Agent CLI
- Orchestron Patch Creator Skill
- Orchestron Performance Creator Skill
- Orchestron Composition Creator Skill
- Sung Melody Transcription Skill

Use a Skill in a Coding Agent

A skill is a directory containing a `SKILL.md` instruction file, supporting references, and executable tooling. It guides the agent through the task; the command-line utility performs the operations.

1. Open the Orchestron repository in a coding agent with access to its files and a terminal.
2. Ask the agent to read the chosen `SKILL.md` under `integrations/skills/` and follow its linked references. Keep the whole skill directory available, including its `src`, `references`, and project files.
3. Give a musical brief, the desired instrument or performance name, and the backend URL. State whether you want a new item or an update to an existing one, and whether playback is wanted.
4. Ask the agent to validate the result and report what it saved, compiled, rendered, and actually listened to. Review the resulting instrument or performance in Orchestron.

Explicitly asking an agent to read the file does not require a product-specific skill installation. Automatic skill discovery and shortcut invocation depend on the agent; the prompts in the following chapters use ordinary language and repository paths.

For installing the packaged tools, follow the plugin installation guide.

The utilities use `uv` and require Python 3.13. Library operations, compilation, performance validation and live sessions need a reachable Orchestron backend, normally `http://localhost:8000/api`. Run the backend using the normal setup procedure (`make run` from the repository root), or provide an existing server's URL. Local patch-spec validation and graph generation can run without the backend. Composition generation, audition and score export also run locally without the backend. The utilities access the library through the API, not by editing SQLite directly.

Explore the architecture

The optional interactive architecture guide shows how an authored score becomes sound. Open the downloaded HTML in a browser; repository previews show its source. This is background reading, not a prerequisite for using Orchestron.

Navigation: Up | Previous | Next

Agent CLI

Navigation: Up | Previous | Next

Two utilities under `integrations/skills/` support repeatable authoring from YAML/JSON specs.

Utility	Purpose
`orchestron_patch_cli`	Generate instrument graphs, validate patch specs, and create, update or compile library patches.
`orchestron_cli`	Discover instruments, stage and save performances, configure routing and mixers, import/export, and manage optional live sessions.

Connect and Discover

Run these read-only checks from the repository root:

```
uv run --project integrations/skills/orchestron-patch-creator \
  orchestron_patch_cli --json health
uv run --project integrations/skills/orchestron-performance-creator \
  orchestron_cli --json patches list
```

Both default to `http://localhost:8000/api`; use `--api-url http://HOST:PORT/api` for another backend. `--json` returns structured results and retry hints; `--help` lists commands. Use `performances list` to discover saved performances, and inspect actual names, IDs and ports before editing.

Validate, Save and Optionally Play

For patches, `spec validate` checks a spec locally; `graph render` produces inspectable JSON. `patch create` and `patch update` write through the backend. Add `--compile` for compile preflight before the final write, then follow the patch skill's audio checks.

For performances, begin a staged edit, add instruments, routes and music, then run `edit validate` before `edit commit`. Validation checks references and routes with the backend; commit saves without starting playback.

For live testing, `edit create-runtime --start` creates and starts a CLI-owned session. `edit push-runtime` updates sequencer, arpeggiator, mixer and controller values with a matching compiled rack; `edit rebuild-runtime` handles instrument, patch or route changes. Continuous instruments need a restart for I-rate controller changes. Audible browser-clock playback requires a connected browser audio client.

Capabilities and Detailed References

Use matching current versions of the app and authoring tools. The performance utility preserves routing, mixer state, instance overrides, arrangement metadata and existing insert chains. Use the app to create or reorder inserts.

- **Sound settings:** discover and set per-instance values with `edit performance-controllers list/set/reset`. Patch specs define the controls; performance edits set their overrides. See controller authoring and instance settings.
- **Routing and mixing:** `edit instruments list`, `edit routes` and `edit mixer` expose ports, main/send paths and strip controls. Quote shell identifiers such as `'$master'`. The optional standard-effects preset requires suitable effect patches. See routing, mixer controls and the preset.
- **Patterns and expression:** discover tracks with `edit sequencers list`; use `edit step-timing` for early/late notes and `edit ratchets` for drum rolls and velocity ramps. Pads are numbered 1–8; steps start at 0. Seetiming, ratchets and score specs.

Use `edit add-score-sequencer --spec PATH` for dedicated Score Sequencers. Use the song workflow to place reusable Score patterns in clips.

Navigation: Up | Previous | Next

Orchestron Patch Creator Skill

Navigation: Up | Previous | Next

Use Orchestron Patch Creator to turn a sound description into an Instrument Design patch, or to revise an existing instrument or effect. The skill combines synthesis guidance with the `orchestron_patch_cli` utility.

What It Does

The agent chooses a suitable template family, writes a structured patch spec, validates it, and creates or updates the graph through the backend. Specs describe sources, envelopes, processing, output, input formulas and optional `perf_controller` settings for later adjustment in the rack. Generated instruments expose a named Stereo Output; route it through the performance mixer to Master when using the instrument in a performance.

For an existing graph, preserve stable node IDs and its control-flow structure. The high-level layer syntax does not create If/Switch branches automatically. See patch specs and graph compatibility for these details.

Example Agent Prompt

> Read `integrations/skills/orchestron-patch-creator/SKILL.md` and follow its references. Use the backend at `http://localhost:8000/api`. Create a new MIDI instrument named “Glass Pad” with a soft attack, a clear bell-like tone, gentle stereo movement and a long release. Expose attack and release as rack controls. Preserve existing instruments. Validate and compile the patch, render representative notes and release tails, and inspect both Mel and log-frequency STFT spectrograms. Report the saved name and which audio checks you performed. Do not start live playback.

The agent should check finite samples, clipping, both channels, velocity response and relevant controller settings. Spectrogram inspection supports verification; it does not replace listening or justify claiming that listening occurred.

Local Validation and Library Creation

After the agent has written `patch.yaml`, these commands illustrate the boundary between local validation and a backend write. Run them from the repository root:

```
uv run --project integrations/skills/orchestron-patch-creator \
  orchestron_patch_cli --json spec validate patch.yaml
uv run --project integrations/skills/orchestron-patch-creator \
  orchestron_patch_cli --json patch create patch.yaml --compile
```

Use `patch update PATCH_ID patch.yaml --compile` for an explicitly selected existing patch. A failed validation or compile preflight must be corrected before updating it. See the authoring workflow and audio validation procedure for the full process.

Navigation: Up | Previous | Next

Orchestron Performance Creator Skill

Navigation: Up | Previous | Next

Use Orchestron Performance Creator to compose, edit, import, validate and save performances using instruments in the library. It guides an agent through the `orchestron_cli` utility, with references for musical form, harmony, rhythm, routing and mixing.

What It Does

The agent discovers available patches, starts a staged edit, assigns rack instruments, connects their audio paths and adds melodic, drummer, controller or arpeggiator material. It can write a YAML/JSON score for longer arrangements, adjust per-instance sound settings, program note timing and drum rolls, and configure mixer levels and sends. The built-in Master needs no extra speaker instrument.

Before saving, the agent validates the staged performance against the backend. Commit saves the result; creating or updating a live session is a separate, optional step. Existing insert chains are retained, while insert creation and reordering use the app.

The skill selects a device for each part: Drummer Sequencers for drum-kit grooves and fills, Melodic Sequencers for repeating bass, riffs and chord patterns, and Score Sequencers for long passages or independently timed overlapping notes. Song length alone does not require Score Sequencers: reusable pads and arranger groups build longer sections. Explicit user choices and existing devices are preserved unless a change is requested.

Custom Drummer patterns support individual MIDI keys, step positions and velocities, with different patterns on each pad. Use `edit add-drummer --hits "35@0:115 38@4:98" --pad-hits "2=35@0:120 54@7:65"` or `hits` lists in a YAML/JSON score. Steps start at zero, velocities range from 1 to 127, and new MIDI keys add rows automatically. An empty hits list creates silence. See custom drum patterns for examples, timing and rolls.

Example Agent Prompt

> Read `integrations/skills/orchestron-performance-creator/SKILL.md` and follow its references. Use the backend at `http://localhost:8000/api`. First list available instruments and choose a drum kit, bass and pad from the library. Create a new 32-bar performance named “Night Drive” at 120 BPM in 4/4 and C minor: four bars of drums, an evolving bass section, a pad-led middle, and a short ending. Keep the bass and pad harmonically compatible, route the instruments to Master, and set balanced levels. Validate and save it without changing existing performances or starting playback. Report the instrument choices and arrangement.

If suitable instruments are missing, the agent should report that and use the Patch Creator skill when creating the missing sounds is part of the request. Avoid assuming that example patch names exist in every library.

Stage and Save

From the repository root, the agent starts a new edit with:

```
uv run --project integrations/skills/orchestron-performance-creator \
  orchestron_cli --json edit begin --new --name "Night Drive"
```

After populating the staged rack, routes and music, validate and save:

```
uv run --project integrations/skills/orchestron-performance-creator \
  orchestron_cli --json edit validate
uv run --project integrations/skills/orchestron-performance-creator \
  orchestron_cli --json edit commit
```

See composition workflow and score specs for detailed authoring instructions.

Use an Accepted Composition

Composition Creator develops melodic candidates through listening and selection, then exports `orchestron-score.json`. Performance Creator imports score versions 1 and 2, including per-note velocity and local meter/subdivision/beat-rate timing. Version 2 checks explicit event overlap and pad boundaries. Polyphonic parts become coordinated voice tracks on the instrument's MIDI channel, preserving normal editable note/chord cells.

Start a new staged performance, bind the score's channels to suitable instruments, connect their outputs, then use `edit apply-score /absolute/path/orchestron-score.json`. Apply once: the command appends tracks. Validate and save using the commands above. The score is separate from a native `.orch.json` performance bundle.

Combine Both Skills

Ask the agent to create and validate a new sound with Patch Creator, then read Performance Creator and discover the saved patch in the backend library. It can add that patch to a new performance, connect its Stereo Output to Master, write complementary parts, validate and save. Include the desired names, musical constraints and playback preference in the same request. Creating the patch, saving the performance and starting playback are separate operations.

Navigation: Up | Previous | Next

Orchestron Composition Creator Skill

Navigation: Up | Previous | Next

Use Orchestron Composition Creator to explore melodies by listening, choosing and developing related ideas. It uses music21 and an arvo compatibility copy, runs independently of the Orchestron backend, and exports a machine-readable score for Performance Creator.

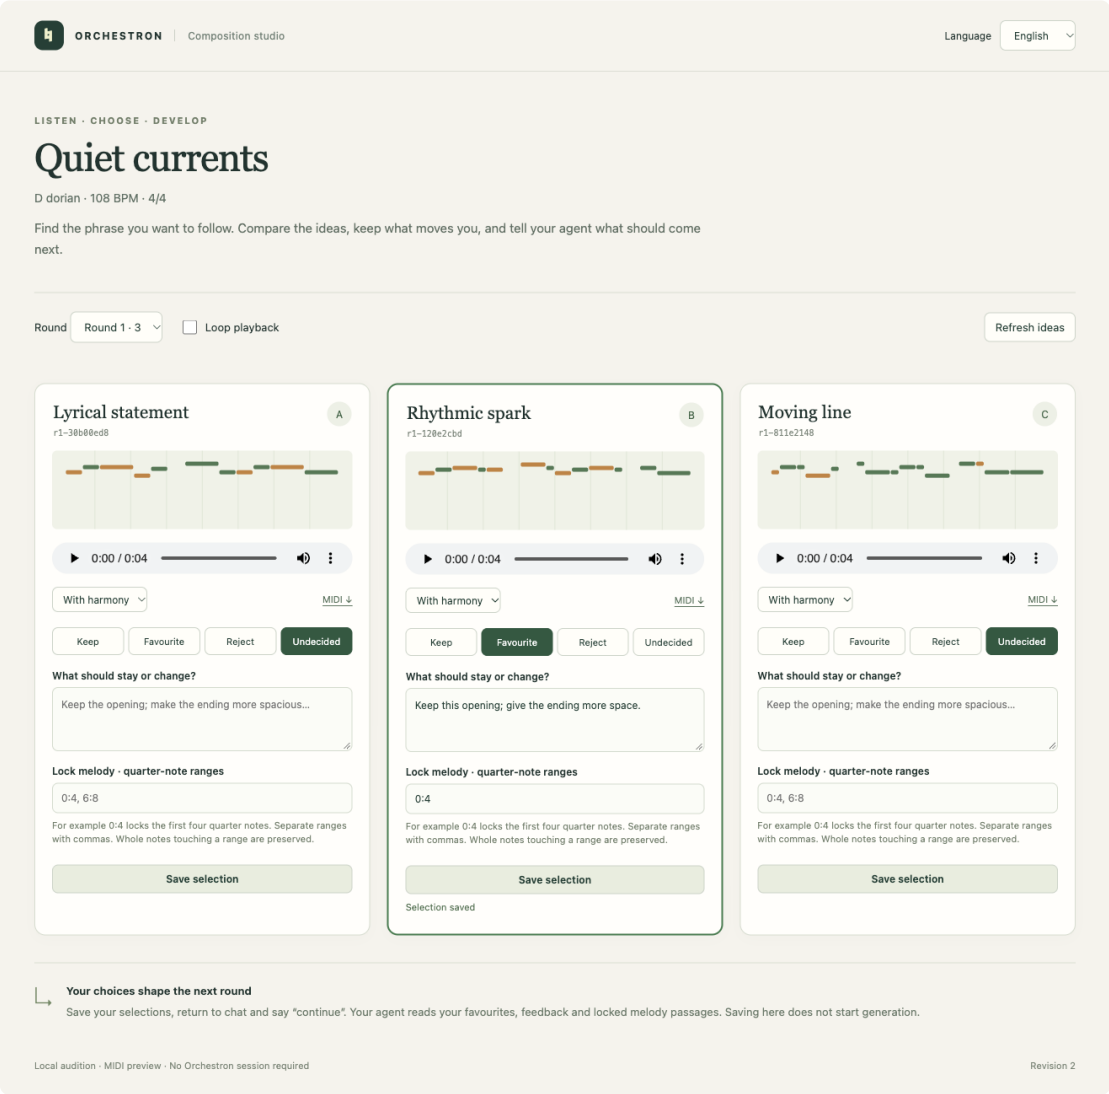

Try It With an Agent

> Read `integrations/skills/orchestron-composition-creator/SKILL.md`. Propose three two-bar melodies in D dorian at 108 BPM, with a calm but slightly syncopated character. Include quiet harmony for comparison. Open the audition page and wait for my choice. Develop the melody I select, preserving any passages I lock. When I accept the phrases, export a score for Orchestron Performance Creator.

The agent creates a session and generates a small shortlist. Each candidate has a contour display and a player. Compare **With harmony** and **Melody only**, loop the phrase, or download its MIDI. All candidates use consistent preview sound and gain. The default simple MIDI synthesizer needs no extra sound installation; FluidSynth with a local SoundFont is optional. Preview timbre differs from the instrument eventually chosen in Orchestron.

Choose and Continue

Mark **Keep**, **Favourite**, **Reject** or **Undecided**, add feedback, and click **Save selection**. To retain an opening, enter a melody lock such as `0:4` for the first four quarter notes. Multiple ranges use commas. Complete notes touching a locked range remain unchanged. Locks protect the melody; ordinary development also leaves its existing harmony intact.

Return to chat and say **“continue”**. The agent reads the saved selections and creates related ideas. It can transpose or invert a motif, develop cyclic or additive material, combine one accepted idea's pitches with another's rhythm, or append a continuation. Feedback guides the agent; the command-line generator does not interpret prose by itself. Saving in the page does not start generation.

Use **Refresh ideas** to see the next round. Earlier rounds, comments and choices remain available. The page works in English, German, French and Spanish. Unsaved changes are marked. A conflicting newer revision must be refreshed before saving. The session is stored locally and can be resumed after closing the page.

Compare Endings and Harmonic Settings

When you like an opening but are unsure of its ending, ask the agent to preserve that opening and compare four alternatives: **Settled ending**, **Suspended ending**, **Higher peak** and **Spacious ending**. The first descends to a held tonic; the second holds scale degree two; the third reaches a new high point; the fourth arrives early and leaves silence. These are intentions to test by listening.

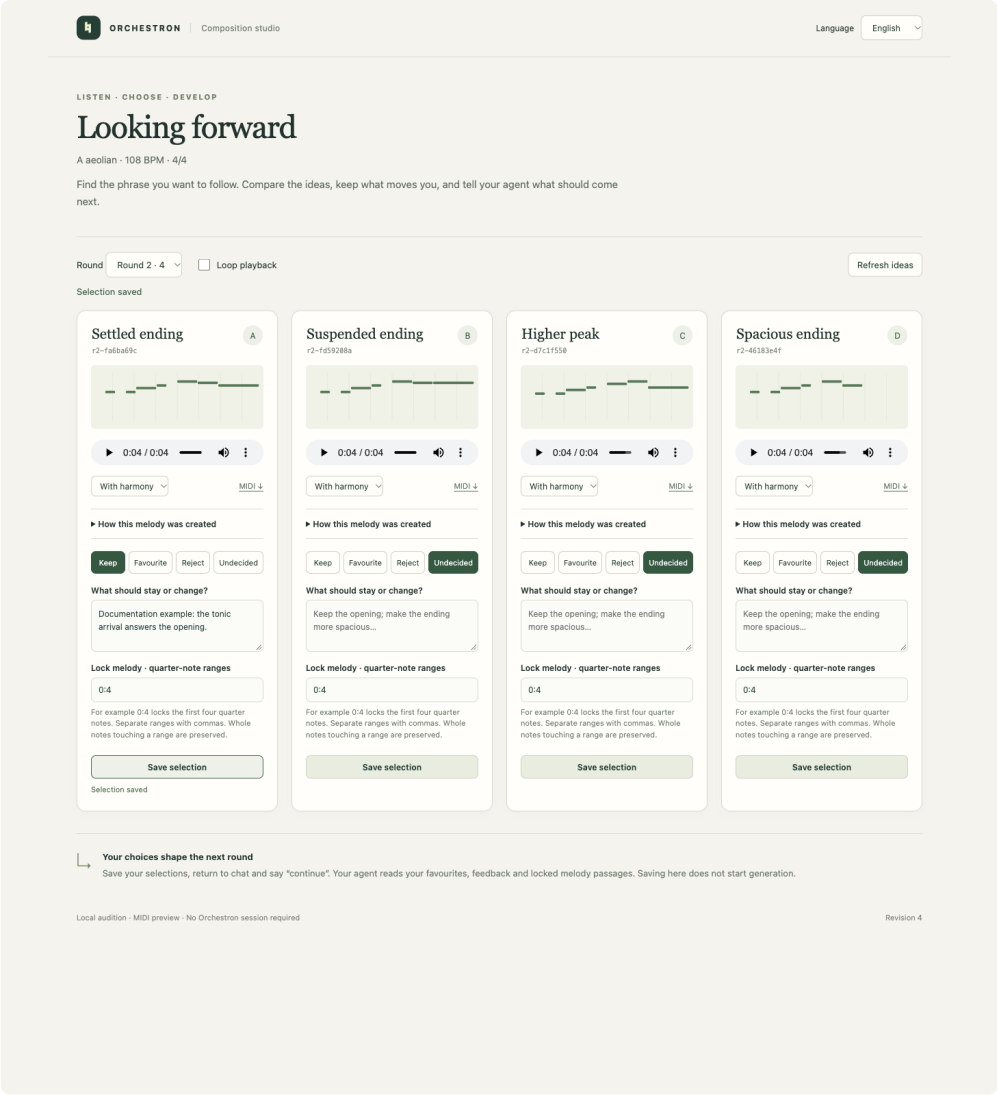

The opening's notes, dynamics and rests stay fixed. A sustained note crossing the chosen boundary is retained whole. Existing locks also apply to the tail; a conflict stops the complete comparison without removing those locks. All four

candidates are saved in one undoable round and wait for your selection. Compare melody alone and with harmony, then record what works and what you want to change.

After you choose an ending, ask to compare two harmonic settings of it. The complete melody remains identical, while the selected accompaniment part changes. The default compares tonic and mediant triads with their bass notes; in the included A-minor study these are Am and C. Solo previews are identical. Use the mix previews to hear the difference, and keep listening levels consistent.

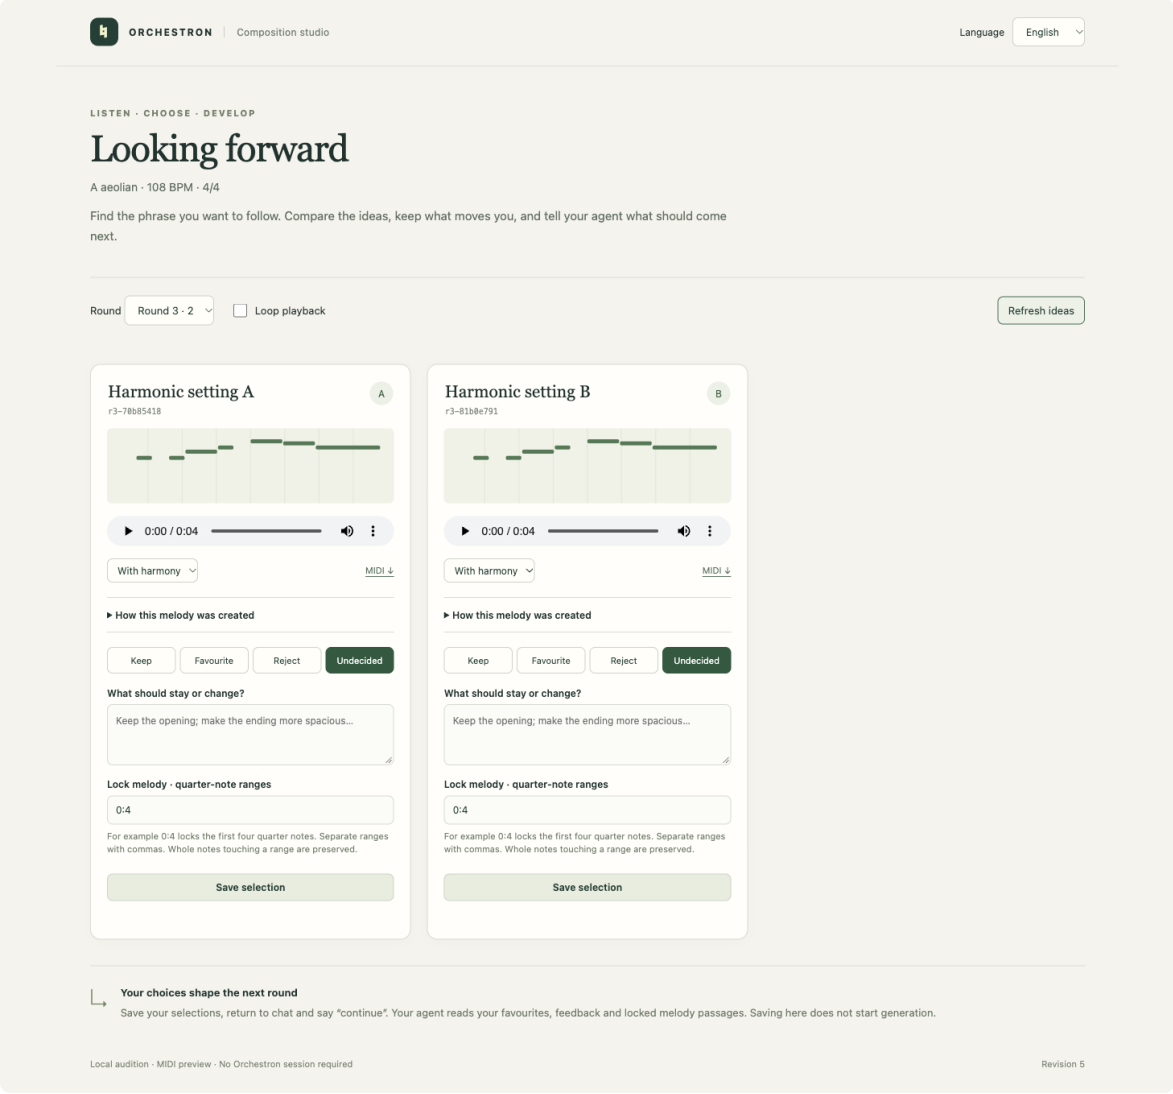

For an original two-bar starting motif, run the study helper from the repository root:

```
uv run --locked --project integrations/skills/orchestron-composition-creator \
  python integrations/skills/orchestron-composition-creator/scripts/start_motif_study.py \
  /absolute/path/to/my-study
```

It imports and renders one pending seed, then waits for your choice. The agent uses `compare-endings SESSION --from ACCEPTED_ID --start 4` , followed after selection by `compare-harmony SESSION --from ACCEPTED_ENDING_ID` . Read the comparison reference for custom harmonic settings, exact timing limits and preservation guarantees. Export only the chosen setting; the alternatives are comparisons, not song sections.

Follow a Melody's Creation

Open **How this melody was created** beneath a candidate's player. The panel shows its original material, earlier versions and contributing candidates in creation order. Expand an ancestor or a step to inspect arvo function names, resolved arguments and input/output material. It also records pitch fitting, rhythm combinations, accompaniment and continuation, so the history

explains the final melody rather than just the first transformation.

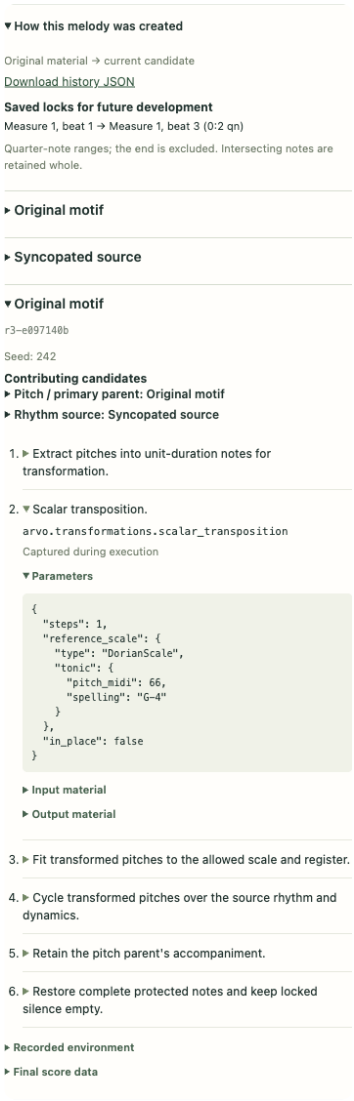

Parent roles identify where pitches and rhythm came from. Saved feedback and reviews are captured when a child is created. **Locks applied during creation** show the ranges used for that operation, with measure/beat positions and exact quarter-note values. Retained notes, protected silence and discarded generated notes are available in the details. **Saved locks for future development** are the current candidate's review and can differ from those historical locks.

The panel loads on demand without interrupting playback or unsaved reviews. **Download history JSON** saves the complete record. CLI users can read the same record with `orchestron_compose trace SESSION CANDIDATE_ID .`

Built-in generation records its operations automatically. Custom authoring scripts use the trace recorder and import their records with `import-score --trace FILE` . Authored descriptions are distinguished from captured function calls. Earlier candidates and unrecorded imports show gaps explicitly; the tool does not reconstruct their missing history.

Local Commands

From the repository root:

```
uv run --locked --project integrations/skills/orchestron-composition-creator \
  orchestron_compose init output/composition/quiet-currents \
  --title "Quiet currents" --key D --mode dorian --tempo 108 --bars 2
uv run --locked --project integrations/skills/orchestron-composition-creator \
  orchestron_compose generate output/composition/quiet-currents
uv run --locked --project integrations/skills/orchestron-composition-creator \
  orchestron_compose serve output/composition/quiet-currents
```

Open the printed loopback URL. Stop the server with Ctrl-C when finished; saved choices remain on disk. The tools require Python 3.13 and `uv` on macOS or Linux. Dependencies are downloaded on first use and locked for repeatability.

Export and Realise the Song

Ask the agent to arrange accepted phrases in a specified order. Export includes `composition.json`, `orchestron-score.json`, MIDI, MusicXML and playable WAV files. The canonical score retains exact rational timing, note dynamics, phrase sections and accepted-candidate provenance. The exporter finds exact supported grids and creates editable voice tracks and pads; it reports unsupported timing or capacity instead of silently changing the music. Initial scope is fixed tempo/meter and standard MIDI pitches, with straight/triplet grids and supported subdivisions. The canonical score also retains full recorded creation history and phrase placements, including repetitions, when imported into another composition session. Performance Creator's score format and native performance bundles are unchanged.

Performance Creator then discovers instruments, maps the score's channels, sets routing and mixer levels, applies the score, validates and saves the performance. These score files are not native Orchestron bundles and should not be passed to the application's Import dialog.

Navigation: Up | Previous | Next

Sung Melody Transcription Skill

Navigation: Up | Previous

This plugin skill transcribes a sung or hummed monophonic melody into MIDI, note JSON and a playable piano WAV. Rhythmic claps can establish pulse timing; they are treated separately from vocal attacks. You can supply tempo, meter, bar count and an ending preference. The default meter is 4/4 with a sixteenth-note grid, and the recorded ending is preserved.

The skill searches all major and natural-minor pitch collections while fitting one tuning offset within ±50 cents. It reports tied collections, including relative major/minor, rather than claiming a unique key. Conservative cleanup merges false splits and removes brief unstable onset scoops and release dips; the original pitch track and every editorial change remain available for review.

Run and review

Ask a coding agent to read the skill source. For example:

```
Transcribe this sung melody as four bars in 4/4. Use the claps for tempo and
leave a short rest at the end. Return a piano preview, MIDI, notes and an edit
summary so I can check the repeated notes and phrasing.
```

The skill requires Python 3.11–3.13, `uv` , the RVC RMVPE checkpoint, FluidSynth and a piano SoundFont. The checkpoint downloader verifies its hash. The setup and command reference explains asset configuration and optional inputs. Python dependencies are locked within the skill; transcription runs locally.

Review `preview.wav` against the recording, particularly repeated attacks, short gestures, rests and the ending. `melody.mid` and `notes.json` contain the quantized result; `pitch.json` retains continuous pitch and voicing evidence; `edits.txt` summarizes the changes. Detailed before/after edits are in `notes.json` . Requested bar counts define the exact MIDI phrase length, including silence. Clipped or excluded notes are reported. FluidSynth's acoustic release tail may continue after the final MIDI note-off.

Claps do not identify a musical downbeat, meter or count-in by themselves. Supply beat zero when the phrase begins after a count-in. Without supplied tempo or a reliable clap pulse, 120 BPM is explicitly marked as a review default. Dense or ambiguous material may need a finer grid or manual review. The output is a draft transcription, not an automatic confirmation of the intended melody.

Plugin packaging and independent previews

The source lives under `integrations/skills/sung-melody-transcription/` and is included in the Orchestron plugin ZIP. Install the plugin or ask the agent to read the skill directly. Transcription and FluidSynth piano previews need no Orchestron backend and do not create instruments, performances or native bundles. Any later Orchestron import or instrument assignment is a separate workflow after reviewing the transcription.

Navigation: Up | Previous